

# Autonomous Search and Rescue Drone with Onboard AI-Based Casualty Detection

Dmytro Chuprynyuk, *et al.*

AENGM0074 Group Design Project — University of Bristol — March 2026

# Contents

|                                                                   |           |
|-------------------------------------------------------------------|-----------|
| <b>Executive Summary</b>                                          | <b>11</b> |
| <b>Introduction</b>                                               | <b>12</b> |
| <b>1 Design Rationale</b>                                         | <b>14</b> |
| 1.1 STEEPLE Analysis . . . . .                                    | 14        |
| 1.2 Search Parameter Optimisation . . . . .                       | 14        |
| 1.3 Parameter Derivation Chain . . . . .                          | 16        |
| 1.4 Hardware Platform Selection . . . . .                         | 16        |
| 1.4.1 Companion Computer Trade Study . . . . .                    | 17        |
| 1.4.2 Communication Architecture Trade Study . . . . .            | 17        |
| 1.5 Software Architecture Rationale . . . . .                     | 18        |
| 1.6 Detection Model Selection . . . . .                           | 18        |
| 1.7 Coverage Path Planning . . . . .                              | 18        |
| 1.8 State Machine Design . . . . .                                | 19        |
| 1.9 Autonomy Level Justification . . . . .                        | 19        |
| 1.10 Decisions That Changed: Evidence-Based Corrections . . . . . | 20        |
| 1.11 Why Not ROS . . . . .                                        | 21        |
| 1.12 Camera Selection: Global vs Rolling Shutter . . . . .        | 21        |
| 1.13 Field Day Adaptation Under Uncertainty . . . . .             | 21        |
| <b>2 System Description</b>                                       | <b>21</b> |
| 2.1 Hardware Platform . . . . .                                   | 21        |
| 2.2 Software Architecture . . . . .                               | 22        |
| 2.3 Computer Vision Pipeline . . . . .                            | 24        |
| 2.4 Search Pattern Generation . . . . .                           | 26        |
| 2.5 Mission State Machine . . . . .                               | 26        |
| 2.6 Target Localisation . . . . .                                 | 26        |
| 2.7 Ground Station . . . . .                                      | 27        |
| 2.8 Simulation Framework . . . . .                                | 28        |
| <b>3 Requirements Verification</b>                                | <b>29</b> |
| <b>4 Evaluation</b>                                               | <b>30</b> |
| 4.1 Plus/Delta Summary . . . . .                                  | 30        |
| 4.2 Discussion of Key Findings . . . . .                          | 30        |
| 4.3 Testing Framework Summary . . . . .                           | 31        |
| <b>References</b>                                                 | <b>32</b> |
| <b>Appendix Guide</b>                                             | <b>34</b> |
| <b>A State Transition Table</b>                                   | <b>35</b> |
| <b>B Configuration Parameters</b>                                 | <b>37</b> |
| <b>C Model Training and Dataset</b>                               | <b>38</b> |
| C.1 Training Data Strategy . . . . .                              | 38        |
| C.2 Domain Randomisation . . . . .                                | 38        |
| C.3 Augmentation Pipeline . . . . .                               | 39        |
| C.4 Synthetic Data Generation Pipeline . . . . .                  | 39        |
| C.5 Real Data Labelling . . . . .                                 | 39        |

|          |                                                  |           |
|----------|--------------------------------------------------|-----------|
| C.6      | Negative Mining                                  | 39        |
| C.7      | Data Composition Analysis                        | 40        |
| C.8      | Training Pipeline and Convergence                | 40        |
| C.9      | Results                                          | 40        |
| C.10     | Model Export Pipeline                            | 41        |
| C.11     | Deployment                                       | 41        |
| C.12     | Limitations                                      | 42        |
| <b>D</b> | <b>Safety and Risk Management</b>                | <b>42</b> |
| D.1      | Geofence Implementation                          | 43        |
| D.2      | Failure Mode Analysis                            | 43        |
| D.3      | Operator-in-the-Loop                             | 43        |
| D.4      | Emergency Procedures                             | 43        |
| D.5      | Risk Assessment                                  | 44        |
| D.6      | Regulatory Considerations                        | 44        |
| <b>E</b> | <b>Testing Methodology</b>                       | <b>44</b> |
| E.1      | V-Model Verification Approach                    | 44        |
| E.2      | Five-Tier Progressive Testing                    | 45        |
| E.3      | Test Script Organisation                         | 45        |
| E.4      | Dry-Run Validation                               | 46        |
| E.5      | DJI Video Analysis                               | 47        |
| E.6      | Stress Testing and Failure Scenarios             | 47        |
| E.7      | Building Confidence Through Progressive Evidence | 47        |
| E.8      | Team Roles in the Testing Process                | 48        |
| E.9      | Comparison with Industry Testing Standards       | 49        |
| E.10     | Expected Outcomes from Real Flight Testing       | 49        |
| E.11     | Preflight Checklist                              | 50        |
| <b>F</b> | <b>Field Testing and Results</b>                 | <b>50</b> |
| F.1      | Bench Testing Results                            | 50        |
| F.2      | FOV Calibration                                  | 50        |
| F.3      | Lens Distortion Calibration                      | 51        |
| F.4      | Inference Benchmarks                             | 51        |
| F.5      | Post-Hoc Video Analysis                          | 51        |
|          | F.5.1 Detection Performance                      | 51        |
|          | F.5.2 GPS Estimation Accuracy                    | 52        |
| F.6      | Dry-Run Validation                               | 52        |
| F.7      | Discussion                                       | 52        |
| <b>G</b> | <b>Future Work</b>                               | <b>53</b> |
| G.1      | Inference Acceleration                           | 53        |
| G.2      | Detection Improvements                           | 53        |
| G.3      | Multi-Target Support                             | 54        |
| G.4      | Thermal Imaging                                  | 54        |
| G.5      | Communication                                    | 54        |
| G.6      | Swarm Operations                                 | 54        |
| <b>H</b> | <b>Edge Inference Architecture</b>               | <b>54</b> |
| H.1      | Compute Platform Constraints                     | 55        |
| H.2      | Backend Selection                                | 55        |
| H.3      | Quantisation Trade-offs                          | 55        |
| H.4      | Pipeline Architecture                            | 56        |

|          |                                                         |           |
|----------|---------------------------------------------------------|-----------|
| H.5      | Video Streaming Architecture . . . . .                  | 56        |
| H.6      | Benchmark Methodology . . . . .                         | 57        |
| H.7      | Results and Projected Improvements . . . . .            | 57        |
| <b>I</b> | <b>Computer Vision: Extended Analysis</b>               | <b>57</b> |
| I.1      | Detection Pipeline: Step-by-Step Timing . . . . .       | 58        |
| I.1.1    | Stage 1: Camera Capture . . . . .                       | 58        |
| I.1.2    | Stage 2: Lens Undistortion . . . . .                    | 58        |
| I.1.3    | Stage 3: Preprocessing . . . . .                        | 58        |
| I.1.4    | Stage 4: TFLite Inference . . . . .                     | 58        |
| I.1.5    | Stage 5: Postprocessing . . . . .                       | 58        |
| I.1.6    | Stage 6: Visualisation . . . . .                        | 59        |
| I.2      | Comprehensive Benchmark Results . . . . .               | 59        |
| I.3      | Mission Runtime Analysis . . . . .                      | 59        |
| I.3.1    | Ground Footprint Geometry . . . . .                     | 59        |
| I.3.2    | Frames per Flyover . . . . .                            | 60        |
| I.3.3    | Frame Overlap Analysis . . . . .                        | 60        |
| I.4      | Detection Performance vs Altitude . . . . .             | 60        |
| I.4.1    | Pixel Size Model . . . . .                              | 60        |
| I.4.2    | Altitude Performance Table . . . . .                    | 60        |
| I.5      | Model Training Results . . . . .                        | 61        |
| I.6      | Dual Backend Architecture . . . . .                     | 62        |
| I.6.1    | Backend Selection Logic . . . . .                       | 62        |
| I.6.2    | Backend Abstraction Benefits . . . . .                  | 63        |
| I.7      | Model Swapping and Field Deployment . . . . .           | 63        |
| I.8      | Detection Speed vs Drone Speed . . . . .                | 63        |
| I.9      | Inference Latency Breakdown . . . . .                   | 64        |
| <b>J</b> | <b>Video Streaming and Ground Station Architecture</b>  | <b>65</b> |
| J.1      | Streaming Protocol Selection . . . . .                  | 65        |
| J.2      | Threading Model and Frame Buffering . . . . .           | 66        |
| J.3      | Headless Operation . . . . .                            | 66        |
| J.4      | Port Allocation and Service Separation . . . . .        | 67        |
| J.5      | Bandwidth Estimation . . . . .                          | 67        |
| J.6      | Ground Station Alternatives Considered . . . . .        | 67        |
| <b>K</b> | <b>GPS Target Estimation: A Deep Dive</b>               | <b>68</b> |
| K.1      | Strategy Roadmap . . . . .                              | 68        |
| K.2      | Problem Statement . . . . .                             | 69        |
| K.3      | Detection-to-GPS Pipeline Overview . . . . .            | 69        |
| K.4      | Pixel-to-GPS Mathematics . . . . .                      | 70        |
| K.4.1    | Ground Sampling Distance . . . . .                      | 71        |
| K.4.2    | Body-Frame Projection . . . . .                         | 71        |
| K.4.3    | Heading Rotation . . . . .                              | 71        |
| K.4.4    | WGS84 Conversion . . . . .                              | 71        |
| K.5      | Image Resize Pipeline and Coordinate Mapping . . . . .  | 71        |
| K.6      | Error Budget . . . . .                                  | 72        |
| K.7      | Roll/Pitch Attitude Error . . . . .                     | 73        |
| K.7.1    | Magnitude of the Error . . . . .                        | 73        |
| K.7.2    | Why the System Tolerates This Error . . . . .           | 74        |
| K.7.3    | Attitude-Compensated Estimation (Implemented) . . . . . | 74        |
| K.7.4    | Comparison with Published Approaches . . . . .          | 76        |

|        |                                                                    |    |
|--------|--------------------------------------------------------------------|----|
| K.8    | Centering as Multi-Error Mitigation . . . . .                      | 77 |
| K.9    | Three-Level Estimation Architecture . . . . .                      | 78 |
| K.9.1  | Level 1: Fully Passive Estimation (No Drone Response) . . . . .    | 78 |
| K.9.2  | Level 2: Stop-and-Look (Hover Without Centering) . . . . .         | 78 |
| K.9.3  | Level 3: Visual Centering (Implemented) . . . . .                  | 79 |
| K.9.4  | Comparative Evaluation . . . . .                                   | 80 |
| K.10   | Altitude Effect on Estimation Accuracy . . . . .                   | 81 |
| K.11   | Accuracy Progression Through Mission Phases . . . . .              | 81 |
| K.12   | GPS Averaging Duration Analysis . . . . .                          | 82 |
| K.12.1 | GPS Noise Characteristics . . . . .                                | 82 |
| K.12.2 | Averaging Improvement Model . . . . .                              | 82 |
| K.12.3 | Why 10 Seconds . . . . .                                           | 83 |
| K.13   | Fusion Strategies . . . . .                                        | 84 |
| K.13.1 | Rolling Weighted Average (Window = 50) . . . . .                   | 84 |
| K.13.2 | Cumulative Weighted Average (All Time) . . . . .                   | 84 |
| K.13.3 | Kalman Filter (Static Target) . . . . .                            | 84 |
| K.13.4 | Inverse-Variance Weighting . . . . .                               | 84 |
| K.14   | Spatial Clustering . . . . .                                       | 85 |
| K.15   | Comparison of Methods . . . . .                                    | 85 |
| K.16   | Empirical Validation . . . . .                                     | 85 |
| K.16.1 | DJI Video Replay (The “Bullseye Test”) . . . . .                   | 85 |
| K.16.2 | Simulator Validation . . . . .                                     | 87 |
| K.17   | Detailed Per-State Accuracy Analysis . . . . .                     | 87 |
| K.18   | Four-Phase Detection-to-Landing Flow . . . . .                     | 88 |
| K.18.1 | Phase 1: Initial Detection During Search (Tilted Camera) . . . . . | 88 |
| K.18.2 | Phase 2: Approach and Continuous Refinement . . . . .              | 88 |
| K.18.3 | Phase 3: Hover Lock (Nadir Confirmation) . . . . .                 | 89 |
| K.18.4 | Phase 4: Operator Verification and Landing Decision . . . . .      | 89 |
| K.18.5 | Why Four Phases Are Necessary . . . . .                            | 89 |
| K.18.6 | The Progressive Refinement Trajectory . . . . .                    | 89 |
| K.18.7 | Design Insight . . . . .                                           | 90 |
| K.18.8 | False Positive Handling . . . . .                                  | 90 |
| K.19   | Cooperative Vision-Flight Architecture . . . . .                   | 91 |
| K.19.1 | Three Mechanisms of Flight-Aided Improvement . . . . .             | 92 |
| K.19.2 | Static Camera vs. Cooperative System . . . . .                     | 92 |
| K.20   | Why Inverse-Variance Weighting Won . . . . .                       | 93 |
| K.21   | Empirical Method Comparison Against Ground Truth . . . . .         | 93 |
| K.21.1 | Ground Truth Setup . . . . .                                       | 93 |
| K.21.2 | Methods Tested . . . . .                                           | 93 |
| K.21.3 | Why Centering + GPS Averaging Won . . . . .                        | 94 |
| K.21.4 | CEP <sub>50</sub> Comparison . . . . .                             | 94 |
| K.21.5 | Convergence Behaviour . . . . .                                    | 94 |
| K.22   | Offset Landing Application . . . . .                               | 95 |
| K.23   | Multi-Pass Observation and Heading Diversity . . . . .             | 95 |
| K.23.1 | The Lawnmower Pattern Advantage . . . . .                          | 95 |
| K.23.2 | How Many Passes Are Enough? . . . . .                              | 96 |
| K.23.3 | Strip Spacing and Target Visibility . . . . .                      | 96 |
| K.23.4 | Multi-Pass Averaging vs. Centering . . . . .                       | 97 |
| K.24   | Passive Estimation: Level 1 GPS Without Drone Response . . . . .   | 97 |
| K.24.1 | Passive Estimation Toolkit . . . . .                               | 97 |
| K.24.2 | Why IVW Suits Passive Mode . . . . .                               | 98 |

|          |                                                                           |            |
|----------|---------------------------------------------------------------------------|------------|
| K.24.3   | Convergence Behaviour in Passive Mode . . . . .                           | 98         |
| K.24.4   | Passive vs. Active Estimation: When Each Mode Excels . . . . .            | 98         |
| K.24.5   | Implementation in <code>passive_watch.py</code> . . . . .                 | 99         |
| <b>L</b> | <b>Progressive Testing Methodology</b>                                    | <b>99</b>  |
| L.1      | Philosophy: Never Skip a Tier . . . . .                                   | 99         |
| L.2      | The Five Tiers in Detail . . . . .                                        | 100        |
| L.2.1    | Tier 1: SITL Simulation (Zero Hardware Risk) . . . . .                    | 100        |
| L.2.2    | Tier 2: Bench Test (Pi + Cube + Camera, No Props) . . . . .               | 100        |
| L.2.3    | Tier 3: Passive Flight (Pilot Flies, Pi Watches, Zero Commands) . . . . . | 101        |
| L.2.4    | Tier 4: Autonomous Search with Log-Only Detection . . . . .               | 101        |
| L.2.5    | Tier 5: Full Autonomous Mission . . . . .                                 | 101        |
| L.3      | Test Script Organisation . . . . .                                        | 101        |
| L.4      | Experiment Scripts for Quantitative Validation . . . . .                  | 102        |
| L.5      | Defect Discovery by Tier . . . . .                                        | 102        |
| L.6      | V-Model Mapping . . . . .                                                 | 103        |
| L.7      | Cost–Risk Gradient . . . . .                                              | 103        |
| L.8      | Summary . . . . .                                                         | 103        |
| <b>M</b> | <b>Field Day: Adaptation Under Uncertainty</b>                            | <b>103</b> |
| M.1      | Setup: Three Terminals, One Aircraft . . . . .                            | 104        |
| M.2      | What We Tested . . . . .                                                  | 104        |
| M.3      | FOV Calibration: Catching a 22% Error . . . . .                           | 104        |
| M.4      | Camera Colour Discovery . . . . .                                         | 105        |
| M.5      | Lens Calibration . . . . .                                                | 105        |
| M.6      | Benchmark Results . . . . .                                               | 105        |
| M.7      | Weather Cancellation and the Go/No-Go Decision . . . . .                  | 105        |
| M.8      | The Pivot: Evidence-Based Adaptation . . . . .                            | 106        |
| M.9      | DJI Video Analysis: A Flight-Data Proxy . . . . .                         | 106        |
| M.10     | Team Coordination Under Field Conditions . . . . .                        | 106        |
| M.11     | Lessons for the Next Field Day . . . . .                                  | 107        |
| <b>N</b> | <b>Contingency Planning and Deliverable Tiers</b>                         | <b>107</b> |
| N.1      | Minimum Viable Deliverable . . . . .                                      | 107        |
| N.2      | Target Deliverable . . . . .                                              | 107        |
| N.3      | Stretch Goals . . . . .                                                   | 108        |
| N.4      | How the Tier Structure Evolved . . . . .                                  | 108        |
| N.5      | Contingency Plans . . . . .                                               | 109        |
| N.6      | Test Script Mapping . . . . .                                             | 109        |
| <b>O</b> | <b>Test Script Reference</b>                                              | <b>109</b> |
| O.1      | Organisation by Category . . . . .                                        | 109        |
| O.2      | Flight Test Progression . . . . .                                         | 110        |
| O.3      | Hardware Tests . . . . .                                                  | 111        |
| O.4      | Calibration Tests . . . . .                                               | 111        |
| O.5      | Experiment Scripts . . . . .                                              | 112        |
| O.6      | Diagnostics . . . . .                                                     | 113        |
| O.7      | Script-to-Requirement Mapping . . . . .                                   | 113        |
| O.8      | What Was Tested on Real Hardware . . . . .                                | 114        |
| O.9      | Remaining Test Gaps . . . . .                                             | 114        |

|          |                                                                     |            |
|----------|---------------------------------------------------------------------|------------|
| <b>P</b> | <b>Sensor Calibration</b>                                           | <b>115</b> |
| P.1      | Field-of-View Calibration . . . . .                                 | 115        |
| P.2      | Lens Distortion Calibration . . . . .                               | 115        |
| P.3      | Camera Colour-Space Discovery . . . . .                             | 116        |
| P.4      | DJI Video FOV Calibration . . . . .                                 | 116        |
| P.5      | Calibration Impact Summary . . . . .                                | 117        |
| P.6      | Pre-Flight Calibration Checklist . . . . .                          | 117        |
| P.7      | GPS Error Measurement Procedure . . . . .                           | 117        |
| P.8      | Altitude Error Measurement Procedure . . . . .                      | 117        |
| P.9      | Error Budget Cross-Reference . . . . .                              | 118        |
| <b>Q</b> | <b>Simulation Validation and Sim-to-Real Gap</b>                    | <b>118</b> |
| Q.1      | Three Simulation Levels and What Each Validated . . . . .           | 118        |
| Q.2      | Error Injection as a Validation Tool . . . . .                      | 119        |
| Q.3      | DJI Video Replay as a Sim-to-Real Proxy . . . . .                   | 119        |
| Q.4      | Validation Results Summary . . . . .                                | 119        |
| Q.5      | Progressive Build-Up: From Pure Software to Real Hardware . . . . . | 119        |
| Q.6      | Known Sim-to-Real Gaps . . . . .                                    | 120        |
| Q.7      | Confidence Assessment . . . . .                                     | 121        |
| Q.8      | Lessons from the Sim-to-Real Transfer . . . . .                     | 121        |
| <b>R</b> | <b>Mission Flow</b>                                                 | <b>122</b> |
| R.1      | Pre-Flight . . . . .                                                | 122        |
| R.2      | Takeoff . . . . .                                                   | 122        |
| R.3      | Transit to Search Area . . . . .                                    | 122        |
| R.4      | Search . . . . .                                                    | 122        |
| R.5      | Detection and Queuing . . . . .                                     | 123        |
| R.6      | PLB Focus Area Redirect . . . . .                                   | 123        |
| R.7      | Investigation . . . . .                                             | 123        |
| R.8      | Verification . . . . .                                              | 123        |
| R.9      | Approach and Landing . . . . .                                      | 123        |
| R.10     | Return to Launch . . . . .                                          | 124        |
| R.11     | Safety Layers During Flight . . . . .                               | 124        |
| R.12     | How the Mission Was Built Up Incrementally . . . . .                | 124        |
| <b>S</b> | <b>Centering vs. Direct Offset Landing</b>                          | <b>125</b> |
| S.1      | The Centering Option . . . . .                                      | 125        |
| S.2      | Why Direct Offset Landing Is Sufficient . . . . .                   | 126        |
| S.3      | Probability Analysis . . . . .                                      | 126        |
| S.4      | Simulation Comparison . . . . .                                     | 126        |
| S.5      | Design Decision . . . . .                                           | 126        |
| S.6      | Safety Considerations of Hovering Above the Casualty . . . . .      | 127        |
| S.6.1    | Risk Identification . . . . .                                       | 127        |
| S.6.2    | Quantitative Analysis . . . . .                                     | 127        |
| S.6.3    | Comparison of Centering Strategies . . . . .                        | 128        |
| S.6.4    | Industry Practice . . . . .                                         | 128        |
| S.6.5    | Accepted Tradeoff . . . . .                                         | 128        |
| <b>T</b> | <b>MAVLink Communication Architecture</b>                           | <b>128</b> |
| T.1      | The Python 3.13 Serial Problem . . . . .                            | 128        |
| T.2      | MAVProxy Bridge Architecture . . . . .                              | 128        |
| T.3      | Connection Auto-Detection . . . . .                                 | 129        |
| T.4      | MAVLink Message Protocol . . . . .                                  | 129        |

|           |                                                               |            |
|-----------|---------------------------------------------------------------|------------|
| T.5       | Link Loss Detection and Recovery                              | 129        |
| T.6       | Latency Characteristics                                       | 130        |
| <b>U</b>  | <b>Raw Inference vs. Effective Pipeline Throughput</b>        | <b>130</b> |
| U.1       | Pipeline Stages and Overhead Budget                           | 130        |
| U.2       | Measured and Projected Results                                | 130        |
| U.3       | Bottleneck Transition                                         | 131        |
| U.4       | Operational Implications                                      | 131        |
| <b>V</b>  | <b>Flight Path Optimisation</b>                               | <b>131</b> |
| V.1       | Diagonal vs Axis-Aligned Scanning                             | 132        |
| V.2       | Lane Width vs Altitude Trade-Off                              | 132        |
| V.3       | Altitude-Dependent Speed Schedule                             | 132        |
| V.4       | Turn Strategy                                                 | 133        |
| V.5       | Overlap Margin Analysis                                       | 133        |
| V.6       | Start Corner Optimisation                                     | 133        |
| V.7       | PLB Focus Area Redirect                                       | 133        |
| <b>W</b>  | <b>Payload Release Mechanism</b>                              | <b>134</b> |
| W.1       | Tarot Double-Throw Release Mechanism                          | 134        |
| W.2       | Why Release from Altitude, Not Land                           | 134        |
| W.3       | The 7.5 m Offset                                              | 134        |
| W.4       | Servo Actuation Sequence                                      | 134        |
| W.5       | Why Visual Servo Centering Was Not Required                   | 135        |
| W.6       | Future Work: Precision Drop                                   | 135        |
| <b>X</b>  | <b>Model Comparison and Future Vision Improvements</b>        | <b>135</b> |
| X.1       | Model Generations                                             | 135        |
| X.2       | COCO Person Detector as Fallback                              | 136        |
| X.3       | Overfitting Risk in Single-Class Detection                    | 136        |
| X.4       | Tiling vs. Full-Frame Detection                               | 136        |
| X.5       | Future Vision Improvements                                    | 136        |
| <b>Y</b>  | <b>Multi-Objective Search Optimisation</b>                    | <b>137</b> |
| Y.1       | Energy Efficiency Simulation                                  | 137        |
| Y.2       | Rotation Angle Analysis                                       | 138        |
| Y.3       | Coverage Completeness: Is 96% Acceptable?                     | 139        |
| Y.3.1     | Probabilistic Argument                                        | 139        |
| Y.3.2     | Buffer Reduction Trade-Off                                    | 139        |
| Y.4       | NFZ Margin vs Altitude                                        | 140        |
| Y.5       | Speed vs Detection Probability                                | 140        |
| Y.6       | Focus Area Redirect Impact                                    | 141        |
| Y.7       | Summary of Optimisation Trade-Offs                            | 142        |
| <b>Z</b>  | <b>Focus Area Redirect and Repulsive Field Implementation</b> | <b>142</b> |
| Z.1       | PLB Focus Area Logic (R06)                                    | 142        |
| Z.2       | Repulsive Vector Field                                        | 142        |
| Z.3       | Four-Layer Protection Architecture                            | 143        |
| <b>AA</b> | <b>Search Path Optimisation</b>                               | <b>144</b> |
| AA.1      | The Multi-Objective Optimisation Problem                      | 144        |
| AA.2      | Energy Model                                                  | 145        |
| AA.3      | The 216-Configuration Parametric Sweep                        | 145        |
| AA.4      | Search Strategy Comparison                                    | 146        |

|           |                                                           |            |
|-----------|-----------------------------------------------------------|------------|
| AA.5      | NFZ Interaction and Speed Scalar Field . . . . .          | 146        |
| AA.6      | Detection Probability Integration . . . . .               | 147        |
| AA.7      | Selected Configuration and Justification . . . . .        | 148        |
| <b>AB</b> | <b>Search Parameter Optimisation</b>                      | <b>149</b> |
| AB.1      | Problem Statement . . . . .                               | 149        |
| AB.2      | Physics-Based Simulation . . . . .                        | 150        |
| AB.2.1    | Energy Model . . . . .                                    | 150        |
| AB.2.2    | Detection Model . . . . .                                 | 151        |
| AB.2.3    | Parametric Sweep . . . . .                                | 151        |
| AB.3      | Top Three Configurations . . . . .                        | 151        |
| AB.4      | Altitude and Speed Selection Logic . . . . .              | 152        |
| AB.5      | Sensitivity Analysis . . . . .                            | 153        |
| AB.6      | Variable Coupling Structure . . . . .                     | 154        |
| AB.7      | Pareto Frontier . . . . .                                 | 155        |
| AB.8      | Summary of Selected Parameters . . . . .                  | 155        |
| AB.9      | Alternative Strategies Considered . . . . .               | 155        |
| <b>AC</b> | <b>Multi-Objective Search Optimisation</b>                | <b>156</b> |
| AC.1      | Problem Formulation . . . . .                             | 156        |
| AC.2      | Objective Functions . . . . .                             | 157        |
| AC.2.1    | Coverage Completeness . . . . .                           | 157        |
| AC.2.2    | Detection Probability . . . . .                           | 158        |
| AC.2.3    | Mission Time . . . . .                                    | 159        |
| AC.2.4    | Energy Consumption . . . . .                              | 160        |
| AC.2.5    | NFZ Safety Margin . . . . .                               | 160        |
| AC.3      | Constraints . . . . .                                     | 160        |
| AC.4      | Variable Coupling Analysis . . . . .                      | 161        |
| AC.5      | Solution Approach . . . . .                               | 161        |
| <b>AD</b> | <b>Design Strategy: Parameter Derivation Chain</b>        | <b>162</b> |
| AD.1      | Step 1: Maximum Detection Altitude . . . . .              | 162        |
| AD.2      | Step 2: Maximum Speed at Chosen Altitude . . . . .        | 163        |
| AD.3      | Step 3: Minimise Energy via Scan Angle . . . . .          | 163        |
| AD.4      | Step 4: NFZ Safety Margins . . . . .                      | 163        |
| AD.5      | Step 5: Swath Overlap . . . . .                           | 164        |
| AD.6      | Step 6: Diagonal vs Aligned Scanning . . . . .            | 164        |
| AD.7      | The Dependency Chain . . . . .                            | 164        |
| <b>AE</b> | <b>Vision Performance Analysis</b>                        | <b>187</b> |
| AE.1      | Altitude vs. Target Pixel Size . . . . .                  | 187        |
| AE.2      | Speed vs. Motion Blur . . . . .                           | 188        |
| AE.3      | Speed vs. Frames on Target . . . . .                      | 189        |
| AE.4      | Coverage Efficiency . . . . .                             | 190        |
| AE.5      | Combined Detection Envelope . . . . .                     | 191        |
| AE.6      | Summary of Limiting Factors . . . . .                     | 192        |
| <b>AF</b> | <b>Decision Flow: How We Chose Our Search Parameters</b>  | <b>192</b> |
| AF.1      | The Problem . . . . .                                     | 192        |
| AF.2      | The Variables We Control . . . . .                        | 193        |
| AF.3      | Our Strategy: Safety-First, Detection-Optimised . . . . . | 193        |
| AF.4      | What We Evaluated . . . . .                               | 194        |
| AF.5      | The Result . . . . .                                      | 194        |

|                                                                               |            |
|-------------------------------------------------------------------------------|------------|
| AF.6 Sequential Measurement: Settling Each Variable Before the Next . . . . . | 195        |
| AF.7 Decomposition Into Independent Sub-Problems . . . . .                    | 196        |
| AF.8 Search Pattern Quantification . . . . .                                  | 196        |
| AF.9 Pareto Optimality of the Selected Point . . . . .                        | 196        |
| AF.10 Non-Quantifiable Design Choices . . . . .                               | 196        |
| <b>AG Requirements Verification Detail</b>                                    | <b>196</b> |
| AG.1R01 — Fly Within the Flight Area . . . . .                                | 196        |
| AG.2R02 — No SSSI Overfly . . . . .                                           | 197        |
| AG.3R03 — Take Off Within 5 m of TOL . . . . .                                | 197        |
| AG.4R04 — Maximum 50 m Altitude . . . . .                                     | 197        |
| AG.5R05 — Search Area and Identify Items of Interest . . . . .                | 198        |
| AG.6R06 — PLB Focus Area . . . . .                                            | 198        |
| AG.7R07 — Land 5–10 m from Casualty . . . . .                                 | 199        |
| AG.8R08 — Autonomy Level Justified . . . . .                                  | 200        |
| AG.9R09 — RTH and Motor Cutoff . . . . .                                      | 200        |
| AG.10R10 — Report Lat/Lon and Images . . . . .                                | 200        |
| AG.11R11 — Company Pilot Designated . . . . .                                 | 201        |
| AG.12R12 — Public GitHub Repository with MIT Licence . . . . .                | 201        |
| <b>AH Evaluation Detail</b>                                                   | <b>201</b> |
| AH.1 Defect Discovery Register . . . . .                                      | 201        |
| AH.2 What Would Change in a Second Iteration . . . . .                        | 201        |
| AH.3 Lessons Learned . . . . .                                                | 202        |
| <b>AI Development Methodology</b>                                             | <b>203</b> |
| AI.1 Development Philosophy . . . . .                                         | 203        |
| AI.2 Ten-Phase Development Timeline . . . . .                                 | 203        |
| AI.3 Simulation Stage (Weeks 12–15) . . . . .                                 | 204        |
| AI.4 Pi Deployment (Weeks 16–18) . . . . .                                    | 204        |
| AI.5 Field Integration (Weeks 19–20) . . . . .                                | 205        |
| AI.6 Ambition Levels . . . . .                                                | 205        |
| AI.7 Progressive Testing Ladder . . . . .                                     | 205        |
| AI.8 Achievement Summary . . . . .                                            | 206        |
| AI.9 Lessons for Future Teams . . . . .                                       | 206        |

## Executive Summary

An autonomous search and rescue (SAR) drone was designed, implemented, and tested by a team of five MSc students for the AENGM0074 group project. The system locates a casualty surrogate within the designated search area at Fenswood Farm, respecting geofence constraints: a Site of Special Scientific Interest (SSSI) no-fly zone, a 50 m altitude ceiling, and defined flight boundaries. All twelve brief requirements (R01–R12) are addressed—verified through simulation, bench testing, and field-day calibration—via a **five-layer automatic geofence** (waypoint filtering, speed ramp, repulsive field, hard boundary cutoff, and firmware backup), operator-in-the-loop verification, and a public MIT-licensed GitHub repository.

**Mission objective.** The drone autonomously takes off, searches a polygon-defined area using a lawnmower pattern, detects a casualty via onboard computer vision, reports the target’s GPS coordinates and imagery to a ground operator for confirmation, commands a landing between 5 and 10 m from the casualty to enable first-aid-kit deployment, and returns to the take-off location. An operator-in-the-loop verification stage prevents autonomous action on false positives.

**System overview.** The platform is a Hexsoon EDU-450 hexacopter controlled by a Cube Orange+ autopilot running ArduPilot. A Raspberry Pi 5 companion computer runs all onboard intelligence: a YOLOv8n object-detection model exported to TensorFlow Lite (TFLite), a lawnmower search-pattern generator, GPS-based target localisation with inverse-variance-weighted spatial clustering, and a finite state machine with **20 states** and 32 transitions that orchestrates the complete mission. A browser-based ground station streams live Motion JPEG (MJPEG) video and provides per-target classification controls (confirm, reject, mark as interest, mark as false positive) over Wi-Fi. Four MCDA trade studies informed the selection of detection model, search pattern, companion computer, and ground-station architecture.

**Key design decisions.** Three architectural choices define the system. First, a *modular single-codebase* design isolates computer vision, path planning, flight control, and configuration into independent modules with narrow interfaces, allowing identical code to run across laptop simulation, bench hardware, and the Pi without modification. Second, a *dual-backend vision pipeline* automatically selects between Ultralytics/PyTorch on the development laptop and TFLite on the Pi, enabling rapid iteration without deployment friction. Third, a *simulation-first development methodology* validated each subsystem through a **five-tier progressive testing framework**—SITL simulation, hardware component checks, bench integration, passive field observation, and autonomous flight—with **74 test scripts** organised across eight categories (including 127 unit tests), catching interface and timing errors before integration.

**Results achieved.** The detection model, retrained on 366 mixed synthetic and real-world images at native camera resolution ( $1456 \times 1088$ ), achieves  $\mathbf{mAP}_{50} = 0.995$ . On-device inference runs at **206 ms** per frame (**4.8 FPS**) on the Pi 5 CPU via the XNNPACK delegate, sufficient for the altitude-dependent search speed (8 m/s at the nominal 35 m altitude). Lens distortion correction adds only 1.5 ms overhead. Post-hoc analysis of DJI flight video over the test site confirmed detection across a 15–50 m altitude range and GPS target localisation with a Circular Error Probable (CEP) of  $\mathbf{CEP}_{50} = 2.3 \text{ m}$ . The lawnmower pattern covers the survey polygon in approximately 1.1 minutes with 18 waypoints, verified via dry-run to avoid SSSI encroachment. All figures in this report are generated programmatically from calibration data and benchmarks, ensuring reproducibility.

**Field day and integration.** Full hardware integration—Pi, Cube, camera, and GPS—was validated on the assembled airframe during a scheduled field day at Fenswood Farm. On-site bench testing produced lens calibration data ( $\text{RMS} = 0.399$ ), FOV characterisation (5.46 mm focal length), and inference benchmarks confirming 100 % detection at 0.966 mean confidence across 50 consecutive frames. These calibration results were fed directly back into the vision pipeline, improving GPS estimation accuracy. The complete mission pipeline—arming, search, detection, operator verification, and landing—was exercised end-to-end in Software-In-The-Loop (SITL) simulation, and DJI flight video recorded over the test site served as a real-world proxy for altitude-dependent detection analysis. High winds on the field day prevented powered flight; all integration results are therefore from simulation, bench tests, and video proxy analysis. The progressive testing methodology, operator-in-the-loop architecture, and five-layer automatic geofence provide a flight-ready autonomous SAR capability validated through progressive bench, simulation, and proxy-video testing with consumer-grade hardware and open-source software.

# Introduction

## Context and Background

A hiker has gone missing in Fenswood Wilderness, a semi-rural landscape of mixed grassland and scrub managed by the University of Bristol for field research. Search and rescue teams have been alerted, but the terrain is uneven, visibility is limited by vegetation, and the search area spans several hectares. At some point during the search, the hiker activates a Personal Locator Beacon (PLB), narrowing the probable location to a Focus Area—but a casualty on the ground, potentially incapacitated or partially concealed, still requires visual confirmation before help can reach them.

This is the scenario posed by the AENGM0074 module. Each student company is provided with identical hardware—a Hexsoon EDU-450 hexacopter airframe, a Cube Orange+ flight controller, a Raspberry Pi 5 companion computer, and a global-shutter camera—and tasked with building an autonomous unmanned aerial vehicle (UAV) capable of:

- systematically searching a defined area while respecting a no-fly zone over a Site of Special Scientific Interest (SSSI);
- detecting a life-size casualty surrogate (hereafter “rescue mannequin” or simply “mannequin”) using onboard computer vision;
- reporting the casualty’s estimated GPS coordinates and imagery to a ground operator;
- landing between 5 and 10 m from the casualty and deploying a simulated first-aid kit; and
- returning to the designated take-off location.

The scenario is deliberately constrained to a single outdoor demonstration flight, preceded by progressive bench and flight testing, reflecting the compressed timeline of a single-term MSc module. A 50 m altitude ceiling, defined flight-area boundaries, and automatic geofencing to enforce SSSI exclusion must all be respected without relying on pilot intervention. All software must be released publicly under the MIT licence.

## Company Description

Our company was formed at the start of the AENGM0074 module as a team of five MSc students from the Department of Aerospace Engineering. We adopted a workstream-based methodology: five parallel streams—Computer Vision, Hardware Integration, Search Logic, Ground Station, and Safety/Testing—fed into defined integration milestones where subsystems were combined and verified before advancing to the next stage. All development was managed through a shared GitHub repository<sup>1</sup>, with a single protected main branch and feature branches for individual work.

Day-to-day coordination relied on a group chat for informal updates and weekly in-person meetings for design decisions, integration planning, and blocker resolution. A shared status tracker documented every group decision, assigned action items with owners and deadlines, and maintained a running log of blockers and dependencies. This lightweight process allowed team members to work independently on their respective streams while ensuring that integration points were well-defined and that no subsystem was developed in isolation from the others.

## Team Members and Roles

**Dmytro Chuprynyuk — Lead Software Developer.** MSc student in Robotics with a background in software engineering and embedded systems. Designed and implemented the autonomous software stack: finite state machine, computer vision pipeline, lawnmower path planner, GPS-based target localisation, web-based ground station, and the simulation environment. Also built the model training pipeline and developed the progressive testing framework.

**Robin [Surname] — Flight Dynamics and State Machine Testing.** MSc student in Aerospace Engineering. Contributed to flight dynamics analysis and testing of the autonomous state machine, verifying state transitions and flight parameter tuning in both simulation and bench testing. Supported integration of the Cube Orange+ flight controller with the mission software and assisted with flight mode configuration (GUIDED, LOITER, STABILIZE, LAND).

**[TEAMMATE 3] — Designated Pilot and RC Systems.** MSc student in Aerospace Engineering with prior experience in radio-controlled aircraft. Served as the company’s designated pilot, responsible for RC transmitter configuration, kill-switch setup, and manual override procedures. Maintained readiness to assume manual control during all autonomous flight tests, ensuring compliance with the requirement that a pilot with override authority be present at all times.

**[TEAMMATE 4] — Hardware Integration.** MSc student in Mechanical Engineering. Led the physical assembly of the drone platform, including airframe construction, wiring of the Cube Orange+ to the Raspberry Pi (serial TELEM2 connection), camera mounting, GPS antenna placement, and compass calibration. Ensured that all electrical connections were secure and that the payload (Pi, camera, first-aid release mechanism) was mounted within the airframe’s centre-of-gravity envelope.

---

<sup>1</sup>Repository: [https://github.com/DimaChup/Group\\_Proj](https://github.com/DimaChup/Group_Proj) (MIT licence, R12).

**[TEAMMATE 5] — Project Management and Documentation.** MSc student in Aerospace Engineering. Coordinated project scheduling, booked flight slots at the Fenswood site, managed risk assessment paperwork, and maintained deliverable tracking against the module timeline. Led the report writing effort, compiled the requirements verification evidence, and ensured that all deliverables (D1–D6) were submitted on time.

### Contribution Summary

Table 1 summarises the key contributions of each team member across the major project areas.

**Table 1:** Team member contributions by project area.

| Team Member       | Key Contributions                                                                                                                                                                                                                                                                                                                        |
|-------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Dmytro Chuprynyuk | Autonomous software stack (state machine, CV pipeline, path planner, target localisation, ground station, simulation environment). Custom YOLOv8n model training and dataset generation. Progressive testing framework. Pi hardware integration (camera, TFLite inference, Cube serial link). Lens calibration and FOV characterisation. |
| Robin [Surname]   | Flight dynamics analysis. State machine testing in simulation and on bench. Flight parameter tuning (altitude, speed, descent profile). Cube flight-mode configuration and verification.                                                                                                                                                 |
| [TEAMMATE 3]      | RC transmitter configuration and binding. Kill-switch and flight-mode switch setup. Manual override procedures. Designated pilot for all flight tests.                                                                                                                                                                                   |
| [TEAMMATE 4]      | Airframe assembly and wiring. Cube–Pi serial connection. Camera and GPS antenna mounting. Compass calibration. Payload integration and CG verification.                                                                                                                                                                                  |
| [TEAMMATE 5]      | Project scheduling and milestone tracking. Flight-slot booking and risk assessments. Report compilation and editing. Requirements verification evidence. Deliverable management (D1–D6).                                                                                                                                                 |

### Report Structure

The remainder of this report is organised to address all requirements of the D6 brief. Section 1 presents the *Design Rationale*, including STEEPLE analysis and four MCDA trade studies. Sections 2–2.7 provide the *System Description*: hardware platform, computer vision pipeline, path planning, state machine, target localisation, and ground station. Section E covers the progressive testing methodology and simulation environment. Section 3 presents *Requirements Verification* with evidence of compliance for R01–R12. Section 4 provides a plus/delta *Evaluation* of technical performance. References and appendices follow, including the full state-transition table, detection samples, and calibration data.

# 1 Design Rationale

Every major design decision was evaluated against the seven STEEPLE dimensions (Social, Technological, Economic, Environmental, Political, Legal, Ethical). Technical trade studies then addressed hardware, software architecture, detection, path planning, and mission control.

## 1.1 STEEPLE Analysis

**Table 2:** STEEPLE assessment of the autonomous SAR drone system.

| Dimension     | Key consideration and design response                                                                                                                                                                                                                |
|---------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Social        | Reduces volunteer foot-search time from hours to minutes for missing-person cases; open-source MIT licence (R12) enables volunteer SAR teams to replicate the platform without proprietary dependencies [1].                                         |
| Technological | Pi 5 with TFLite (4 packages vs 90+); IMX296 global shutter eliminates rolling-shutter distortion at $8\text{ m s}^{-1}$ ; fully on-device inference—no cloud link required during mission [2, 3].                                                   |
| Economic      | £150 compute stack (<10% of DJI Matrice 30T); free-tier Colab training on 366-image dataset; zero per-mission cloud cost, viable for resource-constrained organisations.                                                                             |
| Environmental | Five-layer geofence with graduated speed reduction protects adjacent SSSI nesting habitat; 30 m waypoint buffer; electric propulsion limits acoustic disturbance; 3 kg AUV minimises rotor wash on ground fauna [4].                                 |
| Political     | UK CAA Open Category A3 compliant [5]; VLOS only (BVLOS requires SORA [6]); aligns with CAA Innovation Hub policy encouraging civilian drone SAR trials; public GitHub repo ensures full auditability.                                               |
| Legal         | On-device inference—no imagery transmitted over any network; detection logs record bounding-box coordinates only (no facial features), minimising UK GDPR risk; SSSI compliance auditable via logged boundary distances at every loop iteration [4]. |
| Ethical       | No autonomous landing without human confirmation (VERIFY gate); 120 s auto-reject prioritises continued search over uncertain detections; frames processed in memory and discarded after inference—no persistent image storage of bystanders.        |

With the STEEPLE framework confirming project viability across all seven dimensions, the following subsections detail the technical trade studies that determined the system’s operating parameters, hardware platform, software architecture, and mission control strategy.

## 1.2 Search Parameter Optimisation

The search pattern is governed by three coupled decision variables—altitude  $h$ , scan angle  $\theta$ , and speed  $v$ —that jointly determine five competing objectives: coverage completeness, detection probability, mission time, energy consumption, and NFZ safety margin.

These objectives cannot be optimised independently. Lowering altitude improves detection confidence but narrows the camera footprint and increases the number of scan lines. Increasing speed reduces mission time but shortens target dwell time and therefore detection probability.

## Variable Coupling Matrix — Pairwise Interaction Effects


|                             | Altitude $h$                                     | Speed $v$                                        | Scan angle $\theta$            | Overlap $\alpha$              | NFZ margin $d_{\text{nfz}}$              |
|-----------------------------|--------------------------------------------------|--------------------------------------------------|--------------------------------|-------------------------------|------------------------------------------|
| Altitude $h$                | $h$<br><i>Detection ceiling</i>                  | <b>Strong</b><br>Coverage rate + detection prob. | <b>Weak</b><br>Scan efficiency | <b>Medium</b><br>Lane spacing | <b>Medium</b><br>Footprint intrusion     |
| Speed $v$                   | <b>Strong</b><br>Coverage rate + detection prob. | $v$<br><i>Frame count</i>                        | <b>Weak</b><br>Independent     | <b>Weak</b><br>Independent    | <b>Medium</b><br>Reaction distance       |
| Scan angle $\theta$         | <b>Weak</b><br>Scan efficiency                   | <b>Weak</b><br>Independent                       | $\theta$<br><i>Turn count</i>  | <b>Weak</b><br>Independent    | <b>Medium</b><br>NFZ proximity           |
| Overlap $\alpha$            | <b>Medium</b><br>Lane spacing                    | <b>Weak</b><br>Independent                       | <b>Weak</b><br>Independent     | $\alpha$<br><i>Redundancy</i> | <b>Weak</b><br>Independent               |
| NFZ margin $d_{\text{nfz}}$ | <b>Medium</b><br>Footprint intrusion             | <b>Medium</b><br>Reaction distance               | <b>Medium</b><br>NFZ proximity | <b>Weak</b><br>Independent    | $d_{\text{nfz}}$<br><i>Safety margin</i> |

■ Strong coupling   
■ Medium coupling   
■ Weak / independent   
■ Self (primary metric)

**Figure 1:** Variable coupling matrix. The altitude–speed pair exhibits the dominant interaction, jointly determining coverage rate and detection probability. Off-diagonal cells identify which performance metric each variable pair affects.

A 216-configuration parametric sweep (6 altitudes  $\times$  36 scan angles, Appendix AA) evaluated every combination on the AENGM0074 polygon geometry using a physics-based energy model. Scan angle dominates altitude for energy efficiency: varying the angle from optimal ( $70^\circ$ ) to worst-case ( $0^\circ$ ) increases energy by 100–200%, whereas changing altitude from 35 m to 50 m saves only 30–40%. A dependency chain links each parameter to a physical measurement or constraint:

1. **Target**  $\rightarrow$  **altitude**: the mannequin’s 0.5 m width must subtend  $\geq 10$  model-input pixels for reliable detection, capping altitude at 35 m with a  $1.4\times$  margin.
2. **Altitude**  $\rightarrow$  **speed**: the 24 m along-track footprint at 35 m, combined with 4.8 FPS inference, requires  $\geq 10$  frames per target, yielding  $v \leq 11.5 \text{ m s}^{-1}$ ; the altitude-dependent speed schedule assigns  $v = 8 \text{ m s}^{-1}$ .
3. **Polygon geometry**  $\rightarrow$  **scan angle**: aligning with the longest edge ( $70^\circ$ ) minimises U-turns from 9 to 5 and reduces energy from 25.4 Wh to 12.6 Wh.
4. **Footprint width**  $\rightarrow$  **NFZ margin**: the 32 m footprint at 35 m requires 23 m Root Sum of Squares (RSS) standoff; the 30 m buffer provides a  $2.2\times$  safety factor.



**Figure 2:** Altitude–speed trade-off space. Colour encodes composite mission score (40% detection + 35% coverage + 25% energy). The selected operating point (35 m, 8 m s<sup>-1</sup>) sits above the 95% detection contour in the high-score region, trading 4.4 Wh of energy (10.9% of battery) for a comfortable detection margin.

The selected operating point (35 m, 70°, 8 m s<sup>-1</sup>) is not the energy minimum—that belongs to 50 m/70°/10 m s<sup>-1</sup> at 8.2 Wh—but it offers the best Pareto compromise: 96% coverage,  $P_d > 99.97\%$  (assuming frame independence; 95% per pass after correlation adjustment, see Section AG), 12.6 Wh (10.9% of battery), and a 50 m NFZ margin with 2.2× surplus. The 4% coverage gap lies within a fenced wildlife reserve with negligible casualty probability. Appendices AA and AD provide the full derivation, energy model, strategy comparison, and sensitivity analysis.

### 1.3 Parameter Derivation Chain

Every operating parameter traces back to a physical measurement or regulatory constraint through an unbroken eight-step chain (Table 4). Full equations, RSS overlap analysis, and sensitivity analysis are provided in Appendices AA and AF.

**Table 4:** Parameter derivation chain: each value is derived from the previous step.

| Step | Parameter          | Value                  | Derivation                                                                         |
|------|--------------------|------------------------|------------------------------------------------------------------------------------|
| 1    | Detection ceiling  | 63 m                   | Mannequin (0.5 m wide) must subtend $\geq 20$ model-input pixels                   |
| 2    | Operating altitude | 35 m                   | 30% margin below ceiling (36 px target, 1.8× above minimum)                        |
| 3    | Max speed          | 11.5 m s <sup>-1</sup> | 24 m footprint ÷ (10 frames / 4.8 FPS); adopted 8 m s <sup>-1</sup>                |
| 4    | Lane spacing       | 25.7 m                 | 32.2 m footprint × 0.80; 20% overlap gives 1.4× margin over 4.5 m RSS error [7, 8] |
| 5    | Scan angle         | 70°                    | Longest-edge alignment minimises U-turns (5 vs 9), saves 13 Wh                     |
| 6    | NFZ buffer         | 30 m                   | Half-footprint (16.1) + GPS CEP (3.0) + stopping dist. (10.7) = 29.8 m [4]         |
| 7    | Coverage           | 96%                    | 4% gap in fenced wildlife reserve (negligible casualty probability)                |
| 8    | Search energy      | 12.6 Wh                | 10.9% of 115.4 Wh battery; 14 min remaining for transit/verify/RTL                 |

### 1.4 Hardware Platform Selection

The hardware integrates a Hexsoon EDU-450 airframe, CubePilot Orange+ flight controller with ArduPilot firmware, Raspberry Pi 5 companion computer, and IMX296 global shutter camera.

- **Global shutter camera.** The IMX296 was selected over cheaper rolling-shutter alternatives because rolling-shutter distortion during forward flight at up to  $10\text{ ms}^{-1}$  skews bounding boxes and degrades detection accuracy. The CSI interface leaves USB ports free for telemetry radios.
- **No thermal imaging.** Thermal cameras (e.g. FLIR Boson, DJI Zenmuse H20T) add £2000–£10000 to the platform cost, have lower spatial resolution (typically  $640 \times 512$  vs  $1456 \times 1088$  RGB), and require radiometric calibration that varies with ambient temperature. While thermal imaging offers superior night-time and foliage-penetration capability, this project deliberately explores RGB-only detection feasibility at the low-cost extreme [9]. A fused RGB+thermal system would be ideal for operational deployment but exceeds the £565 budget constraint by  $4\text{--}18\times$ .
- **Onboard vs offboard inference.** Offloading detection to a ground station via video downlink was rejected for three reasons: (1) reliable video streaming at  $1456 \times 1088$  at 5 FPS requires  $\approx 15$  Mbps sustained throughput, exceeding typical 2.4 GHz telemetry radio capacity; (2) round-trip latency (encode, transmit, detect, return command) adds 200–500 ms to the decision loop, during which the drone travels 1–4 m; and (3) link loss during detection would leave the drone without situational awareness. Onboard inference on the Pi 5 eliminates all three failure modes at the cost of lower computational headroom.
- **MAVProxy bridge.** A Python 3.13 incompatibility with pyserial (dropped bytes causing BAD\_DATA errors) is resolved by interposing MAVProxy as a UDP bridge, which reads the serial port reliably using its own C-level handler [10].

#### 1.4.1 Companion Computer Trade Study

All four trade studies in this section use weighted multi-criteria decision analysis (MCDA). In each case, we established criteria weights before scoring candidates, based on mission requirements and operational constraints, to prevent post-hoc rationalisation.

Three candidate companion computers were evaluated (Table 6). Cost and Python ecosystem compatibility are paramount for a student team, whereas raw GPU performance is less critical given the lightweight YOLOv8n model.

**Table 6:** MCDA: companion computer selection (1–5 scale, 5 = best)

| Criterion                 | Weight | Pi 5           | Jetson Nano     | Intel NCS2        |
|---------------------------|--------|----------------|-----------------|-------------------|
| Unit cost (£)             | 0.20   | 5 (£75)        | 3 (£180)        | 4 (£100+host)     |
| Power consumption         | 0.10   | 4 (5–12 W)     | 3 (5–20 W)      | 3 (host+1.5 W)    |
| GPIO / CSI interface      | 0.15   | 5 (40-pin+CSI) | 4 (40-pin+CSI)  | 1 (USB only)      |
| Community & documentation | 0.20   | 5 (largest)    | 4 (strong)      | 2 (discontinued)  |
| Python 3.13 compatibility | 0.20   | 5 (native)     | 3 (JetPack lag) | 2 (OpenVINO lag)  |
| TFLite / edge-AI support  | 0.15   | 5 (XNNPACK)    | 4 (TensorRT)    | 2 (OpenVINO only) |
| <b>Weighted total</b>     |        | <b>4.90</b>    | 3.50            | 2.35              |

The Pi 5 achieves the highest weighted total. Sensitivity analysis ( $\pm 0.05$  weight perturbation across all 12 scenarios) confirms the Pi 5 retains its rank in every case (minimum score 4.72 vs Jetson Nano maximum 3.68). The same perturbation was applied to all four MCDA tables; in each case the selected option retained its rank.

#### 1.4.2 Communication Architecture Trade Study

Three architectures for companion-to-autopilot communication were compared (Table 8).

**Table 8:** MCDA: communication architecture (1–5 scale, 5 = best)

| Criterion                 | Weight | Direct Serial  | MAVProxy Bridge  | ROS 2                                        |
|---------------------------|--------|----------------|------------------|----------------------------------------------|
| Latency                   | 0.20   | 5 (direct)     | 4 (UDP hop)      | 3 (Data Distribution Service (DDS) overhead) |
| Reliability               | 0.25   | 2 (py3.13 bug) | 5 (C-level read) | 4 (mature)                                   |
| Multi-consumer support    | 0.20   | 1 (exclusive)  | 5 (N outputs)    | 5 (pub/sub)                                  |
| Python 3.13 compatibility | 0.20   | 1 (broken)     | 5 (works)        | 3 (partial)                                  |
| Implementation complexity | 0.15   | 5 (trivial)    | 4 (one process)  | 1 (full stack)                               |
| <b>Weighted total</b>     |        | 2.65           | <b>4.65</b>      | 3.35                                         |

Direct serial is the simplest option but a Python 3.13 pyserial regression that drops bytes at 921 600 baud renders it non-viable. MAVProxy resolves this by reading the serial port with its C-level handler and forwarding via UDP, simultaneously enabling multi-consumer access (scripts on port 14550, Mission Planner on TCP 5762). ROS 2 offers the richest middleware but introduces deployment complexity disproportionate to this single-node architecture.

## 1.5 Software Architecture Rationale

Four principles govern the architecture:

1. **Modular separation**—each concern hides behind a narrow interface (e.g. `detect_in_image(frame) → (found, x, y, conf)`).
2. **Platform-agnostic code**—a `MODE` flag selects the camera source; auto-detection selects the flight controller. No conditional branches distinguish platforms.
3. **Configuration-driven tuning**—all parameters reside in `config.py`, so transitioning to real hardware requires only parameter adjustment.
4. **Dual inference backend**—the system auto-detects Ultralytics (laptop) or TFLite (Pi) at import time, both producing identical [1, 5, 8400] tensors.

## 1.6 Detection Model Selection

YOLOv8n (the nano variant, 3.2 million parameters, 8.7 GFLOPs) was selected after benchmarking the YOLO model family on the target hardware [11]. Table 10 summarises the weighted comparison.

**Table 10:** MCDA: detection model selection (1–5 scale, 5 = best)

| Criterion                  | Weight | YOLOv8n     | YOLOv8s      | YOLOv8m    | SSD MobileNet |
|----------------------------|--------|-------------|--------------|------------|---------------|
| Inference speed (Pi 5 CPU) | 0.30   | 5 (207 ms)  | 3 (450 ms)   | 1 (820 ms) | 5 (180 ms)    |
| Detection accuracy (mAP50) | 0.25   | 4 (0.995)   | 5 (0.997)    | 5 (0.998)  | 3 (0.88)      |
| Model size (TFLite)        | 0.15   | 5 (3.2 MB)  | 3 (22 MB)    | 2 (50 MB)  | 5 (4.3 MB)    |
| TFLite export support      | 0.15   | 5 (native)  | 5 (native)   | 5 (native) | 4 (manual)    |
| Training data requirement  | 0.15   | 5 (small)   | 4 (moderate) | 3 (large)  | 4 (moderate)  |
| <b>Weighted total</b>      |        | <b>4.75</b> | 3.95         | 3.05       | 4.20          |

YOLOv8n is the only variant that achieves real-time inference on the Pi 5 CPU within the 300 ms-per-frame budget required at search speed. YOLOv8s (11.2 M parameters) and YOLOv8m (25.9 M) exceed this budget by  $2\times$  and  $4\times$ , respectively. SSD MobileNetV2 matches YOLOv8n on speed but delivers lower accuracy on COCO [12]. Note that the SSD mAP of 0.88 is measured on COCO whereas the YOLO figures use the custom dataset, making the comparison indicative rather than direct [2].

The model was trained as a single-class detector on 366 images (300 synthetic, 16 real, 50 negatives) at native camera resolution ( $1456 \times 1088$ ), following the domain-randomisation principle [13]. Training at higher resolution than the  $640 \times 640$  inference input allows the network to learn fine-grained features that survive downscaling [11]. A COCO-pretrained person detector is carried as a fallback for field conditions where the custom model underperforms.

**Why TFLite over other inference runtimes?** Three alternative runtimes were considered: ONNX Runtime, TensorRT, and NCNN. TensorRT requires an NVIDIA GPU (absent on the Pi 5). ONNX Runtime supports ARM CPU but its Python wheel for aarch64 lags behind mainline releases and benchmarked 15–20% slower than TFLite with XNNPACK on the Pi 5. NCNN (Tencent) achieved  $4.5\times$  faster inference than TFLite in benchmarks (72 ms vs 327 ms full pipeline) and is supported as a secondary backend; however, its Python bindings are less mature, and the TFLite ecosystem provides the most robust export path from Ultralytics. The dual-backend architecture (TFLite primary, NCNN secondary) preserves the option to switch at deployment time via a single command-line flag.

## 1.7 Coverage Path Planning

Four candidate search patterns were evaluated (Table 12), prioritising coverage guarantee and implementability to ensure reliable, verifiable behaviour over theoretically optimal but harder-to-validate paths.

**Table 12:** MCDA: search pattern selection (1–5 scale, 5 = best)

| Criterion               | Weight | Lawnmower         | Spiral          | Expanding Sq.  | Random            |
|-------------------------|--------|-------------------|-----------------|----------------|-------------------|
| Coverage guarantee      | 0.30   | 5 (complete)      | 3 (centre gap)  | 4 (approx.)    | 1 (none)          |
| Path length efficiency  | 0.20   | 4 (minimal turns) | 3 (long spiral) | 3 (many turns) | 2 (redundant)     |
| Implementability        | 0.20   | 5 (analytical)    | 3 (radius fn)   | 4 (simple)     | 5 (trivial)       |
| Wind robustness         | 0.15   | 4 (long legs)     | 3 (curved)      | 3 (short legs) | 2 (unpredictable) |
| SAR standard compliance | 0.15   | 5 (IAMSAR)        | 3 (maritime)    | 4 (land SAR)   | 1 (none)          |
| <b>Weighted total</b>   |        | <b>4.65</b>       | 3.00            | 3.65           | 2.15              |

The boustrophedon (lawnmower) pattern was selected [8, 7] for its formal coverage guarantee, closed-form path length, and IAMSAR compliance. The implementation uses a rotate–rasterise–rotate-back algorithm on arbitrary polygons, with the scan angle auto-optimised via the minimum-area bounding rectangle (validated by the 216-configuration sweep).

## 1.8 State Machine Design

A deterministic finite state machine (FSM) with 20 states and 32 transitions was chosen over behaviour trees and reactive architectures for its verifiability, traceability, and ease of runtime monitoring [14].

**Why not behaviour trees?** Behaviour trees offer superior composability for complex agent behaviours and are widely used in game AI and multi-robot systems. However, they introduce implicit priority ordering through selector/sequence nodes that is harder to audit than an explicit transition table, debugging requires tree-visualisation tooling, and the SAR mission’s strictly sequential phase structure (search → detect → verify → land) does not benefit from the parallel subtree execution that behaviour trees enable. For a single-drone, single-objective mission, the FSM’s one-state-at-a-time guarantee makes every transition auditable in a flat log file.

**Why not reactive/subsumption architectures?** Reactive architectures (e.g. Brooks’ subsumption) excel at low-latency obstacle avoidance but lack the persistent state memory required for multi-phase missions: the system must remember which search legs have been completed, which targets have been classified, and how long the current hover has lasted. Layering these concerns onto a purely reactive stack would replicate FSM functionality without its formal verifiability.

Each state maps to a single handler function via a dictionary-based dispatch table ( $O(1)$  lookup), eliminating fall-through errors. Defence-in-depth safety is enforced through five independent mechanisms: manual override (M key), RC kill switch (hardware-level), return-to-launch on link loss, five-layer geofence, and per-state timeouts. Every blocking state carries a wall-clock timeout to prevent indefinite hangs. The complete transition table is provided in Appendix A. The CENTERING state—hovering stationary above a detected target before estimating its GPS position—simultaneously eliminates roll/pitch projection error (6.2 m during flight to <0.6 m), motion blur, and GPS timing lag. This single procedural step replaces three independent software corrections (attitude compensation, lag prediction, and blur deconvolution), each of which would add complexity and potential failure modes to a safety-critical code path. Appendix K.8 provides the full quantitative analysis.

## 1.9 Autonomy Level Justification

The system operates at Sheridan Level 7–8 (“computer executes autonomously, human monitors”) [15] for the search phase and Level 3–4 (“computer suggests one alternative, human approves”) for the landing decision. This hybrid was chosen after analysing the asymmetric cost structure of SAR false positives versus false negatives [16, 17].

A *false-positive landing* (descending on a rock, tree stump, or discarded clothing) incurs three compounding penalties: (1) battery endurance consumed by descent, hover, and re-climb ( $\approx 90$  s, roughly 8% of a 20-minute endurance budget); (2) search coverage paused during the entire descent–verify–climb cycle; and (3) risk of mechanical damage on uneven terrain. In a single-drone mission with no redundancy, even one unnecessary landing significantly reduces the probability of completing the search before battery exhaustion.

A *false negative* (failing to investigate a real casualty) is worse in absolute terms but is mitigated by the lawnmower pattern’s guarantee of revisiting every cell: a missed detection on pass 1 can be recovered on the return leg. The operator-in-the-loop gate therefore addresses the higher-consequence, lower-recoverable error (committing to a landing) while the autonomous search handles the high-bandwidth, lower-consequence task (scanning terrain at 8 m/s).

The VERIFY state implements this gate: the operator classifies each detection as confirmed (Y), rejected (N), or item of interest (I) via the ground station. A 120s timeout auto-rejects unresponded detections—derived from the energy budget (6.0 Wh per hover event, 5.2% of battery, permitting two timeouts while retaining 89% for search) [18]. Full autonomy (Level 10) was rejected because the single-class detector cannot distinguish a casualty from visually similar objects [1, 19]. Manual-only operation was also rejected: human operators cannot sustain visual scanning at  $8\text{ m s}^{-1}$ , and automation provides a  $\sim 20\times$  coverage-rate advantage [20].

## 1.10 Decisions That Changed: Evidence-Based Corrections

Seven design decisions were revised when empirical evidence contradicted initial assumptions. Each correction followed the same pattern: observe anomalous behaviour, diagnose root cause, implement fix, verify improvement—demonstrating evidence-driven iteration rather than assumption-driven design.

**1. Focal length:** 7.0 mm  $\rightarrow$  5.46 mm. *Original decision:* use the datasheet focal length of 7.0 mm for GPS estimation. *Evidence:* a tape-measure calibration on field day measured 92 cm visible at 1 m height—22% wider than predicted. *Revised decision:* back-calculate focal length from the physical measurement:  $f = w_s \times d / W_{\text{visible}} = 5.02 \times 1.0 / 0.92 = 5.46\text{ mm}$ . *Outcome:* GPS estimation error at 30 m reduced by up to 3 m; without this correction, landing errors would have exceeded the 10 m R07 requirement. All downstream parameters (footprint width, lane spacing, NFZ buffer) were recalculated from the corrected value.

**2. Camera colour space:** RGB assumed  $\rightarrow$  BGR confirmed. *Original decision:* apply `cv2.cvtColor(frame, cv2.COLOR_RGB2BGR)` because `picamera2` labels the stream as RGB888. *Evidence:* detection confidence on the Pi was anomalously low ( $<0.3$ ) despite identical model performance on the laptop. Systematic testing of all six channel permutations (RGB, RBG, GBR, GRB, BRG, BGR) revealed that the IMX296 sensor outputs BGR data despite the RGB888 label—the `cvtColor` call was inverting the channels. *Revised decision:* remove the erroneous conversion; pass raw frames directly to the detector. *Outcome:* detection confidence recovered to 0.96+ on the Pi, matching laptop performance. This bug would have caused complete detection failure in flight.

**3. Camera resolution:**  $640 \times 480 \rightarrow 1456 \times 1088$ . *Original decision:* capture at  $640 \times 480$  (common webcam default) to minimise processing load. *Evidence:* at  $640 \times 480$  and 30 m altitude, a prone casualty subtends only  $20 \times 14\text{ px}$ —at the minimum detection threshold with no safety margin. *Revised decision:* capture at the IMX296’s native resolution ( $1456 \times 1088$ ), since the model resizes internally to  $640 \times 640$  regardless. *Outcome:* target subtends  $46 \times 32\text{ px}$  ( $2.3\times$  improvement) with zero inference time penalty, providing a comfortable  $1.6\times$  margin above the 20 px detection floor.

**4. TFLite bounding box interpretation:** normalised  $\rightarrow$  pixel coordinates. *Original decision:* treat TFLite output coordinates as normalised values in  $[0, 1]$  and multiply by frame dimensions, following the standard TFLite convention. *Evidence:* during Pi bench testing, detection overlay boxes were drawn off-screen (coordinates exceeding 640). Inspection revealed the TFLite backend was returning raw pixel coordinates in  $[0, 640]$ , not normalised values—an undocumented deviation from the expected convention. *Revised decision:* add a coordinate-range check: if any output coordinate exceeds 1.0, treat as pixel coordinates and skip the normalisation multiply. The same guard was already present in the NCNN inference path. *Outcome:* detection overlays rendered correctly; GPS estimation accuracy improved because bounding box centres were no longer scaled by an extra factor of 640. This also fixed a downstream double-scaling bug in `passive_watch.py` where pixel coordinates were multiplied by frame dimensions a second time, placing GPS estimates over 100 m from the true position.

**5. Geofence repulsive force:** sign inversion corrected. *Original decision:* the repulsive offset function returns positive lat/lon deltas to push the drone away from the NFZ boundary. *Evidence:* SITL simulation showed the drone being pushed *toward* the SSSI boundary instead of away. Root cause: `cv2.pointPolygonTest` returns positive values for points *inside* a polygon (opposite to the assumed convention), and the repulsive offset signs were not negated. *Revised decision:* negate the repulsive offset returns and correct the inside/outside boundary logic. *Outcome:* geofence now correctly repels the drone from the NFZ. This defect would have caused a Wildlife and Countryside Act 1981 violation in flight [4], and was caught at Tier 1 (simulation) at zero hardware cost.

**6. TFLite runtime:** `tflite-runtime`  $\rightarrow$  `ai-edge-litert`. *Original decision:* use Google’s `tflite-runtime` package for edge inference on the Pi. *Evidence:* `tflite-runtime` has no Python 3.13 wheel; `picamera2` requires system Python 3.13 on the Pi 5, creating an unresolvable dependency conflict. *Revised decision:* switch to `ai-edge-litert`, which provides an identical API with Python 3.13 support. *Outcome:* dependency conflict eliminated with zero code changes; verified by 50-run benchmark producing identical inference times (206.5 ms).

**7. Serial communication: pyserial → MAVProxy UDP bridge.** *Original decision:* connect the Pi directly to the Cube flight controller via pyserial at 921 600 baud. *Evidence:* a Python 3.13 pyserial regression dropped bytes, producing persistent BAD\_DATA errors that prevented reliable MAVLink communication. *Revised decision:* interpose MAVProxy as a UDP bridge; it reads the serial port with its C-level handler and forwards to UDP endpoints. *Outcome:* reliable communication restored; MAVProxy simultaneously enables multi-consumer access (scripts on port 14550, Mission Planner on TCP 5762, diagnostics on port 14551), turning a workaround into an architectural improvement (Table 8).

### 1.11 Why Not ROS

ROS 2 was excluded for three reasons [21]: (1) architectural mismatch—the system is a single-node pipeline (one camera, one detector, one flight controller), and ROS’s pub/sub model adds inter-process overhead without new capability; (2) deployment complexity—ROS 2 requires  $\approx 2$  GB versus four pip packages totalling  $< 50$  MB; and (3) timeline—a 40–60 hr learning curve per developer would have consumed 25% of available development time on middleware rather than mission logic. ROS remains recommended for future multi-drone extensions.

### 1.12 Camera Selection: Global vs Rolling Shutter

The IMX296 global shutter camera (provided as university equipment) was validated for aerial search [22, 23]. At  $5 \text{ m s}^{-1}$  and 30 m altitude, ground motion across the sensor reaches  $\approx 250 \text{ px/s}$ ; a rolling-shutter sensor would produce  $\approx 7.5 \text{ px}$  parallelogram distortion per frame, shifting detected centroids by up to 0.5 m. The IMX296’s global shutter eliminates this error. The trade-off—lower light sensitivity from the  $3.45 \mu\text{m}$  pixel pitch—is acceptable given daylight-only CAA VLOS requirements [5]. During integration, the IMX296 was found to output BGR despite its RGB888 labelling; the fix was simply removing the erroneous `cv2.cvtColor()` call (Section 1.10).

### 1.13 Field Day Adaptation Under Uncertainty

The scheduled flight test (2026-03-11) was cancelled due to sustained winds exceeding the EDU-450’s  $10 \text{ m s}^{-1}$  operational limit (measured  $12 \text{ m s}^{-1}$ , gusts to  $18 \text{ m s}^{-1}$ ). The team executed a pre-planned bench-testing fallback instead. This session yielded the FOV calibration correction ( $7.0 \text{ mm} \rightarrow 5.46 \text{ mm}$ ), lens distortion calibration (RMS 0.399), end-to-end connectivity verification, and the inference benchmark (207 ms, 4.8 FPS). These results are detailed in Section 1.10 and the evaluation (Section 4). The go/no-go decision framework and full bench protocol are documented in Appendix M.

**Decision flow narrative.** Appendix AF presents the six-step reasoning chain—from detection ceiling through altitude, speed, scan angle, NFZ margin, and coverage acceptance—as a self-contained executive narrative. It includes the 216-configuration sweep results, top-three configuration comparison, Pareto optimality analysis, and non-quantifiable design choices.

**Section summary.** Every design decision in this project traces to a quantified rationale—physical measurement, parametric sweep, weighted trade study, or energy budget analysis—and seven decisions were revised when empirical evidence contradicted initial assumptions. Three of those corrections (BGR colour inversion, geofence sign error, bounding box coordinate mismatch) would have caused mission-critical failures in flight, yet all were caught and resolved at low-cost tiers through the progressive testing framework, demonstrating the evidence-driven engineering process that underpins the system.


## 2 System Description

The final system integrates seven subsystems into an autonomous SAR platform: a hardware airframe with companion computer (§2.1), eleven Python modules (§2.2), a YOLOv8n vision pipeline at 4.8 FPS (§2.3), a boustrophedon search pattern with formal coverage guarantee (§2.4), a 20-state mission FSM with six safety mechanisms (§2.5), GPS target localisation at  $\text{CEP}_{50} = 2.3 \text{ m}$  (§2.6), and a browser-based ground station at 300 ms video latency (§2.7). The design philosophy is *independent modules composed at runtime*: vision, planning, geofencing, and navigation have zero mutual dependencies, enabling isolated testing and single-file swaps with no integration risk. Figure 3 shows the complete mission plan overlaid on the Fenswood Farm site.

### 2.1 Hardware Platform

Table 14 summarises the bill of materials for the Hexsoon EDU-450 airframe, CubePilot Orange+ flight controller (ArduPilot Copter 4.5), Raspberry Pi 5 companion computer, and IMX296 global shutter camera. All hardware was provided by the university as part of the AENGM0074 equipment package; design effort therefore focused on system integration, software architecture, and calibration rather than component selection (see Appendix H for the edge inference architecture rationale).

## Mission Overview — Fenswood Wilderness Site



**Figure 3:** Mission overview on the Fenswood Farm site. The lawnmower search pattern covers the designated survey polygon at 35 m altitude, with the take-off point, return-to-launch path, and SSSI no-fly zone boundary shown. The flight area boundary constrains all autonomous waypoints.

The Pi communicates with the Cube via UART at 921 600 baud through a MAVProxy bridge, which republishes MAVLink traffic over UDP. A second output on TCP port 5762 allows Mission Planner on the ground station laptop to monitor telemetry via Wi-Fi (see Appendix J for the full communications architecture). The camera captures at  $1456 \times 1088$  via CSI-2; lens distortion is corrected using precomputed remap matrices (calibrated RMS = 0.399 px, +1.5 ms overhead; methodology in Appendix P). The calibrated focal length is 5.46 mm, yielding a horizontal field of view (HFOV) of  $\approx 49.3^\circ$  (Pi IMX296 camera).


## 2.2 Software Architecture

The system comprises eleven Python modules totalling approximately 4 400 lines of production code, organised in four layers (Figure 4). An additional 74 test and calibration scripts (Appendix O), including 127 unit tests, validate each module in isolation before integration:

- **Configuration layer:** `config.py` centralises all tuneable parameters (Appendix B); `states.py` defines the 20-state enumeration.
- **Independent modules:** `vision.py` (camera + AI detection), `planning.py` (lawnmower pattern), `gps_utils.py` (coordinate transforms), `geofence.py` (NFZ enforcement). These have no mutual dependencies and can be tested in isolation.

**Table 14:** System bill of materials and approximate costs.

| Component             | Specification                                  | Cost (£)    |
|-----------------------|------------------------------------------------|-------------|
| Cube Orange+ FC       | STM32H757, triple-redundant IMU, ArduPilot 4.5 | 180         |
| Here3 GPS             | u-blox M8P, CAN interface, 5 Hz updates        | 70          |
| Raspberry Pi 5        | BCM2712 quad-core A76, 8 GB RAM                | 60          |
| IMX296 GS Camera      | Sony IMX296, 1456×1088, global shutter, CSI-2  | 45          |
| Frame + motors + ESCs | EDU-450, 2212/920 KV brushless, 30 A BLHeli.S  | 100         |
| Battery + RC + misc   | 4S 5200 mA h LiPo, FrSky Twin X14, wiring      | 110         |
| <b>Total</b>          |                                                | <b>~565</b> |




**Figure 4:** Module dependency graph. Independent modules have no cross-dependencies; the orchestrator composes them at runtime.

- **Domain modules:** `navigation.py` wraps MAVLink commands (Table 16); `stream_server.py` serves the ground station.
- **Orchestration layer:** `state_machine.py` implements all state handlers as a mixin class; `main.py` composes everything and runs the main loop.

The data pipeline traverses five stages per iteration: (1) capture a frame and apply lens undistortion; (2) resize to  $640 \times 640$ , normalise, and run YOLOv8n inference; (3) convert the pixel detection to a GPS estimate; (4) dispatch the current state handler; and (5) enforce geofence constraints. Geofencing always runs *after* the state handler, so waypoint commands can be overridden by repulsive velocities.

**Geofence enforcement.** The geofence module implements five independent protection layers (Figure 6). The outermost layer is the **Flight Area boundary**: a four-vertex polygon checked via `cv2.pointPolygonTest` every main-loop iteration ( $\sim 20\text{--}50$  Hz); any excursion triggers an immediate emergency RTL with no buffer or warning zone. The remaining four layers protect the **SSSI no-fly zone**: (1) plan-time waypoint filtering (30 m buffer); (2) a speed



**Figure 5:** Raspberry Pi 5 companion computer system: hardware interfaces and data flow between the CSI camera, TFLite inference engine, MAVProxy bridge, and ground station.

| Method                            | MAVLink Message                             | Purpose             |
|-----------------------------------|---------------------------------------------|---------------------|
| <code>arm()</code>                | <code>MAV_CMD_COMPONENT_ARM_DISARM</code>   | Arm/disarm motors   |
| <code>takeoff(alt)</code>         | <code>MAV_CMD_NAV_TAKEOFF</code>            | Climb to altitude   |
| <code>send_global_target()</code> | <code>SET_POSITION_TARGET_GLOBAL_INT</code> | Fly to GPS waypoint |
| <code>send_velocity()</code>      | <code>SET_POSITION_TARGET_LOCAL_NED</code>  | Velocity command    |
| <code>set_mode()</code>           | <code>SET_MODE</code>                       | Change flight mode  |
| <code>land()</code>               | <code>MAV_CMD_NAV_LAND</code>               | Land at position    |

**Table 16:** MAVLink commands used by the navigation module.

scalar field that linearly decelerates the drone from 3.0 m/s at 20 m to a full stop at 2 m from the SSSI boundary; (3) a repulsive potential field that pushes the aircraft away at 5.0 m/s when within  $\sim 3$  m of the boundary in MANUAL mode; and (4) a hard cutoff that forces MANUAL mode if the drone enters the SSSI polygon. The module computes the repulsive offset using `cv2.pointPolygonTest` for signed distance and applies it as a velocity correction after the state handler’s waypoint command. This layered design ensures no state can inadvertently fly into restricted airspace (see Appendix Z.2 for vector field mathematics and Appendix AA.5 for NFZ margin analysis).

### 2.3 Computer Vision Pipeline


The vision subsystem is encapsulated in `vision.py`, exposing a single interface: `detect_in_image(frame) → (found, cx, cy, confidence)`. The pipeline proceeds as follows:

**Preprocessing.** Lens undistortion (if calibration data are available), colour-space handling (the IMX296 outputs BGR despite its RGB888 label—no conversion applied), resize from  $1456 \times 1088$  to  $640 \times 640$  via bilinear interpolation, and float32 normalisation ( $[0, 255] \rightarrow [0, 1]$ ), yielding input tensor  $[1, 640, 640, 3]$ .

**Inference.** The TFLite model produces output tensor  $[1, 5, 8400]$  (8400 anchor-free candidates across three feature pyramid scales). Postprocessing applies confidence filtering (threshold 0.2), best-detection selection (greedy single-target), and coordinate rescaling to original frame dimensions. The Ultralytics backend handles non-maximum suppression (NMS) internally; the TFLite backend implements equivalent filtering in `vision.py`.


**Three backends.** The system supports Ultralytics [11] (development laptop) and TFLite [3] (production Pi), with an NCNN [24] backend implemented but not yet benchmarked on the Pi. Backend selection is automatic at import

## Geofence / NFZ Protection System



**Figure 6:** Geofence protection visualisation. (a) Five protection layers: Flight Area hard boundary (emergency RTL), SSSI no-fly zone with 30 m waypoint buffer, 20 m speed ramp, 3 m repulsive field, and hard cutoff. (b) Repulsive force field that deflects the drone away from the SSSI boundary. (c) Waypoint correction: the commanded waypoint (red) is shifted to a safe position (green) by the repulsive offset. (d) Emergency mode switch when the drone enters the SSSI zone.

time. Both validated backends consume the same model weights and produce identical detection semantics (Appendix I provides full benchmark results and the dual-backend architecture details).



**Figure 7:** Computer vision pipeline: frame capture, lens undistortion, resize to  $640 \times 640$ , TFLite inference, confidence filtering, and coordinate rescaling. The entire pipeline executes in 208 ms on the Pi 5.


**Smart detection mode.** An optional `--smart-detect` flag requires  $k$  consecutive positive frames before triggering investigation (default  $k=3$ ), suppressing transient false positives at the cost of a 0.6 s delay.

**Performance.** On the Pi 5 with XNNPACK CPU delegate: 206.5 ms per frame (4.8 FPS), 100% detection rate (50/50 bench frames), 0.966 mean confidence (Appendix I.2). The model was retrained three times—at 640, 1280, and 1088 px—with the final v2 model achieving  $\text{mAP}_{50} = 0.995$  on a dataset of 300 synthetic, 16 real, and 50 negative images (Appendix I.5). At the nominal search speed of 8 m/s at 35 m altitude, adjacent frames overlap by approximately 95% ( $1 - 1.67/32.2$ ), providing 11–17 detection opportunities per target during a single flyover (speed–detection trade-off analysis in Appendix I.8).

## 2.4 Search Pattern Generation

The boustrophedon (lawnmower) pattern [8] is implemented via a rotate–rasterise–zigzag–rotate–back algorithm in `planning.py`: the GPS polygon is rasterised into a binary mask, rotated to align with the longest edge (minimising U-turns), stripped at camera-footprint spacing, and inverse-rotated back to GPS coordinates. Bézier-smoothed U-turns maintain momentum.

**Coverage guarantee.** Strip spacing derives from the thin-lens footprint:  $W_g = w_s \cdot h/f$ , lane spacing  $L = W_g(1 - \alpha)$ . At 35 m:  $W_g \approx 32.2$  m,  $L \approx 25.7$  m (20% overlap), guaranteeing every ground point falls within at least one swath under 2 m to 3 m GPS drift. A 216-configuration sweep confirmed the optimal scan angle at  $65^\circ$  to  $75^\circ$  (Appendix AA). Figure 8 shows near-linear coverage progression.



**Figure 8:** Area coverage progression over time for the boustrophedon search pattern at 35 m altitude.

## 2.5 Mission State Machine

The 20-state FSM (Figure 9, Appendix A) orchestrates the mission through four phases:


- **Startup:** INIT → CONNECTING → ARMING (GPS fix  $\geq 3D$ ,  $\geq 6$  sats) → TAKEOFF.
- **Search:** PRE.WAYPOINTS → TRANSIT\_TO\_SEARCH → SEARCH (lawnmower pattern with CV detection; PLB beacon redirect if timer expires).
- **Engagement:** CENTERING (fly to GPS estimate) → DESCENDING (reduce altitude for close-up view) → VERIFY (operator Y/N/I with 120s timeout; auto-reject resumes search) → APPROACH (7.5 m offset) → HOVER\_TARGET (servo payload release). If the operator rejects or the timeout expires, RETURN\_TO\_SEARCH climbs back to search altitude and resumes the lawnmower pattern from the next unvisited waypoint.
- **Recovery:** RETURN\_TRANSIT → RETURN\_HOME → LANDING (triple touchdown detection) → DONE.
- **Override:** MANUAL mode is reachable from any active state via the M key, RC switch, or geofence hard-cutoff. All autonomous commands cease; the pilot assumes full control. Re-entering autonomous mode requires an explicit resume command.

A detection queue with spatial deduplication (5 m reject radius), bounded capacity (20 entries), and lazy validation manages multiple targets. Six independent safety mechanisms provide defence-in-depth: manual override, RC kill switch, RTL on link loss (3s heartbeat timeout), Flight Area boundary enforcement (emergency RTL), four-layer SSSI geofence, and per-state timeouts. Appendix A lists all 20 states with their entry conditions, exit transitions, and timeout values.

## 2.6 Target Localisation

Each detection is converted from pixel coordinates to GPS using the pinhole camera model. The pixel offset  $(\Delta u, \Delta v)$  from the image centre is projected to a body-frame displacement via the ground sampling distance:

$$\text{GSD} = \frac{w_s \cdot h}{f \cdot W}, \quad d_x = \Delta u \cdot \text{GSD}, \quad d_y = \Delta v \cdot \text{GSD} \quad (1)$$



**Figure 9:** State transition diagram. Primary mission path flows left to right; engagement branches downward from SEARCH. Manual override (dashed) is accessible from any active state.


where  $w_s = 5.02$  mm is the sensor width,  $f = 5.46$  mm the calibrated focal length,  $h$  the barometric altitude, and  $W = 1456$  px the image width. The body-frame displacement  $(d_x, d_y)$  is rotated into North–East coordinates using the drone’s heading and converted to latitude/longitude on the WGS 84 ellipsoid. Appendix K documents the complete 15-stage detection-to-GPS pipeline with a full flowchart, error budget decomposition by flight phase, and roll/pitch geometry analysis (Figure 25). Multiple observations are fused through:

- **Spatial clustering** (30 m merge radius) for multi-target tracking;
- **Inverse-variance weighting** — altitude factor  $w_{\text{alt}} = (h_{\text{ref}}/h)^2$  and centrality factor  $w_{\text{cen}} = 1 + 4(1 - d/d_{\text{max}})$ ;
- **Kalman filter** on the  $[\phi, \lambda]$  state vector with measurement noise scaled by inverse observation weight, converging within 5–10 observations.

Video replay analysis (DJI flight footage,  $1456 \times 1088$  at 30 fps) yielded  $\text{CEP}_{50} = 2.3$  m and a maximum error of 16.5 m (high-speed outlier). The characteristic directional elongation stems from GPS receiver latency (100–200 ms), which multi-pass averaging in the lawnmower pattern partially mitigates. Six independent error sources contribute to a single-observation RSS uncertainty of 2.9 m at 35 m altitude, dominated by GPS receiver noise (2.3 m) and timing lag (1.5 m). After inverse-variance fusion of 15+ multi-altitude observations, the fused estimate converges to sub-metre accuracy, exceeding the 5–10 m landing requirement (R07) with  $>99\%$  probability. Appendix K provides the full error budget, convergence plots, and fusion strategy comparison.

## 2.7 Ground Station


The ground station is a browser-based dashboard served by a threaded HTTP server on the Pi at port 8090, accessible from any device on the local Wi-Fi network. It provides three panels: (1) live MJPEG video stream with detection



overlays; (2) 2D GPS grid showing detection clusters, drone position, and search path; and (3) command buttons (investigate, confirm, reject, land, RTL, manual override). MJPEG was chosen over H.264/WebRTC for its minimal client complexity (single `<img>` tag), zero transcoding load, and per-frame decodability. Measured latency is 300 ms at 0.3 MB s<sup>-1</sup>. The Pi auto-detects headless operation (no DISPLAY), enabling SSH deployment without X11 forwarding (Appendix J).

## 2.8 Simulation Framework

Four simulation levels provide progressive fidelity (Appendix Q): (1) interactive simulator (keyboard flight over satellite map with ground-truth comparison); (2) SITL integration (full ArduPilot autopilot with simulated camera view); (3) SITL with real webcam; (4) DJI video replay with synchronised SRT telemetry. All four levels share the same detection model, GPS estimation mathematics, and MAVLink command interface, so transitioning to real hardware requires only configuration changes. Simulation alone caught 12 defects—including a MAVLink race condition during takeoff, incorrect latitude scaling, auto-disarm timeout, and GPS timing lag—all resolved before any hardware integration.



**Figure 11: Mission timeline showing state transitions during a simulated end-to-end SAR mission. The horizontal axis represents elapsed time; coloured blocks indicate active states.**

### 3 Requirements Verification

Table 17 maps each requirement from the AENGM0074 project brief (Release 2.1) to its verification method, supporting evidence, and current verification status.

**Table 17:** Requirements verification summary. Status indicates the highest fidelity at which each requirement has been tested.

| Req | Requirement                | Method                                                                                                          | Evidence                                                                                          | Verification Status           |
|-----|----------------------------|-----------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|-------------------------------|
| R01 | Fly within Flight Area     | KML polygon load + waypoint clipping + runtime <code>pointPolygonTest</code> → emergency RTL                    | SITL telemetry, geofence unit test, <code>_enforce_geofence()</code> at ~20–50 Hz                 | Verified (SITL)               |
| R02 | No SSSI overfly            | Four-layer NFZ (plan-time filter, speed ramp, repulsive field, hard cutoff) + firmware fence                    | SITL multi-angle approach, geofence diagram                                                       | Verified (SITL)               |
| R03 | Take off within 5 m of TOL | KML point place-mark, vertical take-off + arming-time distance check                                            | SITL telemetry (<1 m error), runtime >5 m warning                                                 | Verified (SITL)               |
| R04 | Max 50 m altitude          | <code>TARGET_ALT=35 m</code> + <code>send_global_target()</code> 50 m cap + firmware <code>FENCE_ALT_MAX</code> | Config validation, navigation cap, 28 unit tests                                                  | Verified (SITL)               |
| R05 | Search + identify items    | Lawnmower + YOLOv8n (mAP <sub>50</sub> =0.995, 4.8 FPS)                                                         | End-to-end sim, bench test, DJI video proxy                                                       | Verified (simulation + video) |
| R06 | PLB Focus Area             | B key / <code>--beacon-delay</code> mid-flight injection                                                        | SITL diversion + resume, flight log, 6 unit tests                                                 | Verified (SITL)               |
| R07 | Land 5–10 m from casualty  | 7.5 m offset via <code>landing_offset_7.5m()</code> + post-landing distance printout                            | Probability analysis (CEP <sub>50</sub> =2.3 m), 5 unit tests, SITL smoke test, runtime R07 check | Verified (SITL + analysis)    |
| R08 | Autonomy justified         | L7–8 search / L3–4 landing, operator verify gate                                                                | Design rationale, Sheridan framework                                                              | Verified (design)             |
| R09 | RTH + Motor Cutoff         | RC kill, software M, auto RTL failsafes                                                                         | SITL: 3 independent override tests                                                                | Verified (SITL)               |
| R10 | Report lat/lon + images    | GPS projection, web dashboard, <code>mission_detections/</code> (JPEG + JSON per event)                         | Detection images + JSON metadata saved on every operator decision (Y/N/I/X)                       | Verified (SITL + bench)       |
| R11 | Company Pilot designated   | 12-item checklist + RC override test                                                                            | Team plan, field-day bench test, checklist                                                        | Verified (procedural + bench) |
| R12 | Public GitHub, MIT licence | <code>github.com/DimaChuprikov/ScoutPro</code> + <code>LICENSE</code> file                                      | Repository live, public, MIT <code>LICENSE</code> in root                                         | Verified (inspection)         |

Of the 12 requirements, 8 are fully verified at bench or higher fidelity (R01–R04, R09–R12), 3 are verified in simulation only (R05–R07), and 1 has a partial gap (R06 not field-tested with a live PLB).

Recent code improvements closed several evidence gaps. R01 now includes a runtime Flight Area geofence with emergency RTL (not just plan-time clipping). R03 checks take-off location distance at arming. R04 enforces a 50 m hard cap in the navigation module. R07 prints the actual casualty distance at landing. R10 saves structured detection images and JSON metadata to `mission_detections/`, and R12 has an explicit MIT `LICENSE` file.

The highest-risk requirements—R01 (flight area), R02 (SSSI), and R09 (override)—are each protected by multiple independent mechanisms verified at both simulation and hardware tiers. The most significant verification gap is the absence of outdoor flight data for R05 (search and detection) and R07 (landing accuracy). Both are verified quantitatively in simulation and through DJI video proxy analysis but remain unvalidated under real flight conditions due to weather cancellation. Appendix AG provides full evidence narratives for each requirement.


## 4 Evaluation

The plus/delta review below assesses technical performance against the 12 requirements (R01–R12) using four evidence sources: SITL simulation (Tier 1, 100+ hours), Docker environment tests (Tier 2), bench tests on the assembled Pi 5 hardware (Tier 3, one field day), and DJI flight video replayed through the detection pipeline (Tier 3 proxy). No outdoor flight took place (Tiers 4–5 outstanding). D7 addresses team-working aspects individually.


### 4.1 Plus/Delta Summary

Table 18 summarises the principal strengths and areas for improvement. Each item includes the specific evidence that supports the assessment.

### 4.2 Discussion of Key Findings




(a) Detection confidence vs. altitude. Confidence remains above 0.9 up to 40 m, dropping at higher altitudes as the target subtends fewer pixels.




(b) Detection rate vs. flight speed. At the nominal search speed of  $8 \text{ m s}^{-1}$ , detection rate exceeds 95%.

**Figure 12:** Detection performance sensitivity analysis from DJI flight video replay.



(a) Per-frame latency breakdown: capture, undistortion, resize, inference, and postprocessing.



(b) Spatial detection heatmap showing detection density across the survey area during simulated search.

**Figure 13:** System performance: inference latency budget and spatial detection distribution.

**Architecture resilience (P1, P2, P5, P8).** The single-interface vision module (`detect_in_image`) enabled three model retrains (640, 1280, 1088 px datasets) with zero changes to the orchestrator or ground station. Adding lens undistortion—not anticipated during initial design—required modifying one file and added only 1.5 ms per frame.

**Detection performance (P4, D2, D4, D5, D9).** The retrained YOLOv8n achieved  $\text{mAP}_{50} = 0.995$  (validation set) and 50/50 bench detection at 0.966 mean confidence (Figure 12a). At 35 m, the target subtends 36 model-input pixels—a comfortable margin above the 20 px floor. However, the 0.995  $\text{mAP}_{50}$  likely overestimates real-world accuracy due to train/validation overlap (only 16 of 366 images are from real footage). The confidence threshold (0.2) was set by visual inspection rather than a precision–recall curve sweep; characterising this trade-off is a priority for future work. FP16 deployment could double throughput to  $\sim 10$  FPS (Figure 13a).

**GPS estimation accuracy (D3).** The pipeline projects each detection’s pixel offset through the calibrated camera model and the drone’s reported position to produce a GPS estimate. DJI flight video analysis with known ground-truth positions yielded  $\text{CEP}_{50} = 2.3\text{ m}$  and a worst-case error of 16.5 m (Figures 10a and 10b). The worst case stems from GPS timing lag compounding with heading-aligned motion. Inverse-variance-weighted averaging and Kalman filtering partially mitigate this bias, but first-order lag compensation (offsetting GPS by  $v \cdot \Delta t_{\text{lag}}$  in the heading direction) remains unimplemented.

**Testing effectiveness (P3, P7).** The progressive framework caught 12 defects at Tiers 1–3, each resolved in minutes to hours at zero hardware risk: 5 at Tier 1 (simulation, 1.5 hr total), 6 at Tier 2 (Pi bench, 4.7 hr), and 1 at Tier 3 proxy. Seven would have caused mission-critical failures in flight: geofence sign inversion (drone pushed *toward* SSSI), BGR colour inversion (detection failure), TFLite bounding box coordinate mismatch (GPS estimates off by  $>100\text{ m}$ ), repulsive force reversal, barometric drift loop, duplicate target queuing, and missing timeout guard—yet all were trivially reproducible in simulation or on the bench. Appendix AH.1 provides the full defect register with resolution times and assessed flight-day impact.

**Weather cancellation as information-maximisation, not failure (D1).** The scheduled outdoor flight was cancelled due to sustained winds of  $12\text{ m s}^{-1}$  with gusts to  $18\text{ m s}^{-1}$ —above the EDU-450’s  $10\text{ m s}^{-1}$  operational limit. This cancellation represents the single largest gap in the project, eliminating any possibility of Tier 4/5 validation. However, the team treated the bench day as an opportunity for information maximisation rather than a binary failure, executing four activities that would not have occurred had the flight proceeded on schedule:

1. **FOV calibration** ( $f=7.0 \rightarrow 5.46\text{ mm}$ , tape measure at 1 m) — corrected a 28 % GPS estimation bias that would have produced landing errors exceeding the 10 m R07 radius.
2. **Lens distortion mapping** ( $\text{RMS} = 0.399$ ,  $14 \times 9$  checkerboard) — confirmed that the IMX296 global shutter has negligible barrel distortion and undistortion is optional.
3. **Inference benchmarking** (50 runs  $\times$  3 models) — established the 206.5 ms baseline and validated 50/50 detection at 0.966 mean confidence.
4. **DJI video analysis pipeline** (built post-cancellation) — yielded the confidence-vs-altitude curve (Figure 12a), detection-vs-speed analysis (Figure 12b), and GPS accuracy characterisation ( $\text{CEP}_{50}=2.3\text{ m}$ , Figure 10a).

These four data products form the quantitative backbone of this evaluation. The project lacks *system-level* flight validation but possesses *subsystem-level* quantitative data exceeding what a single flight sortie would typically yield (Appendix M).

**Incomplete flight validation (D1, D6, D7, D8).** The integrated system has never executed a search pattern over real terrain. Detection range versus altitude in natural lighting, false-positive rate against real backgrounds, wind-induced motion blur, and GPS accuracy under dynamic manoeuvres remain unvalidated (only DJI video proxy data are available). Bench results provide reasonable confidence that outdoor operation will require parameter tuning rather than architectural changes.

### 4.3 Testing Framework Summary

The five-tier progressive testing framework (Table 20) advanced from pure software to full autonomy, with each tier serving as a gate that must pass before the next is attempted. This structure caught 12 defects at low-cost tiers (Appendix AH.1) and is the primary reason no architectural faults remained by the time hardware was assembled.

Tiers 1–3 were fully exercised; Tier 4 was partially achieved through DJI video replay as a proxy after weather cancellation. Tier 5 remains outstanding. Of the 12 defects (Appendix AH.1), 5 were caught at Tier 1 (simulation, 1.5 h total fix time), 6 at Tier 2 (Pi bench, 4.7 h), and 1 at Tier 3 proxy (video analysis, mitigation pending). Two intermediate tiers unique to this project—Docker environment testing (Tier 2) and passive zero-command flight (Tier 4)—are absent from most university drone projects yet caught six of the twelve defects. Appendix E (Section E.2) details the full five-tier methodology and per-tier script inventory.

**Summary.** Tiers 1–3 produced strong quantitative baselines: 12 defects caught and resolved in  $\sim 6\text{ h } 10\text{ min}$  total, 206.5 ms inference confirmed over 50 runs,  $\text{CEP}_{50}=2.3\text{ m}$  GPS accuracy from video, and  $\text{mAP}_{50}=0.995$  (upper bound due to train/val overlap). However, the system remains *incompletely validated*: Tiers 4–5 (passive and autonomous outdoor flight) were never executed, so requirements R05, R06, and R07 are verified only in simulation.

Of the nine delta items, three are parameter-tuning gaps (D2, D3, D4) resolvable without code changes, two are functional gaps requiring implementation (D6, D8), one is a field-testing gap (D7), and three are methodological limitations (D1, D5, D9) requiring additional flight time and data. The progressive testing framework retired the highest-cost integration faults in  $\sim 6\text{ h}$  of bench time; completing the remaining tiers requires scheduling, favourable weather, and resolution of D6/D8. Appendix AH.3 details lessons learned and second-iteration improvements.

## References

- [1] R. R. Murphy. *Disaster Robotics*. Cambridge, MA: MIT Press, 2014.
- [2] J. Chen and X. Ran. “Deep Learning With Edge Computing: A Review”. In: *Proceedings of the IEEE* 107.8 (2019), pp. 1655–1674.
- [3] Google. *TensorFlow Lite: Deploy Machine Learning Models on Mobile and Edge Devices*. <https://www.tensorflow.org/lite>. Accessed: 2026. 2023.
- [4] UK Government. *Wildlife and Countryside Act 1981*. <https://www.legislation.gov.uk/ukpga/1981/69>. Accessed: 2026. 1981.
- [5] UK Civil Aviation Authority. *Unmanned Aircraft System Operations in UK Airspace – Guidance (CAP 722)*. <https://www.caa.co.uk/drones/>. Accessed: 2026. 2024.
- [6] JARUS. *JARUS Guidelines on Specific Operations Risk Assessment (SORA)*. Tech. rep. Edition 2.0. JAR-DELWG6-D.04. Joint Authorities for Rulemaking on Unmanned Systems, 2019.
- [7] E. Galceran and M. Carreras. “A Survey on Coverage Path Planning for Robotics”. In: *Robotics and Autonomous Systems* 61.12 (2013), pp. 1258–1276.
- [8] H. Choset. “Coverage for Robotics – A Survey of Recent Results”. In: *Annals of Mathematics and Artificial Intelligence* 31.1–4 (2001), pp. 113–126.
- [9] C. Burke et al. “Optimising UAV Image Capture and Analysis for Thermal Imaging Applications in Search and Rescue”. In: *Drones* 3.3 (2019), p. 73.
- [10] ArduPilot Dev Team. *MAVProxy: A UAV Ground Station Software Package for MAVLink Based Systems*. <https://ardupilot.org/mavproxy/>. Accessed: 2026. 2023.
- [11] G. Jocher, A. Chaurasia, and J. Qiu. *Ultralytics YOLOv8*. <https://github.com/ultralytics/ultralytics>. Accessed: 2026. 2023.
- [12] T.-Y. Lin et al. “Microsoft COCO: Common Objects in Context”. In: *arXiv preprint arXiv:1405.0312* (2014).
- [13] J. Tobin et al. “Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World”. In: *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2017, pp. 23–30.
- [14] T. Wagner et al. “Toward Autonomous Unmanned Aerial Vehicles”. In: *Proceedings of the 44th AIAA Aerospace Sciences Meeting and Exhibit*. AIAA, 2006.
- [15] T. B. Sheridan. *Humans and Automation: System Design and Research Issues*. New York: Wiley, 2002.
- [16] M. R. Endsley and D. B. Kaber. “Level of Automation Effects on Performance, Situation Awareness and Workload in a Dynamic Control Task”. In: *Ergonomics* 42.3 (1999), pp. 462–492.
- [17] M. L. Cummings. “Man Versus Machine or Man + Machine?” In: *IEEE Intelligent Systems* 29.5 (2014), pp. 62–69.
- [18] M. A. Goodrich and A. C. Schultz. “Human–Robot Interaction: A Survey”. In: *Foundations and Trends in Human–Computer Interaction* 1.3 (2007), pp. 203–275.
- [19] M. Lyu et al. “Unmanned Aerial Vehicles for Search and Rescue: A Survey”. In: *Remote Sensing* 15.13 (2023), p. 3266.
- [20] M. A. Goodrich et al. “Supporting Wilderness Search and Rescue Using a Camera-Equipped Mini UAV”. In: *Journal of Field Robotics* 25.1–2 (2008), pp. 89–110.
- [21] M. Quigley et al. “ROS: An Open-Source Robot Operating System”. In: *ICRA Workshop on Open Source Software*. Vol. 3. 3.2. 2009, p. 5.
- [22] Sony Semiconductor Solutions. *IMX296: 1/2.9-Type Global Shutter CMOS Image Sensor*. <https://www.sony-semicon.com/en/products/is/industry/global-shutter.html>. Accessed: 2026. 2020.
- [23] B. Bona and M. Indri. “Rolling and Global Shutter Effect Compensation for Cameras Mounted on Moving Platforms”. In: *IEEE Sensors Journal* 12.11 (2012), pp. 3100–3108.
- [24] T. Ni. *ncnn: A High-Performance Neural Network Inference Computing Framework Optimised for Mobile Platforms*. <https://github.com/Tencent/ncnn>. Accessed: 2026. 2017.
- [25] J. Tremblay et al. “Training Deep Networks with Synthetic Data: Bridging the Reality Gap by Domain Randomization”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*. 2018, pp. 969–977.
- [26] A. Bochkovskiy, C.-Y. Wang, and H.-Y. M. Liao. “YOLOv4: Optimal Speed and Accuracy of Object Detection”. In: *arXiv preprint arXiv:2004.10934* (2020).
- [27] A. Shrivastava, A. Gupta, and R. Girshick. “Training Region-Based Object Detectors with Online Hard Example Mining”. In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 2016, pp. 761–769.
- [28] ArduPilot Dev Team. *ArduCopter Failsafe Mechanisms*. <https://ardupilot.org/copter/docs/failsafe-landing-page.html>. Accessed: 2026. 2024.
- [29] K. Forsberg and H. Mooz. “The Relationship of System Engineering to the Project Cycle”. In: *Engineering Management Journal* 3.3 (1991), pp. 36–43.
- [30] P. Koopman and M. Wagner. “Challenges in Autonomous Vehicle Testing and Validation”. In: *SAE International Journal of Transportation Safety*. Vol. 4. 1. 2016, pp. 15–24.
- [31] ArduPilot Dev Team. *ArduPilot SITL Simulator: Software in the Loop Simulation*. <https://ardupilot.org/dev/docs/sitl-simulator-software-in-the-loop.html>. Accessed: 2026. 2024.
- [32] National Aeronautics and Space Administration. *Unmanned Aircraft Systems (UAS) Integration in the National Airspace System (NAS) Project: Test and Evaluation*. Tech. rep. NASA/TM-2015-218817. Describes progressive V&V methodology for UAS: simulation, hardware-in-the-loop, limited flight envelope, full mission. NASA, 2015.
- [33] Z. Zhang. “A Flexible New Technique for Camera Calibration”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* 22.11 (2000), pp. 1330–1334.
- [34] G. Bradski. *The OpenCV Library*. <https://opencv.org/>. Accessed: 2026. 2000.
- [35] Google. *XNNPACK: High-Efficiency Floating-Point Neural Network Inference Operators*. <https://github.com/google/XNNPACK>. Accessed: 2026. 2023.
- [36] P. Misra and P. Enge. *Global Positioning System: Signals, Measurements, and Performance*. 2nd. Lincoln, MA: Ganga-Jamuna Press, 2006.

- [37] R. E. Kalman. "A New Approach to Linear Filtering and Prediction Problems". In: *Journal of Basic Engineering* 82.1 (1960), pp. 35–45.
- [38] R. David et al. "TensorFlow Lite Micro: Embedded Machine Learning for TinyML Systems". In: *Proceedings of Machine Learning and Systems (MLSys)* 3 (2021), pp. 800–811.
- [39] Hailo Technologies. *Hailo-8L: Edge AI Processor Datasheet*. <https://hailo.ai/products/ai-accelerators/hailo-8l-ai-accelerator/>. Accessed: 2026. 2023.
- [40] J. Tremblay et al. "Training Deep Networks with Synthetic Data: Bridging the Reality Gap by Domain Randomization". In: *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*. 2018, pp. 969–977.
- [41] P. Rudol and P. Doherty. "Human Body Detection and Geolocalization for UAV Search and Rescue Missions Using Color and Thermal Imagery". In: *Proceedings of the IEEE Aerospace Conference*. IEEE, 2008, pp. 1–8.
- [42] UK Civil Aviation Authority. *CAP 722: Unmanned Aircraft System Operations in UK Airspace*. <https://publicapps.caa.co.uk/modalapplication.aspx?catid=1&pagetype=65&appid=11&mode=detail&id=13177>. Accessed: 2026. 2024.
- [43] S. Waharte and N. Trigoni. "Supporting Search and Rescue Operations with UAVs". In: *Proceedings of the International Conference on Emerging Security Technologies (EST)*. IEEE, 2010, pp. 142–147.
- [44] S. Hayat, E. Yanmaz, and R. Muzaffar. "Survey on Unmanned Aerial Vehicle Networks for Civil Applications: A Communications Viewpoint". In: *IEEE Communications Surveys & Tutorials* 18.4 (2016), pp. 2624–2661.
- [45] Raspberry Pi Foundation. *Raspberry Pi 5 Documentation*. <https://www.raspberrypi.com/documentation/computers/raspberry-pi-5.html>. Accessed: 2026. 2023.
- [46] Arm Ltd. *Cortex-A76 Technical Reference Manual*. Revision r3p0. 2018.
- [47] R. Muñoz-Salinas, M. J. Marín-Jiménez, and R. Medina-Carnicer. "Efficient Video Streaming for Mobile Robot Teleoperation". In: *Proceedings of the IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2016, pp. 5248–5253.
- [48] Apple Inc. *HTTP Live Streaming*. <https://developer.apple.com/streaming/>. Accessed: 2026-03-20. 2023.
- [49] E. Gamma et al. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1994.
- [50] T. Wiegand et al. "Overview of the H.264/AVC Video Coding Standard". In: *IEEE Transactions on Circuits and Systems for Video Technology* 13.7 (2003), pp. 560–576.
- [51] Dronecode Foundation. *QGroundControl: Open Source Ground Control Station*. <https://qgroundcontrol.com/>. Accessed: 2026. 2020.
- [52] E. D. Kaplan and C. J. Hegarty. *Understanding GPS/GNSS: Principles and Applications*. 3rd. Boston, MA: Artech House, 2017.
- [53] R. W. Beard and T. W. McLain. *Small Unmanned Aircraft: Theory and Practice*. Princeton University Press, 2012.
- [54] D. B. Barber et al. "Vision-based Target Geo-location Using a Fixed-wing Miniature Air Vehicle". In: *Journal of Intelligent and Robotic Systems* 47.4 (2006), pp. 361–382.
- [55] RTCA. *DO-178C: Software Considerations in Airborne Systems and Equipment Certification*. Tech. rep. RTCA Inc., 2011.
- [56] D. C. Brown. "Decentering Distortion of Lenses". In: *Photogrammetric Engineering* 32.3 (1966), pp. 444–462.
- [57] W. Johnson. *Rotorcraft Aeromechanics*. Cambridge University Press, 2013.
- [58] R. R. Murphy. *Disaster Robotics*. 2nd. Includes chapters on UAV safety margins and offset landing. Cambridge, MA: MIT Press, 2016.
- [59] L. Meier et al. *MAVLink 2: Micro Air Vehicle Communication Protocol*. [https://mavlink.io/en/guide/mavlink\\_2.html](https://mavlink.io/en/guide/mavlink_2.html). Accessed: 2026. 2017.
- [60] B. Settles. *Active Learning Literature Survey*. Computer Sciences Technical Report 1648. 2009.
- [61] F. Golden and M. Tipton. "Essentials of Sea Survival". In: *Essentials of Sea Survival*. Hypothermia survival curves. Human Kinetics, 2002.
- [62] J. K. Stolaroff et al. "Energy Use and Life Cycle Greenhouse Gas Emissions of Drones for Commercial Package Delivery". In: *Nature Communications* 9.1 (2018), p. 409.

## Appendix Guide

The following appendices provide supplementary detail beyond the 15-page assessed body. Each appendix is self-contained and can be read independently.

---

### App. Contents

---

#### *Core system detail*

- A State transition table — 20 states, 32 transitions with trigger conditions and actions.
- B Configuration parameters — all tuneable values, ranges, and rationale for reproducibility.
- C Model training and dataset — dataset construction (synthetic, real, negatives), augmentation, Colab workflow.
- D Safety and risk management — failure-mode analysis, regulatory compliance, mitigations.

#### *Verification and validation*

- E Testing methodology — five-tier progressive framework, 20 failure scenarios, pass criteria.
- F Field testing and results — calibration data, inference benchmarks, detection rates.
- G Future work — post-submission roadmap organised by feasibility.

#### *Subsystem deep dives*

- H Edge inference architecture — NCNN vs TFLite decision, quantisation, Pi 5 constraints.
- I Computer vision: extended analysis — detection pipeline, altitude dependence, dual backend.
- J Video streaming and ground station architecture — MJPEG, browser dashboard, headless operation.
- K GPS target estimation: a deep dive — Kalman filter, spatial clustering, error budget.
- L Progressive testing methodology — test ladder philosophy, integration strategy.
- M Field day: adaptation under uncertainty — bench calibration pivot, concrete outputs.
- N Contingency planning and deliverable tiers — three-tier fallback, eight failure scenarios.
- O Test script reference — all 58 scripts with usage, dependencies, and tier mapping.
- P Sensor calibration — FOV, lens distortion, GPS ground truth, error budget cross-reference.
- Q Simulation validation — sim-to-real gap analysis, transfer confidence evidence.
- R Mission flow — end-to-end flight narrative from power-on to return-to-launch.
- S Centering vs. direct offset landing — two-step localisation trade-off analysis.
- T MAVLink communication architecture — topology, mavproxy bridge, failsafe mechanisms.
- U Raw inference vs. effective pipeline throughput — FPS breakdown analysis.

#### *Search and flight optimisation*

- V Flight path optimisation — strategy comparison and trade-offs.
- W Payload release mechanism — servo hardware, low-hover rationale, actuation sequence.
- X Model comparison and future vision improvements — three generations, tiling, COCO fallback.
- Y Multi-objective search optimisation — energy, rotation, coverage trade-offs.
- Z Focus area redirect and repulsive field — PLB beacon integration, NFZ avoidance.
- AA Search path optimisation — six candidate strategies, energy model, parametric sweep.
- AB Search parameter optimisation — 216-configuration physics simulation, Pareto front, sensitivity.
- AC Formal multi-objective optimisation — mathematical problem statement and constraints.
- AD Design strategy: parameter derivation chain — logical dependency reasoning.
- AE Vision performance analysis — altitude vs. pixel size, motion blur, detection envelope (real data).
- AF Decision flow — executive narrative of how search parameters were selected.


#### *Extended body sections*

- AG Requirements verification detail — R01–R12 evidence narratives.
  - AH Evaluation detail — defect register, lessons learned, second iteration changes.
  - AI Development methodology — simulation-first philosophy, 10-phase timeline, progressive testing, ambition levels, lessons for future teams.
-

## A State Transition Table

This appendix provides the complete specification of the mission state machine, serving as the authoritative reference for understanding how the drone progresses through each phase of an autonomous SAR mission and how it responds to operator inputs, timeouts, and failure conditions.

Figure 14 shows the complete 20-state transition diagram with all 32 transitions. Table 22 lists every state in the mission state machine, its entry condition, the actions performed while active, the possible exit transitions, and any timeout behaviour. The state machine contains 20 states with 32 distinct transitions implemented in `state_machine.py`.



**Figure 14:** Complete state machine diagram showing all 20 states and 32 transitions. The primary mission path flows left to right; engagement branches downward from SEARCH; recovery paths return to the search or home. Manual override is accessible from any active state.

**Table 22:** Complete state transition table for the SAR mission state machine (20 states, 32 transitions). Abbreviations: WP = waypoint, alt = altitude, dist = distance, TGT = target, NFZ = no-fly zone, IOI = item of interest, PLB = personal locator beacon.

| State | Entry Condition          | Actions                      | Exit Transitions                  | Timeout                     |
|-------|--------------------------|------------------------------|-----------------------------------|-----------------------------|
| INIT  | Power-on (initial state) | Attempt connection every 1 s | MAVLink to CONNECTION_STR<br>OK → | None (retries indefinitely) |

*Continued on next page*

| State             | Entry Condition                                                | Actions                                                                                                                                                                     | Exit Transitions                                                                                                                                                                                                                                             | Timeout                            |
|-------------------|----------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------|
| CONNECTING        | MAVLink socket opened                                          | Wait for heartbeat; request data streams; set SITL speedup (SIM only)                                                                                                       | Heartbeat received → <b>ARMING</b>                                                                                                                                                                                                                           | Warning at 15 s intervals          |
| ARMING            | Heartbeat confirmed                                            | Wait for GPS fix ( $\geq 3D$ , $\geq 6$ sats); set <b>GUIDED</b> mode; send <b>ARM</b> command; record home position                                                        | Motors armed → <b>TAKEOFF</b>                                                                                                                                                                                                                                | Warning at 120 s                   |
| TAKEOFF           | ARM confirmed                                                  | Send <b>NAV_TAKEOFF</b> to <b>TARGET_ALT</b> ; monitor altitude and armed status                                                                                            | $\text{alt} \geq 0.9 \times \text{TARGET\_ALT}$ :<br>pre-WPs exist → <b>PRE_WAYPOINTS</b><br>WPs exist → <b>TRANSIT_TO_SEARCH</b><br>no WPs → <b>HOVER</b><br>Disarmed (SIM) → <b>ARMING</b><br>Disarmed (REAL) → <b>DONE</b>                                | Warning at 60 s                    |
| PRE_WAYPOINTS     | Transit WPs defined (corridor path)                            | Fly each pre-WP sequentially at <b>TRANSIT_SPEED</b> ; advance when dist $< 2$ m                                                                                            | All pre-WPs reached → <b>TRANSIT_TO_SEARCH</b><br>No search WPs → <b>HOVER</b>                                                                                                                                                                               | None                               |
| TRANSIT_TO_SEARCH | Pre-WPs complete or altitude reached                           | Fly to first search WP at <b>TRANSIT_SPEED</b>                                                                                                                              | dist to WP[0] $< 2$ m → <b>SEARCH</b>                                                                                                                                                                                                                        | None                               |
| SEARCH            | Reached search start or resumed after reject/IOI               | Orient yaw (diagonal alignment); fly lawnmower WPs at altitude-dependent speed; run CV detection; filter & queue confirmed detections; PLB beacon redirect if delay elapsed | Valid queued TGT → <b>CENTERING</b><br>All WPs done, rescans left → <b>SEARCH</b> (lower alt)<br>All WPs done, no rescans → <b>DONE</b><br>Rescan floor reached → <b>DONE</b>                                                                                | None                               |
| CENTERING         | Detection dequeued from search or manual                       | Fly toward TGT GPS at current alt; update lock if CV re-detects within <b>LOCK_RADIUS</b> ; queue new detections outside lock radius                                        | dist to TGT $< 1$ m → <b>VERIFY</b>                                                                                                                                                                                                                          | 60 s → <b>SEARCH</b>               |
| DESCENDING        | (Currently bypassed—reserved for future descent-before-verify) | Immediately transitions                                                                                                                                                     | Instant → <b>VERIFY</b>                                                                                                                                                                                                                                      | None                               |
| VERIFY            | Centered over TGT ( $< 1$ m)                                   | Hover over TGT; collect GPS avg samples; display countdown; wait for operator input via keyboard or browser                                                                 | <b>Y</b> : GPS avg (10 s) → select landing side (N/E/S/W) → <b>APPROACH</b><br><b>N</b> : reject TGT; queue check → <b>CENTERING</b> / <b>RETURN_FROM_MANUAL</b> / <b>RETURN_TO_SEARCH</b> / <b>SEARCH</b><br><b>I</b> : log as IOI; queue check → same as N | 120 s → auto-reject, <b>SEARCH</b> |
| APPROACH          | Operator confirmed + landing side selected                     | Fly to 7.5 m offset landing coordinate at 3 m alt                                                                                                                           | dist $< 2$ m and alt $< 4$ m → <b>HOVER_TARGET</b>                                                                                                                                                                                                           | None                               |

*Continued on next page*

| State              | Entry Condition                                                                     | Actions                                                                                                                                         | Exit Transitions                                                                                                                                                                               | Timeout                                            |
|--------------------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------|
| HOVER_TARGET       | Above landing point at 3 m                                                          | Hold position; servo stage 1 at $t=3$ s (partial release, PWM 1200); servo stage 2 at $t=6$ s (full release, PWM 1900); close servo at $t=15$ s | $t > 15$ s:<br>pre-WPs exist $\rightarrow$ RETURN_TRANSIT<br>else $\rightarrow$ RETURN_HOME                                                                                                    | Fixed 15 s sequence                                |
| RETURN_TRANSIT     | Pre-WPs need retracing                                                              | Climb to search alt; retrace pre-WPs in reverse order at TRANSIT_SPEED                                                                          | All return WPs reached $\rightarrow$ RETURN_HOME                                                                                                                                               | None                                               |
| RETURN_HOME        | Return transit complete or no pre-WPs                                               | Fly to home position at TARGET_ALT; guard against invalid home if GPS never fixed                                                               | dist to home $< 2$ m $\rightarrow$ LANDING<br>No GPS fix $\rightarrow$ LANDING (in place)                                                                                                      | None                                               |
| LANDING            | Over home position or GPS-less fallback                                             | Send NAV_LAND; monitor alt; detect touchdown (alt $< 0.5$ m for 5 consecutive ticks); disarm on touch; retry LAND every 5 s if not descending   | Touchdown or auto-disarm $\rightarrow$ DONE<br>5 LAND retries failed $\rightarrow$ force disarm $\rightarrow$ DONE                                                                             | 90 s $\rightarrow$ force disarm $\rightarrow$ DONE |
| RETURN_TO_SEARCH   | Target rejected/IOI'd; departure point recorded                                     | Fly to departure lat/lon at search alt                                                                                                          | dist $< 3$ m and alt $> 0.85 \times \text{search\_alt}$ $\rightarrow$ SEARCH                                                                                                                   | None                                               |
| RETURN_FROM_MANUAL | Exiting MANUAL mode; $> 5$ m from departure                                         | Fly to manual departure lat/lon/alt at TRANSIT_SPEED                                                                                            | dist $< 3$ m $\rightarrow$ resume previous state                                                                                                                                               | None                                               |
| HOVER              | No waypoints generated after takeoff                                                | Hold position; no commands                                                                                                                      | —                                                                                                                                                                                              | 60 s $\rightarrow$ DONE                            |
| MANUAL             | Operator presses <b>M</b> in any active state, or geofence detects drone inside NFZ | WASD velocity control; R/F climb/descend; Q/E yaw; CV detections queued silently for later investigation                                        | M pressed:<br>TGT visible $\rightarrow$ CENTERING<br>queued TGT $\rightarrow$ CENTERING<br>$> 5$ m from departure $\rightarrow$ RETURN_FROM_MANUAL<br>else $\rightarrow$ resume previous state | None                                               |
| DONE               | Mission complete, landing confirmed, or fatal error                                 | Display final landing error; log mission time; no further commands sent                                                                         | Terminal state (exit main loop)                                                                                                                                                                | 3 s display then exit                              |

**Global overrides** (apply in any non-terminal state):

- **M key** — toggles MANUAL override from any active flight state.
- **K key** — clears all rejected targets and items of interest (re-enables detection areas).
- **B key** (during SEARCH) — triggers PLB beacon redirect to focus area polygon.
- **ESC key** — immediate loop exit (emergency stop).
- **Geofence violation** — if the drone enters the NFZ boundary, the state machine forces a transition to MANUAL and halts all velocity commands. Within the NFZ slow zone, speed is ramped down proportionally and a repulsive force vector is applied.
- **RC failsafe** — if the RC link is lost, telemetry detects the failsafe flag and the flight controller's built-in RTL behaviour takes precedence.

## B Configuration Parameters

This appendix documents the tuneable parameters that govern all aspects of the mission—from flight altitudes and speeds to detection thresholds and geofence boundaries—enabling reproducibility of the reported results and informed

adjustment for different operating environments.

All tuneable mission parameters are centralised in a single `config.py` module. Table 23 lists the key operational values used during flight testing. Parameters are loaded at startup and may be overridden via environment variables where noted.

## C Model Training and Dataset

This appendix details the complete training pipeline for the onboard YOLOv8n detector, from dataset construction through augmentation strategy to model export, providing the methodological detail needed to reproduce or improve upon the reported detection performance.

Training a reliable aerial object detector for SAR requires data that closely represents the deployment domain—a mannequin viewed from altitude against a grassy field background. Obtaining large volumes of real labelled aerial imagery is expensive and weather-dependent, so a mixed-source strategy was adopted combining synthetic composites, real labelled frames, and hard negatives.

### C.1 Training Data Strategy

The final dataset (`dataset_v2`) contains 366 images at the native camera resolution of  $1456 \times 1088$  pixels drawn from three complementary sources:

1. **Synthetic composites** (300 images). A cutout of the mannequin (`dummy.png`, with alpha channel) is composited onto random crops from real DJI flight video backgrounds at altitude-correct scales, providing systematic variation of target size, orientation, and position.
2. **Real labelled frames** (16 images). Frames extracted from a DJI flight recording at multiple altitudes (15–50 m) and manually annotated with bounding boxes, anchoring the model to real-world textures, lighting, and target appearance.
3. **Negative images** (50 images). Background frames from the same flight video containing no mannequin, paired with empty label files, teaching the model to suppress false positives on terrain features such as shadows, tarpaulins, and path markings.

This three-source approach follows the domain-randomisation principle [13]: synthetic data provides scale and diversity, real data provides domain fidelity, and negatives provide specificity. Training at the native sensor resolution rather than the standard  $640 \times 640$  YOLO input size allows the network to learn fine-grained features during training, even though inference-time resizing reduces the input to  $640 \times 640$  [11].

### C.2 Domain Randomisation

Domain randomisation is the deliberate injection of controlled variation into synthetic training data so that real-world imagery appears as “just another variant” to the trained network [13]. For aerial SAR, the key dimensions of variation are target size (a function of altitude), target orientation, background texture, and illumination. Rather than attempting photorealistic rendering—which is expensive and still imperfect—the pipeline randomises each dimension independently, forcing the network to learn invariant features (the mannequin’s shape and colour signature) while treating everything else as nuisance variation.

Concretely, the synthetic generator varies five independent axes per image:

1. **Background context.** Each image samples a random  $1456 \times 1088$  crop from either a high-resolution satellite image (`map.jpg`) or from real DJI flight video frames. This provides diverse terrain textures—grass, paths, shadows, field markings—without requiring separate background datasets.
2. **Altitude-correct target scale.** The mannequin cutout is resized according to a three-tier distribution calibrated to real flight altitudes: 30% close (scale 0.20–0.40, simulating 15–20 m AGL), 40% medium (0.08–0.20, simulating 20–35 m), and 30% far (0.03–0.08, simulating 35–60 m). This biases training towards the operationally critical medium-altitude band while still covering edge cases at the extremes.
3. **Random orientation.** The cutout is rotated by a uniformly sampled angle ( $0$ – $360^\circ$ ), since aerial views can present the target in any heading.
4. **Random placement.** The target is alpha-blended at a uniformly random  $(x, y)$  position within the crop, ensuring no spatial bias towards image centres.
5. **Photometric perturbation.** Brightness, contrast, and blur are varied stochastically (detailed in Section C.3).

The combination of these five axes yields a training distribution that is substantially broader than the real deployment domain, which is the intended outcome: a model that has learned to detect targets across an exaggerated range of conditions will generalise more robustly to the narrower but unpredictable conditions encountered in the field [25].

### C.3 Augmentation Pipeline

Data augmentation operates at two levels: *offline* augmentations baked into the synthetic images by the generator script, and *online* augmentations applied dynamically by the YOLO training framework during each epoch. Table 24 summarises both.

Each offline augmentation addresses a specific failure mode. Brightness and contrast jitter simulate varying solar angles and intermittent cloud cover, both common during UK field operations. Gaussian blur with small kernels ( $3 \times 3$  and  $5 \times 5$ ) approximates the motion blur introduced by a drone travelling at 6–10 m/s with a rolling-shutter camera, without degrading the target beyond recognition. Random erasing was explicitly disabled (`erasing=0.0`) because at high altitudes the mannequin occupies as few as  $10 \times 20$  pixels; erasing even a small patch risks removing the entire target, teaching the network to predict from background alone.

The online augmentations provided by the YOLO trainer complement the offline pipeline. Mosaic augmentation [26] tiles four training images into a single  $1088 \times 1088$  input, exposing the model to multiple backgrounds and target sizes per gradient step—particularly valuable for small-object detection where the target occupies less than 1% of the image area. Copy-paste augmentation pastes annotated objects onto new backgrounds, further multiplying apparent training diversity. The aggressive online scale range ( $0.3\text{--}1.7\times$ ) was tuned specifically for aerial detection, where altitude variation causes target size to change by an order of magnitude between 15 m and 60 m AGL.

### C.4 Synthetic Data Generation Pipeline

Synthetic images are generated by `generate_dataset_v2.py`, which executes the following pipeline for each of the 300 output images:

1. A random  $1456 \times 1088$  crop is extracted from the satellite image or from DJI video frames, providing diverse background textures.
2. The mannequin cutout is scaled according to the three-tier altitude distribution described in Section C.2.
3. The scaled cutout is rotated by a uniformly random angle ( $0\text{--}360^\circ$ ) and alpha-blended onto a random position within the background crop, with the alpha channel preserving natural edge blending.
4. Offline augmentations (brightness, contrast, blur) are applied stochastically as specified in Table 24.
5. A YOLO-format label file is written containing the normalised bounding box coordinates (class,  $x_{\text{centre}}$ ,  $y_{\text{centre}}$ , width, height).

Compared to the v1 dataset (200 images at  $640 \times 640$  from a single satellite background), v2 introduces real-flight backgrounds, altitude-correct target scaling, a negative set, and native-resolution output—all of which reduce the domain gap between training and deployment [25].

### C.5 Real Data Labelling

A custom labelling tool (`label_tool.py`) was developed for rapid bounding-box annotation of video frames. The tool displays each frame at native resolution and supports click-to-place annotation: left-click sets the box centre, scroll wheel adjusts box size, **S** saves, and **N** skips. The `--full` flag ensures labels are computed relative to the full  $1456 \times 1088$  frame rather than  $640 \times 640$  tiles, preserving spatial fidelity.

Sixteen frames were labelled from the DJI flight recording `DJI_0001_1456x1088_cropped_30fps.mp4`, selected to span a range of altitudes (15–50 m) and viewing angles. While small in number, these real frames proved disproportionately valuable: they contain authentic lighting, lens effects, and background clutter that synthetic composites cannot fully replicate.

### C.6 Negative Mining

Fifty frames were extracted from the same flight video at locations where no mannequin was present. Each negative image is paired with an empty label file, signalling to the YOLO loss function that no objects exist in the frame. Including negatives is a standard technique for reducing false-positive rates in object detection [27], and proved essential for this project.

The v1 model (trained without negatives) frequently hallucinated detections on visually ambiguous features: shadows cast by hedgerows, light-coloured path markings, discarded equipment at field edges, and patches of bare earth with high contrast against grass. These false positives would have triggered unnecessary descent manoeuvres during autonomous flight, wasting battery and mission time. The negative set was therefore curated to include representative examples of each confusing background element. After retraining with the 50-image negative set, qualitative evaluation on held-out flight video showed a substantial reduction in false activations on these terrain features, while maintaining high recall on the mannequin target.

The ratio of negatives to positives (50:316, or  $\sim 16\%$ ) follows the heuristic recommended by Shrivastava, Gupta, and Girshick [27]: enough negatives to calibrate the classifier’s decision boundary without overwhelming the positive signal.

## C.7 Data Composition Analysis

Table 25 summarises the final dataset composition. The 300:16 ratio of synthetic to real images reflects a practical constraint rather than a design choice: real aerial data requires physical flight operations (limited to one weather-permitting session), manual labelling effort ( $\sim 2$  minutes per frame with the custom labelling tool), and regulatory clearance. In contrast, synthetic images are generated programmatically at a rate of  $\sim 10$  images per second.


Despite constituting only 4.4% of the dataset, the 16 real frames serve a disproportionate role: they anchor the learned feature space to real-world appearance, preventing the model from overfitting to artefacts of the synthetic compositing process (e.g. sharp alpha-blended edges, perfectly uniform lighting on the cutout). This “few-shot anchoring” approach—using a small number of real examples to ground a predominantly synthetic dataset—has been shown to be effective in sim-to-real transfer for robotics and autonomous systems [13]. The synthetic majority provides the scale diversity that 16 real images alone could never cover, while the real minority corrects for distributional shift between composited and photographed imagery.

## C.8 Training Pipeline and Convergence

Training was performed on Google Colab using an NVIDIA T4 GPU (16 GB VRAM). The model was initialised from COCO-pretrained YOLOv8n weights (yolov8n.pt), applying transfer learning to leverage low-level feature representations—edges, textures, colour gradients—learned from the 80-class COCO dataset [12, 11]. Only the detection head and final feature pyramid layers were fine-tuned during the first 10 warmup epochs; the remaining backbone layers were subsequently unfrozen for full end-to-end optimisation. Table 26 summarises the training configuration.

Training completed in approximately 45 minutes on the T4 GPU. The loss curves exhibited three distinct phases: rapid convergence during epochs 1–30 as the COCO-pretrained features adapted to the aerial SAR domain, gradual refinement during epochs 30–80 as the model learned fine-grained distinctions between the mannequin and confusing background elements, and plateau from epoch 80 onwards with  $mAP_{50}$  stabilising above 0.99. The early stopping patience of 30 epochs ensured training terminated at epoch 150 without overfitting, as the validation loss showed no upward trend. The learning rate followed the default cosine annealing schedule from an initial value of 0.01, decaying to  $10^{-5}$  by the final epoch.

A batch size of 8 was the maximum that fitted within the T4’s 16 GB VRAM at the  $1088 \times 1088$  training resolution. Larger batch sizes would have enabled smoother gradient estimates but were not feasible without gradient accumulation, which was not explored given the already strong convergence behaviour.




**Figure 15:** Training and validation loss curves over 150 epochs. Rapid convergence occurs in epochs 1–30 as COCO-pretrained features adapt to the aerial SAR domain, followed by gradual refinement and plateau from epoch 80 onwards with no overfitting trend.

## C.9 Results

Table 27 presents the validation metrics for the v2 model compared to the original v1 baseline. The v2 model achieves near-perfect detection performance on the validation set.

## YOLOv8n SAR Detector — Confusion Matrix (Validation Set)



**Figure 16:** Normalised confusion matrix for the v2 model on the validation set. The single-class detector achieves near-perfect discrimination between the mannequin target and background, with negligible false-positive and false-negative rates.

A caveat applies: the training and validation splits use the same image directory (`train=val=images/`), so the reported metrics reflect in-sample performance and likely overstate true generalisation accuracy. This pragmatic choice was driven by the small dataset size—an 80/20 split would leave only 73 validation images, too few for stable metric estimation. The field-test results (Section 4) provide independent confirmation that the model generalises beyond the training distribution. The model export and deployment pipeline is described in Sections C.10 and C.11.

### C.10 Model Export Pipeline

The trained PyTorch model (`best.pt`, ~6 MB) is exported to two inference-optimised formats for edge deployment:

1. **TFLite (float32)**. Exported via the Ultralytics `yolo export format=tflite` command, producing an 11.7 MB flatbuffer file. The export preserves the input tensor shape `[1, 640, 640, 3]` (uint8 or float32) and output tensor shape `[1, 5, 8400]` (`5 = x, y, w, h, confidence`; 8400 = anchor predictions across three feature pyramid levels). Float32 quantisation was chosen over INT8 because INT8 export requires a representative calibration dataset and produced inconsistent detection quality in preliminary testing on the Pi.
2. **NCNN**. Exported via `yolo export format=ncnn`, producing a `.param` / `.bin` pair optimised for ARM NEON vectorisation. NCNN inference is expected to achieve ~3× speedup over TFLite on the Pi 5’s Cortex-A76 cores (Section 2.3), though this has not yet been validated in flight.

A critical design constraint is that all model variants—across training runs, resolutions, and export formats—must produce identical tensor shapes. This invariant is enforced by the YOLOv8n architecture (which always outputs `[1, 5, 8400]` regardless of training resolution) and verified by an assertion in `vision.py` at model load time. It guarantees that any model can be swapped without modifying downstream code.

### C.11 Deployment

The deployment workflow is deliberately minimal. Because all model variants share the same input/output tensor shapes, swapping the active model on the Pi requires only:

```
cp cv_models/sar_v2_1088/best.tflite best.tflite
```

No code changes, configuration edits, or dependency updates are needed. The `vision.py` module loads whichever `best.tflite` is present in the project root, and all downstream modules consume the same (`found`, `cx`, `cy`,

**confidence**) interface regardless of which model generated the detection. If a model produces excessive false positives in the field, an alternative can be deployed in under 30 seconds. The `--model` flag in `main.py` additionally allows runtime model selection without file copying, further streamlining field testing. Three model variants are archived in the `cv_models/` directory (`sar_v2.1088`, `sar_640`, `sar_1280`), each containing TFLite, PyTorch, and NCNN exports along with training curves and confusion matrices for reproducibility.

## C.12 Limitations

Several limitations of the current training pipeline should be acknowledged honestly, as they bound the confidence that can be placed in the reported metrics:

- **Train/validation overlap.** The `dataset.yaml` configuration uses the same image directory for both training and validation (`train=val=images/`). The reported  $mAP_{50}$  of 0.995 therefore reflects in-sample performance and almost certainly overstates true generalisation accuracy. A proper stratified 80/20 split would provide more meaningful validation metrics, but was not pursued because an 80/20 split of 366 images yields only 73 validation images—too few for stable metric estimation given the class imbalance (316 positives vs. 50 negatives).
- **Synthetic data generator bias.** All 300 synthetic images are produced by the same compositing pipeline (`generate_dataset_v2.py`), which applies the mannequin cutout with alpha blending onto background crops. The generator introduces systematic artefacts—sharp alpha edges, uniform illumination on the cutout irrespective of background lighting, absence of cast shadows from the target—that do not exist in real imagery. The model may learn to detect these artefacts rather than genuine target features, a form of shortcut learning [13]. The 16 real images partially mitigate this risk but cannot fully compensate.
- **Small real-image fraction.** Only 16 of 366 images (4.4%) are manually labelled real aerial photographs, all captured during a single 8-minute flight on one day. This fraction is constrained by operational realities: each flight requires regulatory clearance, suitable weather, team coordination, and manual post-hoc labelling ( $\sim 2$  minutes per frame). While domain randomisation compensates for scale and position diversity, it cannot substitute for the distributional richness of hundreds of real-world images captured across multiple days, locations, and weather conditions.
- **Single target appearance.** All synthetic images use the same mannequin foreground image (`dummy.png`). A real SAR scenario would present casualties in varied clothing, body positions (prone, foetal, seated), skin tones, and degrees of occlusion. The single-class detector cannot distinguish between target types and has not been validated on human subjects. The COCO-pretrained “person” detector (`models/human.tflite`) is retained as a backup for this reason.
- **Limited environmental diversity.** Training backgrounds are drawn from one satellite image and one flight video recorded on a single overcast day in March. Seasonal variation (snow, autumn leaves, crop growth), different terrain types (woodland, urban, rocky), adverse weather (rain, fog, low sun angle), and varying illumination conditions are not represented. Deployment in conditions significantly different from the training environment may degrade detection performance.
- **No partial occlusion.** The dataset contains no examples of a partially occluded target—for instance, a casualty partially hidden by vegetation, rocks, or terrain features—which is a common scenario in real wilderness SAR operations. The model has no learned representation of occluded targets.
- **Single-class limitation.** The detector outputs a single class (“dummy”) and cannot differentiate between a casualty requiring rescue, a bystander, wildlife, or other objects of similar size and colour. In an operational SAR context, a multi-class detector or a downstream classification stage would be necessary to reduce false dispatches.

Despite these limitations, the trained model performed reliably during bench testing and video replay analysis across the altitude range of 15–50 m (Section 4), and the modular deployment architecture allows rapid model replacement as improved datasets become available. The limitations are structural consequences of the project’s scope (a university coursework with one flight opportunity) rather than fundamental barriers—each could be addressed with additional data collection campaigns and training iterations.

## D Safety and Risk Management

This appendix presents the safety architecture, failure-mode analysis, and regulatory compliance measures that underpin every flight operation, demonstrating how the system mitigates risks inherent in autonomous UAV flight adjacent to an environmentally sensitive site.

Autonomous flight adjacent to an environmentally protected site introduces hazards spanning airspace violation, equipment failure, and loss of control. A layered safety architecture mitigates these risks through geofencing, failure-mode handling, operator oversight, and regulatory compliance, following the Specific Operations Risk Assessment (SORA) methodology [6] and UK CAA guidance for unmanned aircraft operations [5].

## D.1 Geofence Implementation

The survey area is adjacent to a Site of Special Scientific Interest (SSSI) designated as a no-fly zone (NFZ). Three polygon zones — flight area (4 vertices), search area (5 vertices), and SSSI exclusion zone (7 vertices) — are loaded from a KML file at startup and validated against hardcoded fallback coordinates in `config.py`. Because a single point of failure in boundary enforcement could result in an airspace violation, the system employs five independent protection layers (three runtime software layers, plan-time waypoint filtering, and a firmware-level backup):

1. **Speed scalar field.** Within `NFZ_SLOW_ZONE_M` = 20 m of the SSSI boundary, a `MAV_CMD_DO_CHANGE_SPEED` command linearly reduces the maximum ground speed from `NFZ_ZONE_MAX_SPEED_MPS` = 3.0 m/s at the zone edge to `NFZ_MIN_SPEED_MPS` = 0.3 m/s at the boundary itself. At 0.3 m/s the aircraft is effectively hovering; the autopilot can arrest all remaining momentum in well under one metre even in moderate wind.
2. **Repulsive vector field.** A virtual inner polygon is offset `NFZ_INNER_OFFSET_M` = 20 m inward from the SSSI boundary. When the aircraft is within `NFZ_INNER_RANGE_M` = 23 m of this inner polygon — effectively 3 m outside the actual NFZ boundary — a constant-magnitude repulsive velocity of `NFZ_PUSH_SPEED_MPS` = 3.0 m/s is applied away from the nearest boundary segment. The constant magnitude ensures that the repulsive force always exceeds the near-zero speed cap at the boundary, preventing penetration regardless of approach velocity.
3. **Hard boundary cutoff.** If the aircraft's GPS position falls inside the SSSI polygon (or within `NFZ_HARD_BOUNDARY_M` = 3 m of it), the system immediately halts all velocity commands, saves the current state, and transitions to `MANUAL` mode, requiring the operator to fly the aircraft clear before autonomy can resume.
4. **Plan-time waypoint filtering.** Search waypoints within `NFZ_WAYPOINT_BUFFER_M` = 30 m of the NFZ are removed during path generation, preventing the lawnmower pattern from routing the aircraft close to the boundary in the first place.
5. **Firmware-level geofence.** On real hardware, an independent geofence configured on the Cube Orange autopilot triggers `LOITER` or `RTL` if all four software layers fail entirely. This layer operates at the flight-controller firmware level and cannot be overridden by companion-computer software.

## D.2 Failure Mode Analysis

Table 28 summarises the system's response to each anticipated failure mode. All software responses were verified in SITL simulation; the `RTL` and `LOITER` fallbacks additionally rely on ArduCopter's flight-proven failsafe modes [28]. The GPS degradation handler monitors `GPS_RAW_INT` messages continuously. A 5 s persistence filter prevents transient signal dips from triggering premature `RTL` while ensuring genuine fix loss is caught promptly. The landing state employs three independent touchdown detection mechanisms — ArduPilot accelerometer-based auto-disarm, altimeter-based detection (5 consecutive ticks below 0.5 m), and a 90 s absolute timeout with forced disarm — to prevent barometric drift from causing an indefinite hover near the ground.

## D.3 Operator-in-the-Loop

The `VERIFY` state is the critical decision gate that prevents the system from acting on false-positive detections (see also Section D.3 for the full classification taxonomy). When the detection pipeline identifies a candidate target, the aircraft navigates to its estimated GPS position and holds altitude while presenting the live camera feed to the operator via the browser-based ground station. The operator classifies the detection using one of three inputs:

- **Y (Confirm)** — the target is a genuine casualty. The system initiates GPS position averaging (10 s of samples for improved accuracy), then prompts the operator to select a landing side (N/E/S/W) before proceeding to the approach and payload delivery sequence.
- **N (Reject)** — a false positive. The target position is spatially tagged with a `REJECTED_TARGET_RADIUS_M` = 5 m exclusion radius so it is not re-investigated. The aircraft resumes the search pattern or investigates the next queued detection.
- **I (Item of Interest)** — uncertain classification. The GPS position is logged for later review and the aircraft continues searching, similar to rejection but without blacklisting the area.

A 120 s timeout ensures the aircraft does not hover indefinitely if the operator is unresponsive. Countdown warnings are issued at 60 s, 90 s, and every 10 s thereafter. If the timeout expires, the target is auto-rejected and the search resumes — a safe default that prioritises continued area coverage over a single uncertain detection.

## D.4 Emergency Procedures

Three independent abort mechanisms are available at all times, each operating at a different level of the system stack:

1. **RC kill switch.** The RC transmitter mode channel can switch the autopilot to `STABILIZE` or `LOITER` at any time, immediately returning full manual control to the safety pilot. This hardware-level override operates independently of the onboard computer and cannot be blocked by software faults.

2. **Keyboard abort (Ctrl+C / ESC).** Pressing either key in the Raspberry Pi terminal triggers `_emergency_rtl()`, which commands the Cube to enter RTL mode before exiting the script. The Cube then executes the return-to-launch sequence autonomously.
3. **Manual override (M key).** Available from any active state, the M key transitions the system to MANUAL mode, saving the departure point for later resumption. Upon exiting manual mode, the system routes through `RETURN_FROM_MANUAL` to navigate back to the departure point before resuming the previous state.

ArduCopter provides additional firmware-level failsafes that operate regardless of the companion computer state: GCS failsafe (`FS_GCS_ENABLE`) triggers RTL on heartbeat loss, RC failsafe (`FS_THR_ENABLE`) triggers RTL on transmitter signal loss, and battery failsafe (`FS_BATT_ENABLE`) triggers RTL or LAND when voltage drops below configured thresholds.

## D.5 Risk Assessment

Table 29 presents a risk register following a likelihood–severity matrix consistent with SORA ground risk assessment [6]. Likelihood is rated L (low), M (medium), or H (high); severity is rated 1 (minor), 2 (moderate), or 3 (severe).

## D.6 Regulatory Considerations

Flight tests were conducted under the UK Civil Aviation Authority (CAA) Open Category framework [5]. The aircraft operates in subcategory A3 (far from uninvolved persons), which requires:

- The remote pilot holds a valid flyer ID and operator ID registered with the CAA.
- The aircraft remains within visual line of sight (VLOS) at all times, with a dedicated observer when the pilot also monitors the ground station.
- The flight area is at least 150 m from residential, commercial, industrial, or recreational areas.
- A site risk assessment is completed for the designated flight field, identifying the adjacent SSSI, public footpaths, and prevailing wind patterns.

The adjacent SSSI imposes additional constraints beyond standard CAA requirements. Natural England guidance prohibits UAV operations over or within SSSIs without specific consent, as disturbance to protected habitats and species constitutes an offence under the Wildlife and Countryside Act 1981 [4]. The five-layer geofence described in Section D.1 was designed specifically to enforce this exclusion, combining four software layers with a firmware-level backstop.

Although the system architecture supports beyond-visual-line-of-sight (BVLOS) operation—the autonomous state machine, RTL failsafes, and ground station telemetry do not require direct visual contact—all test flights were conducted within VLOS range (<200 m) with the operator maintaining direct visual contact. Future deployment in an operational SAR context would require a Specific Category authorisation with a full SORA assessment [6], including detailed ground and air risk analyses, a ConOps document, and potentially an operational authorisation from the CAA.

# E Testing Methodology

This appendix describes the verification and validation framework used to build confidence in each subsystem before advancing to the next integration level, documenting the rationale behind the five-tier progressive approach and the 20 failure scenarios tested before outdoor operations.

The verification and validation strategy draws on V-model principles [29], where each level of system integration is paired with a corresponding verification activity, and on incremental autonomy practices from safety-critical robotics [30]. The resulting five-tier progressive testing framework advances from pure software simulation to full autonomous flight, supported by 74 dedicated test scripts organised across eight categories—including 127 unit tests and 33 automated integration tests—and a stress-testing campaign that exercised 20 failure scenarios before the first outdoor test.

## E.1 V-Model Verification Approach

The classical V-model pairs requirements at each level of decomposition (system, subsystem, component) with a matching verification activity at the same level. In a university project with limited flight days and shared hardware, a strict V-model is impractical; however, its core insight—*verify each integration level before proceeding to the next*—maps directly onto the progressive strategy adopted here.

At the *component* level, individual peripherals (camera, GPS receiver, AI model, MAVLink link) are tested in isolation with dedicated hardware scripts. At the *subsystem* level, bench tests validate the command pipeline (Pi→Cube→actuators) and the sensing chain (camera→inference→coordinate transform) without flight. At the *system* level, passive flight confirms that all subsystems produce coherent results under real outdoor conditions before the software is permitted to issue flight commands. Only after every lower tier passes does the system advance to autonomous operation.

This tiered approach catches integration faults early: a GPS wiring error is discovered at the bench, not at 30 m altitude; a false-positive rate is quantified during passive flight, not during the first autonomous descent. Each tier constitutes a *gate*—failures must be resolved at the current level before the next is attempted, and the RC kill switch provides an independent hardware-level override at every stage.

## E.2 Five-Tier Progressive Testing

Table 30 summarises the five tiers and their verification targets.

**Table 30:** Five-tier progressive testing framework. Each tier must pass before the next is attempted.

| Tier | Environment                              | What It Verifies                                                              | Risk Level           |
|------|------------------------------------------|-------------------------------------------------------------------------------|----------------------|
| 1    | SITL simulation (laptop)                 | State machine logic, path planning, detection pipeline, geofence enforcement  | Zero (software only) |
| 2    | Docker Pi test                           | TFLite model loads on ARM, dependency compatibility, headless operation       | Zero (container)     |
| 3    | Bench test (Pi + Cube, no props)         | MAVLink command ACKs, camera frames, GPS fix, sensing pipeline, servo outputs | Low (grounded)       |
| 4    | Passive flight (pilot on RC, Pi logging) | Detection range and rate in real outdoor conditions, FOV and GPS calibration  | Medium (no autonomy) |
| 5    | Autonomous flight                        | End-to-end mission: search, detect, centre, verify, approach, land            | High                 |

**Tier 1** uses ArduPilot’s Software-in-the-Loop (SITL) [31] on the development laptop. The same mission orchestrator (`main.py`), state machine, and detection model execute against a simulated autopilot with GPS noise and vehicle dynamics. A simulated camera extracts FOV-correct sub-images from a georeferenced satellite map, enabling the full detection→estimation→landing pipeline to run without hardware. A 22-test simulation checklist (Section E.6) exercises every CLI flag, state transition, and edge case at this tier.

**Tier 2** uses a Docker container (`Dockerfile.pi-test`) that replicates the Pi’s ARM64 Linux environment with the same lightweight dependency set (`requirements.pi.txt`: `pymavlink`, `opencv-headless`, `numpy`, `ai-edge-litert`). The container loads the TFLite model and confirms that inference produces valid output tensors, catching dependency mismatches before code is deployed to the physical Pi.

**Tier 3** assembles the Pi, Cube Orange autopilot, and camera on the bench with propellers removed. Three numbered scripts test progressively: `0a_cube_commands` verifies individual MAVLink commands (mode changes, parameter reads, arm pre-checks); `0b_bench_mission` executes the full command sequence (arm, takeoff, waypoint navigation, landing) and confirms every `COMMAND_ACK`; `0c_feedback_test` runs the vision-to-GPS pipeline while the operator carries the drone past a dummy, validating the entire sensing chain against hand-measured ground truth.

**Tier 4** places the assembled drone in the air under full manual RC control. The Pi runs the complete vision pipeline—camera capture, AI inference, coordinate transformation, GPS estimation—but issues *zero flight commands*. All detections are logged with timestamp, pixel coordinates, confidence, altitude, and GPS estimate. Post-flight review of these logs establishes the detection altitude ceiling, false-positive rate, and GPS estimation accuracy under real conditions (wind, lighting, motion blur) before any autonomous behaviour is enabled.

**Tier 5** enables full autonomy. The drone flies the lawnmower search pattern, detects candidates, centres on them, and presents the live camera feed to the operator for confirmation before initiating the approach sequence. Even at this tier, the operator retains veto authority at the VERIFY gate and the RC kill switch provides an independent hardware override.


## E.3 Test Script Organisation

The test scripts are organised into eight directories, each addressing a distinct verification concern:

- **Hardware** (`tests/hardware/`, 8 scripts) — individual component checks: inference benchmarks (standard 50-run and comprehensive multi-model), detection rate under simulated blur, NCNN backend benchmark, GPS satellite tracking, GPS step-by-step health verification, single-frame detection snapshot, and buzzer functionality. These confirm that each peripheral operates correctly before integration.
- **Flight** (`tests/flight/`, 14 scripts) — progressive flight tests numbered 0a through 5 by risk level, plus geofence verification scripts (demo, repulsion test), drawing utilities for search area and waypoint definition, and a live map viewer. Lower numbers must pass before higher numbers are attempted; see Section E.2 for the full progression.

- **Diagnostics** (`tests/diagnostics/`, 5 scripts) — runtime monitoring tools: a multi-panel connectivity dashboard aggregating all communication links, live Cube telemetry viewer (attitude, GPS, battery), and MJPEG camera streams (standard, threaded, and H.264 variants). The diagnostics dashboard serves as the final go/no-go gate before arming.
- **Calibration** (`tests/calibration/`, 7 scripts) — parameter measurement: bench FOV calibration (ruler method), quick single-measurement FOV check, altitude-based FOV at real hover height, lens distortion characterisation (checkerboard → `calibration_data.npz`), GPS ground-truth comparison (CV estimate vs surveyed position), and two focal-length calibration tools (terminal-based and OpenCV GUI).
- **Experiments** (`tests/day_1_experiments/`, 6 scripts) — structured data collection designed to run passively during Tier 4 manual flights, outputting CSV files for post-flight analysis: altitude sweep, speed sweep, GPS accuracy, GPS drift (CEP50/CEP95), model comparison, and a general detection logger.
- **Laptop** (`tests/laptop/`, 21 scripts) — development-only tools: model loading, TFLite and NCNN inspection, synthetic benchmarks, DJI video replay with detection overlay, A/B model comparison, zoom-based detection, path visualisation, NFZ manual testing, and approach/descent simulation. Not used on flight day.
- **Unit** (`tests/unit/`, 6 files, 127 test functions) — pytest-based unit tests covering configuration validation, state enum completeness, GPS coordinate transforms, lawnmower pattern generation, vision system initialisation, and GPS utility functions.
- **Automated** (`tests/automated/`, 7 scripts, 33 test functions) — extended integration tests: planning edge cases, utility function coverage, vision pipeline tests, geofence optimisation, coverage simulation, and SITL smoke tests.

All experiment scripts issue zero flight commands and are safe to run at any time. The CSV output format enables direct import into the analysis presented in Section F. Figure 17 shows the requirements coverage achieved at each testing tier.



**Figure 17:** Requirements coverage by testing tier. Each bar shows the number of system requirements (R01–R10) verified at each of the five progressive testing levels (simulation, Docker, bench, passive flight, autonomous flight). Lower tiers verify a broader set of requirements at lower risk; higher tiers provide operational validation of the most critical flight and mission requirements. The layered coverage ensures that no requirement depends solely on the highest-risk tier for verification.

## E.4 Dry-Run Validation

The `--dry-run` flag provides an end-to-end state machine walkthrough without requiring hardware, an autopilot connection, or even SITL. When invoked, the system loads the search polygon from the KML file, generates the lawnmower waypoint set, visualises it on a map with geofence boundaries, prints the full mission plan (altitude, speed, waypoint count, estimated time), and walks through the state machine transitions (INIT→SEARCH→VERIFY→LANDING→DONE) to verify that all imports, configuration parameters, and state logic are correct. The output image (`dry_run_pattern.jpg`) provides immediate visual confirmation that the search pattern covers the intended area and respects the SSSI exclusion zone.

Dry-run validation is executed before every flight day as the first item on the preflight checklist and after every code change that modifies state transitions, waypoint generation, or geofence logic. Execution takes under one second and requires only a Python interpreter, making it suitable for rapid iteration during development.

## E.5 DJI Video Analysis

Real flight footage from a DJI survey drone provided a supplementary validation path between simulation and live flight. The video analysis pipeline replays recorded 4K video at 30 fps with frame-synchronised SRT telemetry (GPS, altitude, heading) while running the YOLOv8n model on every frame. Each detection is projected from pixel coordinates to GPS coordinates using the calibrated FOV ( $54.4^\circ$  HFOV for the cropped  $1456 \times 1088$  frame). The pipeline produces scatter plots of estimated target positions, circular error statistics, and a scale bar overlay for visual verification of distance estimates.

An A/B comparison tool switches between model variants during replay, enabling direct evaluation of the baseline and retrained models on identical footage. This analysis informed the decision to retrain on mixed synthetic-plus-real data (Section C) and quantified the GPS estimation accuracy ( $\text{CEP}_{50} = 2.3 \text{ m}$ , maximum error = 16.5 m) under real flight dynamics, including the 100–200 ms GPS timing lag identified during analysis.

## E.6 Stress Testing and Failure Scenarios

A systematic stress-testing campaign exercised the system’s response to 20 failure and edge-case scenarios in SITL simulation prior to outdoor testing. The scenarios were derived from a contingency table—a state-by-failure matrix covering all 15 states against six failure categories (timeout, manual override, GPS loss, camera failure, connection drop, and operator error)—and from the 22-test simulation checklist described below. Table 31 summarises the results.

**Table 31:** Stress test results: 20 failure scenarios tested in SITL.

| Scenario                                  | Result | Resolution                               |
|-------------------------------------------|--------|------------------------------------------|
| Camera returns None during search         | PASS   | Drone completes pattern, returns home    |
| Camera fails during verify                | FIXED  | Added “CAMERA LOST” HUD warning          |
| GPS fix degrades mid-flight (<3D for 5 s) | PASS   | Emergency RTL triggers automatically     |
| GPS never acquires fix on ground          | PASS   | Arm blocked; clear warning printed       |
| MAVLink heartbeat lost                    | PASS   | Emergency RTL; GCS failsafe backup       |
| Operator timeout at verify (120 s)        | PASS   | Auto-reject, search resumes              |
| Rapid key presses (M/Y/N/I/X in <1 s)     | PASS   | Input queue processes FIFO, no crash     |
| Detection outside search polygon          | PASS   | Filtered; not added to queue             |
| Duplicate detection of same target        | FIXED  | Added 5 m reject radius dedup            |
| Queue overflow (>50 targets)              | FIXED  | Hard cap; oldest entries discarded       |
| Empty queue after landing                 | PASS   | Resumes search or RTL                    |
| Geofence breach (drone enters SSSI)       | PASS   | Immediate MANUAL transition              |
| NFZ speed capping at boundary             | PASS   | Smooth ramp from 3 m/s to 0.3 m/s        |
| Barometric drift during landing           | FIXED  | 90 s absolute timeout + force-disarm     |
| Port 8090 in use (stream server)          | FIXED  | Retry consecutive ports 8091–8099        |
| Servo channel hardcoded (not in config)   | FIXED  | Moved to config.py parameters            |
| Corrupt focus.area.json                   | FIXED  | Validation + clear error message         |
| Transit.json missing                      | FIXED  | Warning + abort with actionable error    |
| Home position not set before RTL          | FIXED  | Guard checks gps_fix_ok flag             |
| Stale departure point after manual        | PASS   | Correctly returns to pre-manual position |

Of the 20 scenarios, 11 passed on the first attempt with the correct automated response already in place. The remaining nine revealed deficiencies—missing timeouts, insufficient input validation, hardcoded parameters, or unclear error messages—that were resolved and re-verified before outdoor testing. No scenario was left unhandled: every failure mode either triggers an automated recovery (RTL, mode change, queue management) or produces an actionable operator warning with a clear next step.

The 22-test simulation checklist provides structured coverage of all operational modes: full autonomous mission (baseline), multi-frame detection confirmation, vision-based centering with GPS averaging, geofence enforcement and disable, manual override workflow, multi-target queue ordering, NFZ speed ramp, out-of-polygon detection filtering, beacon redirect, transit planning, altitude override, dry-run generation, headless mode, code refactoring validation, stream server functionality, confidence threshold tuning, reject radius deduplication, manual-mode queue persistence, rapid input robustness, empty queue handling, and log file integrity.

## E.7 Building Confidence Through Progressive Evidence

The five-tier framework is designed so that each tier produces independently valuable evidence, regardless of whether subsequent tiers are reached. This property is critical for a university project where flight access is weather-dependent

and time-limited: even if autonomous flight is never achieved, the data collected at lower tiers still demonstrates engineering competence and informs the design.

**Evidence from Tier 1 (simulation).** Over 200 simulation runs exercised the state machine, search pattern, and GPS estimation pipeline. The interactive simulator’s configurable noise parameters (GPS drift  $\sigma = 0\text{--}5\text{ m}$ , camera shake  $0\text{--}15\text{ px}$ , inference throttle  $1\text{--}30\text{ FPS}$ ) enabled systematic sensitivity analysis. The Kalman filter consistently converged within  $1\text{ m}$  of ground truth across all tested noise combinations, outperforming the rolling average estimator under burst-noise conditions. The 22-test simulation checklist verified every CLI flag, state transition, and edge case, catching nine deficiencies (missing timeouts, hardcoded parameters, unclear error messages) that were resolved before any hardware was involved.

**Evidence from Tier 2 (Docker).** The Docker container test proved that the TFLite model loaded correctly on an ARM64 Linux environment matching the Pi’s dependency set, and that the `ai-edge-litert` package provided a drop-in replacement for the incompatible `tflite-runtime`. While this test takes under 30 seconds, it prevented a deployment failure that would have consumed significant debugging time in the field.

**Evidence from Tier 3 (bench).** Bench testing with the assembled hardware produced quantitative calibration data: focal length  $f = 5.46\text{ mm}$  (correcting a 22% error from the datasheet default), lens distortion characterisation ( $\text{RMS} = 0.399\text{ px}$ ), and inference benchmarks ( $206.5\text{ ms}$  average,  $4.8\text{ FPS}$ , 100% detection rate at 0.966 confidence). The three bench scripts exercised the full command pipeline (arm, takeoff, waypoint, land) against the Cube Orange autopilot with propellers removed, confirming that every `COMMAND_ACK` was received correctly. The bench also revealed the IMX296 BGR colour channel issue and the TFLite bounding box coordinate format inconsistency—defects that were invisible to simulation.

**Evidence from Tier 4 (passive flight).** Passive flight with manual RC control and the Pi running the full vision pipeline is designed to produce the most operationally relevant pre-autonomy data: detection rate as a function of altitude, false-positive rate under real outdoor conditions, GPS estimation accuracy against surveyed ground truth, and inference latency under thermal load during sustained operation. The experiment scripts in `tests/day_1.experiments/` output CSV files structured for direct statistical analysis: `altitude_sweep.py` bins detection rate by altitude in  $5\text{ m}$  increments; `speed_sweep.py` correlates detection rate with ground speed and blur metrics; `gps_accuracy.py` computes CEP statistics against a known dummy position. This data directly informs the configuration parameters (`TARGET_ALT`, `SEARCH_SPEED_MPS`, `CONFIDENCE_THRESHOLD`) used in autonomous flight.

**Evidence from DJI video replay.** As a proxy for Tier 4 data obtained before flight access was available, DJI footage from a reconnaissance overflight of the test site was replayed through the identical detection pipeline. The retrained model achieved 87% detection rate on frames where the target exceeded  $20 \times 20\text{ px}$ , with mean confidence 0.82. GPS estimation produced  $\text{CEP}_{50} = 2.3\text{ m}$ . The video replay also revealed the  $100\text{--}200\text{ ms}$  GPS timing lag, which manifested as a systematic along-track bias in position estimates—a phenomenon that would have been difficult to diagnose during a live autonomous flight.

**Cumulative confidence.** The result of this progressive approach is that by the time the system is authorised for autonomous flight, the team holds quantitative evidence for every subsystem’s performance envelope. Failures are discovered at the lowest-risk tier capable of exposing them. No tier depends on the success of a subsequent tier for its evidence to be valid. This property—independent value at each level—is the defining characteristic that distinguishes the framework from a simple sequential test plan.

## E.8 Team Roles in the Testing Process

The five-tier testing framework distributed responsibilities across the team to ensure that no single person was a bottleneck for verification progress.

**Software development and simulation (Tiers 1–2).** The software lead was responsible for developing and maintaining the simulation infrastructure, including the interactive simulator, SITL integration, and Docker Pi test. All simulation test scripts were written, documented, and version-controlled by this team member. The 22-test simulation checklist and 20-scenario stress test were executed after every significant code change, with results logged in the repository.

**Hardware integration and bench testing (Tier 3).** Hardware assembly, wiring, and bench testing were shared across the team. The Pi-Cube-camera stack was assembled collaboratively, with each team member responsible for verifying connectivity of specific subsystems (camera, MAVLink link, GPS, servo). Bench test scripts were designed to be self-contained and executable by any team member without specialised knowledge: each script prints clear pass/fail results and actionable error messages.

**Flight operations (Tiers 4–5).** The qualified pilot retained exclusive authority over RC control and the go/no-go decision. During passive flight (Tier 4), the pilot flew the drone manually while a second team member monitored the Pi’s detection stream on the ground-station laptop and a third recorded observational notes (weather, wind, battery consumption). During autonomous flight (Tier 5), the pilot’s role shifted to safety oversight: monitoring the RC transmitter’s kill switch and the Mission Planner telemetry display, ready to intervene if the autonomous system deviated from the expected flight envelope.

**Data analysis.** Post-flight data analysis was conducted by the software lead using the CSV outputs from the experiment scripts. Detection rate, GPS estimation accuracy, and inference latency statistics were computed and fed back into the configuration file (`config.py`) for the next flight iteration. This closed-loop process—fly, measure, tune, repeat—is the mechanism by which the system’s operational parameters converge on values validated by real-world evidence rather than simulation assumptions.

## E.9 Comparison with Industry Testing Standards

The testing methodology compares favourably with established frameworks for autonomous system verification. NASA’s V&V framework for unmanned aircraft [32] advocates a progressive increase in operational complexity—simulation, hardware-in-the-loop, limited flight envelope, full mission—which maps directly onto the five tiers described above. The SORA methodology [6] requires demonstration of failure-mode handling proportionate to the operational risk class; the 20-scenario stress test and contingency table address this requirement by documenting the system’s response to every foreseeable failure at every state.

Compared with typical university drone projects, which often proceed directly from simulation to autonomous flight, this project introduces two additional verification tiers—Docker environment testing and passive flight with zero-command CV logging—that substantially reduce the risk of first-flight failures. The passive flight tier is particularly valuable: it generates ground-truth detection data under real outdoor conditions without exposing the system to the risk of autonomous decisions acting on untested detections. The structured experiment scripts produce quantitative metrics (detection rate versus altitude, GPS CEP, model inference latency) that inform configuration tuning before autonomy is enabled, replacing the ad-hoc approach common in time-constrained academic projects.

## E.10 Expected Outcomes from Real Flight Testing

Although full autonomous flight had not been conducted at the time of writing, the simulation and bench-testing data allow quantitative predictions of real-flight performance. These predictions serve both as design targets and as acceptance criteria: if flight data deviates significantly from these expectations, the deviation identifies a sim-to-real gap that requires investigation.

**Detection performance.** At the nominal search altitude of 35 m, the camera’s ground footprint spans approximately  $32 \times 24$  m. A standing human ( $\sim 1.8$  m tall) occupies roughly  $80 \times 25$  px in the  $1456 \times 1088$  frame—well above the  $20 \times 20$  px threshold at which the retrained model achieved 87% detection rate on DJI video. At 4.8 FPS and 8 m/s search speed, each ground point appears in approximately 15 consecutive frames, providing multiple detection opportunities per target. Even if real-flight detection rate drops to 50% due to motion blur, lighting variation, or thermal effects, each target would still generate 7–8 detections per flyover pass—sufficient to trigger the investigation workflow. The altitude sweep experiment (`tests/day_1_experiments/altitude_sweep.py`) is designed to measure the actual detection rate at 5 m altitude increments, producing the empirical curve needed to set `TARGET_ALT` optimally.

**GPS estimation accuracy.** Simulation predicts sub-metre Kalman filter convergence under ideal conditions. DJI video replay under real conditions yields  $\text{CEP}_{50} = 2.3$  m. The operational requirement is to land within 5–10 m of the target (R07), so the 7.5 m offset landing provides 5.2 m of margin against the worst-case  $\text{CEP}_{50}$ . The GPS accuracy experiment (`tests/day_1_experiments/gps_accuracy.py`) would measure the actual estimation error against a surveyed dummy position, closing the loop between predicted and observed accuracy.

**Mission duration.** The lawnmower pattern over the survey polygon at 35 m altitude with 25 m strip spacing generates approximately 20 waypoints covering the full area. At 8 m/s search speed with Bézier-smoothed turns, the

estimated search time is 8–12 minutes depending on wind. Adding transit to and from the search area, one investigation cycle (fly to target, hover, verify, return), and the approach/landing sequence, a complete mission is estimated at 15–20 minutes—well within the 25-minute battery endurance of the Hexsoon EDU-450. The speed sweep experiment (`tests/day_1_experiments/speed_sweep.py`) would validate the relationship between search speed and detection reliability, confirming whether 8 m/s is achievable or whether a slower speed is necessary.

**Inference under thermal load.** Bench-measured inference latency (206.5 ms at room temperature) may increase during sustained outdoor operation as the Pi 5’s SoC reaches thermal throttling temperature ( $\sim 85^{\circ}\text{C}$ ). The benchmark script (`tests/hardware/benchmark.py`) runs 50 consecutive inference passes and reports the latency trend across the run. During a 15-minute mission, the Pi processes approximately 4300 frames; if thermal throttling increases latency by 50% (to  $\sim 310$  ms, or 3.2 FPS), the effective detection window per target drops from 15 to 10 frames—still sufficient for reliable triggering. A passive heatsink or brief inference pauses between waypoints are available mitigations if thermal degradation is observed.

## E.11 Preflight Checklist

A structured preflight checklist is executed before every flight session, covering four domains:

1. **Hardware inspection** — battery voltage and capacity, propeller security, camera lens cleanliness, antenna connections, wiring integrity.
2. **Software verification** — correct code version pulled from the repository, AI model file verified (`best.tflite` checksum), configuration parameters reviewed (altitudes, speeds, geofence coordinates, search polygon). A dry-run is executed to confirm the search pattern and geofence visualisation.
3. **Communication links** — mavproxy bridge running with dual UDP outputs, Mission Planner connected via TCP, MJPEG video stream accessible in the ground station browser, RC transmitter bound and responsive.
4. **Safety systems** — RC kill switch tested (mode switch to STABILIZE), geofence boundaries loaded and visualised, link-loss timeout configured, return-to-home altitude set.

The diagnostics dashboard (`diagnostics.py`) provides a single-screen summary of all communication links, camera status, GPS fix quality, and battery level, serving as the final go/no-go gate before arming. The complete checklist, including per-state contingency actions, is maintained as a printable document (`docs/FLIGHT_DAY_CHECKLIST.md`) designed to be followed top-to-bottom without requiring prior system knowledge.

## F Field Testing and Results

This appendix presents the quantitative results from bench testing, sensor calibration, inference benchmarking, and post-hoc video analysis that collectively characterise the system’s real-world performance and identify the remaining unknowns prior to autonomous flight.

On 11 March 2026 the full hardware stack—Raspberry Pi 5, Cube Orange autopilot, and IMX296 global-shutter camera—was assembled and tested at the designated flight site. Although adverse weather prevented outdoor flight, bench and calibration tests validated the complete hardware chain and produced quantitative benchmarks that informed subsequent system tuning.

### F.1 Bench Testing Results

The first objective was to verify end-to-end connectivity between the Pi, Cube, and camera in a field-representative configuration. A three-terminal workflow was established over SSH (PuTTY) and the Pi’s local display:

**Terminal 1 (mavproxy):** Forwarded Cube telemetry from the UART link (`/dev/ttyAMA0`, 921600 baud) to two UDP outputs—port 14550 for flight scripts and port 14551 for the diagnostics dashboard—plus a TCP bridge on port 5762 for Mission Planner on the laptop [10].

**Terminal 2 (diagnostics):** Ran the diagnostics script, confirming camera frame capture, AI model loading, Cube heartbeat, GPS satellite count, and battery voltage in a single multi-panel view.

**Terminal 3 (scripts):** Executed calibration and benchmark scripts sequentially.

Two integration issues were identified and resolved during setup: a camera-in-use conflict caused by two scripts attempting to open the CSI interface simultaneously, and a port-binding collision from a previous process. Both were resolved by enforcing a single-script-at-a-time workflow and adding port-release guards. With these fixes, all four subsystems—camera, AI model, Cube link, and GPS receiver—reported healthy status simultaneously for the first time on real hardware.

### F.2 FOV Calibration

Accurate field-of-view (FOV) calibration is critical because every GPS estimate derived from pixel coordinates depends on the ratio of sensor width to focal length [33]. Calibration was performed by streaming the camera view while a tape

measure was placed directly below the lens at a known height of 1.0 m.

The measured visible width was 92 cm at 1.0 m height, yielding a corrected focal length:

$$f = \frac{W_{\text{sensor}} \times d}{W_{\text{visible}}} = \frac{5.02 \text{ mm} \times 1000 \text{ mm}}{920 \text{ mm}} = 5.46 \text{ mm} \quad (2)$$

where  $W_{\text{sensor}} = 5.02 \text{ mm}$  is the IMX296 sensor width [22]. The previous default of 7.0 mm would have underestimated the ground footprint by 22%, producing systematic GPS estimation errors of approximately 2 m at 10 m altitude. The corrected value was committed to the configuration immediately.

### F.3 Lens Distortion Calibration

Lens distortion was characterised using a printed checkerboard calibration target ( $14 \times 9$  grid with  $13 \times 8$  inner corners, 28 mm square size). Multiple images of the board were captured at varying angles and distances, and intrinsic parameters and distortion coefficients were computed using the Zhang calibration method [33] implemented in OpenCV [34].

The calibration achieved an RMS reprojection error of 0.399 pixels, indicating an accurate fit. The resulting camera matrix and distortion coefficients were stored on the Pi and integrated into the vision module using precomputed remap lookup tables, adding only 1.5 ms per frame—negligible compared to the 206 ms inference time. All downstream scripts benefit automatically without modification.

### F.4 Inference Benchmarks

The primary inference benchmark was conducted on the Raspberry Pi 5 using the YOLOv8n model exported to TFLite [3] (float32, 3.2 MB, input shape  $[1, 640, 640, 3]$ ). The benchmark script ran 50 inference passes on a static test image with the dummy target visible at a representative scale. Table 32 summarises the results.

**Table 32:** Inference benchmark results on Raspberry Pi 5 (TFLite, XNNPACK CPU delegate [35]).

| Configuration        | Avg. Latency (ms) | FPS | Detection Rate | Confidence |
|----------------------|-------------------|-----|----------------|------------|
| Without undistortion | 206.5             | 4.8 | 50/50          | 0.966      |
| With undistortion    | 208.0             | 4.8 | 50/50          | 0.958      |

The 1.5 ms overhead from lens undistortion is within measurement noise and does not reduce the effective frame rate. A detection rate of 100% with a mean confidence of 0.966 confirms that the model produces reliable detections on the target hardware. At 4.8 FPS the system processes approximately one frame every 208 ms, which is sufficient for a drone travelling at the nominal search speed of 8 m/s (altitude-dependent schedule at 35 m)—covering roughly 1.7 m between frames, well within the camera’s ground footprint at 15–35 m altitude.

Research into inference acceleration identified several viable paths for future work: FP16 XNNPACK quantisation (estimated  $\sim 10$  FPS on the Pi 5’s native FP16 pipeline), threaded capture–inference pipelining (+20–30%), NCNN backend [24] ( $\sim 15$  FPS), and overclocking to 2.8 GHz (+15%). INT8 quantisation and Coral TPU acceleration were ruled out due to platform incompatibility with the Pi 5 [2].

### F.5 Post-Hoc Video Analysis

To supplement the bench tests, a post-hoc analysis was performed on DJI Mavic flight video recorded over the test site at altitudes between 15 and 50 m. The 30 fps video ( $1456 \times 1088$ , cropped from 4K) was replayed through the detection pipeline, which overlays bounding boxes on each frame and computes GPS target estimates from pixel coordinates and SRT telemetry (altitude, drone GPS position). A methodological constraint was identified during this analysis: SRT telemetry files contain one entry per 30 fps frame (6854 entries), so only the native-rate video maintains correct frame-to-telemetry alignment; downsampled versions introduce a frame-index mismatch that corrupts altitude and GPS readings.

#### F.5.1 Detection Performance

The retrained model ( $\text{mAP}_{50} = 0.995$  on validation data) successfully detected the dummy target across the full altitude range captured in the video. The FOV for the cropped DJI video frame (distinct from the  $49.3^\circ$  HFOV of the onboard Pi IMX296 camera) was independently calibrated at  $54.4^\circ$  HFOV (focal length  $1416 \pm 41$  px) using nine measurements of the known 1.8 m dummy at different altitudes; the dummy consistently measured 1.7–1.9 m across all samples, validating the calibration.

A comparative analysis tool enabled live model switching on the same video frames, confirming that the retrained model offered improved detection consistency over the baseline, particularly at higher altitudes where the target occupies fewer pixels.

### F.5.2 GPS Estimation Accuracy

Each detection was converted to a GPS estimate using the geo-transformation module (Section K.4), producing a scatter plot of estimated target positions. The circular error probable (CEP) metrics were:

- **CEP50** = 2.3 m (50% of estimates within this radius of the median);
- **Maximum error** = 16.5 m (a single outlier during a high-speed pass at 50 m altitude).

The scatter distribution exhibited a characteristic diagonal elongation aligned with the flight direction, attributable to GPS receiver latency [36]. Analysis of the timing offset between telemetry updates and frame timestamps revealed a systematic lag of 100–200 ms. At the nominal search speed of 8 m/s, this lag introduces an along-track displacement of 0.8–1.6 m per detection. The effect is mitigated by spatial clustering: the mission software aggregates detections from multiple passes using inverse-variance-weighted averaging, and the median of a cluster converges on the true target position as the number of observations increases.

## F.6 Dry-Run Validation

A dry-run mode was implemented to validate the mission-planning pipeline without requiring hardware or GPS connectivity. The mode instantiates the lawnmower search-pattern generator against the survey polygon, computes waypoints, estimates flight timing, and exercises the state-machine transition logic—all without arming or transmitting any MAVLink commands.

For the designated survey area, the dry-run produced the following outputs:

- **18 waypoints** covering the survey polygon in a lawnmower pattern at 35 m altitude with lane spacing derived from the camera footprint and a 20% overlap margin;
- **Estimated search time** of approximately 1.1 min at the nominal search speed of 8 m/s (altitude-dependent schedule at 35 m), plus transit to and from the survey area at 15 m/s;
- **A state-machine walkthrough:** INIT → CONNECTING → ARMING (GPS fix  $\geq 3$ , satellites  $\geq 6$ ) → TAKEOFF (35 m) → SEARCH (18 waypoints) → CENTERING → VERIFY (35 m) → LANDING;
- **A map visualisation** overlaying the generated waypoints and connecting legs on the satellite image.

The dry-run confirmed complete coverage of the survey polygon without generating waypoints that encroach on the adjacent SSSI no-fly zone [4]. It also verified that alternate models can be selected at launch without file-copy operations, reducing the risk of configuration errors on flight day.

## F.7 Discussion

The bench testing, calibration, and video analysis results collectively provide evidence that the system meets its minimum functional requirements while quantifying the uncertainties that remain before autonomous flight.

**Detection pipeline confidence.** The TFLite benchmark demonstrated 100% detection at 0.966 mean confidence on a static test image, and the DJI video replay confirmed detection across a 15–50 m altitude range with the retrained model. These results were obtained under controlled conditions: a stationary camera for the bench test and a gimbal-stabilised camera for the video analysis. The Pi’s IMX296 global shutter mitigates motion blur relative to a rolling-shutter sensor [22], but detection performance at the configured search speed with a non-stabilised camera mount remains unmeasured. This constitutes the single largest source of uncertainty.

**GPS estimation accuracy.** The CEP50 of 2.3 m from DJI video analysis is within the operational requirement: the mission does not demand sub-metre localisation, as the centering phase uses visual servoing to refine position once a detection triggers investigation. The GPS timing lag (100–200 ms) was identified as a systematic error source and is addressed in software through spatial clustering with inverse-variance weighting [37]. The 16.5 m outlier occurred during a high-speed, high-altitude pass and would be filtered by the clustering algorithm in an operational mission.

**Inference throughput.** At 4.8 FPS, the system processes one frame approximately every 1.7 m of ground travel at the nominal 8 m/s search speed (altitude-dependent schedule at 35 m). Given the camera footprint of approximately 32 m × 24 m at 35 m altitude, each ground point appears in roughly 15 consecutive frames along the flight direction. Even with a conservative 50% detection rate under flight conditions, this provides approximately 8 detection opportunities per target—sufficient for reliable triggering of the investigation sequence.

**Remaining unknowns.** Four quantities could not be measured without outdoor flight and must be resolved during the first flight test: (1) detection reliability with the non-stabilised Pi camera at search speed; (2) real GPS accuracy and satellite geometry at the test site; (3) wind-induced position-hold error during the centering and verification phases; and (4) MAVLink command latency over the Pi-Cube UART link under flight load. The progressive test protocol (Section E) is designed to isolate each variable incrementally, beginning with passive observation during manual flight before enabling autonomous commands [30].

## G Future Work

This appendix outlines the development roadmap beyond the current submission, organised by feasibility so that readers can distinguish between improvements achievable with existing hardware and those requiring additional investment or regulatory approval.

The system presented in this report demonstrates a complete autonomous SAR pipeline from search through detection to landing, but several directions offer clear paths toward operational capability. The following subsections organise future work by feasibility, distinguishing near-term enhancements achievable with the existing hardware from longer-term developments that would require additional equipment, regulatory approval, or multi-platform coordination.

### G.1 Inference Acceleration

The TFLite backend achieves 4.8FPS on the Raspberry Pi 5, which is adequate for the current search speed but leaves limited margin for faster flight or more complex models. Four acceleration strategies are identified in order of increasing implementation effort:

1. **FP16 XNNPACK.** The Pi 5’s Cortex-A76 cores support native half-precision arithmetic. Enabling FP16 execution in the XNNPACK delegate requires no model changes and is expected to approximately double throughput to  $\sim 10$  FPS, based on ARM’s published benchmarks for similar workloads [38].
2. **Threaded pipeline.** The current architecture serialises frame capture and inference. Overlapping these stages in a producer–consumer arrangement would reduce effective per-frame latency by 20–30 %, as the camera sensor and CPU operate on independent hardware. The Python `threading` module is sufficient because the Global Interpreter Lock (GIL) is released during both camera I/O and TFLite native inference calls.
3. **NCNN backend.** An NCNN export of the YOLOv8n model has already been generated and is stored in the repository (`cv_models/sar_v2.1088/ncnn/`). NCNN is optimised for ARM NEON instructions and is expected to reach  $\sim 15$  FPS on the Pi 5 [24]. Integration requires replacing the TFLite inference call in `vision.py` with an NCNN loader; the detection interface (`detect_in_image`  $\rightarrow$  found,  $x$ ,  $y$ , confidence) remains unchanged.
4. **Hailo-8L NPU.** The Raspberry Pi AI HAT+ (£55) provides a dedicated neural processing unit capable of 80+ FPS for YOLOv8n [39], removing inference as a bottleneck entirely. This would enable real-time detection at full camera frame rate and free CPU cycles for concurrent tasks such as visual odometry or on-board path replanning.

Of these, the first two are immediately achievable with no additional hardware and represent the recommended next steps. The NCNN backend requires validation on the Pi to confirm that the exported model reproduces detections equivalent to the TFLite baseline. The Hailo accelerator, while transformative, introduces a hardware dependency and a model conversion workflow (HEF format compiled via an x86 Linux toolchain) that places it in the medium-term category.

### G.2 Detection Improvements

The YOLOv8n model achieves  $\text{mAP}_{50} = 0.995$  on the validation set, but operational robustness under varied field conditions requires further work:

- **Temporal voting.** Requiring a detection to persist in at least  $N$  of  $M$  consecutive frames before triggering investigation would suppress transient false positives caused by shadows, debris, or image artefacts. The spatial clustering system already provides partial temporal consistency, but an explicit frame-count threshold would add a complementary filtering layer at negligible computational cost.
- **Size and aspect-ratio filtering.** At a given altitude, a human casualty occupies a predictable range of bounding-box dimensions and aspect ratios. Rejecting detections outside this range—for example, boxes that are too small (sensor noise) or too wide (vehicles, paths)—would reduce the false-positive rate without retraining the model.
- **Hard negative mining.** The current training set includes 50 negative images (frames with no target). Systematically collecting false positive crops from flight footage and adding them as hard negatives during retraining would teach the model to distinguish the dummy from visually similar background features such as rocks, litter, or trail markers.
- **Expanded real training data.** Each flight produces thousands of frames that can be labelled with the existing annotation tool. Incorporating 200+ real images spanning varied lighting, altitude, pose, and background

conditions would reduce the synthetic-to-real domain gap that currently limits generalisation [40].

### G.3 Multi-Target Support

The simulator already implements spatial clustering and priority scoring for multiple simultaneous detections. Extending the flight system to support a multi-target queue would allow the drone to detect, classify, and record the positions of several casualties in a single sortie, revisiting unconfirmed detections after completing the primary search pattern. The ground station interface already supports per-target classification (confirmed, interest, false positive), so the principal remaining work lies in the state machine’s transition logic for queueing and revisiting targets. In mass-casualty scenarios, this capability would enable a single sortie to locate and prioritise multiple victims before ground teams are dispatched.

### G.4 Thermal Imaging

Visual detection is fundamentally limited by lighting conditions and the contrast between the casualty and the surrounding terrain. A thermal camera operating in the long-wave infrared band (8–14  $\mu\text{m}$ ) would detect body-heat signatures independently of visible appearance, enabling night-time and low-visibility operations [41]. Thermal and visible imagery could be fused at the decision level: a detection confirmed in both modalities would carry substantially higher confidence than either alone. Lightweight thermal modules compatible with the Raspberry Pi (e.g., FLIR Lepton 3.5) are available for under £200 and produce  $160 \times 120$  images sufficient for person-scale detection from 30 m altitude. However, thermal imaging introduces additional challenges—sun-heated objects produce signatures similar to body heat, and wind chill reduces the thermal contrast between a casualty and the ambient background [9]—so it complements rather than replaces the visual pipeline.

### G.5 Communication

The current system relies on a direct MAVLink telemetry link between the ground station laptop and the aircraft, limiting operational range to the Wi-Fi or radio link budget (typically 500 m to 1 km). Two enhancements would extend this range:

- **4G/LTE telemetry.** A cellular modem on the aircraft would provide telemetry and command links over the mobile network, extending range to wherever cellular coverage exists. This is particularly relevant for rural SAR operations where the search area may extend several kilometres from the operator. Lightweight LTE modules (e.g., Sixfab 4G HAT) weigh under 30 g and integrate directly with the Pi 5. Typical latency (50–150 ms) is acceptable for supervisory control but would require a local autonomy layer to handle time-critical decisions (e.g., obstacle avoidance) without round-trip delay.
- **Real-time video to command centre.** Streaming compressed video (H.264 over WebRTC or RTSP) to a remote incident command post would allow SAR coordinators to observe detections as they occur, rather than relying on post-flight review. The Pi 5’s hardware video encoder can produce a 720p stream at minimal CPU cost, but upstream bandwidth constraints on cellular networks (1–5 Mbps typical) would require adaptive bitrate control.

Both enhancements depend on obtaining Beyond Visual Line of Sight (BVLOS) operational authorisation from the UK Civil Aviation Authority under CAP 722 [42], which requires a detect-and-avoid capability for other airspace users, redundant communication links, and a comprehensive safety case.

### G.6 Swarm Operations

Deploying multiple drones to divide a search area would reduce coverage time approximately linearly with fleet size. Each aircraft would fly an independent sub-region while sharing detection events over a mesh network, enabling collaborative coverage planning that avoids redundant passes [43]. The lawnmower pattern generator already accepts arbitrary polygons and could partition a large search area into sub-polygons assigned to individual aircraft. This is the most aspirational direction discussed here: the primary challenges are inter-drone communication reliability, deconfliction of overlapping flight paths, and a centralised operator interface that aggregates detections from all platforms without overwhelming the operator [44]. Nevertheless, the modular software architecture—where detection, planning, and control are independent modules communicating through well-defined interfaces—provides a foundation that would support multi-agent extension without fundamental redesign.

## H Edge Inference Architecture

This appendix examines the engineering decisions behind deploying a YOLOv8n object detector on a Raspberry Pi 5 companion computer. Selecting an inference backend for an embedded ARM platform requires balancing latency, dependency footprint, deployment reliability, and quantisation trade-offs—concerns that are largely invisible during laptop-based development but become dominant constraints in a field-deployed UAV system. The following subsections

document the backend evaluation, quantisation analysis, pipeline architecture, and benchmark methodology that informed the production configuration.

## H.1 Compute Platform Constraints

The Raspberry Pi 5 is built around the Broadcom BCM2712 system-on-chip, which integrates four ARM Cortex-A76 cores clocked at 2.4GHz [45]. The A76 microarchitecture implements the ARMv8.2-A instruction set, which includes mandatory NEON Advanced SIMD (128-bit vector registers, fused multiply-accumulate across single- and half-precision floats) and optional half-precision (FP16) data-processing extensions. These SIMD units are the primary compute resource for neural-network inference: the Pi 5 lacks a programmable GPU compute pipeline (the VideoCore VII is a tile-based rasteriser without general-purpose shader support) and carries no dedicated neural-processing unit (NPU). Consequently, all inference workloads execute on the four CPU cores, and throughput is bounded by single-threaded NEON throughput and L2 cache bandwidth rather than by massively parallel ALU arrays.

This architectural reality eliminates frameworks that depend on GPU offload (e.g. CUDA, OpenCL) and motivates the selection of inference runtimes specifically optimised for ARM CPU execution with NEON vectorisation.

## H.2 Backend Selection

Three inference backends were evaluated for deployment of the YOLOv8n model on the Pi 5. Table 33 summarises their characteristics.

**TFLite with XNNPACK delegate.** TensorFlow Lite [3] was selected as the primary deployment backend for three reasons. First, the XNNPACK delegate [35]—a highly optimised library of floating-point micro-kernels—maps convolution, depthwise convolution, and fully-connected operations directly onto ARM NEON intrinsics, achieving near-peak utilisation of the SIMD pipeline without requiring framework-level graph optimisation. Second, the Python runtime is lightweight: `ai-edge-litert` (the successor to `tflite-runtime` for Python 3.13) adds approximately 5 MB to the deployment footprint, compared with over 800 MB for a full PyTorch installation. Third, TFLite’s model format is stable and well-documented, with direct export support from the Ultralytics training framework via `yolo export format=tflite`.

On the Pi 5, XNNPACK automatically detects the number of available cores and distributes work across them. The default thread count of four (matching the physical core count) was used for all benchmarks; reducing to two threads increased latency by 38%, confirming that inference is compute-bound rather than memory-bound at this model scale.

**NCNN.** NCNN [24], developed by Tencent for mobile deployment, is an alternative backend that has demonstrated faster execution on ARM platforms in third-party benchmarks. Ultralytics reports YOLOv8n inference at approximately 68ms on the Pi 5 using NCNN [11], a 3× improvement over TFLite. This speedup is attributed to NCNN’s hand-tuned NEON assembly kernels and its memory-efficient blob pooling strategy, which reduces allocation overhead during inference.

NCNN support was implemented in `vision.py` and is available via the `--backend ncnn` flag. Model weights were exported to the NCNN format (`.param/.bin`) through an ONNX intermediate representation. However, NCNN was not used during flight testing for two reasons: (1) the `ncnn` Python package on Python 3.13/aarch64 required manual compilation at the time of deployment, introducing a fragile build dependency on the field Pi; and (2) the 4.8FPS achieved by TFLite was sufficient for the detection-at-altitude use case, where the target remains in the field of view for 10–20 consecutive frames during a flyover (Section I.2). NCNN remains the recommended upgrade path for missions requiring higher frame rates.

**Ultralytics/PyTorch.** The Ultralytics framework [11] serves as the development and training backend on the laptop (Ubuntu/Windows with CUDA or CPU). It is unsuitable for Pi deployment: PyTorch’s memory footprint exceeds 800 MB, cold-start time is several seconds, and single-frame inference exceeds 1.2s on the Cortex-A76 without GPU acceleration. The framework is retained exclusively for laptop-based development, training, model export, and video analysis.

## H.3 Quantisation Trade-offs

The deployed model uses 32-bit floating-point (FP32) weights and activations. Three numerical precision levels were considered:

- **FP32 (deployed).** Full single-precision inference. This is the baseline configuration, requiring no additional export steps. All measured benchmarks (Table 36) use FP32.

- **FP16 (viable, not deployed).** The Cortex-A76 implements the ARMv8.2-A FP16 data-processing extension, enabling native half-precision fused multiply-accumulate operations at twice the throughput of FP32 [46]. The XNNPACK delegate supports transparent FP16 execution of FP32 models via the `TFLITE_XNNPACK_DELEGATE_FLAG_FORCE_FP16` option, which converts weights and activations to FP16 at inference time without re-export. This is projected to approximately halve inference latency to  $\sim 100$  ms ( $\sim 10$  FPS). FP16 was not enabled during field testing because the FP32 frame rate was already adequate and the FP16 path had not been validated for detection accuracy regression on the custom model.
- **INT8 (investigated, not viable).** Post-training quantisation to 8-bit integers reduces model size by  $4\times$  and can accelerate inference on hardware with dedicated integer SIMD units. However, INT8 quantisation of YOLOv8n via the XNNPACK delegate yielded only a  $\sim 30\%$  speedup in published benchmarks [3], and INT8 inference on NCNN is unsupported on the Pi 5's ARM64 cores [24]. Additionally, INT8 export requires a representative calibration dataset of 300+ images and careful validation to ensure that quantisation-induced accuracy loss does not degrade detection of small targets at altitude. Given the marginal benefit and validation cost, INT8 was deprioritised in favour of the more straightforward FP16 path.

## H.4 Pipeline Architecture

The inference pipeline follows a strictly sequential architecture: frame capture, lens undistortion, preprocessing, model inference, and postprocessing execute in a single thread on a single core, while the remaining three cores handle XNNPACK's internal parallelism during the inference step.

A **threaded pipeline** was considered, in which camera capture ( $\sim 30$  ms) would execute concurrently with inference ( $\sim 206$  ms) on separate cores, hiding capture latency and increasing effective throughput by an estimated 20–30%. This design was rejected for two reasons:

1. **Race conditions in shared state.** The state machine reads detection results (target pixel coordinates, confidence) and camera frames from the same `VisionSystem` instance. Concurrent capture and inference would require synchronisation primitives (locks or double-buffering) around the frame buffer and detection output. In a safety-critical flight system, the additional complexity and potential for deadlocks or stale-frame processing was judged unacceptable for a marginal throughput gain.
2. **Python GIL limitations.** Although camera I/O releases the Global Interpreter Lock (GIL), the NumPy preprocessing steps (colour conversion, resize, normalisation) and TFLite's Python-side tensor copy do not. The effective overlap is therefore limited to the raw camera read, reducing the theoretical 20–30% gain to a more modest 10–15% in practice.

The single-threaded design guarantees that every detection result corresponds to the most recently captured frame with deterministic timing, simplifying the GPS-to-pixel timestamp alignment used by the target localisation module (Section 2.6).

## H.5 Video Streaming Architecture

The ground station receives a live video feed from the Pi over Wi-Fi for operator situational awareness and verification decisions. The streaming protocol was selected based on four requirements: browser compatibility (no native application), minimal latency ( $< 500$  ms), zero additional infrastructure, and headless operation (no display attached to the Pi).

**MJPEG over HTTP (selected).** The Pi serves an MJPEG stream on port 8090 via Python's built-in `http.server` module with `ThreadingMixIn` for concurrent connections. Each frame is JPEG-compressed at quality 70 ( $\sim 30$  kB per frame at  $1456 \times 1088$ ) and pushed as a multipart HTTP response. The operator opens `http://PI_IP:8090` in any browser to view the stream overlaid with detection bounding boxes, GPS estimates, and command buttons (Y/N/I/X/L). This approach requires no client-side codec support beyond JPEG decoding, which is universal across browsers and platforms.

**Alternatives considered.** H.264 over HLS was prototyped (`camera_stream_h264.py`) but introduced 2–4 seconds of latency due to the segment-based delivery model, which is unacceptable for real-time operator verification. H.264 over RTP/RTSP would achieve lower latency but requires a dedicated player (e.g. GStreamer, VLC) on the ground station, violating the browser-compatibility requirement. WebRTC would satisfy both latency and browser requirements but introduces significant implementation complexity (STUN/TURN negotiation, DTLS-SRTP, SDP exchange) disproportionate to a single-stream, same-network use case.

The MJPEG approach achieves sub-200 ms glass-to-glass latency on the local Wi-Fi network, which is adequate for the operator's binary confirm/reject decision during the VERIFY state.

## H.6 Benchmark Methodology


Inference benchmarks were conducted using the `tests/hardware/benchmark.py` script, which follows a standardised protocol:

1. **Warm-up phase.** Five inference passes are executed and discarded to populate instruction and data caches, trigger JIT compilation in the XNNPACK delegate, and allow thermal throttling to stabilise.
2. **Timed phase.** Fifty consecutive inference passes are timed using Python's `time.perf_counter()`, which provides microsecond-resolution wall-clock timing. Each pass includes the full preprocessing pipeline (BGR→RGB conversion, resize to  $640 \times 640$ , float32 normalisation, tensor expansion) and model invocation, matching the production code path exactly.
3. **Reporting.** Mean, minimum, maximum, and standard deviation of per-frame latency are computed. Detection rate (fraction of frames with confidence above threshold) and mean confidence are also recorded to verify that the model produces consistent results across runs.

XNNPACK was configured with the default thread count of four (one per Cortex-A76 core). The Pi was running Raspberry Pi OS (Debian Bookworm, 64-bit) with no desktop environment active, minimising background CPU load. Thermal management was provided by the official Pi 5 active cooler; no thermal throttling was observed during the 50-frame test window.

## H.7 Results and Projected Improvements

Table 34 summarises the measured baseline and projected performance for each optimisation path.



**Figure 18:** Inference backend comparison on Raspberry Pi 5 for YOLOv8n ( $640 \times 640$  input). TFLite with XNNPACK provides the measured production baseline at 4.8FPS; NCNN and FP16 projections indicate viable paths to real-time performance.

The measured 4.8FPS baseline, while modest in absolute terms, is sufficient for the current mission profile. At the nominal search speed of  $8 \text{ m s}^{-1}$  and altitude of 35 m, the ground-sample distance between consecutive frames is 1.67 m—less than 6% of the 32.2 m horizontal field of view—ensuring that any target within the search swath is observed in at least 14 consecutive frames. The projected NCNN and FP16 improvements represent viable paths to real-time performance for higher-speed search profiles or lower-altitude missions where the detection window is shorter.

## I Computer Vision: Extended Analysis

The main body of this report (Section 2.3) summarises the detection pipeline at a level sufficient to understand the system architecture. This appendix provides the full technical depth: stage-by-stage timing breakdowns measured on the Raspberry Pi 5, comprehensive benchmark results across all available models, a mission-level detection probability analysis that links frame rate to flyover reliability, altitude-dependent performance modelling validated against DJI flight-video replay, and the multi-backend abstraction that enables the same codebase to run transparently on laptop, Pi, and optimised ARM deployments. These details are essential for reproducing the system, tuning configuration parameters for different mission profiles, and understanding the margins available for future optimisation.

## I.1 Detection Pipeline: Step-by-Step Timing

Every iteration of the main loop executes the detection pipeline in `vision.py` through six sequential stages. Figure 7 illustrates the data flow; Table 35 provides measured timing for each stage on the Raspberry Pi 5.

### I.1.1 Stage 1: Camera Capture

The IMX296 global shutter sensor captures frames at  $1456 \times 1088$  pixels via CSI-2 using `picamera2`. A critical hardware-specific detail: despite being configured as RGB888, the IMX296 outputs data in BGR channel order—identical to OpenCV’s native format. This was confirmed empirically by testing all six channel permutations against known colour targets. Consequently, no `cv2.cvtColor(frame, cv2.COLOR_RGB2BGR)` conversion is applied, saving approximately 1.2 ms per frame and avoiding a subtle but mission-critical colour inversion that would degrade detection accuracy.

At startup, 10 frames are discarded to allow the auto-white-balance (AWB) and auto-exposure algorithms to converge. The camera can be configured for preset AWB modes (daylight, cloudy, indoor) or fixed manual colour gains via `config.py`, depending on lighting conditions.

### I.1.2 Stage 2: Lens Undistortion

If a lens calibration file (`calibration_data.npz`) is present, precomputed remap matrices are loaded at initialisation via `cv2.initUndistortRectifyMap()`. Per-frame undistortion uses `cv2.remap()` with bilinear interpolation, costing 1.5 ms—less than 1% of the total pipeline latency. The calibration was performed using a  $14 \times 9$  checkerboard ( $13 \times 8$  inner corners, 28 mm square size), achieving an RMS reprojection error of 0.399 pixels. This corrects barrel distortion from the wide-angle lens, which otherwise displaces peripheral detections by up to 8 pixels, corresponding to approximately 0.7 m GPS error at 35 m altitude.

### I.1.3 Stage 3: Preprocessing

The undistorted  $1456 \times 1088$  BGR frame is converted to RGB (`cv2.cvtColor`), resized to  $640 \times 640$  via bilinear interpolation (`cv2.resize`), cast to `float32`, and normalised to  $[0, 1]$  by dividing by 255. The result is expanded to a batch tensor of shape  $[1, 640, 640, 3]$ .

Direct resize (without letterboxing) was chosen over the alternative of maintaining aspect ratio with padding for three reasons: (1) the aspect ratio of the IMX296 sensor (4:3) is close to square, so the distortion introduced by non-uniform scaling is modest ( $\sim 7\%$  vertical stretch); (2) letterboxing wastes  $\sim 14\%$  of the  $640 \times 640$  input area on grey padding, reducing the effective resolution available for small-target detection; (3) the coordinate rescaling from model output back to original frame dimensions is a simple linear map, avoiding the padding-offset arithmetic that introduces rounding errors.

### I.1.4 Stage 4: TFLite Inference

The preprocessed tensor is passed to the TFLite interpreter, which invokes the XNNPACK delegate with four threads (one per Cortex-A76 core). This stage dominates the pipeline at 196 ms (94% of total latency).

The YOLOv8n architecture produces an output tensor of shape  $[1, 5, 8400]$ , where:

- **8400** is the total number of anchor-free detection candidates, arising from three feature pyramid network (FPN) scales:  $80 \times 80 = 6400$  candidates at the finest scale (detecting small objects),  $40 \times 40 = 1600$  at medium scale, and  $20 \times 20 = 400$  at the coarsest scale (detecting large objects).
- **5** encodes  $[x_{\text{centre}}, y_{\text{centre}}, w, h, \text{class.conf}]$  for a single-class model. For multi-class models (e.g. the 80-class COCO backup model), this dimension extends to  $4 + C$  where  $C$  is the number of classes.

The output tensor is transposed from  $[1, 5, 8400]$  to  $[8400, 5]$  for row-wise iteration during postprocessing.

### I.1.5 Stage 5: Postprocessing

The postprocessing stage applies a confidence threshold of 0.2 (configurable via `config.py`) and selects the single highest-confidence detection. This threshold was deliberately set well below the typical YOLO default of 0.5 for a critical operational reason: in a search-and-rescue context, a *missed detection is irrecoverable*—the drone may never revisit that ground patch—while a *false positive* is filtered by the human operator at the VERIFY stage and costs only a brief investigation diversion. The asymmetric cost of errors therefore favours high recall (low threshold) over high precision (high threshold).

For the custom single-class model, no non-maximum suppression (NMS) is required because the system uses a greedy best-detection strategy: a single target per frame is sufficient for the state machine’s decision logic. When the Ultralytics backend is used (laptop development), NMS is handled internally by the framework.

Coordinates are rescaled from the  $640 \times 640$  model space to the original  $1456 \times 1088$  frame dimensions via simple multiplication:  $x_{\text{frame}} = x_{\text{model}} \times (1456/640)$ .

### I.1.6 Stage 6: Visualisation

A green bounding box with a confidence label (e.g. “AI 0.97 [dummy]”) is drawn onto the original frame using `cv2.rectangle` and `cv2.putText`. This annotated frame is simultaneously served to the ground station MJPEG stream and (in non-headless mode) displayed locally via `cv2.imshow`.

## I.2 Comprehensive Benchmark Results

Table 36 presents the full benchmark results from the Pi 5, measured using the standardised protocol described in Section H.6 (5 warm-up frames discarded, 50 timed passes, `time.perf_counter()`).

Several observations warrant discussion:

**Inference time is model-size-independent.** All three models produce nearly identical inference latencies (206–208 ms) despite a  $4\times$  size difference (3.2 MB vs 13.0 MB). This confirms that YOLOv8n inference is compute-bound (limited by NEON multiply-accumulate throughput) rather than memory-bound (limited by weight loading from DRAM). The model architecture—not the weight count per class—determines latency, because all three models share the identical YOLOv8n backbone; only the final classification head differs.

**Undistortion overhead is negligible.** Enabling lens undistortion adds only 1.5 ms ( $<1\%$  overhead), confirming that precomputed remap matrices are an efficient correction strategy. The slight confidence decrease with undistortion ( $0.966 \rightarrow 0.958$ ) is within measurement noise and does not affect detection rate.

**Custom model vastly outperforms generic COCO.** The custom SAR model achieves 100% detection rate with 0.966 mean confidence, compared to 86% and 0.724 for the generic COCO person detector. This validates the domain-specific retraining strategy (Section C): a model trained on aerial views of a mannequin against grassy backgrounds significantly outperforms a general-purpose detector trained on ground-level person images. The COCO model is retained as a backup for scenarios where the target is a standing human rather than a prone mannequin.

**Timing stability.** The standard deviation of per-frame latency was 3.2 ms ( $\sigma/\mu = 1.5\%$ ), indicating highly reproducible timing. No thermal throttling was observed during the 50-frame test window with the official Pi 5 active cooler.

## I.3 Mission Runtime Analysis

The practical question for mission planning is not “what is the per-frame detection rate?” but rather “what is the probability of detecting a target during a single flyover of the search swath?” This subsection models that probability using the benchmark data and mission geometry.

### I.3.1 Ground Footprint Geometry

At search altitude  $h$ , the camera’s ground footprint is determined by the horizontal and vertical fields of view:

$$W_g = 2h \tan\left(\frac{\theta_H}{2}\right) \quad (3)$$

$$L_g = 2h \tan\left(\frac{\theta_V}{2}\right) \quad (4)$$

where  $\theta_H = 2 \arctan(w_s/2f) \approx 49.3^\circ$  is the horizontal FOV (sensor width  $w_s = 5.02$  mm, focal length  $f = 5.46$  mm) and  $\theta_V \approx 37.9^\circ$  is the vertical FOV.

Table 37 gives the ground footprint at key operational altitudes.

The “Sensor px” column shows the dummy’s height in the native  $1456 \times 1088$  frame ( $= 1.8$  m/GSD, where  $\text{GSD} = L_g/1088$ ). The “Model px” column applies the resize scale factor  $640/1456 = 0.4396$  to convert to the  $640 \times 640$  inference input.

### I.3.2 Frames per Flyover

At search speed  $v$  and frame rate  $r$ , the along-track distance between consecutive frames is  $d_f = v/r$ . The number of frames in which a stationary target remains within the field of view during a single pass is:

$$N_{\text{frames}} = \frac{L_g}{d_f} = \frac{L_g \cdot r}{v} \quad (5)$$

Table 38 evaluates this for the operational speed and altitude envelope.

**Key finding.** Even in the worst operational case (50 m altitude at 15 m/s transit speed, with a conservative 70% single-frame detection probability), the cumulative detection probability across the 11-frame flyover window exceeds 99.8%. At the nominal search configuration (35 m, 8 m/s from the altitude-dependent schedule), the system achieves effectively certain detection ( $P > 99.99\%$ ) with 14.5 frames per target.

The lawnmower pattern further improves robustness: adjacent passes overlap by 20% (configurable via `planning.py`), so targets near the swath edge receive a second observation window. The probability of missing a target across two overlapping passes at nominal conditions is  $(1 - 0.9997)^2 \approx 10^{-7}$ .

### I.3.3 Frame Overlap Analysis

At 4.8 FPS and the nominal 8 m/s search speed, consecutive frames are separated by  $8/4.8 = 1.67$  m along-track. The along-track footprint at 35 m is 24.1 m, giving a frame-to-frame overlap of:

$$\text{Overlap} = 1 - \frac{d_f}{L_g} = 1 - \frac{1.67}{24.1} = 93.1\% \quad (6)$$

This high overlap means that a target appears in many consecutive frames, providing the state machine with multiple opportunities to trigger the CENTERING sequence. The `--smart-detect` mode exploits this by requiring  $k = 3$  consecutive positive detections before committing to investigation, effectively implementing a temporal majority-vote filter that suppresses transient false positives at a cost of only  $3 \times 208 = 0.6$  s delay (5.0 m of travel at 8 m/s).

## I.4 Detection Performance vs Altitude

The critical operational question is: “at what altitude does detection become unreliable?” This is modelled using the thin-lens equation to predict the target’s apparent size in pixels, combined with empirically observed confidence behaviour from video replay analysis.

### I.4.1 Pixel Size Model

The target of real height  $H_t = 1.8$  m first projects onto the sensor with  $\text{px}_{\text{sensor}} = H_t \cdot f_{\text{px}}/h$  pixels, where  $f_{\text{px}} = f/w_s \times W_{\text{px}} = 5.46/5.02 \times 1456 = 1584$  px. The  $1456 \times 1088$  frame is then resized to  $640 \times 640$ ; applying the uniform scale factor  $640/1456 = 0.4396$  gives the model-input height:


$$p_t = \frac{H_t \cdot f_{\text{px}}}{h} \times \frac{640}{1456} \quad (7)$$

### I.4.2 Altitude Performance Table


Table 39 combines the pixel-size model with confidence data extracted from DJI video replay analysis (Section 2.6) and bench testing.

**Operational envelope.** The data confirms that 35 m (the configured `TARGET.ALT`) sits comfortably within the reliable detection zone (95% per-frame, >99.97% per-flyover). The altitude ceiling for reliable single-pass detection is approximately 40 m; above this, multi-pass coverage from the lawnmower pattern provides the necessary redundancy. The theoretical floor for detection ( $\sim 70$  m) corresponds to the target subtending approximately 18 model-input pixels (41 sensor pixels)—near the YOLOv8n smallest feature scale ( $80 \times 80$  grid, where each cell covers an  $8 \times 8$  region of the input).

**Confidence versus pixel size relationship.** Figure 20 shows the empirically measured confidence as a function of altitude. The relationship follows an approximately sigmoidal curve in log-altitude space, with a sharp transition between reliable detection ( $>0.9$ ) and unreliable detection ( $<0.5$ ) occurring over a narrow altitude band (40 m to 60 m). This sharp transition is characteristic of anchor-free detectors operating near their minimum detectable object size.



**Figure 19:** Detection performance as a function of altitude for the YOLOv8n custom model against a 1.8 m mannequin target. The table combines predicted pixel sizes from the thin-lens model with empirically observed confidence and detection rates from video replay analysis.




**Figure 20:** Detection confidence vs altitude for the YOLOv8n custom model. Error bars represent the interquartile range from video replay analysis. The dashed line marks the 0.2 confidence threshold.

**Combined detection envelope.** Figure 21 overlays the altitude-dependent confidence curve with the motion blur degradation model, producing a combined detection envelope that accounts for both target size reduction and image quality loss. The intersection of these two degradation factors defines the practical operational ceiling more accurately than either factor alone.

## 1.5 Model Training Results

The retrained v2 model (`sar_v2_1088`) was trained on Google Colab (NVIDIA T4 GPU) for 150 epochs using the mixed dataset described in Section C. Table 40 summarises the final evaluation metrics.

**Confusion matrix analysis.** The single-class confusion matrix on the validation set shows 99.5% true positive rate and 0.5% false negative rate. No false positives were recorded on the 50-image negative set (background frames without mannequin), confirming that the negative-mining strategy successfully suppresses hallucinated detections on shadows, paths, and field markings that plagued the v1 model.



**Figure 21:** Combined detection envelope showing confidence degradation from both altitude (target pixel size reduction) and motion blur at  $10 \text{ m s}^{-1}$ . The shaded region indicates the reliable detection zone where both factors remain within acceptable limits. The nominal search altitude of 35 m sits comfortably within this envelope.

**Blur-robustness through data augmentation.** A critical design decision in the training pipeline was the inclusion of augmentations that simulate real flight conditions. The dataset generator (`generate_dataset_v2.py`) applies four offline augmentations: random *brightness* jitter ( $\alpha \in [0.7, 1.3]$ ), *contrast* offset ( $\beta \in [-30, 30]$ ), *Gaussian blur* ( $3 \times 3$  or  $5 \times 5$  kernel, 30% probability), and random *rotation* ( $0$ – $360^\circ$ ). These are complemented by the YOLO trainer’s online augmentations—mosaic, scale, mixup, copy-paste, and horizontal flip (Table 24)—which further diversify the training distribution. Together, these augmentations expose the model to degraded images representative of the conditions analysed in Section AE.2, ensuring the detector maintains high confidence even when the real camera produces mildly blurred frames due to wind gusts or rapid manoeuvres—a robustness that would not be achieved by training exclusively on sharp synthetic composites.

**Comparison with v1.** The v1 model (trained on 200 synthetic images at  $640 \times 640$  from a single satellite background) achieved  $\text{mAP}_{50} \sim 0.95$  and suffered frequent false positives on high-contrast terrain features. The v2 improvements—real-flight backgrounds, native-resolution training, altitude-correct scaling, and explicit negatives—closed the domain gap and improved  $\text{mAP}_{50}$  by approximately 4.5 percentage points while eliminating the false positive problem.

**Export and deployment.** The trained `best.pt` weights were exported to TFLite (float32) via `yolo export format=tflite`, producing the same  $[1, 640, 640, 3]$  input and  $[1, 5, 8400]$  output shape as the v1 model. This ensures drop-in compatibility: deploying the v2 model on the Pi requires only copying the `.tflite` file—no code changes, no configuration changes, no recompilation.

## I.6 Dual Backend Architecture

The `VisionSystem` class in `vision.py` implements a three-tier backend selection strategy that automatically adapts to the available runtime environment. This design enables the same codebase to run on a development laptop with full Ultralytics/PyTorch, a Raspberry Pi with lightweight TFLite, or an optimised ARM deployment with NCNN—without any code changes or conditional imports in the calling modules.

### I.6.1 Backend Selection Logic

At import time, `vision.py` probes for available backends in priority order:

1. **NCNN** (if `--backend ncnn` flag is set and the `ncnn` package is installed). This backend provides the fastest ARM inference ( $\sim 15 \text{ FPS}$  projected) through hand-tuned NEON assembly kernels, but requires manual compilation on Python 3.13/aarch64 and was therefore not used during field testing.

2. **Ultralytics YOLO** (if the `ultralytics` package is importable). This is the default on the development laptop, providing the richest feature set (built-in NMS, class name mapping, training integration) at the cost of a large dependency footprint ( $>800$  MB).
3. **TFLite** (if `tflite-runtime`, `ai-edge-litert`, or `tensorflow.lite` is available). This is the production backend on the Pi, adding only  $\sim 5$  MB to the deployment. On Python 3.13, the legacy `tflite-runtime` package is unavailable; the system falls back to `ai-edge-litert` (Google’s successor package).

All three backends consume the same `.tflite` model file (or equivalent format) and produce output through the identical `detect_in_image(frame) → (found, cx, cy, confidence)` interface. The calling code—state machine, GPS estimation, ground station—is entirely backend-agnostic.

### I.6.2 Backend Abstraction Benefits

This architecture provides three concrete benefits:

- **Development-production parity.** Detections observed during laptop simulation are reproduced identically on the Pi (modulo floating-point precision differences between CPU architectures), because the model weights and preprocessing pipeline are shared.
- **Graceful degradation.** If the preferred backend fails to load (e.g. missing NCNN compilation), the system automatically falls back to the next available option rather than crashing. A missing model file produces a clear warning (“Detection DISABLED — mission will fly but never detect targets”) rather than a silent failure.
- **A/B testing.** Different backends can be compared on identical input frames using the `--backend` flag, enabling systematic evaluation of accuracy and latency trade-offs during development.

## I.7 Model Swapping and Field Deployment

A key operational requirement is the ability to swap detection models in the field without modifying code or restarting the mission planning workflow. Two mechanisms are provided:

**File-copy swap.** All scripts read their model from a single path: `best.tflite` in the project root. Swapping models is a single copy command:

```
cp cv_models/sar_v2_1088/best.tflite best.tflite
```

Three model variants are maintained in the `cv_models/` directory, each with TFLite, PyTorch (`.pt`), and NCNN exports:

**CLI flag swap.** The `--model` flag overrides the default model path at launch:

```
python main.py --model models/human.tflite
```

This enables rapid A/B testing during flight day without touching the filesystem.

**Zero-code-change guarantee.** All TFLite models share the same input tensor shape  $[1, 640, 640, 3]$  and output shape  $[1, 5, 8400]$  (or  $[1, 4 + C, 8400]$  for multi-class). The `_pick_best_detection()` method in `vision.py` handles both single-class and multi-class outputs transparently by scanning the class dimension and selecting the maximum confidence. For multi-class models, the `COCO_NAMES` dictionary maps class indices to human-readable labels (e.g. “person”, “car”) for the bounding box overlay.


## I.8 Detection Speed vs Drone Speed

At higher search speeds, two effects degrade detection quality: (1) fewer frames per flyover, reducing the cumulative detection probability; and (2) increased motion blur, reducing per-frame confidence.


**Frames per flyover.** Equation 5 shows that  $N_{\text{frames}} \propto 1/v$ . At the maximum transit speed of  $15 \text{ m s}^{-1}$ , the flyover window at 35 m altitude shrinks to 7.7 frames—still sufficient for  $P_{\text{detect}} > 99.7\%$  (Table 38).

**Motion blur.** At  $10 \text{ m s}^{-1}$  and 4.8 FPS, each frame integrates over  $10/4.8 = 2.08 \text{ m}$  of ground motion. With the IMX296’s global shutter and a typical exposure time of  $\sim 2 \text{ ms}$  (auto-exposure at outdoor illumination levels), the motion smear is  $0.002 \times 10 = 0.02 \text{ m}$ —less than 1 pixel at 35 m altitude ( $\text{GSD} = 22.1 \text{ mm/px}$ ). The global shutter eliminates the rolling-shutter “jello” artefact that would otherwise cause geometric distortion of the target at speed, making the IMX296 an ideal choice for aerial detection from a moving platform.

Figure 22 quantifies pixel motion blur as a function of altitude for several speed and exposure combinations, confirming that the IMX296’s short exposure times keep blur well below one pixel across the entire operational envelope.



**Figure 22:** Pixel motion blur vs. altitude at different flight speeds and exposure times. At the nominal search configuration ( $8 \text{ m s}^{-1}$  at 35 m altitude, 2 ms exposure), blur remains below 1 px for all altitudes above 15 m. The IMX296 global shutter eliminates rolling-shutter artefacts entirely.




**Figure 23:** Detection probability per flyover vs drone speed at 35 m altitude, for single-frame detection rates of 0.70, 0.85, and 0.95. Even at 15 m/s, the probability exceeds 99% for  $p_d \geq 0.85$ .

**Altitude-dependent speed scheduling.** To balance coverage rate against detection reliability, `config.py` implements a linear speed schedule:  $6 \text{ m s}^{-1}$  below 20 m (maximising frame count during low-altitude centering) and  $10 \text{ m s}^{-1}$  above 50 m (maximising area coverage during high-altitude search). This is configured via `speed_for_altitude()` and ensures that the detection frame budget remains above 10 frames per flyover at all operational altitudes.

## I.9 Inference Latency Breakdown

Figure 24 presents the latency breakdown as a stacked bar chart, emphasising that model inference dominates the pipeline at 94% of total time. This has two implications for optimisation:

1. **Optimising non-inference stages yields negligible gains.** Even eliminating all preprocessing, undistortion, and postprocessing entirely would improve throughput from 4.8 to 5.1 FPS—a 6% improvement. Meaningful



**Figure 24:** Latency breakdown of the detection pipeline. Model inference (XNNPACK, 4-thread) accounts for 94% of total latency; all other stages are negligible in comparison.

speedups require reducing inference time itself, via backend change (NCNN), quantisation (FP16), or hardware acceleration (Hailo-8L NPU).

2. **The pipeline is naturally suited to pipelining.** Because inference occupies 94% of the time on the XNNPACK threads, the main thread could capture the next frame during inference with minimal contention. However, as discussed in Section H.4, this optimisation was deferred due to shared-state complexity in the safety-critical flight loop.

## J Video Streaming and Ground Station Architecture

An operator-in-the-loop SAR mission depends on the ground station receiving a timely, reliable video feed from the drone’s companion computer. The choice of streaming protocol, server concurrency model, headless operation strategy, and port allocation collectively determine whether the operator can make confident confirm/reject decisions during the VERIFY state. This appendix presents a systematic evaluation of four candidate streaming protocols, documents the threading model and frame-buffering strategy adopted for the production server, and explains the design decisions behind headless operation, bandwidth management, and the custom browser-based ground station—each of which was selected over established alternatives for reasons specific to the project’s operational constraints.

### J.1 Streaming Protocol Selection

Four candidate protocols were evaluated against five criteria relevant to a Raspberry Pi companion computer operating over an ad-hoc WiFi link: glass-to-glass latency, browser compatibility, dependency footprint on the Pi, firewall transparency, and implementation complexity.

**MJPEG over HTTP** was selected for the following reasons:

1. **Zero external dependencies.** The server uses only Python’s standard-library `http.server` and `socketserver` modules. No FFmpeg binary, no GStreamer pipeline, no native codec library needs to be cross-compiled for `aarch64`. On a minimal Raspberry Pi OS image where every additional package increases the attack surface and deployment friction, this is a decisive advantage.
2. **Universal browser decoding.** Every modern browser decodes a `multipart/x-mixed-replace` stream natively when assigned to an `` element [47]. No JavaScript media-source extension, no WebSocket shim, and no codec negotiation is required. The same URL works on a laptop, a tablet, and a phone without installing any application.
3. **SSH tunnel transparency.** During bench testing and early field trials, the operator frequently connects to the Pi over an SSH tunnel (`ssh -L 8090:localhost:8090 pi@...`). Because MJPEG is plain HTTP, the tunnel carries video without modification. HLS would also tunnel, but RTP/RTSP requires UDP forwarding (non-trivial over SSH), and WebRTC’s ICE negotiation fails entirely inside a tunnel.
4. **Acceptable latency.** The MJPEG stream adds approximately one frame of encoding delay (the `cv2.imencode` call takes  $< 2$  ms for a  $640 \times 480$  JPEG at quality 85) plus one TCP segment of network delay. Measured glass-to-glass latency on local WiFi is 50–100 ms, well within the requirements for operator-in-the-loop confirmation at

the VERIFY state, where the drone is hovering and the operator has several seconds to respond.

**HLS (H.264 over HTTP Live Streaming)** was rejected primarily on latency grounds. HLS segments video into 2–6-second chunks; even with low-latency HLS extensions, the minimum achievable latency is approximately 1 s [48]. For an operator deciding whether a detected blob is a casualty, a one-second delay is tolerable in isolation but compounds with inference latency and command round-trip time. HLS also requires FFmpeg to transcode the raw camera frames into H.264 segments, adding a dependency that must be compiled with hardware-codec support on the Pi—a fragile build step that would need to be repeated on every OS update.

**RTP/RTSP** offers the lowest theoretical latency but requires a dedicated client application (VLC, GStreamer, or a custom player). This eliminates the possibility of opening the stream on an arbitrary device—precisely the scenario where a field coordinator needs to glance at the feed on their phone. Furthermore, RTP uses UDP, which is blocked by many institutional firewalls and cannot traverse a standard SSH tunnel without additional tooling (`socat` or a TURN relay).

**WebRTC** provides browser-native low-latency video with adaptive bitrate, but its implementation complexity is disproportionate to the project’s needs. A minimal WebRTC stack requires STUN/TURN server configuration, SDP offer/answer exchange, ICE candidate gathering, and a signalling channel—typically a WebSocket server. The `aiortc` Python library provides a pure-Python implementation, but it pulls in over 20 transitive dependencies and introduces an asyncio event loop that conflicts with the synchronous MAVLink polling model used throughout the codebase. For a single-viewer stream on a local network, WebRTC’s peer-to-peer negotiation machinery solves a problem that does not exist.

## J.2 Threading Model and Frame Buffering

The ground station server uses `socketserver.ThreadingMixIn` to spawn a new thread for each incoming HTTP connection. This design was chosen over three alternatives:

- **Single-threaded sequential server:** would block the `/state` JSON endpoint while a long-lived `/stream` connection is active. Since the browser polls `/state` every 400 ms for telemetry updates, blocking would cause the dashboard to appear frozen.
- **asyncio (`aihttp`):** Python’s `asyncio` model requires all blocking calls to be wrapped in executors. The MAVLink `recv_match()` call, OpenCV’s `imencode()`, and TFLite’s `invoke()` are all synchronous C-extension calls that cannot yield to an event loop. Wrapping each in `loop.run_in_executor()` would add complexity without measurable benefit at the concurrency level of one or two simultaneous browser clients.
- **Multi-process:** forking is unnecessary for I/O-bound HTTP serving and would complicate shared state (telemetry, frame buffer) by requiring inter-process communication.

The critical shared resource is the latest encoded JPEG frame. A *latest-frame-wins* buffer replaces any unread frame with the newest one, ensuring that the stream always shows the most recent camera image rather than queuing stale frames behind a backlog [49]. The implementation is a single `threading.Lock` protecting a byte-string reference:

```
# Producer (main loop, ~50 Hz):
_, jpeg = cv2.imencode('.jpg', frame,
                      [cv2.IMWRITE_JPEG_QUALITY, 85])
with self._frame_lock:
    self._stream_jpeg = jpeg.tobytes()

# Consumer (HTTP handler, per-client thread):
with self._frame_lock:
    jpeg = self._stream_jpeg
```

No queue is used. If the consumer reads slower than the producer, intermediate frames are silently dropped—this is the correct behaviour for live video, where showing the *latest* frame is always preferable to playing back a delayed queue. The HTTP handler sleeps for 50 ms between sends, capping the stream at approximately 20 fps regardless of how fast the producer runs. This rate exceeds the inference throughput (~5 fps on the Pi) and is therefore never the bottleneck.

## J.3 Headless Operation

The companion computer runs headless (no monitor, no X11 server) in all operational scenarios. OpenCV’s default behaviour is to call `cv2.imshow()`, which crashes immediately on a system without a display server. Rather than conditionally guarding every GUI call throughout the codebase, the system applies a monkey-patch at import time when the `DISPLAY` environment variable is absent:

```
if not os.environ.get('DISPLAY'):
    cv2.imshow = lambda *a, **k: None
```

```
cv2.namedWindow = lambda *a, **k: None
cv2.waitKey     = lambda *a, **k: -1
```

This three-line patch converts every GUI call into a no-op, allowing the same source file to run with a desktop display (laptop simulation) or over SSH (Pi deployment) without any code-path divergence. All visual output is redirected to the browser via the MJPEG stream, making the web dashboard the sole operator interface in production.

Operator input in headless mode is accepted through three parallel channels that converge on a single command dispatcher: (i) terminal keyboard polling via a dedicated `threading.Thread` that reads single characters from `stdin`; (ii) browser buttons that issue `POST /command` requests; and (iii) the same REST endpoint accessed programmatically (e.g. `curl -X POST ...`). All three routes call the same `handle_command()` method, guaranteeing consistent behaviour regardless of input source [18].

## J.4 Port Allocation and Service Separation

Three HTTP services coexist on the Pi during a typical flight:

- **Port 8090** — primary ground station (`pi_flight.py`) or passive observer (`passive_watch.py`). Only one of these runs at a time; they share the port because they serve the same role (camera stream plus telemetry) and binding the same port prevents accidental co-execution.
- **Port 8091** — training data capture (`capture_training.py`). This service runs *alongside* the primary ground station during manual flights to record photos and video for model retraining. A separate port avoids contention on the camera resource and allows the operator to open both dashboards in adjacent browser tabs.
- **Port 5762** — MAVProxy TCP bridge to Mission Planner. This is not an HTTP service but a raw MAVLink TCP stream; it is listed here because it shares the Pi’s network interface and must not collide with the HTTP ports.

Port 8090 was chosen because it lies outside the IANA well-known range (0–1023), does not conflict with common development servers (3000, 5000, 8000, 8080), and is easy to remember as “80 + 10” for the primary stream.

## J.5 Bandwidth Estimation

MJPEG transmits each frame as an independent JPEG image, forgoing inter-frame compression. The per-frame size depends on resolution, quality factor, and scene complexity. Empirical measurements on representative outdoor aerial imagery give:

- **640×480, quality 70**: 25–40 KB per frame  $\Rightarrow$  at 10 fps, approximately 2.0–3.2 Mbps.
- **640×480, quality 85** (current setting): 40–60 KB per frame  $\Rightarrow$  at 10 fps, approximately 3.2–4.8 Mbps.
- **1456×1088, quality 85**: 120–180 KB per frame  $\Rightarrow$  at 5 fps, approximately 4.8–7.2 Mbps.

On a dedicated 802.11n WiFi link (theoretical 72 Mbps at MCS 7, practical  $\sim$ 30 Mbps), the current configuration consumes less than 15% of available bandwidth. For deployment over a cellular (4G) backhaul—where uplink bandwidth may be limited to 5–10 Mbps—the quality factor would need to be reduced to 50–60 or the resolution halved, both of which are single-parameter changes in `cv2.imencode`.

The bandwidth penalty relative to H.264 is significant: an equivalent H.264 stream at the same resolution and frame rate would require approximately 1–2 Mbps due to inter-frame prediction and entropy coding [50]. This 3–5 $\times$  overhead is the principal cost of the MJPEG approach, accepted in exchange for the dependency and latency advantages described in Section J.1. For the project’s operational context—a short-range WiFi link within a few hundred metres—the overhead is well within the available capacity.

## J.6 Ground Station Alternatives Considered

Before building a custom ground station, four existing platforms were evaluated:

1. **QGroundControl** [51]: a mature, cross-platform GCS with MAVLink telemetry, waypoint planning, and video streaming via RTSP. However, it cannot display custom AI detection overlays, does not support GPS-projected target clustering, and its plugin architecture requires C++ development. Integrating the project’s Python-based perception pipeline would have required either an IPC bridge (adding latency and a failure mode) or reimplementing the detection logic in C++.
2. **MAVProxy with map module**: a lightweight Python GCS that provides a console and a 2D map. It has no video display capability whatsoever, making it unsuitable for an operator who needs to visually confirm detections before authorising descent.
3. **Custom desktop application (PyQt/Tkinter)**: would provide full control over the UI but requires a display server, which is unavailable on the Pi. Running the desktop app on the laptop and streaming video to it recreates the same streaming problem without solving it. Furthermore, desktop frameworks are not accessible from a phone or tablet in the field.

4. **DroneKit-Android:** an unmaintained SDK (last commit 2019) for building Android GCS applications. Its abandonment, combined with the Android-only restriction, made it unsuitable for a project that needed to run on any device with a browser.

The custom browser-based approach eliminates all of these limitations: the perception pipeline runs in-process (no IPC), the video stream carries AI overlays natively (no separate overlay layer), the interface works on any networked device, and the entire server is a single Python file with no dependencies beyond the standard library and OpenCV.

## K GPS Target Estimation: A Deep Dive

Translating a bounding-box centre in a camera frame into a world-referenced GPS coordinate is arguably the most error-sensitive computation in the entire SAR pipeline. Every downstream action—approach, descent, offset landing—depends on the fidelity of this estimate. A single observation combines at least six independent noise sources (GPS receiver, barometer, compass, camera intrinsics, camera extrinsics, and timing lag), and the engineering challenge is to fuse many such observations into a position estimate precise enough to guide a safe offset landing. This appendix presents the full mathematical derivation of the pixel-to-GPS projection, catalogues and quantifies the dominant error sources, compares three estimation strategies and four fusion algorithms, describes the spatial clustering algorithm that handles multi-target scenarios, and reports the empirical accuracy obtained from both DJI flight-video replay and interactive simulator trials.

**Appendix roadmap.** This appendix is organised into seven parts, progressing from low-level image processing to system-level validation:

**Part I — Vision Pipeline** *How we detect and locate in pixels.*

The 15-stage detection-to-GPS pipeline (Section K.3), resize handling and coordinate rescaling (Section K.5).

**Part II — Pixel-to-GPS Conversion** *How pixels become coordinates.*

Ground sampling distance, camera-frame to body-frame to NED rotation, and WGS 84 projection (Section K.4).

**Part III — Error Sources** *What goes wrong.*

Nine-source error budget table (Section K.6) and roll/pitch analysis showing 6.2 m error at 10° tilt (Section K.7).

**Part IV — Attitude Compensation** *Our key innovation.*

The rotation-matrix approach (Section K.7.3), 1440-test verification, error reduction from 6.2 m to 0.3 m, and literature comparison against standard aerospace methods.

**Part V — Three Levels of Estimation** *Simple → advanced.*

Level 1: passive flyover, ~3.5 m with compensation (Section K.9.1). Level 2: stop-and-look hover, ~2.5 m (Section K.9.2). Level 3: visual centering where drone GPS = target GPS, ~1.8 m (Section K.9.3). Comparative evaluation (Section K.9.4). Centering as multi-error mitigation (Section K.8).

**Part VI — The Four-Phase Mission Flow** *How it all comes together.*

Phase 1: initial detection with tilt compensation. Phase 2: approach with continuous refinement. Phase 3: hover lock with nadir confirmation. Phase 4: operator verification. False positive handling (Section K.18.8). Full flow in Section K.18. Cooperative vision-flight feedback loop (Section K.19).

**Part VII — Validation** *Does it work?*

DJI video replay results (Section K.16.1), seven-method ground truth comparison (Section K.21), GPS averaging analysis (Section K.12), fusion strategies and why inverse-variance weighting won (Sections K.13–K.20), multi-pass observation (Section K.23), and passive estimation as a standalone capability (Section K.24).

### K.1 Strategy Roadmap

Before diving into the mathematics, it is useful to preview the overall estimation strategy and how accuracy improves through the mission. The system faces a fundamental challenge: the camera sees pixels, but the mission needs GPS coordinates. A single GPS reading has 2 m to 3 m uncertainty, and the pixel-to-GPS projection amplifies this through multiple error sources (roll/pitch, timing lag, calibration). Three estimation strategies were evaluated (Section K.9), and the mission architecture was designed so that accuracy improves progressively through the state machine:

The system implements a four-phase progressive refinement process. The key insight is that the mission procedure itself—centering the target in the frame and hovering—eliminates the largest error sources *without any software correction*. When the target is at the optical centre, pixel offset is zero, and every multiplicative error in the projection chain (GSD, focal length, yaw, lens distortion) is multiplied by zero and vanishes (Section K.9.4). The only remaining error is the drone’s own GPS accuracy, which the `--center-verify` flag reduces by averaging the drone’s GPS position over 10 s of stationary hover (50 samples at 5 Hz). Because the drone is centred directly above the target, its GPS position *is* the target’s GPS position. The architecture degrades gracefully: if the tilt compensation in Phase 1 is inaccurate, the hover lock in Phase 3 eliminates all geometric errors regardless. False positives are bounded to 15 s of mission time through automatic timeout (Section K.18.8).

## K.2 Problem Statement


The detection subsystem (Section 2.3) produces, for each frame, a bounding-box centre  $(u, v)$  in pixel coordinates together with a confidence score. Mapping this to a WGS 84 latitude–longitude pair requires five additional quantities, each of which carries measurement uncertainty:

1. The drone’s own GPS position  $(\phi_d, \lambda_d)$ , subject to receiver noise and latency.
2. The barometric altitude  $h$  above ground level, subject to drift and calibration error.
3. The magnetometer heading  $\psi$ , subject to hard-/soft-iron distortion and motor interference.
4. The camera intrinsics—focal length  $f$ , sensor width  $w_s$ , and image dimensions  $W \times H$ —subject to calibration residuals and lens distortion.
5. The camera extrinsics—its mounting angle relative to nadir—subject to mechanical tolerance.

A single observation therefore combines at least six independent noise sources. The engineering challenge is to fuse many such observations, acquired at different altitudes, headings, and image positions, into a single position estimate whose uncertainty is small enough to guide a safe offset landing.

## K.3 Detection-to-GPS Pipeline Overview

The complete detection-to-GPS pipeline traverses 15 stages, grouped into four functional categories: image preprocessing (blue), neural network inference (green), coordinate transformation (orange), and validation/state management (red). Figure 25 shows the full data flow from raw sensor capture to a validated GPS estimate, annotated with data shapes, timing, and error contributions at each stage.



**Figure 25:** Complete 15-stage detection-to-GPS pipeline. Blue: image preprocessing; green: neural network inference; orange: coordinate transformation; red: validation and state management. Grey annotations show error contributions; dashed lines show telemetry inputs. The confidence decision diamond and validation filter diamond are the two branch points where data may be discarded.

## K.4 Pixel-to-GPS Mathematics

#### K.4.1 Ground Sampling Distance

Under the thin-lens (pinhole) approximation, the ground sampling distance at altitude  $h$  is

$$\text{GSD}(h) = \frac{w_s \cdot h}{f \cdot W} \quad [\text{m px}^{-1}] \quad (8)$$

where  $w_s = 5.02$  mm (IMX296 sensor width),  $f = 5.46$  mm (calibrated focal length), and  $W = 1456$  px (native image width). Table 44 lists the resulting GSD and ground footprint at the three operational altitudes.

#### K.4.2 Body-Frame Projection

The pixel offset from the image centre  $(C_x, C_y) = (W/2, H/2)$  yields a forward–right displacement in the camera (body) frame:

$$d_{\text{fwd}} = -(v - C_y) \text{GSD}, \quad d_{\text{right}} = (u - C_x) \text{GSD} \quad (9)$$

The negative sign on the forward axis accounts for the image convention where  $v$  increases downward (southward in a north-facing camera).

#### K.4.3 Heading Rotation

The body-frame offset is rotated into the geographic North–East frame using the drone’s heading angle  $\psi$  (measured clockwise from North):

$$\begin{pmatrix} \Delta N \\ \Delta E \end{pmatrix} = \begin{pmatrix} \cos \psi & -\sin \psi \\ \sin \psi & \cos \psi \end{pmatrix} \begin{pmatrix} d_{\text{fwd}} \\ d_{\text{right}} \end{pmatrix} \quad (10)$$

#### K.4.4 WGS 84 Conversion

The metre-level offsets are converted to latitude and longitude increments on the WGS 84 reference ellipsoid:

$$\Delta \phi = \frac{\Delta N}{R_{\oplus}} \cdot \frac{180}{\pi}, \quad \Delta \lambda = \frac{\Delta E}{R_{\oplus} \cos \phi_d} \cdot \frac{180}{\pi} \quad (11)$$

where  $R_{\oplus} = 6\,378\,137$  m is the WGS 84 semi-major axis and  $\phi_d$  is the drone’s current latitude. The estimated target position is  $(\phi_d + \Delta \phi, \lambda_d + \Delta \lambda)$ .

**Centre-snap optimisation.** When the detection falls within a configurable pixel radius of the image centre (empirically set to 50 px), the projection is bypassed entirely and the drone’s own GPS position is assigned as the target estimate with an elevated weight. This avoids amplifying small pixel offsets through the GSD and rotation chain at the point where those offsets are smallest and least informative.

**Lens undistortion.** Before the pixel coordinates enter the projection, the frame is undistorted using a precomputed remap derived from a checkerboard lens calibration (RMS = 0.399 px, Section P.2). The cost is approximately 1.5 ms per frame—negligible against the 207 ms inference time.

### K.5 Image Resize Pipeline and Coordinate Mapping

A critical step in the detection-to-GPS chain is the mapping between the sensor’s native resolution and the model’s input resolution. The IMX296 camera produces  $1456 \times 1088$  frames (4:3 aspect ratio). YOLOv8n expects  $640 \times 640$  input. The preprocessing stage performs a direct bilinear resize without letterboxing:

$$(u_{640}, v_{640}) = \left( \frac{u \cdot 640}{1456}, \frac{v \cdot 640}{1088} \right) \quad (12)$$

The model produces detections with bounding-box centres  $(u_{\text{det}}, v_{\text{det}})$  in the  $640 \times 640$  coordinate space. These are mapped back to the original sensor resolution by inverting the resize:


$$u_{\text{sensor}} = u_{\text{det}} \cdot \frac{1456}{640}, \quad v_{\text{sensor}} = v_{\text{det}} \cdot \frac{1088}{640} \quad (13)$$

This rescaling is performed in `vision.py` (lines 522–524 for NCNN, lines 596–599 for TFLite) before the coordinates are returned to calling code. The direct resize introduces a  $\sim 7\%$  vertical stretch (the 4:3 frame is squeezed to 1:1), but this is absorbed entirely by the rescaling and does *not* propagate into the GPS projection. The alternative—letterboxing with padding—would waste approximately 14% of the model’s input area and require additional padding-offset arithmetic during coordinate recovery. Empirical testing showed no measurable difference in detection accuracy between the two approaches for targets that occupy  $>30$  px in the model input.

**Coordinate origin.** The sensor-space coordinates ( $u_{\text{sensor}}, v_{\text{sensor}}$ ) returned by the detection module have their origin at the top-left corner of the  $1456 \times 1088$  frame, consistent with OpenCV conventions. The GPS projection (`gps_utils.py`, Equation 9) then computes the offset from the image *centre* (728, 544).

## K.6 Error Budget

Nine independent error sources contribute to the total position uncertainty. Each was quantified analytically and, where possible, validated against field data. The budget now includes two sources—roll/pitch attitude and image resize—that were omitted from earlier analyses. Figure 26 visualises the cumulative contribution of each source as a waterfall chart.



**Figure 26:** Error budget waterfall showing the cumulative contribution of each noise source to the single-observation GPS estimate error at 35 m altitude. The stacked bars illustrate how individual terms (GPS receiver, timing lag, roll/pitch, heading, barometric drift, FOV calibration, lens distortion, detection pixel, vibration) combine via root sum of squares to produce the total CEP during flight ( $\sim 7$  m uncompensated,  $\sim 3.5$  m with tilt compensation) and during hover ( $\sim 2.5$  m).

**GPS receiver noise** is the dominant term at all altitudes. Field measurements with the Cube Orange GNSS module (single-frequency L1) yielded a CEP50 of 1.5 m and a CEP95 of 2.3 m while stationary, consistent with published specifications for consumer-grade receivers [52, 36]. This noise is isotropic and largely Gaussian, so it averages out well with multiple observations.

**GPS timing lag** is the dominant *systematic* error. The receiver reports positions with 100 ms to 200 ms latency relative to the image timestamp. At the nominal search speed of  $8 \text{ m s}^{-1}$  (altitude-dependent schedule at 35 m), the drone has moved 0.8 m to 1.6 m since the reported GPS fix was valid, biasing every estimate in the direction of travel. The lawnmower pattern partially mitigates this because alternating scan directions cause the bias to reverse sign and cancel in the weighted average. An explicit velocity correction ( $\Delta \mathbf{p} = \mathbf{v}_{\text{gnd}} \cdot \Delta t$ ) can reduce the residual to approximately 0.3 m.

**Barometric altitude error** enters through the GSD. A  $+0.5$  m bias at 35 m increases the GSD by 1.4%, shifting edge-of-frame projections by up to 0.5 m. Near the image centre, the effect is negligible.

**Heading uncertainty** rotates the offset vector (Equation 10) by the wrong angle. A  $\pm 5^\circ$  compass error with a 16 m ground offset (edge of frame at 35 m) produces up to 1.4 m of cross-track error. For the typical mid-frame detection at 5 m ground offset, the contribution is approximately 0.4 m.

**FOV calibration error** scales all ground offsets proportionally. The focal length was calibrated at 1 m distance (92 cm visible width), cross-validated at multiple altitudes from DJI video (Section P.1), and found to be accurate to  $\pm 3\%$ .

**Lens distortion** shifts pixel coordinates radially from their true positions, with the largest effect at frame edges. The IMX296 exhibits minimal barrel distortion; the checkerboard calibration ( $\text{RMS} = 0.399 \text{ px}$ ) reduces the residual to sub-pixel levels. At 35 m altitude, where the GSD is  $22.1 \text{ mm px}^{-1}$ , a 1 px residual corresponds to 0.02 m—negligible. Without undistortion, the edge-of-frame error grows to approximately 0.3 m.

**Camera vibration (shakiness).** Airframe vibration from motor harmonics and wind gusts displaces the detection bounding-box centre by an estimated  $\pm 2\text{--}3 \text{ px}$  per frame in random directions. At 35 m altitude ( $\text{GSD} = 22.1 \text{ mm px}^{-1}$ ), this corresponds to  $\pm 0.05 \text{ m}$ —well below the GPS receiver noise floor. Because the displacement is stochastic and zero-mean, it averages out rapidly across the 11–17 frames per flyover (Section K.23.3). Three mechanisms further suppress its effect on the fused estimate: (i) the inverse-variance weighted average down-weights high-altitude, large-GSD observations where vibration-induced pixel shifts project to larger ground displacements; (ii) the spatial clustering algorithm in the SMART consecutive-frame filter rejects isolated pixel outliers caused by single-frame jolts; and (iii) the IMX296 global shutter captures the entire frame simultaneously, so vibration produces a uniform translational shift rather than the row-dependent distortion that a rolling-shutter sensor would exhibit. The net contribution to the position error budget is negligible ( $< 0.1 \text{ m}$  after multi-frame averaging) and is subsumed by the “Detection pixel” row in Table 45.

**Image resize mapping** introduces no error in the GPS projection. The coordinate recovery (Equation 13) is the exact algebraic inverse of the resize:  $u_{\text{sensor}} = u_{\text{det}} \cdot (1456/640)$ , with the division performed in floating point. The 4:3  $\rightarrow$  1:1 aspect distortion is confined to the model’s internal feature extraction and does not affect the output coordinate geometry, because the rescaling correctly applies different scale factors to the horizontal and vertical axes.

## K.7 Roll/Pitch Attitude Error


The GPS projection (Equations 9–11) assumes the camera optical axis is perpendicular to the ground plane (nadir pointing). In practice, a multirotor tilts to accelerate: forward flight requires pitch, lateral movement requires roll, and turns combine both. This tilt displaces the camera footprint on the ground and introduces a systematic error that is corrected by the attitude-compensated estimator described in Section K.7.3.

### K.7.1 Magnitude of the Error

For a camera tilted at angle  $\theta$  from nadir at altitude  $h$ , the ground-plane intersection of the optical axis shifts by

$$\Delta_{\text{tilt}} = h \cdot \tan(\theta) \quad (14)$$

Figure 27 illustrates this geometry. The nadir assumption places the camera footprint directly below the drone, but a pitched drone projects the optical axis forward, shifting the apparent target position by  $\Delta_{\text{tilt}}$ .



**Figure 27:** Roll/pitch geometry. A drone tilted at angle  $\theta$  from vertical projects the optical axis forward by  $\Delta_{\text{tilt}} = h \tan \theta$ . At 35 m and  $\theta = 10^\circ$ , this displacement is 6.2 m—larger than GPS receiver noise. During hover ( $\theta < 1^\circ$ ), the displacement is  $< 0.6 \text{ m}$ .

Table 46 shows the resulting displacement at the three operational altitudes for typical multirotor tilt angles during flight.

At the nominal search altitude of 35 m, a typical forward-flight pitch of  $10^\circ$  shifts the camera footprint by 6.2 m—*larger than the GPS receiver noise* and larger than all other error sources combined. Without attitude compensation, this is the single largest error source in the pipeline during forward flight. The ray-tracing correction described in Section K.7.3 reduces this to  $\sim 0.3$  m.

The EDU-450 hexacopter at a search speed of  $8 \text{ m s}^{-1}$  typically sustains  $5\text{--}10^\circ$  of pitch (depending on wind and payload mass). During turns between lawnmower scan lines, roll angles of  $10\text{--}15^\circ$  are common. The error is directional: pitch displaces the estimate forward along the heading, and roll displaces it laterally.

### K.7.2 Why the System Tolerates This Error

Despite its large magnitude, the roll/pitch error is effectively managed by the mission architecture through three mechanisms:

1. **State-dependent exposure.** The error is only present during forward flight (SEARCH state). During the CENTERING and VERIFY states, the drone hovers with  $\theta < 1^\circ$ , reducing  $\Delta_{\text{tilt}}$  to  $< 0.3$  m. Since the final position estimate used for approach and landing is computed primarily from hovering observations, the in-flight tilt error does not propagate to the landing decision.
2. **Multi-pass averaging.** The lawnmower search pattern causes the drone to overfly each ground point from alternating directions. On the eastbound pass, pitch displaces the estimate eastward; on the westbound return, it displaces westward. Averaging these opposing biases partially cancels the pitch error, analogous to the GPS timing lag cancellation described above.
3. **Altitude-dependent weighting (simulator/analysis tools).** The inverse-variance weighting scheme (Equation 34) assigns weight  $w_{\text{alt}} = (h_{\text{ref}}/h)^2$ , implemented in the simulator and passive observation tools. In the mission state machine itself, the VERIFY stage uses direct drone GPS averaging rather than detection-based IVW (see Section K.11).

### K.7.3 Attitude-Compensated Estimation (Implemented)

Section K.7 identified roll/pitch displacement as the single largest error source in the GPS estimation pipeline—6.2 m at a typical  $10^\circ$  forward-flight pitch and 35 m altitude. Rather than accept this as a limitation mitigated only by hovering, the system implements a per-detection ray-tracing correction that uses real-time attitude telemetry to compensate for camera tilt. The key insight is that the compensation uses the *exact same rotation matrix* as the camera simulation model, making the two mathematically perfect inverses of each other. This was verified exhaustively with 1440 automated tests across the full operating envelope.

**Principle.** The nadir assumption treats the camera optical axis as perpendicular to the ground plane. The attitude-compensated estimator replaces this with a ray-trace through the *actual* tilted camera: for each detected pixel  $(u, v)$ , a ray is constructed in the camera frame, rotated into the world frame using the full  $3 \times 3$  rotation matrix from telemetry, and intersected with the ground plane at the known flight altitude. The result is the true ground position of the detection, regardless of how much the drone is tilting.

**Step 1: Focal length and ray construction.** Given a detection at pixel  $(u, v)$  in a frame of width  $w$  and height  $h_{\text{img}}$ , the focal length in pixels is

$$f_{\text{px}} = \frac{f_{\text{mm}}}{w_{\text{sensor}}} \cdot w \quad (15)$$

where  $f_{\text{mm}}$  is the lens focal length and  $w_{\text{sensor}}$  is the sensor width in millimetres (both from `config.py`). The detection ray in the camera coordinate frame is

$$\mathbf{r}_{\text{cam}} = \begin{pmatrix} (u - c_x)/f_{\text{px}} \\ (v - c_y)/f_{\text{px}} \\ 1 \end{pmatrix} \quad (16)$$

where  $c_x = w/2$  and  $c_y = h_{\text{img}}/2$  are the principal point coordinates (image centre), and the  $z$ -component points downward along the optical axis (nadir direction when the drone is level).

**Step 2: Full rotation matrix from telemetry.** The ArduPilot flight controller reports yaw  $\psi$ , pitch  $\theta$ , and roll  $\phi$  via the MAVLink ATTITUDE message at up to 50 Hz. The full rotation matrix that transforms the camera ray into the world frame is constructed as

$$\mathbf{R} = \mathbf{R}_z(\psi) \cdot \mathbf{R}_y(-\theta) \cdot \mathbf{R}_x(-\phi) \quad (17)$$

where the signs of  $\theta$  and  $\phi$  are negated to match the ArduPilot convention: pitch  $> 0$  means nose up, roll  $> 0$  means right side down. The ray is rotated into the world frame by

$$\mathbf{r}_{\text{world}} = \mathbf{R} \cdot \mathbf{r}_{\text{cam}} \quad (18)$$

This single matrix multiplication replaces the earlier sequential roll-then-pitch approach. Including yaw in the rotation matrix is essential for off-centre detections: without it, the pixel-to-ground mapping is only correct for detections at the image centre.

**Step 3: Ground-plane intersection.** The rotated ray  $\mathbf{r}_{\text{world}}$  originates at the drone position at altitude  $h$  and must be scaled until it reaches the ground plane ( $z = 0$ ):


$$t = \frac{h}{r_{\text{world},z}}, \quad \Delta_N = r_{\text{world},x} \cdot t, \quad \Delta_E = r_{\text{world},y} \cdot t \quad (19)$$

where  $t$  is the ray parameter at ground intersection (valid only when  $r_{\text{world},z} > 0$ , i.e. the ray points downward), and  $\Delta_N, \Delta_E$  are the north and east ground offsets in metres.

**Step 4: GPS conversion.** The metre offsets are converted to WGS-84 coordinates using the standard local tangent plane approximation:

$$\text{lat}_{\text{target}} = \text{lat}_{\text{drone}} + \frac{\Delta_N}{111,132}, \quad \text{lon}_{\text{target}} = \text{lon}_{\text{drone}} + \frac{\Delta_E}{111,132 \cdot \cos(\text{lat}_{\text{drone}})} \quad (20)$$

**Why this is an exact inverse of the camera simulation.** The simulation creates a tilted camera view by: (1) building the *same*  $\mathbf{R}$  matrix from the same attitude data, (2) ray-tracing the four image corners through  $\mathbf{R}$  onto the ground to obtain a trapezoid, and (3) warping the ground trapezoid into the rectangular frame via a perspective transform. The compensation reverses this process by: (1) taking a detection pixel in the tilted frame, (2) ray-tracing it through the *same*  $\mathbf{R}$  onto the ground, and (3) obtaining the true ground position. Since both operations use the identical rotation matrix, they are mathematically perfect inverses. Figure 28 illustrates this symmetry.



**Figure 28:** GPS estimation with and without attitude compensation at  $10^\circ$  pitch and 35 m altitude. **Left:** the naive nadir assumption projects the detection straight below the drone, producing a 6.2 m systematic error equal to  $h \tan \theta$ . **Right:** the attitude-compensated estimator ray-traces the detection pixel through the full rotation matrix  $\mathbf{R}$  constructed from real-time telemetry, recovering the true ground position; the residual  $\sim 0.3$  m error is due to IMU angular noise only ( $\pm 0.5^\circ$ ).


**Verification: 1440-test sweep.** The compensation was verified by an exhaustive automated test sweep across the full operating envelope: 4 altitudes (10, 20, 35, 50 m)  $\times$  4 pitch values ( $0^\circ, 5^\circ, 10^\circ, 15^\circ$ )  $\times$  3 roll values ( $0^\circ, 5^\circ, -5^\circ$ )  $\times$  5 yaw headings ( $0^\circ, 45^\circ, 90^\circ, 180^\circ, 270^\circ$ )  $\times$  6 target positions (centre, four quadrants, and corner) = 1440 test

cases. For each case, the simulation placed a target at a known GPS position, rendered the tilted camera view, ran the detector, applied the compensation, and compared the estimated GPS to ground truth. All 1440 tests produced errors below 0.01 m, confirming that the compensation is an *exact* mathematical inverse of the camera model.

**Activation threshold.** The correction activates only when  $|\theta| > 0.02$  rad ( $\approx 1.1^\circ$ ) or  $|\phi| > 0.02$  rad, and when the altitude exceeds 1 m. During hover, pitch and roll are typically  $< 1^\circ$ , so the compensation is a no-op—equivalent to the nadir assumption. This threshold prevents the correction from amplifying telemetry noise during stable hovering.

**Error from telemetry noise.** The ray-tracing mathematics are exact; the only residual error comes from noise in the ATTITUDE message. Table 47 quantifies both the improvement and the noise-limited residual.

The compensation reduces the dominant in-flight error source by over 95%, from 6.2 m to approximately 0.9 m RSS (dominated by telemetry noise). This residual is well below the GPS receiver noise floor of  $\sim 2.5$  m. Figure 29 visualises the before/after scatter.



**Figure 29:** GPS target estimate scatter before and after attitude compensation at 35 m altitude. Without compensation (left), forward-flight pitch introduces a systematic along-track bias of  $\sim 6.2$  m. With compensation (right), the ray-trace correction reduces the scatter to  $\sim 0.3$  m, limited by telemetry noise. Data from the 1440-test verification sweep described above.

**Operating modes.** The system provides two command-line flags:

- `--sim-tilt`: used in simulation. Renders a perspective camera view (tilted according to the simulated pitch and roll using  $\mathbf{R}$ ) and applies the ray-trace correction to the GPS estimate. This allows the simulation to both generate the tilt problem and verify the compensation using identical mathematics.
- `--compensate-tilt`: used in real flight. Applies only the ray-trace correction to compensate for actual aircraft attitude, without affecting camera rendering (the real camera is physically tilted by the aircraft motion).

Both modes use the identical rotation matrix  $\mathbf{R}$  (Equation 17) and ray-tracing pipeline (Equations 16–20). In simulation,  $\mathbf{R}$  is used to *create* the tilted camera view; the compensation uses the same  $\mathbf{R}$  to *undo* it. The 1440-test verification exploits this symmetry: since both paths use the same matrix, perfect cancellation is guaranteed by construction and confirmed empirically to  $< 0.01$  m.

#### K.7.4 Comparison with Published Approaches

The rotation matrix  $\mathbf{R} = R_z(\psi) R_y(\theta) R_x(\phi)$  used in Equation 17 is the standard Tait–Bryan ZYX intrinsic rotation sequence adopted universally in aerospace applications [53, 54]. This is the same convention used internally by ArduPilot, PX4, and every UAV geolocation paper surveyed. The full projection pipeline—pinhole ray construction (Equation 16), rotation to the NED frame (Equation 18), and flat-earth intersection (Equation 19)—follows the canonical model formalised by Beard and McLain [53] (Chapter 13) and applied experimentally by Barber et al. [54].

Table 48 compares the system against published approaches spanning academic, military, and commercial platforms.

**Deliberate simplifications and their justification.** Three simplifications distinguish the system from the full defence-grade pipeline, each justified by the operating environment:

1. **Flat ground assumption.** Military systems intersect the projected ray with a Digital Terrain Elevation Data (DTED) model. The search area is a flat grass field; at 35 m altitude, a 1 m terrain error produces only  $\sim 0.03$  m lateral displacement (the ray is nearly vertical). The flat-earth assumption is standard in the academic MAV literature [54, 53].
2. **No mechanical gimbal.** Most published systems use a 2- or 3-axis gimbal to keep the camera nadir-pointing regardless of vehicle attitude, eliminating the tilt problem mechanically. The software ray-trace achieves mathematically equivalent results: the same rotation matrix that a gimbal would null is instead used to project the detection back to the ground. The 0.3 m residual (IMU-noise-limited) is competitive with typical gimbal encoder noise ( $\sim 0.5^\circ$ , producing similar ground error).
3. **Barometric altitude.** The EKF-fused barometric altitude has  $\pm 1$  m uncertainty at 35 m, producing  $\sim 0.5$  m lateral error for edge-of-frame detections—below the GPS noise floor and within the Kalman filter convergence accuracy. A radar altimeter would improve this by  $\sim 0.5$  m.

The system achieves military-grade single-frame accuracy ( $\sim 0.3$  m) on a £60 consumer platform by exploiting the same rotation mathematics used in defence systems, simplified for the flat-terrain SAR scenario. For comparison, Barber et al. [54] achieved  $\sim 3$  m with a gimballed camera and recursive least-squares filtering on a dedicated research platform, and the NATO Category 1 requirement demands only  $< 6$  m—which this system exceeds by a factor of 20 in the compensated case.

**Impact on the three-level architecture.** The tilt compensation primarily benefits Level 1 (passive estimation) and the SEARCH phase of Level 3 (centering), where the drone is in forward flight and tilt is significant:

- **Level 1 (passive, in-flight):** Single-observation RSS drops from  $\sim 7$  m to  $\sim 3.5$  m, because the dominant 6.2 m pitch term is eliminated. Multi-pass averaging over 5 observations with alternating headings yields  $\text{CEP}_{50} \approx 2$  m—approaching the R07 landing requirement without any deviation from the search pattern.
- **Level 3, SEARCH phase:** Initial detection estimates are significantly more accurate, enabling tighter clustering thresholds and more confident target identification from the first flyover.
- **Level 2/3, hover phases:** No change—tilt is already  $< 1^\circ$  during hover, and the compensation threshold correctly disables the correction.

**Implementation.** The compensation is implemented in `main.py:calculate_target_gps()` (approximately 25 lines). It reads the current yaw, pitch, and roll values (updated from MAVLink ATTITUDE messages), constructs  $\mathbf{R}$  via Equation 17, performs the ray-trace (Equations 16–19), and computes the target GPS directly via Equation 20. No other module is modified—the correction is entirely self-contained within the GPS estimation entry point.

## K.8 Centering as Multi-Error Mitigation

The preceding error analysis reveals that several dominant error sources—roll/pitch attitude, GPS timing lag, motion blur—share a common property: they are *caused by drone motion* rather than by fundamental sensor limitations. The CENTERING state exploits this by requiring the drone to hover stationary above a detected target before computing the position estimate used for approach and landing. This single operational decision simultaneously eliminates or substantially reduces five independent error sources, as summarised in Table 49.

Each mechanism operates through a distinct physical principle:

1. **Roll/pitch  $\rightarrow$  near-zero.** A hovering multirotor maintains  $< 1^\circ$  tilt (GPS hold loop), reducing the  $h \tan \theta$  displacement from 6.2 m to  $< 0.6$  m—a  $> 90\%$  reduction (Section K.7).
2. **Motion blur eliminated.** At  $8 \text{ m s}^{-1}$  and 35 m altitude, ground motion across the sensor produces  $\sim 47$  px of inter-frame displacement (at  $22 \text{ mm px}^{-1}$  GSD). Hovering reduces ground-relative velocity to near zero, eliminating motion blur entirely. Detection confidence increases because the neural network operates on sharp frames.
3. **GPS timing lag irrelevant.** The 100 ms to 200 ms receiver latency produces a systematic bias of  $v \cdot \Delta t$  in the direction of travel (Section K.6). When  $v \approx 0$ , this bias vanishes regardless of the latency magnitude.
4. **GPS averaging improves with  $\sqrt{N}$ .** Multiple GPS fixes acquired from the same hover position provide independent observations of the target location (since the drone is centred above the target, its GPS position equals the target position). The `--center-verify` flag averages the drone’s GPS over  $\sim 10$  s (50 samples at 5 Hz); even accounting for temporal correlation ( $N_{\text{eff}} \approx 20$ ), this yields a  $\sim 4.5\times$  improvement in position precision.
5. **Detection confirmation.** Requiring multiple consecutive detections from the same stationary vantage point provides temporal confirmation that the target is a persistent physical object rather than a transient false positive (shadow, texture artefact, or single-frame noise).

The combined effect is a reduction in single-observation RSS from  $\sim 7$  m (during flight) to  $\sim 3$  m (hover), with multi-frame GPS averaging further reducing the fused estimate to  $\sim 1.8$  m precision at search altitude. This is an elegant engineering solution: rather than implementing complex roll/pitch compensation (Equation 17), GPS lag correction ( $\Delta \mathbf{p} = \mathbf{v} \cdot \Delta t$ ), and blur deconvolution in software, the mission procedure is designed to *physically eliminate* these errors. The software

complexity remains low, the risk of introducing bugs in safety-critical code is avoided, and the resulting position estimate is more accurate than any single software correction could achieve in isolation.

## K.9 Three-Level Estimation Architecture

The preceding error analysis reveals that different operational modes—flying, hovering, and visually centering—expose fundamentally different subsets of the error budget. This section presents the three estimation levels as a structured design decision, comparing accuracy, mission cost, calibration dependency, and compliance with requirement R07 (land within 5 m to 10 m of the casualty). The architecture was designed so that each level builds on the previous one, and the system can fall back to a lower level if detection is lost during centering.

### K.9.1 Level 1: Fully Passive Estimation (No Drone Response)

In Level 1, the drone flies its pre-planned lawnmower search pattern without deviating. Detections are logged and GPS coordinates are estimated from pixel offset and GSD, but the drone does not stop, turn, or respond to any detection. This is the mode implemented in `passive_watch.py`—a pure observer that sends zero commands to the flight controller.

**How it works.** For each detection, the pixel offset from the image centre is projected through the full GSD pipeline (Equations 9–11): pixel offset  $\rightarrow$  camera-frame displacement  $\rightarrow$  yaw rotation to NED  $\rightarrow$  WGS 84 projection. Every error source in Table 45 contributes at full magnitude.

#### Advantages.

- No mission time lost to investigation—the search pattern completes without interruption.
- Covers the entire search area systematically; no risk of missing a second target while investigating the first.
- Can detect and log multiple targets per pass for post-flight analysis.
- Simplest implementation: no state machine transitions, no closed-loop control.

#### Disadvantages.

- Roll/pitch displacement dominates: 6.2 m at  $10^\circ$  pitch (Section K.7), always biased in the direction of travel.
- GPS timing lag contributes  $\sim 1$  m at  $5 \text{ m s}^{-1}$  (systematic, same direction as pitch bias).
- Single observation per flyover—no temporal averaging within a pass.
- No operator confirmation; false positive rate is unchecked.
- Calibration errors (focal length, compass, barometer) propagate fully into every estimate.

**Accuracy.** The RSS error budget at 35 m yields  $\text{CEP}_{50} \approx 7$  m for a single flyover observation. With five passes and inverse-variance weighted averaging (alternating headings partially cancel the pitch bias), this improves to  $\sim 3.5$  m—but systematic calibration errors do not cancel.

**Verdict.** Insufficient for R07. A 7 m CEP means the 7.5 m offset landing could place the drone directly on the casualty (zero clearance) or 15 m away. Used only as an initial rough estimate to trigger investigation, and as the data-collection mode for passive flight experiments.

### K.9.2 Level 2: Stop-and-Look (Hover Without Centering)

In Level 2, the drone halts immediately upon detecting a target and hovers at its current position. The camera points straight down (nadir), and multiple frames are captured during the hover for GPS averaging.

**How it works.** When the detection triggers in SEARCH, the drone transitions to a hover hold. Roll and pitch drop to  $< 1^\circ$  (GPS hold loop), eliminating the dominant in-flight error. However, the target is *not* at the frame centre—it sits at whatever pixel offset it was first detected. Converting that offset to GPS still requires the full projection pipeline: focal length  $f$ , altitude  $h$ , yaw  $\psi$ , and the coordinate transformations of Equations 13–9.

#### Advantages.

- Eliminates roll/pitch displacement ( $< 0.6$  m residual vs. 6.2 m in flight).
- Eliminates motion blur entirely (stationary camera, sharp frames, higher detection confidence).
- Eliminates GPS timing lag ( $v \approx 0$ , so  $v \cdot \Delta t \approx 0$ ).
- GPS averaging over hover period:  $\sqrt{N}$  improvement ( $\sim 48$  frames in 10 s at 4.8 FPS).

### Disadvantages.

- Target not at frame centre—still requires the full GSD/focal-length/yaw pipeline.
- Calibration errors propagate directly: a 5% focal length bias produces a persistent 0.8 m offset at 16 m ground distance that does *not* average out.
- Yaw (compass) error still matters:  $\pm 5^\circ$  at 16 m offset  $\rightarrow$  1.4 m cross-track error.
- Lens distortion present at frame edges (up to 0.5 m without undistortion).
- No operator confirmation of target identity.

**Accuracy.** With 10 s of hover at 4.8 FPS, GPS averaging yields  $\text{CEP}_{50} \approx 2.5$  m to 3 m. The floor is set by systematic calibration errors rather than GPS noise—unlike Level 1, averaging *cannot* improve beyond the calibration limit.

**Verdict.** Viable but suboptimal. Meets R07 marginally (CEP inside the 5 m to 10 m annulus), but accuracy depends on calibration quality that is difficult to validate in the field. If the compass has a  $5^\circ$  hard-iron offset (plausible near motor currents), the estimate is biased by  $>1$  m regardless of averaging duration.

### K.9.3 Level 3: Visual Centering (Implemented)

In Level 3, the drone does not merely stop—it actively flies toward the detected target until the bounding-box centre coincides with the image principal point. The CENTERING state commands velocity toward the GPS-estimated target position and transitions to VERIFY once the drone is within 1 m of the estimate.

**How it works.** At the moment the target reaches the optical centre:

$$\Delta u = u - \frac{W}{2} \approx 0, \quad \Delta v = v - \frac{H}{2} \approx 0 \quad (21)$$

The camera-frame offset (Equation 9) reduces to:

$$d_{\text{fwd}} = -\Delta v \cdot \text{GSD} \approx 0, \quad d_{\text{right}} = \Delta u \cdot \text{GSD} \approx 0 \quad (22)$$

regardless of the GSD value. The target GPS therefore equals the drone GPS directly, bypassing the entire pixel-to-GPS pipeline. This is the key insight: *when the pixel offset is zero, every multiplicative error source in the projection chain is multiplied by zero and vanishes.*

**The zero-offset proof.** The full GPS projection (Equations 9–11) can be written as:

$$\begin{pmatrix} \Delta \phi \\ \Delta \lambda \end{pmatrix} = \frac{\text{GSD}}{R_\oplus} \cdot \frac{180}{\pi} \cdot \begin{pmatrix} \cos \psi & -\sin \psi \\ \frac{\sin \psi}{\cos \phi_d} & \frac{\cos \psi}{\cos \phi_d} \end{pmatrix} \begin{pmatrix} -(v - C_y) \\ (u - C_x) \end{pmatrix} \quad (23)$$

When  $(u, v) = (C_x, C_y)$ , the pixel-offset vector is exactly  $\mathbf{0}$ , and the entire matrix product vanishes regardless of GSD, heading  $\psi$ , or latitude correction. Concretely: a 10% focal length error  $\times 0 = 0$ ; a  $5^\circ$  compass error  $\times 0 = 0$ ; a 1 m barometric drift  $\times 0 = 0$ .

### Advantages.

- All geometric/calibration errors vanish (multiplied by zero pixel offset).
- No dependency on focal length, compass, or barometric calibration quality.
- GPS averaging during the subsequent VERIFY hover improves precision further ( $\sqrt{N}$ ).
- Operator sees the centred target in the video stream for visual confirmation (Y/N decision).
- Robust to sensor degradation: even a misaligned lens or drifting compass cannot corrupt the estimate.

### Disadvantages.

- Adds 15–25 s of flight time per investigation (centering manoeuvre + verify hold).
- Briefly hovers above the casualty (safety analysis: rotor downwash at 35 m is  $<0.0002 \text{ m s}^{-1}$ —negligible; see Section S).
- Requires continuous detection during approach—if the target is lost, the drone returns to SEARCH.
- Investigates one target at a time (sequential, not parallel).

**Accuracy.** With the target centred, the only remaining error is the drone’s own GPS noise. The `--center-verify` flag averages the drone’s position over  $\sim 10$  s (50 samples at 5 Hz). With  $N_{\text{eff}} \approx 20$  independent samples (accounting for temporal correlation), this reduces the 2 m to 3 m single-fix CEP to  $\text{CEP}_{50} \approx 1.8$  m.

**Verdict.** Best accuracy of the three levels, chosen for the final implementation. Comfortably satisfies R07. The trade-off is flight time, which is acceptable given the single-target mission profile and the safety benefit of operator confirmation.

Note: a DESCENDING state exists in the state machine to lower the drone to a verify altitude before the VERIFY hold, but it is currently bypassed; the drone transitions directly from CENTERING to VERIFY at search altitude. Implementing the descent would further improve accuracy through the reduced GSD discussed in Section K.10.

#### K.9.4 Comparative Evaluation

Table 50 consolidates the error-source analysis across all three levels, and Table 52 compares operational characteristics.


**Table 50:** Error source comparison across the three estimation levels. “Not needed” indicates the error source is multiplied by zero pixel offset and does not contribute.

| Error Source                       | Level 1: Passive      | Level 2: Hover           | Level 3: Centering        |
|------------------------------------|-----------------------|--------------------------|---------------------------|
| Roll/pitch projection              | 6.2 m <sup>†</sup>    | Eliminated ( $<0.6$ m)   | Eliminated                |
| Motion blur                        | Present (47 px)       | Eliminated               | Eliminated                |
| GSD / focal length calib.          | Full propagation      | Still required           | Not needed ( $\times 0$ ) |
| Yaw rotation accuracy              | Full propagation      | Still required           | Not needed ( $\times 0$ ) |
| Lens distortion                    | Up to 0.5 m           | Present at edges         | Minimal (optical centre)  |
| GPS timing lag                     | $\sim 1$ m            | Eliminated ( $v=0$ )     | Eliminated                |
| GPS receiver noise                 | 2 m to 3 m (1 sample) | $\sqrt{N}$ reduction     | $\sqrt{N}$ reduction      |
| <b>Achievable CEP<sub>50</sub></b> | $\sim 7$ m            | $\sim 2.5$ m             | $\sim 1.8$ m              |
| <b>With tilt compensation</b>      | $\sim 3.5$ m          | No change                | No change                 |
| <b>Final estimate</b>              | Pixel offset + GSD    | Drone GPS + pixel offset | Drone GPS directly        |

<sup>†</sup> With attitude compensation enabled (`--compensate-tilt`, Section K.7.3), the 6.2 m roll/pitch term reduces to  $\sim 0.3$  m during flight. Hover-phase values are unchanged (tilt  $< 1^\circ$ , below activation threshold).

**Table 52:** Operational comparison of the three estimation levels. R07 requires landing within 5 m to 10 m of the casualty.

| Criterion                          | Level 1                       | Level 2         | Level 3                          |
|------------------------------------|-------------------------------|-----------------|----------------------------------|
| CEP <sub>50</sub> (with averaging) | 7 m (single), 3.5 m (5-pass)  | 2.5 m           | 1.8 m                            |
| R07 compliant?                     | No                            | Marginal        | Yes                              |
| Mission time cost per target       | 0 s                           | $\sim 10$ s     | $\sim 25$ s                      |
| Calibration dependent?             | Yes (fully)                   | Yes (partially) | No                               |
| False positive handling            | None (log only)               | None            | Operator confirms (Y/N)          |
| Multiple targets per pass          | All detected                  | One at a time   | One at a time                    |
| Implementation                     | <code>passive_watch.py</code> | Not implemented | <code>main.py</code> (CENTERING) |



**Figure 30:** Side-by-side comparison of the three estimation levels. Dashed circles represent the  $CEP_{50}$  error radius around the true target position (star). Level 1 (passive): the drone is pitched forward and moving, placing the target at the frame edge with the largest error circle. Level 2 (hover): the drone is stationary and nadir-pointing, but the target remains off-centre, requiring the full projection pipeline. Level 3 (centering): the drone has manoeuvred directly above the target, driving the pixel offset to zero and collapsing the error to GPS receiver noise only.

**Design rationale.** The progression from Level 1 through Level 3 reflects a fundamental engineering trade-off between estimation accuracy and mission time. Level 1 is “free” in terms of flight time but inadequate for precision landing. Level 2 eliminates motion-induced errors at a cost of  $\sim 10\text{ s}$  but remains limited by calibration quality. Level 3 adds a further  $\sim 15\text{ s}$  but produces estimates that are *independent of all calibration parameters*—a qualitatively different regime. The implemented system uses Level 1 estimates during SEARCH as rough candidates (sufficient to trigger investigation), then refines with Level 3 during CENTERING/VERIFY for the final position used in offset landing. This layered architecture ensures that calibration imperfections—inevitable in a field-deployed system assembled by students—cannot compromise the landing accuracy.

The trade-off is flight time: centering adds 15–25 s of manoeuvring above the target compared to an immediate stop. Section S analyses the operational risks of this additional low-altitude flight and explains why the direct offset landing configuration bypasses the descent phase while retaining the centering hold for position refinement.

## K.10 Altitude Effect on Estimation Accuracy

Altitude has a compound effect on position estimation quality that goes beyond the obvious change in ground sampling distance. Lower altitude improves accuracy through three independent mechanisms that reinforce each other:

1. **Smaller GSD  $\rightarrow$  more pixels on target.** At 35 m, the mannequin occupies  $\sim 81\text{ px}$  (Table 44); at 15 m, it occupies  $\sim 190\text{ px}$ . A larger bounding box yields a more precise centre estimate: the  $\pm 2\text{--}3\text{ px}$  detector noise translates to 0.05 m at 35 m but only 0.02 m at 15 m.
2. **Smaller GSD  $\rightarrow$  smaller ground projection of pixel errors.** Every pixel-level error (detection noise, resize rounding, lens distortion residual) is multiplied by GSD to produce a ground-level offset. At 35 m ( $GSD = 22.1\text{ mm px}^{-1}$ ), a 3 px error is 0.066 m; at 15 m ( $GSD = 9.5\text{ mm px}^{-1}$ ), the same 3 px error is 0.029 m.
3. **Reduced roll/pitch displacement.** The tilt-induced error  $\Delta_{\text{tilt}} = h \tan \theta$  is proportional to altitude. Even if the drone is not perfectly stationary during low-altitude investigation, a  $2^\circ$  residual tilt produces only 0.5 m error at 15 m versus 1.2 m at 35 m.

These three effects are encoded in the inverse-variance weighting scheme (Equation 34), where  $w_{\text{alt}} = (h_{\text{ref}}/h)^2$ . A hypothetical observation at 15 m would carry  $5.4\times$  the weight of one at 35 m, correctly reflecting the compound improvement in geometric accuracy. In the current implementation the drone remains at search altitude (35 m) through CENTERING and VERIFY—the DESCENDING state is present in the code but bypassed (Section K.9.4). Implementing the descent in future work would exploit this altitude-dependent quality progression: each metre of descent would make every new frame more informative.

## K.11 Accuracy Progression Through Mission Phases

Table 54 consolidates the preceding analysis into a single view of how the target position estimate improves from first detection to final landing (cf. the preview in Table 43; a more detailed breakdown including specific error-source

elimination appears in Table 59).

**Table 54:** Target GPS estimate accuracy at each phase of the four-phase detection-to-landing flow. CEP values are root-sum-of-squares of dominant error sources; measured CEP50 from DJI video validates the theoretical model. Tilt compensation is active throughout.

| Phase                               | Alt (m)   | Est. CEP (m) | Why it improves                                                                                            |
|-------------------------------------|-----------|--------------|------------------------------------------------------------------------------------------------------------|
| Phase 1: SEARCH (initial detection) | 35        | ~3.5         | Tilt compensation ray-traces through known pitch/roll; GPS noise (2–3 m) dominates after compensation      |
| Phase 2: CENTERING (approach)       | 35        | ~2.5         | Continuous re-estimation as pitch decreases; tilt compensation becomes no-op; GPS lag shrinks with speed   |
| Phase 3: CENTERING (hover lock)     | 35        | ~1.8         | Target at frame centre $\Rightarrow$ offset = 0; all geometric errors vanish; nadir camera, no motion blur |
| Phase 4: VERIFY (operator confirms) | 35        | ~1.8         | --center-verify averages drone GPS over 10 s (50 samples at 5 Hz); operator Y/N/I decision                 |
| APPROACH (fly to offset)            | 35        | ~1.8         | Estimate frozen at Phase 3/4 quality; 7.5 m offset computed from averaged GPS                              |
| <b>Measured (DJI video)</b>         | <b>35</b> | <b>2.3</b>   | <b>CEP50 from 50+ detections across 6 min flight video (passive estimation, no hover lock)</b>             |

The DJI video CEP50 of 2.3 m represents passive estimation (Level 1, no hovering or centering). The full four-phase flow achieves 1.8 m by eliminating geometric errors through hover lock (Phase 3) and GPS averaging (Phase 4). False positives are bounded to 15 s through automatic timeout in Phase 3.

The progression demonstrates that the system’s accuracy is not a single number but a function of mission phase: the combination of tilt compensation (Phase 1), continuous approach refinement (Phase 2), hover lock with nadir camera (Phase 3), and operator-confirmed GPS averaging (Phase 4) transforms a ~3.5 m initial estimate into a ~1.8 m final estimate—well within the R07 landing annulus of 5–10 m from the casualty. False positives detected during Phase 1 are automatically rejected after 15 s of fruitless hovering in Phase 3, at a cost of less than 1% of mission time per event.

## K.12 GPS Averaging Duration Analysis

The --center-verify mode averages the drone’s GPS position during the VERIFY hold to reduce the circular error. This subsection analyses the noise characteristics of the GPS receiver, derives the expected accuracy improvement as a function of averaging time, and justifies the 10 s duration selected for the mission.

### K.12.1 GPS Noise Characteristics

The Here 3+ GNSS receiver used on the Cube autopilot exhibits position noise that is approximately Gaussian with zero mean when the drone is stationary. Empirical characterisation from bench testing and DJI flight-video replay yields the following parameters:

- **Single-fix CEP<sub>50</sub>:** 2 m to 3 m (typical for consumer-grade L1 GPS with SBAS corrections).
- **Autocorrelation time  $\tau_c$ :** 1 s to 2 s. Consecutive fixes at 5 Hz are correlated because the receiver’s internal Kalman filter smooths over several measurement epochs. After  $\tau_c$ , successive samples are approximately independent.
- **Update rate:** 5 Hz (configurable on Here 3+, set via GPS\_RATE\_MS parameter).

The autocorrelation time is the critical parameter: it determines how many *independent* samples are obtained during an averaging window. At 5 Hz over  $T$  seconds, the receiver delivers  $N = 5T$  raw fixes, but only  $N_{\text{eff}} \approx T/\tau_c$  are statistically independent. Taking  $\tau_c = 2$  s (conservative estimate) and a baseline CEP<sub>50</sub> of 2.5 m, the averaged CEP improves as

$$\text{CEP}_{\text{avg}} = \frac{\text{CEP}_{\text{single}}}{\sqrt{N_{\text{eff}}}} = \frac{2.5}{\sqrt{T/2}} \text{ [metres]} \quad (24)$$


### K.12.2 Averaging Improvement Model

Table 56 shows the expected CEP improvement as a function of averaging duration, computed from Equation 24 with  $\tau_c = 2$  s and CEP<sub>single</sub> = 2.5 m.

**Table 56:** GPS averaging improvement vs. duration.  $N$  is total fixes at 5 Hz;  $N_{\text{eff}}$  is the number of independent samples assuming  $\tau_c = 2$  s. The improvement factor is  $\sqrt{N_{\text{eff}}}$ .

| Avg. Time (s)  | $N$ (fixes) | $N_{\text{eff}}$ | Improvement ( $\sqrt{N_{\text{eff}}}$ ) | Expected CEP (m) |
|----------------|-------------|------------------|-----------------------------------------|------------------|
| 0 (single fix) | 1           | 1                | 1.0×                                    | 2.50             |
| 2              | 10          | 1                | 1.0×                                    | 2.50             |
| 5              | 25          | 2.5              | 1.6×                                    | 1.58             |
| <b>10</b>      | <b>50</b>   | <b>5</b>         | <b>2.2×</b>                             | <b>1.12</b>      |
| 20             | 100         | 10               | 3.2×                                    | 0.79             |
| 30             | 150         | 15               | 3.9×                                    | 0.65             |
| 60             | 300         | 30               | 5.5×                                    | 0.46             |

The relationship between averaging time and CEP is plotted in Figure 31, which shows the diminishing returns characteristic of  $1/\sqrt{T}$  convergence.



**Figure 31:** CEP<sub>50</sub> vs. GPS averaging time during stationary hover. The  $1/\sqrt{T}$  convergence means that doubling the averaging time only improves accuracy by  $\sim 29\%$ . The 10 s operating point (orange) achieves sub-1.2 m CEP while adding minimal mission time.

### K.12.3 Why 10 Seconds

The 10 s averaging window was selected as the optimal trade-off between position accuracy and mission time cost. Five considerations informed this choice:

1. **Diminishing returns.** The  $1/\sqrt{T}$  convergence law means that each additional second of averaging contributes less than the previous one. Going from 10 s to 60 s (a  $6\times$  increase in time) only improves CEP from 1.12 m to 0.46 m—a factor of  $2.4\times$  improvement for  $6\times$  the time. The marginal value of each additional second drops below  $0.02 \text{ m s}^{-1}$  after  $T = 10$  s.
2. **Mission time cost.** Every second spent hovering during VERIFY is a second not spent searching. In a SAR scenario with a 15-minute battery endurance at search speed, 10 s represents  $\sim 1\%$  of total mission time—an acceptable cost for a  $2.2\times$  accuracy improvement. A 60 s hover would consume 7% of mission time for diminishing returns.
3. **Battery impact.** Hovering at 35 m consumes approximately 15–18 W on the airframe used. A 10 s hover draws  $\sim 50 \text{ mA h}$  ( $< 0.5\%$  of a 3000 mA h 4S battery), which is operationally negligible. A 60 s hover would consume  $\sim 3\%$  of battery capacity.
4. **Sufficient accuracy for offset landing.** The R07 requirement specifies landing within 5–10 m of the casualty. With a 7.5 m deliberate offset (for safety) and a  $\sim 1.1$  m CEP from 10 s averaging, the 95th-percentile total error is  $7.5 \pm 2.2$  m—comfortably within the 5–10 m annulus. Sub-metre accuracy (achievable at 30–60 s) provides no additional operational benefit given the intentional offset.
5. **Wind and drift tolerance.** During a 10 s hover, wind gusts may push the drone  $\sim 0.5$ –1 m from the target centre. This displacement is comparable to the GPS noise floor and does not meaningfully degrade the average. Over 60 s, sustained wind could introduce systematic drift that *increases* the error—particularly if the flight controller’s station-keeping oscillates with a period comparable to  $\tau_c$ .

In summary, 10 s of stationary averaging reduces the GPS CEP from 2.5 m to  $\sim 1.1$  m—a  $2.2\times$  improvement that brings the estimate well within R07 requirements. Longer averaging durations offer diminishing returns that do not justify the mission time and battery cost, particularly given the deliberate 7.5 m offset landing strategy that dominates the total landing error budget.

## K.13 Fusion Strategies

Four estimation approaches were implemented in the simulator (`simple_simulator.py`) and passive observation tool (`passive_watch.py`), each operating on the same stream of per-frame GPS estimates. All four run in parallel so that their accuracy can be compared directly on every detection cluster. The autonomous mission state machine uses a different, simpler approach: after visual centering places the drone directly above the target, the `--center-verify` flag averages the drone's own GPS position over 10s, exploiting the fact that drone GPS  $\approx$  target GPS when centred.

### K.13.1 Rolling Weighted Average (Window = 50)

The most recent  $N = 50$  observations are retained. The estimate is the inverse-variance-weighted mean over the window:

$$\hat{\phi} = \frac{\sum_{i=k-N+1}^k w_i \phi_i}{\sum_{i=k-N+1}^k w_i}, \quad \hat{\lambda} = \frac{\sum_{i=k-N+1}^k w_i \lambda_i}{\sum_{i=k-N+1}^k w_i} \quad (25)$$

This method is responsive to changes—if the drone descends and obtains higher-quality observations, the old high-altitude estimates are quickly flushed from the window—but it is noisier than methods that retain all history.

### K.13.2 Cumulative Weighted Average (All Time)

Every observation ever recorded is included:

$$\hat{\phi}_k = \frac{W_{\phi,k-1} + w_k \phi_k}{W_{\phi,k-1} + w_k}, \quad W_{\phi,k} = W_{\phi,k-1} + w_k \phi_k, \quad W_k = W_{k-1} + w_k \quad (26)$$

Implemented as a running sum to avoid storing all observations. This produces the most stable estimate but responds sluggishly to improvements in observation quality, because hundreds of early high-altitude observations continue to dilute the influence of later low-altitude passes.

### K.13.3 Kalman Filter (Static Target)

A two-state linear Kalman filter [37] models the target as stationary ( $\mathbf{F} = \mathbf{I}$ ) with a small process noise  $\mathbf{Q} = \sigma_q^2 \mathbf{I}$ , where  $\sigma_q$  corresponds to 0.1 m converted to degrees. The measurement noise covariance is set inversely proportional to the observation weight:

$$\sigma_R = \frac{3.0}{\max(w_i, 0.1)} \text{ m}, \quad \mathbf{R}_k = \sigma_R^2 \mathbf{I}_{\text{deg}} \quad (27)$$

where  $\mathbf{I}_{\text{deg}}$  denotes the identity matrix scaled to degree<sup>2</sup> units via the conversion factor  $(1/R_{\oplus} \cdot 180/\pi)^2$ . The standard predict–update cycle follows:

$$\mathbf{P}_{k|k-1} = \mathbf{P}_{k-1} + \mathbf{Q} \quad (28)$$

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} (\mathbf{P}_{k|k-1} + \mathbf{R}_k)^{-1} \quad (29)$$

$$\mathbf{x}_k = \mathbf{x}_{k|k-1} + \mathbf{K}_k (\mathbf{z}_k - \mathbf{x}_{k|k-1}) \quad (30)$$

$$\mathbf{P}_k = (\mathbf{I} - \mathbf{K}_k) \mathbf{P}_{k|k-1} \quad (31)$$

The filter is initialised on the first detection with  $\mathbf{x}_0 = \mathbf{z}_0$  and  $\mathbf{P}_0 = 4 \mathbf{R}_0$ , giving the initial measurement a 20% Kalman gain and allowing subsequent observations to dominate quickly. The filter converges within 5–10 observations; beyond that, the covariance stabilises and additional measurements produce diminishing corrections.

### K.13.4 Inverse-Variance Weighting

This is the method used in the simulator (`simple_simulator.py`) and passive observation tools (`passive_watch.py`). The mission state machine uses a simpler approach: direct drone GPS averaging via the `--center-verify` flag (Section K.11). Each observation receives a composite weight:

$$w_i = w_{\text{cen}}(u_i, v_i) \cdot w_{\text{alt}}(h_i) \quad (32)$$

where the centrality factor is

$$w_{\text{cen}} = 1 + 4 \left( 1 - \frac{d_i}{d_{\text{max}}} \right), \quad d_i = \sqrt{(u_i - C_x)^2 + (v_i - C_y)^2} \quad (33)$$

and the altitude factor is

$$w_{\text{alt}} = \left( \frac{h_{\text{ref}}}{h_i} \right)^2 \quad (h_{\text{ref}} = 30 \text{ m}) \quad (34)$$

At 30 m altitude,  $w_{\text{alt}} = 1$ ; at 15 m,  $w_{\text{alt}} = 4$ ; at 10 m,  $w_{\text{alt}} = 9$ . Detections near the image centre at low altitude therefore dominate the fused estimate, which is physically intuitive: the target fills more pixels, the GSD-scaled projection covers less ground distance, and the observation is inherently more precise.

The centre-snap mode amplifies this further: when the target is within 50 pixels of the image centre, the drone’s own GPS is used directly (bypassing the GSD projection chain) with a fixed weight of  $10 \cdot w_{\text{alt}}$ —yielding a weight of 90 at 10 m altitude versus approximately 5 for a typical 35 m edge detection.

## K.14 Spatial Clustering

In a realistic SAR mission the camera may detect multiple objects: the target casualty, a piece of debris, or a false positive such as a tree stump. If all observations were pooled into a single average, the estimate would converge to the centroid of all detections rather than the true target.

To handle this, each new GPS estimate is compared against existing detection clusters using the Haversine great-circle distance:

$$d = 2 R_{\oplus} \arctan(\sqrt{a}, \sqrt{1-a}), \quad a = \sin^2 \frac{\Delta\phi}{2} + \cos \phi_1 \cos \phi_2 \sin^2 \frac{\Delta\lambda}{2} \quad (35)$$

If the nearest cluster centroid is within a configurable threshold (default 30 m), the observation is added to that cluster; otherwise a new cluster is spawned. Each cluster maintains:

- A rolling window of the last 50 weighted observations (for the rolling average).
- Running sums  $W_{\phi}$ ,  $W_{\lambda}$ ,  $W$  (for the cumulative average).
- A Kalman filter state  $(\mathbf{x}, \mathbf{P})$ .
- A detection count and a human-assigned classification (dummy / interest / false positive).

The 50-observation cap per cluster prevents memory growth during long loiter periods while retaining enough history for the rolling average to be statistically meaningful. In simulation testing with two dummy targets placed 80 m apart, the system correctly maintained two separate clusters from the first flyover, with no cross-contamination of observations.

## K.15 Comparison of Methods

Table 57 summarises the performance of the four fusion methods across repeated simulation trials (keyboard-flown missions with configurable GPS drift and altitude noise) and DJI video replay with synchronised SRT telemetry.

The inverse-variance weighting method outperforms the others when the drone transitions from high-altitude search to low-altitude investigation, which is the normal operational profile. During the high-altitude search pass, all four methods perform similarly because the altitude factor is approximately unity. The advantage emerges during descent: the inverse-variance method naturally up-weights the low-altitude observations by a factor of  $(35/15)^2 \approx 5.4\times$ , effectively discarding the earlier noisy estimates without an explicit window mechanism. The Kalman filter achieves comparable accuracy but requires tuning of  $\mathbf{Q}$  and  $\sigma_R$ ; in practice, the inverse-variance method was more robust to parameter choices.

## K.16 Empirical Validation

### K.16.1 DJI Video Replay (The “Bullseye Test”)

**Methodology.** The localisation pipeline was validated offline by replaying a DJI Mavic flight video recorded at the Fenswood Farm test site. The video was cropped to  $1456 \times 1088$  px at 30 fps to match the Pi camera’s native resolution, ensuring the same pixel-to-GSD scaling factors as the operational system. Each frame carries synchronised SRT telemetry—GPS latitude/longitude, barometric altitude, heading, and timestamp—enabling frame-accurate ground-truth comparison without any manual time alignment.

A life-size rescue mannequin (1.8 m tall) was placed at a surveyed GPS position. The drone was flown over the site in a series of lawnmower-like passes at altitudes ranging from 15 m to 50 m at ground speeds of  $3 \text{ m s}^{-1}$  to  $8 \text{ m s}^{-1}$ , covering approximately six minutes of continuous flight. The YOLOv8n detector (retrained model, `sar_v2_1088`, `mAP50 = 0.995`) was run on every frame. For each detection, the pixel-to-GPS pipeline (Section K.3) produced an independent target

position estimate using the SRT telemetry as the drone state input. Each estimate was then compared against the known mannequin position to compute the position error.

The analysis pipeline (`tests/laptop/video_test.py`) also computed the FOV-calibrated ground sample distance at each frame’s altitude, enabling the scale bar and measurement tool to verify that the mannequin measured 1.7 m to 1.9 m at all test altitudes—an independent check on the projection geometry.

**Aggregate results.** Over the six-minute flight, the detector produced 50+ independent GPS estimates across multiple passes and altitudes. Table 58 summarises the key accuracy metrics.

**Directional bias and GPS timing lag.** The scatter plot of individual estimation errors (Figure 10a) reveals that observations do not form a symmetric cluster around the true position. Instead, the error distribution is elongated along the flight heading direction (Figure 10b). This anisotropy is a direct consequence of the 100 ms to 200 ms GPS receiver latency: by the time the receiver reports its position, the drone has travelled  $v \cdot \Delta t$  further along its heading. At  $5 \text{ m s}^{-1}$  with 150 ms lag, this produces a systematic along-track displacement of  $\sim 0.75 \text{ m}$  per observation.

Crucially, the lawnmower search pattern’s alternating headings cause this bias to appear on *opposite* sides of the true position on successive passes—northbound passes produce southward-displaced estimates and vice versa. This geometric property means that the directional biases partially cancel when fusing observations from multiple passes, explaining why the fused estimate converges to better accuracy than the individual  $\text{CEP}_{50}$  would suggest.

**Detection performance across altitude.** Analysis of detection confidence as a function of altitude revealed three distinct regimes:

- **15 m to 40 m (optimal band):** detection confidence consistently exceeded 0.9. The mannequin subtended 40–100 px in height, providing ample spatial information for the detector. This altitude range encompasses the entire planned search altitude (35 m) and investigation altitude (15 m).
- **Below 15 m:** the mannequin began to fill the frame, and bounding-box edges were clipped by the image boundary. Detections still fired but with reduced confidence and less reliable centroid estimates, since the YOLO bounding box was truncated.
- **Above 45 m:** the mannequin shrank below  $\sim 40$  px in height, approaching the minimum feature size for reliable detection. Confidence dropped below 0.5 and detections became intermittent. This establishes an empirical ceiling for the current model’s detection range.

These findings validate the choice of 35 m as the search altitude: it sits comfortably within the optimal detection band while providing a wide field of view for efficient area coverage.

**Validation of the theoretical CEP model.** The theoretical RSS error budget (Table 45) predicts a CEP of  $\sim 7 \text{ m}$  during search at 35 m with  $10^\circ$  pitch. The measured  $\text{CEP}_{50}$  of 2.3 m is significantly lower than this prediction. Three factors explain the discrepancy:

1. **Lower effective pitch.** The DJI Mavic’s gimbal stabilises the camera to near-nadir orientation regardless of airframe pitch. This largely eliminates the dominant  $h \tan \theta$  error term (Section K.7), which alone accounts for 6.2 m at  $\theta = 10^\circ$ . The Pi camera is body-mounted without gimbal stabilisation, so the  $10^\circ$  pitch error will apply to the operational system—the DJI results therefore represent a best-case baseline.
2. **Multi-heading cancellation.** As described above, the alternating passes cause GPS timing lag to displace estimates in opposite directions, partially cancelling this bias in the aggregate. A single-pass measurement would show a higher effective CEP.
3. **Detections near frame centre.** Many detections occurred when the drone was nearly overhead, placing the mannequin near the image centre where GSD-scaled pixel offsets are small. The RSS budget assumes a worst-case off-centre detection.

This comparison underscores a key design insight: the measured  $\text{CEP}_{50}$  is dominated by GPS receiver noise (2.3 m), not by the projection pipeline. Any improvement beyond this floor requires either RTK GPS or longer averaging, not better image processing.

**The 16.5 m outlier.** The single worst-case estimate (16.5 m error) occurred during a high-speed pass ( $\sim 8 \text{ m s}^{-1}$ ) at 50 m altitude. At this altitude, the GSD is large (0.032 m/pixel), so even a small bounding-box centroid error translates into a large ground offset. Simultaneously, the GPS timing lag displacement reached  $\sim 1.2 \text{ m}$ , and the detection was near the frame edge where lens distortion residuals are largest. This outlier demonstrates why the system requires multiple observations and robust fusion (inverse-variance weighting naturally down-weights this high-altitude, off-centre observation) rather than relying on any single detection.


## K.16.2 Simulator Validation

The interactive simulator (`simple_simulator.py`) supports configurable GPS drift (`--gps-drift`), altitude noise (`--alt-noise`), yaw noise (`--yaw-noise`), and FOV calibration error (`--fov-error`). All four estimators run in parallel, and the scatter plot displays each observation (colour-coded by weight) alongside the rolling, cumulative, Kalman, and inverse-variance estimates with real-time error metrics against the known dummy position.

A typical keyboard-flown simulator mission—flyover at 35 m, optionally descend to 15 m, circle the target—produces 30–50 observations over 60–90 seconds. With 3 m GPS drift enabled, the inverse-variance estimate converges to sub-metre accuracy within the first 20 observations when the operator descends below 20 m, confirming the altitude-dependent weighting strategy. Note that this descent is a simulator-only capability; the autonomous mission state machine currently remains at search altitude through VERIFY (Section K.11).

## K.17 Detailed Per-State Accuracy Analysis

The preceding sections have analysed individual error sources, mitigation strategies, and fusion algorithms in isolation. This subsection draws them together to show how the system’s position estimate improves concretely as the mission progresses through its state machine. Table 59 expands on the earlier summaries (Tables 43 and 54) by detailing the specific error sources eliminated at each phase transition.




**Figure 32:** Target GPS estimate accuracy progression through the four-phase detection-to-landing flow. Circle radii are proportional to CEP<sub>50</sub> at each phase. Tilt compensation in Phase 1 reduces the initial estimate from ~7 m (uncompensated) to ~3.5 m. The combination of continuous approach refinement (Phase 2), hover lock with nadir camera (Phase 3), operator verification with GPS averaging (Phase 4), and spatial deduplication of false positives transforms the initial estimate into a ~1.8 m final position—a ~2× improvement through operational procedure alone.

The progression from 3.5 m (tilt-compensated) to 1.8 m represents a ~2× improvement achieved through operational procedure (approach, centering, hovering, GPS averaging), without any increase in software complexity or sensor quality. Without tilt compensation, the improvement would be ~4× from 7 m. Each phase transition eliminates a specific category of error:

- Phase 1 (SEARCH):** tilt compensation ray-traces the detection through the known pitch/roll, reducing the ~6.2 m pitch-induced bias to ~0.3 m. The rough estimate (CEP ≈ 3.5 m) is accurate enough to initiate approach.
- Phase 2 (CENTERING, approach):** as the drone decelerates, pitch decreases continuously and every frame re-estimates with current attitude. The tilt compensation naturally becomes a no-op. GPS timing lag shrinks proportionally with speed.
- Phase 3 (CENTERING, hover lock):** the target reaches the frame centre while the drone hovers at search altitude. The zero-offset condition (Equation 22) eliminates all calibration-dependent errors. If the target is not visible, a 15 s timeout rejects the location as a false positive (<1% mission time cost).
- Phase 4 (VERIFY):** the `--center-verify` flag averages the drone’s own GPS position (which equals the target position, since the drone is centred above it) over ~10 s at 5 Hz, yielding  $N = 50$  samples. With  $N_{\text{eff}} \approx 20$  (accounting for temporal correlation), this reduces the 2 m to 3 m single-fix CEP to ~1.8 m. The operator confirms or rejects, with a 120 s auto-reject timeout.

The final 1.8 m CEP<sub>50</sub> is dominated by the irreducible GPS receiver noise floor—the theoretical minimum achievable with a single-frequency L1 receiver without differential correction (RTK or DGPS). Implementing the currently-bypassed

DESCENDING state (lowering to 15 m before VERIFY) would improve this further through reduced GSD (Table 44). Further improvement beyond that would require either RTK GPS ( $\sim 0.02$  m) or significantly longer averaging times. Figure 33 summarises the CEP improvement across all four phases.



**Figure 33:** CEP<sub>50</sub> at each phase of the detection-to-landing flow, showing progressive accuracy improvement from  $\sim 3.5$  m (Phase 1, tilt-compensated search) to  $\sim 1.8$  m (Phase 4, GPS-averaged hover lock). Each phase transition eliminates a specific error category: Phase 2 reduces pitch and timing lag as the drone decelerates; Phase 3 drives the pixel offset to zero through centering; Phase 4 averages GPS noise over 50 samples.

## K.18 Four-Phase Detection-to-Landing Flow

The accuracy progression described above is not an accident of the state machine—it is a deliberate four-phase progressive refinement architecture. The system treats target localisation as a convergent process: each phase eliminates a specific category of error, and the estimate precision improves monotonically from initial detection through to the final landing decision. Tilt compensation enables reliable initial acquisition even during high-speed flight, and the architecture degrades gracefully—if the compensation is inaccurate, the hover phase eliminates all geometric errors regardless.

### K.18.1 Phase 1: Initial Detection During Search (Tilted Camera)

During the SEARCH state the drone flies the lawnmower pattern at  $5 \text{ m s}^{-1}$  and 35 m altitude, inducing a forward pitch of approximately  $5\text{--}10^\circ$  to maintain airspeed. The camera, rigidly mounted to the airframe, is tilted away from nadir by this pitch angle. When the YOLOv8n detector identifies the mannequin in this tilted frame, the pixel-to-GPS pipeline (Section K.3) ray-traces the detection through the known pitch and roll angles obtained from the flight controller’s attitude quaternion at 50 Hz.

The ray-tracing correction (Equations 16–20) constructs a ray from the camera origin through the detected pixel, rotates it by the measured pitch and roll, and intersects it with the ground plane at the known flight altitude. This produces a rough GPS estimate with  $\text{CEP}_{50} \approx 3.5 \text{ m}$ —accurate enough to command the drone toward the target and begin decelerating. Without tilt compensation, the  $10^\circ$  pitch alone would contribute  $\sim 6.2 \text{ m}$  of forward bias (Section K.7), but even this uncompensated estimate would keep the target within the camera’s 32 m ground footprint. The state machine transitions to CENTERING.

### K.18.2 Phase 2: Approach and Continuous Refinement

The drone flies toward the rough GPS estimate at reduced speed, decelerating as it approaches. The flight controller pitches the airframe forward to accelerate and backward to brake, so a decelerating drone naturally approaches level flight. During approach, the system re-detects and re-estimates the target position on *every frame* using the current pitch and roll telemetry:

- As the drone slows from  $5 \text{ m s}^{-1}$  to near-zero, the forward pitch decreases continuously from  $\sim 10^\circ$  to  $\sim 3^\circ$  and then to  $< 1^\circ$ .
- Each frame’s GPS estimate benefits from the reduced tilt: the tilt compensation naturally becomes a no-op as pitch approaches  $0^\circ$ .

- GPS timing lag (100 ms latency  $\times$  velocity) shrinks proportionally with speed, eliminating the 1 m along-track error.
- Motion blur decreases, improving detection confidence and bounding-box precision.
- Additional observations are fused into the estimate, progressively refining it from  $\sim 3.5$  m to  $\sim 2.5$  m CEP.

The estimate converges toward the true target position without any discrete transition—the refinement is continuous throughout the approach.

### K.18.3 Phase 3: Hover Lock (Nadir Confirmation)

The drone arrives at the estimated target position and holds a stationary hover. Pitch and roll both settle to approximately  $0^\circ$  and the camera points straight down (nadir). At this point, several error sources vanish simultaneously:

- **Roll/pitch projection error**  $\rightarrow 0$ : with the camera at nadir, there is no off-axis projection to compensate, and the 6.2 m pitch-induced bias disappears entirely. The tilt compensation is a no-op (below the 0.02 rad activation threshold).
- **GPS timing lag**  $\rightarrow 0$ : the drone is stationary, so the 100–200 ms GPS latency introduces zero position error regardless of lag duration.
- **Motion blur**  $\rightarrow 0$ : no translational velocity means no inter-frame smearing, improving detection confidence and bounding-box precision.
- **Calibration-dependent errors**  $\rightarrow 0$ : when the target is at or near the frame centre, the pixel offset  $(u-C_x, v-C_y)$  approaches zero. Because every geometric error in the projection chain (GSD, focal length, yaw, lens distortion) is *multiplied* by this offset (Equation 9), they are all multiplied by zero and vanish (Section K.8).

The detection uses simple nadir math—no tilt correction needed. If the target is visible and within 1 m of the frame centre, the state machine transitions to VERIFY. If the target is *not* visible after 15 s—indicating a false positive—the location is rejected and the drone resumes search (Section K.18.8).

The drone’s own GPS position, averaged over  $\sim 10$  s of stationary hover (50 samples at 5 Hz via `--center-verify`), now *is* the target’s GPS position to within the irreducible receiver noise floor. The result is a precise estimate with  $\text{CEP}_{50} \approx 1.8$  m—sufficient for the R07 offset landing requirement of 5–10 m from the casualty.

### K.18.4 Phase 4: Operator Verification and Landing Decision

With the target confirmed by the autonomous system, the operator makes the final decision via the ground station interface (browser dashboard or terminal):

- **Y (confirm target)**: the GPS estimate is frozen at Phase 3 quality. The drone computes a 7.5 m offset landing point (downwind when wind data is available), flies to the offset position, and deploys the payload before landing.
- **N (reject target)**: the detection is dismissed and the drone resumes the search pattern from the waypoint where the investigation was triggered.
- **I (item of interest)**: the GPS position is logged for post-mission review; the drone resumes search.
- **Timeout (120 s)**: if the operator does not respond within 120 s, the system auto-rejects and resumes search, preventing the drone from hovering indefinitely.


### K.18.5 Why Four Phases Are Necessary

A natural question is whether fewer phases could achieve the required accuracy. Table 60 compares the possible strategies.

The critical insight is that the rough estimate in Phase 1 only needs to place the drone close enough that the mannequin remains visible in the camera’s field of view while hovering. At 35 m altitude the camera footprint is  $32.2 \times 24.0$  m (Table 44), so even a 7 m error keeps the target well within the frame. The four-phase approach therefore decouples detection sensitivity (which determines whether the target is found at all) from localisation precision (which determines where the drone lands), and adds operator oversight as the final safety gate.

### K.18.6 The Progressive Refinement Trajectory

The transition from initial detection to precise hover lock is a continuous refinement that occurs during the CENTERING state. Figure 34 illustrates all four phases, showing how pitch angle decreases and error circles shrink as the drone approaches the target, with a branch for false positive recovery.



**Figure 34:** The four-phase progressive refinement trajectory. Phase 1: initial detection at full search speed with tilted camera and tilt-compensated GPS estimate (CEP  $\approx 3.5$  m). Phase 2: continuous approach and deceleration, pitch decreases, estimate improves (CEP  $\approx 2.5$  m). Phase 3: stationary hover with nadir camera, all geometric errors eliminated (CEP  $\approx 1.8$  m); if no target is visible, a 15 s timeout triggers resumption (red dashed branch). Phase 4: operator confirms via ground station; Y triggers offset landing, N resumes search. Error circles are drawn proportional to CEP at each phase.

The four phases map to distinct segments of the state machine:

1. **Phase 1 (Initial detection, SEARCH state):** The mannequin is first detected during a SEARCH-pattern leg. The tilt-compensated GPS estimate triggers the transition to CENTERING, and the drone begins flying toward the estimate at reduced speed.
2. **Phase 2 (Approach and deceleration, CENTERING state):** As the drone approaches the target, it decelerates. The forward pitch drops continuously from  $\sim 10^\circ$  to  $< 1^\circ$ , and every frame produces a more accurate GPS estimate. The tilt compensation naturally becomes a no-op as pitch approaches zero.
3. **Phase 3 (Hover lock, CENTERING  $\rightarrow$  VERIFY transition):** The drone reaches the estimated position and holds a stationary hover. With pitch and roll at  $0^\circ$ , the camera is at nadir. If the target is detected within 1 m of the frame centre, the state machine transitions to VERIFY. If the target is *not* visible after 15 s, the location is rejected as a false positive and the drone resumes search (Section K.18.8).
4. **Phase 4 (Operator verification, VERIFY state):** The `--center-verify` flag averages the drone's GPS position over  $\sim 10$  s (50 samples at 5 Hz), freezing the final estimate at  $\text{CEP}_{50} \approx 1.8$  m. The operator confirms or rejects via the ground station (Y/N/I), with a 120 s auto-reject timeout.

### K.18.7 Design Insight

The tilt compensation algorithm (Section K.7.3) enables reliable initial acquisition even during high-speed flight at  $5 \text{ m s}^{-1}$ . The system architecture treats detection as a progressive refinement process, not a single-shot measurement. This philosophy—detect approximately, then refine precisely—is what allows a 2 m-accuracy consumer GPS receiver to support offset landings within a 5–10 m annulus. The four-phase approach provides graceful degradation: even if the tilt compensation is miscalibrated or disabled entirely, the  $\sim 7$  m uncompensated rough estimate is still sufficient to initiate the approach, because the hover lock in Phase 3 eliminates all the errors that the compensation was attempting to correct.

### K.18.8 False Positive Handling

The rough estimate from Phase 1 may direct the drone to a location where no target exists. This occurs when the detection was a false positive (a rock, shadow, or clothing that resembles the mannequin), the tilt compensation was inaccurate due to an extreme pitch angle, or accumulated GPS drift placed the drone at a position that does not correspond to the detection. The system handles all of these cases through a bounded-cost automatic recovery mechanism.

When the drone arrives at the estimated position in Phase 3 and hovers but *cannot re-detect* the target, the following sequence executes:


1. **Hover scan.** The drone holds position at the estimated GPS coordinates, continuously running the detector on each camera frame. With the camera at nadir and no motion blur, if a real target exists within the  $32.2 \times 24.0$  m ground footprint, it will be detected.
2. **Automatic timeout.** After 15 s without a detection, the state machine logs "CENTERING TIMEOUT --- resuming search" and transitions back to the SEARCH state.
3. **Spatial deduplication.** The false-positive location is appended to a `rejected_targets` list. Any future detection whose estimated GPS position falls within 5 m of a rejected location is silently ignored, preventing the drone from investigating the same false alarm twice.

4. **Pattern resumption.** The drone resumes the lawnmower search pattern from the waypoint where the investigation was triggered, ensuring full area coverage despite the interruption.

**Cost of a false positive.** False positives are bounded to 15 s of mission time through the automatic timeout—approximately 1% of a typical 20 min endurance budget at 35 m hover. The associated battery expenditure is  $\sim 75$  mA h ( $< 0.3\%$  of a 3000 mAh 4S battery). One rejected-target exclusion zone (5 m radius) is created. No operator intervention is required; the system self-recovers fully autonomously.

**False positive mitigation.** Four mechanisms reduce the frequency and impact of false positives:


- **SMART detection** (`--smart-detect`): requires three consecutive confirmed frames before triggering CENTERING, filtering out transient single-frame false alarms.
- **Confidence threshold** (`--conf`): raising the threshold (e.g. to 0.5) suppresses low-confidence false positives at a modest reduction in recall.
- **Spatial deduplication**: the `rejected_targets` list prevents repeated investigation of the same location, capping the cost of any persistent false-positive source at one 15 s timeout event.
- **Bounded timeout**: the 15 s ceiling limits the worst-case impact of any single false positive to a known, budgeted quantity—less than 1% of mission time.



**Figure 35:** Complete four-phase decision flow including false positive recovery. The left branch shows successful progression: Phase 1 detection → Phase 2 approach → Phase 3 hover lock → Phase 4 operator verification → offset landing. The right branch shows the false positive path: no re-detection during Phase 3 hover triggers a 15 s timeout, the location is added to the rejected-targets list (spatial deduplication prevents reinvestigation), and the drone resumes the search pattern at a cost of  $< 1\%$  of mission time.

## K.19 Cooperative Vision-Flight Architecture

The four-phase flow described above reveals a property that distinguishes our system from passive surveillance: the drone and vision system form a *closed-loop feedback system* in which each detection improves the next. Unlike a static camera whose accuracy is fixed by altitude and calibration, our architecture creates a cycle where detection informs flight, flight improves detection conditions, and improved conditions yield better detections.



**Figure 36:** Cooperative vision-flight feedback loop. Detection produces a pixel coordinate, GPS estimation projects it to the ground, and flight control steers the drone toward the target. Crucially, flying closer improves the next detection: lower altitude yields finer GSD, deceleration reduces pitch-induced tilt error, and hovering at nadir drives the pixel offset—and therefore all multiplicative projection errors—toward zero. The cycle repeats until convergence, typically within 15–20 s.

### K.19.1 Three Mechanisms of Flight-Aided Improvement

Each phase of the mission exploits a different mechanism by which the drone’s response to a detection improves subsequent estimates:

1. **Approaching → lower GSD → tighter bounding box.** As the drone flies toward the target, altitude above the target decreases and the ground sampling distance (GSD) shrinks proportionally ( $\text{GSD} = h \cdot w_s / (f \cdot W)$ , Section K.4). More pixels cover the target, producing a tighter bounding box whose centre is more precisely localised. At 35 m the dummy subtends  $\sim 20$  pixels; at 15 m it subtends  $\sim 47$  pixels—a  $2.3\times$  improvement in spatial resolution without any software change.
2. **Decelerating → less pitch → less tilt error.** During the SEARCH phase the drone flies at  $5 \text{ m s}^{-1}$ , inducing  $\sim 10^\circ$  forward pitch (Section K.7). This pitch introduces 6.2 m of projection error before compensation. As the drone decelerates toward hover, pitch drops toward  $0^\circ$ , and the tilt error vanishes *even without the attitude compensation algorithm*—the flight manoeuvre itself eliminates the error source.
3. **Hovering → nadir camera → maximum accuracy.** In stationary hover the camera points straight down. When the target is centred in the frame, the pixel offset  $(u - C_x, v - C_y) \rightarrow 0$ , and every multiplicative error in the projection chain (GSD, focal length, yaw, lens distortion) is multiplied by zero and vanishes (Section K.8). The only remaining error is the drone’s own GPS noise, which is reduced by temporal averaging.

### K.19.2 Static Camera vs. Cooperative System

Table 61 contrasts a fixed surveillance camera with our cooperative architecture. The key difference is that a static system’s accuracy is determined at installation time, while our system’s accuracy *improves during the mission* as a direct consequence of the drone responding to detections.

**Table 61:** Comparison of static camera estimation vs. cooperative vision-flight estimation. The cooperative system exploits the drone’s ability to reposition, producing accuracy that improves over time rather than remaining fixed.

| Property                  | Static Camera             | Cooperative (Ours)                                          |
|---------------------------|---------------------------|-------------------------------------------------------------|
| Altitude                  | Fixed                     | Decreases toward target                                     |
| GSD                       | Fixed                     | Improves as drone approaches                                |
| Tilt error                | Fixed (mounting angle)    | Eliminated by hover                                         |
| Pixel offset              | Fixed (target position)   | Driven to zero by centering                                 |
| Observation count         | Accumulates over time     | Accumulates <i>and</i> improves per-obs                     |
| Accuracy trend            | $\propto 1/\sqrt{N}$ only | $\propto 1/\sqrt{N}$ <i>and</i> per-obs $\sigma \downarrow$ |
| Typical CEP <sub>50</sub> | 3 m–5 m                   | 3.5 m → 1.8 m (converges)                                   |

This cooperative property is what allows the three estimation levels (Section K.9) to achieve progressively better accuracy: Level 1 (passive) corresponds to a static camera, Level 2 (hover) adds deceleration, and Level 3 (centering) closes the full feedback loop. The four mission phases (Section K.18) are the operational realisation of this feedback loop, with each phase exploiting a different improvement mechanism. The architecture ensures that even if early

estimates are coarse ( $\sim 3.5$  m CEP), the system converges to GPS-limited accuracy ( $\sim 1.8$  m CEP) through the natural progression of approach, decelerate, and hover.

## K.20 Why Inverse-Variance Weighting Won

Among the four fusion methods evaluated, the inverse-variance method is preferred in the simulator and analysis tools for three reasons:

1. **Physics-aligned weighting.** Lower altitude means more pixels on target, smaller GSD-scaled projection errors, and reduced heading-rotation amplification. The  $1/h^2$  weight directly encodes this relationship: it is not a tuned hyperparameter but a consequence of the error model.
2. **No convergence plateau.** The Kalman filter’s covariance  $\mathbf{P}$  contracts monotonically, so later observations have diminishing impact even if they are dramatically better. The inverse-variance average, by contrast, is a weighted sum that always gives full credit to a high-weight observation regardless of how many low-weight observations preceded it.
3. **Simplicity.** The rolling weighted mean requires no matrix algebra, no tuning of process noise  $\mathbf{Q}$ , and no initialisation heuristics. It is computed in three lines of Python and is straightforward to verify.

In practice, the Kalman filter and inverse-variance method converge to similar accuracy after 20+ observations. The key advantage of inverse-variance would appear if the drone descended during investigation: the first few low-altitude observations would immediately dominate the estimate, whereas the Kalman filter takes several update cycles to “catch up” because its covariance was already small from the high-altitude phase. In the current mission implementation the drone remains at search altitude and uses direct drone GPS averaging (`--center-verify`) rather than detection-based IVW; the IVW method is used in the simulator and passive observation tools for post-hoc analysis.

## K.21 Empirical Method Comparison Against Ground Truth

The preceding sections derived error budgets analytically and compared fusion algorithms on observation streams. This subsection takes a different approach: seven fundamentally different *estimation strategies*—not just fusion filters, but entire operational procedures—were evaluated against a known ground-truth target position in SITL simulation, where the dummy coordinates are exact.

### K.21.1 Ground Truth Setup

In the ArduPilot SITL environment, the dummy is placed at a surveyed coordinate on the simulation map (`map.jpg`), and the pixel position is converted to a GPS coordinate using the same cartographic transform used to generate the map overlay. Because SITL provides perfect GPS ground truth (no receiver noise unless explicitly injected via `--gps-drift`), the estimation error can be measured to centimetre precision. For these trials, realistic noise was injected: 3 m GPS drift, 1 m barometric altitude noise,  $2^\circ$  yaw noise, and 3% FOV calibration error.

Twenty simulated missions were run for each of the seven strategies. Each mission followed the standard lawnmower search pattern over the survey area at 35 m altitude. CEP<sub>50</sub> (radius containing 50% of final estimates), CEP<sub>95</sub> (95%), and maximum error were recorded. Convergence time was measured from first detection to the point where the running CEP stabilised to within 10% of its final value.

### K.21.2 Methods Tested

Table 62 summarises the seven strategies in order of decreasing CEP<sub>50</sub>.

**Table 62:** Empirical comparison of GPS estimation strategies against SITL ground truth. Each row represents the median result across 20 simulated missions with 3 m GPS drift, 1 m altitude noise, and 3% FOV calibration error. Convergence time is measured from first detection.

| Strategy                     | CEP <sub>50</sub><br>(m) | CEP <sub>95</sub><br>(m) | Max Err.<br>(m) | Conv. Time<br>(s)           | Notes                               |
|------------------------------|--------------------------|--------------------------|-----------------|-----------------------------|-------------------------------------|
| Single flyover               | 7.2                      | 14.8                     | 22.3            | 0 (instant)                 | Roll/pitch dominated                |
| Multi-pass avg (5)           | 3.8                      | 8.1                      | 12.4            | $\sim 60$                   | Alternating headings reduce bias    |
| Hover in place               | 2.9                      | 5.4                      | 8.2             | $\sim 5$                    | Roll eliminated; calibration limits |
| Visual centering only        | 2.1                      | 4.2                      | 6.8             | $\sim 15$                   | Zero offset; GPS noise remains      |
| IVW (passive watch)          | 1.8                      | 3.8                      | 6.1             | $\sim 30$                   | Needs many obs; altitude-weighted   |
| Centering + Kalman           | 1.6                      | 3.3                      | 5.5             | $\sim 20$                   | Similar to averaging; more complex  |
| <b>Center + 10 s GPS avg</b> | <b>1.5</b>               | <b>3.1</b>               | <b>5.2</b>      | <b><math>\sim 25</math></b> | <b>Best overall; implemented</b>    |

The results confirm the analytical predictions from the error budget (Table 45) and the three-strategy comparison (Table 50). Two findings are particularly notable:

1. The gap between “single flyover” ( $\text{CEP}_{50} = 7.2 \text{ m}$ ) and “centering + GPS avg” ( $\text{CEP}_{50} = 1.5 \text{ m}$ ) represents a  $4.8\times$  improvement—achieved entirely through operational procedure, without any change to the sensor hardware or detection model.
2. The Kalman filter ( $\text{CEP}_{50} = 1.6 \text{ m}$ ) and the simple GPS average ( $\text{CEP}_{50} = 1.5 \text{ m}$ ) are statistically indistinguishable ( $p > 0.3$ , paired  $t$ -test). The extra complexity of the Kalman filter—process noise tuning, matrix algebra, initialisation heuristics—buys no measurable accuracy improvement when combined with the centering procedure.


### K.21.3 Why Centering + GPS Averaging Won

The visual centering strategy combined with 10s GPS averaging was selected for the flight implementation for five reasons:

1. **Simplest of the top performers.** The centering + GPS average requires no tuning parameters beyond the averaging duration. The Kalman filter achieves comparable accuracy but requires selection of process noise  $\mathbf{Q}$  and measurement noise scaling, both of which affect convergence behaviour.
2. **Eliminates all geometric errors by design.** When the target is at the optical centre, every multiplicative term in the projection chain (GSD, focal length, yaw, lens distortion) is multiplied by zero (Equation 22). This means the estimate is correct even if the focal length calibration is 10% wrong or the compass has a  $5^\circ$  hard-iron offset.
3. **GPS averaging is statistically optimal for zero-mean noise.** The drone’s GPS receiver noise is approximately Gaussian and zero-mean when stationary. The arithmetic mean of  $N$  independent samples is the maximum-likelihood estimator for Gaussian noise; no Kalman filter or weighted scheme can improve upon it when the underlying distribution is symmetric and stationary.
4. **Convergence time is acceptable.** The 25s convergence time (15s centering manoeuvre + 10s GPS hold) adds less than 5% to total mission time for a 10-minute search. The IVW passive approach achieves comparable accuracy but requires 30+s of accumulated flyover observations, during which the drone has moved on to other scan lines.
5. **No calibration dependency.** The single flyover, multi-pass, and hover-in-place strategies all require accurate focal length, compass heading, and barometric altitude to project pixel offsets into GPS coordinates. A 5% focal length error at 35 m produces a 1.6 m bias at the frame edge—comparable to the GPS noise floor. The centering strategy bypasses this entire computation, making the estimate robust to calibration drift during flight.

### K.21.4 $\text{CEP}_{50}$ Comparison


Figure 37 shows the  $\text{CEP}_{50}$  for each estimation strategy, ordered by accuracy.



**Figure 37:**  $\text{CEP}_{50}$  comparison across seven GPS estimation strategies (lower is better). The centering + GPS averaging approach achieves 1.5 m—a  $4.8\times$  improvement over a single flyover estimate. Error bars omitted for clarity; inter-trial variation was  $<0.3 \text{ m}$  for all strategies except “single flyover” ( $\pm 1.1 \text{ m}$ ).

### K.21.5 Convergence Behaviour

Figure 38 shows how the position error evolves over time for the three best-performing strategies. The centering approach achieves lower error faster than IVW, and the addition of GPS averaging drives the final error below both alternatives.



**Figure 38:** Convergence of position error over time for the three best estimation strategies. All three begin with the same initial flyover estimate ( $\sim 7$  m). Visual centering converges rapidly to 2.1 m (GPS noise floor for single fixes) by  $t = 15$  s. GPS averaging during the subsequent 10 s hold reduces the error further to 1.5 m. IVW converges more slowly because it relies on accumulating many flyover observations rather than actively centering.

The convergence plot reveals two distinct phases in the centering strategy: (i) a rapid geometric convergence (0–15 s) as the drone manoeuvres above the target and eliminates roll/pitch, calibration, and timing errors; and (ii) a slower statistical convergence (15–25 s) as GPS averaging reduces the residual receiver noise. The IVW approach lacks phase (i) because the drone continues flying; its convergence is purely statistical, driven by the  $1/\sqrt{N}$  law applied to passively collected observations.

## K.22 Offset Landing Application

The GPS estimate feeds directly into the offset landing calculation. The system computes a landing point 7.5 m north of the estimated target position (configurable) to avoid rotor downwash on the casualty. Given a typical fused estimate error of 1 m to 2 m, the probability that the drone lands within 10 m of the actual target (and therefore within visual confirmation range) exceeds 99%. The offset distance was chosen as  $3\times$  the worst-case single-observation CEP, ensuring a safe margin even if the estimation converges poorly.

With the `--center-verify` drone GPS averaging ( $N = 50$  samples over 10 s,  $N_{\text{eff}} \approx 20$ ), the expected landing accuracy is approximately  $\sigma/\sqrt{N_{\text{eff}}} \approx 2.5/\sqrt{20} \approx 0.6$  m, where  $N_{\text{eff}}$  accounts for temporal correlation in GPS noise at 5 Hz. In the simulator and passive observation tools, the inverse-variance weighted average across multiple detection-based estimates yields comparable accuracy when sufficient observations are available.

## K.23 Multi-Pass Observation and Heading Diversity

Sections K.6 and K.7 noted that the lawnmower pattern’s alternating headings partially cancel directional biases. This subsection develops that observation quantitatively and shows that the boustrophedon (zigzag) search pattern is not merely a convenient coverage strategy but an inherent source of heading diversity that improves GPS estimation quality.


### K.23.1 The Lawnmower Pattern Advantage

Two dominant systematic error sources—GPS timing lag and roll/pitch displacement—share a critical property: they bias the GPS estimate *in the direction of travel*. GPS timing lag displaces the estimate forward by  $v \cdot \Delta t$  (Section K.6), while forward-flight pitch displaces the camera footprint forward by  $h \tan \theta$  (Section K.7). Both biases reverse sign when the drone reverses its heading.

In the boustrophedon pattern, consecutive scan lines are flown in opposite directions. At the Bristol field site, scan lines run at headings of approximately  $155^\circ$  (south-eastbound) and  $335^\circ$  (north-westbound). Consider a target detected on both passes:

- **Pass 1** (heading  $155^\circ$ ): GPS lag biases the estimate  $\sim 1$  m along  $155^\circ$ ; pitch bias adds  $\sim 6.2$  m along  $155^\circ$ . Total systematic offset:  $\sim 7.2$  m south-east of the true position.
- **Pass 2** (heading  $335^\circ$ ): both biases reverse direction, displacing the estimate  $\sim 7.2$  m north-west of the true position.
- **Average of Passes 1+2:** the two opposing systematic offsets cancel, leaving only the random GPS receiver noise and minor residuals from asymmetric wind gusts or speed differences between passes.

Figure 39 illustrates this cancellation geometrically.



**Figure 39:** Multi-pass bias cancellation. Pass 1 (heading  $155^\circ$ , blue) biases the GPS estimate south-east; Pass 2 (heading  $335^\circ$ , red) biases it north-west by the same magnitude. The average (green) falls near the true target position (star), because the systematic GPS lag and pitch biases cancel between opposing headings. Only random GPS receiver noise remains.

### K.23.2 How Many Passes Are Enough?

The total position error after  $N$  passes combines two components: the residual systematic bias (which depends on heading diversity) and the random noise (which decreases as  $1/\sqrt{N}$ ). The GPS lag bias for a single pass is approximately  $v \cdot \Delta t \approx 1 \text{ m}$  at  $5 \text{ m s}^{-1}$ ; the pitch bias is  $h \tan \theta \approx 3 \text{ m}$  at  $5^\circ$  (conservative). Both are directional and reverse with heading. The random component is dominated by GPS receiver noise ( $\sigma_{\text{GPS}} \approx 2.5 \text{ m}$  per observation).

For  $N$  passes with varying heading configurations, the expected CEP<sub>50</sub> is:

$$\text{CEP}_{50} \approx \sqrt{\left(\frac{\bar{b}}{N}\right)^2 + \frac{\sigma_{\text{GPS}}^2}{N}} \quad (36)$$

where  $\bar{b}$  is the mean residual bias after heading cancellation. For perfectly opposing headings,  $\bar{b} = 0$ ; for random headings,  $\bar{b}$  decreases as  $1/\sqrt{N}$ . Table 63 shows the expected accuracy for different numbers of passes.

The most significant accuracy improvement occurs between one and two passes: the second pass with an opposing heading eliminates the systematic bias entirely, and the  $\sqrt{2}$  averaging improvement reduces random noise by 29%. Beyond two passes, gains follow the slower  $1/\sqrt{N}$  law.

### K.23.3 Strip Spacing and Target Visibility

The number of observations per pass depends on the camera footprint, search speed, and detection frame rate. At the nominal search altitude of 35 m:

- **Camera footprint:**  $32.2 \times 24.0 \text{ m}$  (Table 44).
- **Strip spacing:** with 20% overlap, the lawnmower spacing is  $0.8 \times 24.0 \approx 19 \text{ m}$  (along the cross-track direction) or  $0.8 \times 32.2 \approx 26 \text{ m}$  (along the flight direction, depending on pattern orientation).
- **Along-track visibility:** a stationary target remains within the 24 m along-track footprint for  $24/5 = 4.8 \text{ s}$  at  $5 \text{ m s}^{-1}$ .
- **Frames per flyover:** at 4.8 FPS (TFLite inference rate on Pi 5), this yields  $4.8 \times 4.8 \approx 23$  frames per pass—though detection only fires when the target is within the frame and above the confidence threshold, so the effective count is approximately 11–17 detections per pass.

With a 10-strip lawnmower pattern, most of the search area receives 1–2 passes (the target may fall within the overlap region of adjacent strips). The typical target therefore generates 11–34 observations across 1–2 passes, providing sufficient data for the inverse-variance weighted average to converge to sub-2 m accuracy even before the centering procedure begins.

### K.23.4 Multi-Pass Averaging vs. Centering

Despite the theoretical effectiveness of multi-pass averaging, the mission pipeline does *not* rely on it as the primary estimation strategy for three reasons:

1. **First detection triggers CENTERING.** The state machine transitions from SEARCH to CENTERING on the first confirmed detection (or after the SMART consecutive-frame filter confirms the target, if `--smart-detect` is enabled). This means the drone rarely completes a second pass over the same target—the centering manoeuvre begins immediately.
2. **Centering provides a fundamentally better estimate.** As shown in Section K.9.4, centering the target at the optical axis eliminates all calibration-dependent errors by driving the pixel offset to zero. Multi-pass averaging can cancel systematic biases but cannot eliminate random calibration errors (focal length, compass, barometric drift). Centering removes these errors entirely.
3. **Spatial clustering limitations.** The spatial clustering algorithm (Section K.14) in `passive_watch.py` correctly merges multiple observations of the same target within a 30 m radius. However, in the mission state machine (`main.py`), the first detection event triggers an immediate state transition—multiple passes would create separate queue entries for the same target, but the state machine processes the first entry and never returns to the search pattern for that target.

Multi-pass averaging therefore serves as a *theoretical safety net*: if the centering procedure were to fail (e.g. due to loss of visual lock during the approach), the system could fall back to the accumulated flyover estimates. In the passive observation tools (`passive_watch.py`, `simple_simulator.py`), where no centering manoeuvre occurs, multi-pass averaging with heading diversity is the *primary* estimation strategy and achieves the accuracies shown in Table 63. The lawnmower pattern thus serves a dual purpose: maximising spatial coverage for detection probability, and providing heading diversity that improves estimation accuracy through systematic bias cancellation.

## K.24 Passive Estimation: Level 1 GPS Without Drone Response

The preceding sections analysed estimation strategies in the context of the autonomous mission, where the drone actively centres on the target and hovers. However, the estimation toolkit implemented in `passive_watch.py` represents a fundamentally different operational mode: *passive estimation*, in which the drone continues its pre-planned flight path and sends zero control commands, yet still produces a usable target GPS estimate purely from accumulated flyover observations. This “Level 1” capability—estimation without response—is valuable both as an independent operational mode and as a fallback when active centering is unavailable.

### K.24.1 Passive Estimation Toolkit

The `DummyEstimator` class in `passive_watch.py` implements the inverse-variance weighting scheme described in Section K.13.4, operating as a cumulative weighted average over all detections. Each detection frame triggers the following pipeline:

1. **Pixel-to-GPS projection:** the detection bounding-box centre  $(u, v)$  is projected to a ground-plane GPS estimate using the same GSD-based projection described in Section K.4, with the drone’s current GPS position, altitude, and heading as inputs.
2. **Composite weight calculation:** each observation receives a weight  $w_i = w_{\text{cen}} \cdot w_{\text{alt}}$ , combining a centrality factor (Equation 32, with a tighter  $0.3 \cdot d_{\text{max}}$  radius for the bonus zone) and an altitude factor  $w_{\text{alt}} = 1/h_i^2$  (equivalent to Equation 34 with  $h_{\text{ref}} = 1$  m). At 10 m altitude with a centred detection, the weight is  $\sim 45\times$  that of a 35 m edge detection.
3. **Running weighted accumulation:** the weighted latitude and longitude sums are updated incrementally (no observation history stored), making the memory footprint  $O(1)$  regardless of mission duration.
4. **Spatial clustering:** new estimates are compared against existing clusters using the Haversine distance (Section K.14). If the nearest cluster centroid is within 30 m, the observation joins that cluster; otherwise a new cluster is spawned. Each cluster maintains its own `DummyEstimator` instance, enabling simultaneous tracking of multiple targets.

In addition to the IVW accumulator, the `SmartEstimator` class implements a consensus-based lock mechanism: once a configurable number of estimates (default 5) cluster within a maximum spread threshold (default 1 m), the estimator declares a “lock” and reports the median latitude and longitude of the tightest cluster as the final target position. This greedy tightest-cluster algorithm is robust to outliers because it selects the densest subset of observations rather than averaging all of them.

## K.24.2 Why IVW Suits Passive Mode

The four fusion algorithms described in Section K.13 were all evaluated in the passive context. IVW emerged as the preferred method for passive estimation for reasons that are distinct from—and complementary to—the reasons it was preferred in the active mission (Section K.20):

1. **Natural downweighting of edge detections.** In passive mode the drone does not manoeuvre to centre the target, so most detections occur at the frame periphery as the target drifts through the field of view. Edge detections carry larger projection errors (GSD $\times$ pixel offset is maximised), and the centrality weight  $w_{\text{cen}}$  correctly assigns them lower influence. The rolling and cumulative averages treat all pixel positions equally, allowing peripheral noise to contaminate the estimate.
2. **Altitude weighting without tuning.** The  $1/h^2$  weight is derived directly from the GSD error model—it is not a tuned hyperparameter. The Kalman filter achieves a similar effect through its measurement noise covariance  $\mathbf{R}_k \propto 1/w_i$ , but requires selection of the process noise  $\mathbf{Q}$  and initialisation covariance  $\mathbf{P}_0$ , both of which affect convergence behaviour and must be tuned empirically.
3. **No convergence saturation.** During a multi-pass search, the drone may descend for a closer look at an unrelated area before returning to search altitude. If it passes over the target during a low-altitude transit, the Kalman filter’s already-contracted covariance  $\mathbf{P}$  limits the Kalman gain, preventing the high-quality low-altitude observation from having its full impact. The IVW average gives full credit to the low-altitude observation regardless of prior history.

## K.24.3 Convergence Behaviour in Passive Mode

In simulation with realistic noise (3 m GPS drift, 1 m altitude noise,  $2^\circ$  yaw noise), the IVW passive estimate converges as follows:

- **5 observations** ( $\sim 1$  pass):  $\text{CEP}_{50} \approx 4.5$  m. Systematic bias from a single heading dominates.
- **15–20 observations** ( $\sim 2$  passes with opposing headings):  $\text{CEP}_{50} \approx 2.0$  m. Heading bias cancellation (Section K.23) removes the dominant systematic error; IVW downweighting of edge detections reduces the random component.
- **30+ observations** ( $\sim 3$ – $5$  passes):  $\text{CEP}_{50} \approx 1.8$  m. Diminishing returns as the estimate approaches the GPS receiver noise floor.

The convergence is slower than the active centering strategy (which reaches  $\text{CEP}_{50} = 1.5$  m in  $\sim 25$  s) because the passive mode lacks the geometric error elimination that centering provides (Section K.9.4). However, the passive estimate requires no interruption of the search pattern and no state machine transitions.

## K.24.4 Passive vs. Active Estimation: When Each Mode Excels

Table 64 compares the two estimation modes across operational metrics. The comparison reveals that passive and active estimation are not competing alternatives but complementary capabilities suited to different mission phases and failure modes.

**Table 64:** Comparison of passive (flyover-only) and active (centering + hover) estimation modes. Passive mode uses IVW fusion over accumulated flyover detections; active mode uses visual centering followed by drone GPS averaging.

| Metric                        | Passive IVW                 | Active Centering            |
|-------------------------------|-----------------------------|-----------------------------|
| $\text{CEP}_{50}$ (converged) | $\sim 1.8$ m                | $\sim 1.5$ m                |
| Time to convergence           | $\sim 60$ s (3–5 passes)    | $\sim 25$ s                 |
| Calibration dependency        | Yes (focal length, compass) | No (zero-offset projection) |
| Interrupts search pattern     | No                          | Yes (15 s manoeuvre)        |
| Works without drone response  | Yes                         | No                          |
| Handles multiple targets      | Yes (spatial clustering)    | One at a time               |
| Fallback if centering fails   | Yes (accumulated estimate)  | N/A                         |
| Memory footprint              | $O(1)$ per cluster          | $O(N)$ GPS samples          |

Three operational scenarios favour passive estimation over active centering:

1. **Multi-target reconnaissance.** In a multi-casualty scenario, the drone may detect several targets during the search phase. Active centering investigates one target at a time, each consuming  $\sim 25$  s of mission time. Passive estimation tracks all detected clusters simultaneously, providing GPS estimates for every target by the end of the search—even targets the drone never actively investigated. The spatial clustering algorithm (Section K.14) maintains independent `DummyEstimator` instances per cluster, so estimates for different targets do not interfere.

2. **Equipment degradation.** If the centering algorithm loses visual lock during approach (e.g. due to sun glare, motion blur, or a false negative streak), the passive estimate accumulated during the search phase provides an immediate fallback position. The mission state machine can use this estimate for an offset landing without repeating the search. In the current implementation, this fallback is not yet automated—the operator would need to command a manual landing—but the passive estimate is always available through the web dashboard.
3. **Covert observation.** In scenarios where the drone should not reveal interest in a specific target (e.g. security applications, wildlife monitoring), passive estimation allows the drone to maintain its pre-planned flight path while accumulating a target position estimate. The target sees a drone flying a routine pattern; the operator sees the GPS estimate converging on their dashboard.

#### K.24.5 Implementation in `passive_watch.py`

The passive estimation mode is implemented as a standalone tool that sends zero commands to the flight controller. It reads GPS telemetry from mavproxy (read-only), runs the camera and AI detection pipeline, and streams results to a web dashboard at `http://PI_IP:8090/`. The operator sees:

- A live MJPEG video stream with detection bounding boxes overlaid.
- A bullseye scatter plot showing all GPS estimates colour-coded by weight, with the IVW fused estimate marked as a crosshair.
- A satellite map panel showing the SSSI no-fly zone (red polygon), the search area (yellow polygon), and the estimated target position (magenta star).
- Detection count, saved snapshot count, and cluster statistics.

The tool can be run during manual RC flight (pilot flies, Pi watches) or during an autonomous mission (as a parallel observer). Because it sends zero commands, it is safe to run at any time without risk of interfering with the flight controller. This makes it the recommended first step in the progressive flight testing sequence (Step 3 in Section O.2): manual flight with passive detection establishes that the CV pipeline works in real outdoor conditions before any autonomous commands are attempted.

## L Progressive Testing Methodology


This appendix documents the five-tier verification framework that governed all testing throughout the project, from desktop simulation through to autonomous flight. The framework is motivated by a fundamental asymmetry in autonomous UAV development: a logic error that manifests as a misplaced pixel in simulation can, on real hardware, command a 2kg aircraft into an unrecoverable trajectory. The cost of discovering a defect rises steeply with integration level—a sign-convention bug caught in simulation costs minutes; the same bug discovered during autonomous flight costs days of rework, potential hardware damage, and a lost flight slot that a five-person university team cannot easily reschedule. The following subsections detail the philosophy, tier definitions, test script organisation, defect discovery data, V-model mapping, and cost-risk gradient that together form the project’s verification and validation strategy.

### L.1 Philosophy: Never Skip a Tier

The governing principle is borrowed from safety-critical avionics: *each verification tier gates the next, and no gate may be bypassed*. A failure at Tier 2 (bench) blocks all flight testing until the root cause is resolved, regardless of schedule pressure. This discipline is counterintuitive on a time-constrained student project, where the temptation to “just try it in the air” is strongest precisely when time is shortest. The justification is empirical: of the 15 defects discovered during this project’s verification campaign, 12 were caught at Tiers 1–3 (Table 66), where the turnaround time from discovery to fix was under one hour. Had those defects propagated to flight, each would have consumed an entire field session.

The framework also enforces a separation of concerns. Each tier tests a *specific integration boundary*: Tier 1 tests software logic; Tier 2 tests the software–hardware interface; Tier 3 tests the hardware–environment interface; Tier 4 tests the perception pipeline under operational conditions; Tier 5 tests closed-loop autonomy. This decomposition ensures that when a failure occurs, its root cause is immediately localisable to the boundary most recently crossed.

## L.2 The Five Tiers in Detail



**Figure 40:** Tier progression. Each arrow represents a gate: the downstream tier is blocked until all upstream tests pass.

### L.2.1 Tier 1: SITL Simulation (Zero Hardware Risk)

ArduPilot’s Software-in-the-Loop (SITL) emulator runs a full flight dynamics model on the development laptop, accepting the same MAVLink commands as the real Cube Orange autopilot. The mission orchestrator, state machine, detection model, geofence logic, and path planner execute unmodified against this simulated vehicle. A synthetic camera view is generated by extracting a field-of-view-correct crop from a georeferenced satellite image (`map.jpg`), enabling the complete detect–estimate–land pipeline to execute without any physical hardware.

At this tier, 22 structured simulation tests exercise every operational mode: baseline autonomous mission, multi-frame detection confirmation, geofence enforcement, manual override, multi-target queue ordering, NFZ speed ramping, beacon redirect, transit planning, altitude override, dry-run generation, headless mode, and 11 further scenarios. The interactive simulator (`simple_simulator.py`) adds keyboard-driven flight with configurable GPS drift, camera shake, and inference throttle, enabling rapid iteration on the GPS estimation and Kalman filtering logic.

A separate `--dry-run` mode walks the state machine through all transitions without connecting to SITL, printing the full waypoint set and generating a pattern visualisation. This mode executes in under one second and is run before every code commit that modifies state logic or path planning.

**Defects caught.** State machine logic errors (incorrect transition guards, missing timeout handlers), geofence sign-convention bug (`cv2.pointPolygonTest` returns positive for *inside*, opposite to the initial assumption), repulsive-force direction inversion at NFZ boundary, barometric drift causing infinite landing loops (resolved with a 90 s absolute timeout), and duplicate-detection queueing of the same target from consecutive frames (resolved with a 5 m reject radius).

### L.2.2 Tier 2: Bench Test (Pi + Cube + Camera, No Props)

The Raspberry Pi 5, Cube Orange, and IMX296 global-shutter camera are assembled on the bench with propellers removed. Three numbered scripts test with increasing integration:

- 0a **Cube commands** — sends individual MAVLink commands (mode switch, parameter read, arm pre-check) and verifies each `COMMAND_ACK`. Confirms the Pi–Cube serial link and mavproxy UDP bridge operate at 921 600 baud.
- 0b **Bench mission** — executes the full command sequence (arm, takeoff, waypoint navigation, land) and logs every `ACK`. On the bench, the flight controller accepts these commands but the vehicle does not move, verifying that the command pipeline is correct end-to-end.
- 0c **Feedback test** — runs the vision-to-GPS pipeline while the operator walks the drone past a printed dummy. The camera captures frames, the TFLite model runs inference, and the coordinate transform produces GPS estimates that are compared against hand-measured ground truth. This validates the entire sensing chain—camera, colour space, lens undistortion, inference, and projection—on the target hardware.

Additional bench scripts include a 50-run inference benchmark (`benchmark.py`), a comprehensive multi-model benchmark with system profiling (`benchmark_full.py`), GPS satellite tracking (`gps_test.py`), and a lens distortion calibration using a checkerboard target (`lens_calibrate.py`).

**Defects caught.** BGR/RGB colour-space mismatch (the IMX296 outputs BGR data despite picamera2’s `RGB888` label, causing inverted colour channels that degraded detection confidence from 0.97 to 0.3), FOV miscalibration (focal length corrected from 7.0 mm to 5.46 mm after ruler-based measurement), Python 3.13 compatibility breakage (pyserial’s serial read is broken on 3.13, requiring a mavproxy UDP bridge as a workaround; `tflite-runtime` package replaced by `ai-edge-litert`), and camera-in-use resource conflicts when multiple scripts attempt concurrent access.

### L.2.3 Tier 3: Passive Flight (Pilot Flies, Pi Watches, Zero Commands)

The drone is flown manually by the RC pilot over the search area containing a casualty dummy. The Pi runs the full vision pipeline—camera capture at  $1456 \times 1088$ , YOLOv8n inference via TFLite, lens undistortion, GPS coordinate estimation—and streams the annotated video to a browser-based ground station. Critically, the Pi sends *zero flight commands* to the Cube. All detections are logged with timestamp, pixel coordinates, bounding box dimensions, confidence, barometric altitude, GPS position, heading, and estimated target GPS.

Post-flight analysis of these logs establishes:

- **Detection altitude ceiling** — the maximum altitude at which the model reliably detects the dummy (target: 30 m, to be confirmed).
- **False-positive rate** — detections triggered by terrain features, shadows, or equipment under real lighting and motion conditions.
- **GPS estimation accuracy** — scatter of estimated target positions versus surveyed ground truth, quantified as CEP50 and CEP95.
- **Inference latency under load** — confirmed at 208 ms (4.8 FPS) on Pi 5 with XNNPACK, including 1.5 ms for lens undistortion.

This tier is the most valuable of the five: it generates real-world performance data with no risk of autonomous misbehaviour. The six experiment scripts in `tests/day_1_experiments/` (Section L.4) are designed specifically for this tier, each producing a CSV file suitable for direct statistical analysis.

### L.2.4 Tier 4: Autonomous Search with Log-Only Detection

The full mission orchestrator runs autonomously—the drone flies the lawnmower search pattern under GUIDED mode control—but the detection pipeline operates in log-only mode. When the model identifies a candidate, the system records the detection and annotates the ground station display but does *not* transition to the CENTERING or DESCENDING states. The drone completes the search pattern and returns to launch.

Post-flight review of the detection log answers the operational question: *would the system have found the target, and would it have acted on the correct detection?* False positives, missed detections, and latency-induced position errors are all visible in the log without having been acted upon. Configuration parameters (confidence threshold, multi-frame confirmation count, reject radius) are tuned at this tier before enabling closed-loop autonomy.

### L.2.5 Tier 5: Full Autonomous Mission

All systems are enabled. The drone executes the complete state machine: INIT, CONNECTING, ARMING, TAKEOFF, SEARCH (lawnmower pattern), CENTERING (visual servo onto target), VERIFY (operator confirms or rejects via ground station), APPROACH (offset landing approach), LANDING, and DONE. A DESCENDING state for altitude reduction before VERIFY exists in the code but is currently bypassed; the drone transitions directly from CENTERING to VERIFY at search altitude. The operator retains veto authority at the VERIFY gate and the RC kill switch provides a hardware-level override independent of all software.

## L.3 Test Script Organisation

Table 65 maps the six test script directories to the verification tiers they support, with representative scripts and total counts.

**Table 65:** Test script categories and their tier coverage.

| Category     | Directory                             | Scripts   | Tiers Served   |
|--------------|---------------------------------------|-----------|----------------|
| Hardware     | <code>tests/hardware/</code>          | 8         | 2 (bench)      |
| Flight       | <code>tests/flight/</code>            | 14        | 2, 3, 4, 5     |
| Diagnostics  | <code>tests/diagnostics/</code>       | 5         | 2, 3, 4, 5     |
| Calibration  | <code>tests/calibration/</code>       | 7         | 2, 3           |
| Experiments  | <code>tests/day_1_experiments/</code> | 6         | 3, 4           |
| Laptop       | <code>tests/laptop/</code>            | 21        | 1 (simulation) |
| Unit         | <code>tests/unit/</code>              | 6         | 1              |
| Automated    | <code>tests/automated/</code>         | 7         | 1              |
| <b>Total</b> |                                       | <b>74</b> |                |

The flight scripts are numbered by risk level (0a through 5), establishing an explicit ordering that prevents a higher-risk test from being executed before its prerequisites have passed. All experiment scripts issue zero flight commands and are safe to execute at any time during manual or autonomous flight.

## L.4 Experiment Scripts for Quantitative Validation

Six experiment scripts are designed for execution during Tier 3 passive flights, each producing a timestamped CSV file with standardised column headers for post-flight analysis:

1. **Altitude sweep** (`altitude_sweep.py`) — the pilot hovers above the dummy at 10, 15, 20, 25, and 30 m for 30 s each. The script bins every frame by barometric altitude into 5 m buckets and computes detection rate, mean confidence, and expected dummy pixel size per bucket. Output directly determines the safe `TARGET_ALT` parameter.
2. **Speed sweep** (`speed_sweep.py`) — the pilot flies repeated passes over the dummy at increasing ground speeds. The script correlates detection rate and confidence with ground speed and a Laplacian-based blur metric, establishing the maximum search speed at which the model maintains acceptable detection performance.
3. **GPS accuracy** (`gps_accuracy.py`) — compares the pixel-to-GPS coordinate estimate against the surveyed ground-truth position of the dummy, logging the error in metres for every detection. Computes an inverse-variance-weighted average across all observations and reports CEP50, CEP95, and maximum error.
4. **GPS drift** (`gps_drift.py`) — the drone hovers stationary above a known point. The script logs raw GPS positions at 10 Hz and computes the noise floor: CEP50, CEP95, and maximum deviation, establishing the baseline GPS uncertainty that the target estimation algorithm must overcome.
5. **Model comparison** (`model_compare.py`) — benchmarks all available `.tflite` models (custom SAR detector, retrained v2, COCO person detector) on identical live frames, reporting inference time, detection rate, and confidence for each. Informs model selection for the mission.
6. **Detection log** (`detection_log.py`) — a general-purpose catch-all logger that records every frame’s detection result, telemetry, and metadata. Serves as the raw data source for any post-hoc analysis not covered by the specialised scripts.

## L.5 Defect Discovery by Tier

Table 66 lists representative defects, the tier at which they were discovered, the resolution time, and the estimated cost had they propagated to flight.

**Table 66:** Defects discovered at each testing tier, with resolution time and estimated cost if undetected.

| Defect                                                                               | Tier | Fix Time | Cost if Missed                                        |
|--------------------------------------------------------------------------------------|------|----------|-------------------------------------------------------|
| Geofence sign convention inverted ( <code>pointPolygonTest</code> positive = inside) | 1    | 15 min   | SSSI incursion; regulatory violation                  |
| Repulsive force direction reversed at NFZ boundary                                   | 1    | 20 min   | Drone pushed <i>into</i> NFZ                          |
| Barometric drift causes infinite landing loop                                        | 1    | 30 min   | Battery depletion in hover                            |
| Duplicate target queued from consecutive frames                                      | 1    | 10 min   | Repeated descent on same target                       |
| State transition missing timeout guard                                               | 1    | 15 min   | Stuck in <code>CENTERING</code> indefinitely          |
| BGR/RGB colour inversion (IMX296)                                                    | 2    | 45 min   | Detection confidence drops to 0.3; mission failure    |
| FOV miscalibration (7.0 vs 5.46 mm)                                                  | 2    | 1 hr     | GPS estimates 28% off; landing outside 10 m radius    |
| Python 3.13 pyserial breakage                                                        | 2    | 2 hr     | No Pi-Cube communication                              |
| <code>tflite-runtime</code> unavailable on 3.13                                      | 2    | 30 min   | No inference on Pi                                    |
| Camera resource conflict (concurrent access)                                         | 2    | 20 min   | Stream or detection crashes on startup                |
| GPS timing lag (100–200 ms)                                                          | 3*   | N/A      | 0.8–1.6 m systematic offset at 8 m/s                  |
| <code>self.investigating</code> uninitialised                                        | 2    | 5 min    | Crash on first detection in <code>pi_flight.py</code> |

\*Discovered via DJI video analysis, a proxy for Tier 3.

The table demonstrates the intended cost gradient: Tier 1 defects (simulation) were resolved in 10–30 minutes with zero hardware involvement. Tier 2 defects (bench) required 20 minutes to 2 hours but were discovered on the ground with propellers removed. The GPS timing lag, identified through video analysis of a DJI survey flight (a Tier 3 proxy), would have caused a systematic 1 m landing offset—detectable only under real flight dynamics.

## L.6 V-Model Mapping

The five-tier framework maps onto the classical V-model [29] as follows:

**Table 67:** Mapping between V-model verification levels and the progressive testing tiers.

| V-Model Level         | Tier | Verification Activity                                                          |
|-----------------------|------|--------------------------------------------------------------------------------|
| Requirements analysis | —    | Mission brief (R01–R12) + KML zones                                            |
| System design         | —    | State machine, architecture, module interfaces                                 |
| Subsystem design      | —    | <code>vision.py</code> , <code>planning.py</code> , <code>utils.py</code> APIs |
| Implementation        | —    | Source code + unit tests                                                       |
| Unit testing          | 1    | Unit tests ( <code>tests/unit/</code> , 6 scripts), TFLite model validation    |
| Integration testing   | 1, 2 | SITL end-to-end, bench command pipeline                                        |
| System testing        | 3, 4 | Passive flight data collection, log-only autonomous                            |
| Acceptance testing    | 5    | Full autonomous mission with operator verification                             |

The left side of the V (decomposition) was completed during the design phase: requirements were traced from the project brief through the system architecture to individual module interfaces. The right side (recomposition and verification) is realised by the five tiers, with each tier verifying a progressively higher level of integration. The critical addition relative to a textbook V-model is the separation of system testing into *passive* (Tier 3) and *log-only autonomous* (Tier 4), which inserts two low-risk verification stages between integration testing and acceptance testing—stages that a strict V-model would collapse into a single system test phase.

## L.7 Cost–Risk Gradient

Table 68 quantifies the cost and risk at each tier, reinforcing the economic argument for catching defects early.

**Table 68:** Cost and risk gradient across testing tiers.

| Tier | Environment      | Hardware at Risk  | Fix Turnaround | Iteration Cost       | People Required               |
|------|------------------|-------------------|----------------|----------------------|-------------------------------|
| 1    | Laptop (SITL)    | None              | Minutes        | Electricity          | 1 (developer)                 |
| 2    | Bench (no props) | None              | 30 min–2 hr    | Travel to lab        | 1–2                           |
| 3    | Passive flight   | Airframe          | Hours          | Flight slot          | 3+ (pilot, operator, spotter) |
| 4    | Auto log-only    | Airframe          | Hours–day      | Flight slot          | 3+                            |
| 5    | Full autonomous  | Airframe + target | Day+           | Flight slot + repair | 3+                            |

At Tier 1, a developer working alone can execute hundreds of test iterations per day. At Tier 5, a single iteration requires a minimum crew of three (pilot, operator, safety observer), a reserved outdoor site, favourable weather, charged batteries, and an assembled airframe—resources that are available for at most two sessions per week in a university setting. The 50:1 ratio in iteration throughput between Tier 1 and Tier 5 makes the case for aggressive front-loading of verification effort at the simulation level.

## L.8 Summary

The progressive methodology treated testing not as a phase that follows development, but as a continuous, tiered process that shapes development decisions. The five-tier gating system, supported by 74 dedicated scripts across eight categories (Table 65), including 127 unit tests and 33 automated integration tests, ensured that every integration boundary was verified before the system was permitted to cross it. Of the 15 significant defects discovered, 80% were caught at Tiers 1 and 2 where the fix cost was measured in minutes, not field sessions. The structured experiment scripts prepared for Tier 3 passive flights provide a repeatable, quantitative basis for tuning configuration parameters before autonomy is enabled—replacing the ad-hoc “fly and see” approach with a data-driven verification pipeline.

## M Field Day: Adaptation Under Uncertainty

This appendix presents a narrative account of the project’s first field day on 11 March 2026, where weather cancellation forced the team to pivot from flight testing to intensive bench calibration. The account is included not as an anecdote but as an empirical case study in progressive testing methodology: it demonstrates that each verification tier produces

independently valuable data, that calibration errors invisible to simulation can dominate real-world accuracy, and that a well-prepared fallback programme can extract maximum information from a session that produces zero flight time. The concrete outputs—a 22% focal length correction, lens calibration, inference benchmarks, and a colour-channel fix—each improved the system’s subsequent performance in ways that would have been difficult to achieve under the time pressure of an active flight session.

On 11 March 2026 the team arrived at the designated test site with the full hardware stack assembled for the first time outside the lab: Raspberry Pi 5, Cube Orange autopilot, IMX296 global-shutter camera, and the airframe with propellers fitted. The objective was to progress through Tiers 2–5 of the progressive testing framework (Section L)—bench verification, passive flight with manual RC, and, conditions permitting, first autonomous waypoint flight. What followed was a field day defined not by flying, but by the calibration data and integration insights that only become visible when the system meets the real world.

## M.1 Setup: Three Terminals, One Aircraft

All interaction with the Pi was conducted remotely from a laptop via SSH (PuTTY), supplemented by the Pi’s own screen for scripts requiring a display. The three-terminal workflow—established within the first twenty minutes and used for the remainder of the day—became the standard operating procedure for all subsequent field sessions:

**Terminal 1 — mavproxy:** Launched first, always. The MAVLink router forwarded Cube telemetry from the UART link (`/dev/ttyAMA0` at 921 600 baud) to two UDP outputs: port 14550 for flight and detection scripts, port 14551 for the live diagnostics dashboard. A TCP bridge on port 5762 allowed Mission Planner on the laptop to connect simultaneously, providing a familiar GCS interface for the safety pilot. The startup sequence was scripted into a single copy-paste block to eliminate typos under time pressure.

**Terminal 2 — diagnostics:** The multi-panel diagnostics script displayed camera frame rate, AI model status, Cube heartbeat, GPS satellite count, and battery voltage on the Pi’s screen. This terminal served as the go/no-go dashboard: all four indicators had to show green before any test script was permitted to run.

**Terminal 3 — scripts:** Calibration, benchmark, and detection scripts were executed sequentially from PuTTY. Browser-based camera streams (`http://PI_IP:8090`) provided live video to the laptop without requiring a display server on the Pi, confirming that the headless architecture designed in simulation translated directly to field use.

Two integration issues surfaced immediately. First, a camera-in-use conflict arose when two scripts attempted to open the CSI interface simultaneously—`picamera2` does not permit shared access to the sensor. The fix was procedural: a strict single-script-at-a-time rule, enforced by a `sudo pkill -f libcamera` guard before each script launch. Second, a port-binding collision left port 14550 occupied by a zombie `mavproxy` process from a previous session. Adding a `sudo pkill -f mavproxy; sleep 2` preamble to the startup sequence resolved this permanently. Neither issue required code changes; both were captured in the field quick-reference sheet for future sessions.

## M.2 What We Tested

With connectivity established, the team worked through the bench-test checklist in order:

1. **Camera verification.** The Pi camera delivered 640×480 frames at the expected rate. The MJPEG stream to the laptop browser confirmed correct colour rendering—a non-trivial validation, given the colour-channel discovery described below.
2. **Cube heartbeat and telemetry.** The diagnostics dashboard confirmed a stable MAVLink heartbeat, GPS satellite count (initially 0 indoors, rising to 7 near a window), and battery voltage within limits. Mode changes issued from Mission Planner propagated to the Pi’s telemetry stream within one heartbeat period.
3. **AI inference.** The TFLite model loaded successfully on the Pi 5 and produced detections on a test image within 210 ms—consistent with Docker-based estimates from Tier 2 testing.
4. **FOV calibration.** See Section M.3.
5. **Lens distortion calibration.** See Section M.5.
6. **Full benchmark.** See Section M.6.

## M.3 FOV Calibration: Catching a 22% Error

The most consequential discovery of the field day was a systematic error in the focal length parameter. The configuration file had carried a default value of  $f = 7.0$  mm since the earliest development phase, inherited from a datasheet approximation for the IMX296 lens assembly. No simulation or Docker test could have caught this: the parameter only matters when converting pixel coordinates to real-world distances, and simulation uses ground-truth geometry.

Calibration was straightforward. The camera was held pointing vertically downward at exactly 1.0 m above a tape measure laid on a flat surface. The browser stream showed 92 cm of the tape measure edge-to-edge. Applying the thin-lens model:

$$f = \frac{W_{\text{sensor}} \times d}{W_{\text{visible}}} = \frac{5.02 \text{ mm} \times 1000 \text{ mm}}{920 \text{ mm}} = 5.46 \text{ mm} \quad (37)$$

The corrected value is 22% lower than the 7.0 mm default. In operational terms, the old value would have *underestimated* the camera’s ground footprint, causing the GPS estimation module to place detected targets closer to the drone’s nadir than they actually were. At 10 m altitude this corresponds to a systematic error of approximately 2 m; at 30 m search altitude, the error would have reached 6 m—large enough to place the drone outside visual contact with the target during the centering phase. The corrected value was committed to the configuration repository within minutes of measurement, and every subsequent GPS estimate benefited immediately.

This single calibration justified the field day. It demonstrated a failure mode that no amount of simulation could have revealed, and it underscored the value of the progressive testing philosophy: measure real hardware before trusting derived quantities.

## M.4 Camera Colour Discovery

A subtler hardware-specific behaviour was identified during camera verification. The IMX296 sensor, when configured for RGB888 output via picamera2, delivers pixel data in BGR channel order—the opposite of what the format label implies. The team discovered this by capturing a frame of a known-colour object and systematically testing all six channel permutations (RGB, RBG, GRB, GBR, BRG, BGR) until the rendered image matched reality. The original codebase had included a `cv2.cvtColor(frame, COLOR_RGB2BGR)` conversion that, counterintuitively, *doubled* the error by swapping channels that were already in the order OpenCV expects.

Removing the unnecessary conversion produced correct colours immediately. The fix was trivial—deleting one line—but the diagnosis required physical hardware and a colour reference target. Like the FOV calibration, this was a defect invisible to simulation.

## M.5 Lens Calibration

Lens distortion was characterised using a printed checkerboard target: a 14×9 grid (13×8 inner corners) with 28 mm squares, sourced from calib.io. Multiple images were captured at varying orientations and distances using the capture-training stream. The calibration script initially failed to detect corners reliably; adding OpenCV’s robust detection flags and a `findChessboardCornersSB` fallback resolved the issue.

The resulting calibration achieved an RMS reprojection error of 0.399 pixels. The camera matrix and distortion coefficients were saved as a NumPy archive on the Pi. Undistortion was integrated into the vision module using precomputed remap lookup tables (`cv2.remap`), adding 1.5 ms per frame—negligible against the 206 ms inference cost. Because all detection scripts call the same vision module, the correction propagated automatically to every downstream tool without modification.

## M.6 Benchmark Results

The inference benchmark ran 50 passes of the YOLOv8n TFLite model (float32, 3.2 MB) on a static test image containing the dummy target. The results confirmed that the Pi 5 with the XNNPACK CPU delegate delivers consistent, operationally adequate performance:

- **Average latency:** 206.5 ms without undistortion, 208.0 ms with (difference within measurement noise).
- **Effective frame rate:** 4.8 FPS.
- **Detection rate:** 50/50 (100%).
- **Mean confidence:** 0.966 without undistortion, 0.958 with.

At 4.8 FPS and a nominal search speed of 8 m/s (altitude-dependent schedule at 35 m), the drone covers approximately 1.7 m between frames. At 35 m altitude the camera footprint spans roughly 32 m along track, so each ground point appears in approximately 15 consecutive frames. Even a pessimistic 50% detection rate under flight conditions would yield 7–8 detection opportunities per target—sufficient for reliable triggering.

## M.7 Weather Cancellation and the Go/No-Go Decision

By mid-afternoon the bench testing programme was complete, but the weather had deteriorated steadily. Wind gusts exceeded the airframe’s rated limit, and intermittent rain made outdoor electronics operation inadvisable. The go/no-go decision was unambiguous: the team’s own flight-day checklist specified wind and precipitation thresholds, and both were exceeded. The decision to stand down was made collectively, without hesitation, and without any attempt to “push through”—a discipline that the progressive testing philosophy explicitly encourages.

The cancellation was disappointing but not damaging. The bench tests had already validated Tier 2 of the testing framework in full. No flight-critical subsystem remained untested at the bench level, and the calibration data obtained during the day would have been impossible to collect during a time-pressured flight session.

## M.8 The Pivot: Evidence-Based Adaptation

Rather than treating the cancelled flight as lost time, the team reframed the day as an intensive calibration and integration session. The concrete outputs were:

1. **FOV correction** ( $f = 5.46$  mm, eliminating a 22% systematic GPS error);
2. **Lens calibration** (RMS = 0.399 px, integrated at 1.5 ms cost);
3. **Inference benchmark** (206.5 ms, 4.8 FPS, 100% detection at 0.966 confidence);
4. **Colour-channel fix** (BGR output confirmed, unnecessary conversion removed);
5. **Connectivity verification** (three-terminal workflow validated, camera/Cube/GPS all healthy);
6. **Field quick-reference sheet** (copy-paste commands for every step, removing internet dependency).

Each of these outputs fed directly into subsequent development. The FOV correction improved every GPS estimate computed thereafter. The lens calibration was automatically applied by every script that calls the vision module. The benchmark data informed the decision to prioritise FP16 quantisation and threaded pipelining as the most cost-effective inference acceleration paths. The field quick-reference sheet reduced setup time at the next session from exploratory debugging to a five-minute copy-paste sequence.

This pivot illustrates a broader principle: in safety-critical robotics development, a day spent measuring is never a day wasted. The progressive testing framework is designed precisely so that each tier produces value independently of whether the next tier is reached. The bench-test data collected on 11 March would have been needed regardless of weather; obtaining it on a day when flight was impossible simply reordered the schedule without losing any information.

## M.9 DJI Video Analysis: A Flight-Data Proxy

To extract flight-representative data despite the cancellation, the team conducted a post-hoc analysis using DJI Mavic footage recorded over the same test site at altitudes between 15 and 50 m during a previous reconnaissance visit. The 30 fps video ( $1456 \times 1088$ , cropped from 4K) was replayed through the full detection pipeline, with SRT telemetry providing per-frame altitude and GPS position.

This analysis yielded two key results. First, the retrained model ( $mAP_{50} = 0.995$ ) detected the dummy across the entire altitude range, confirming that the training dataset—300 synthetic images, 16 real labelled frames, and 50 negatives at native  $1456 \times 1088$  resolution—generalised to real flight conditions. Second, GPS target estimates computed from pixel coordinates and telemetry data produced a CEP50 of 2.3 m, with a maximum outlier of 16.5 m during a high-speed pass at 50 m altitude. The scatter distribution exhibited a characteristic diagonal elongation aligned with the flight direction, which timing analysis attributed to 100–200 ms GPS receiver latency—an error of 1–2 m at 10 m/s search speed. This systematic bias was subsequently addressed in software through spatial clustering with inverse-variance weighting: by averaging estimates from multiple passes, the cluster centroid converges on the true target position regardless of the per-frame lag.

The DJI analysis also revealed a methodological trap. SRT telemetry files contain one entry per native-rate frame (6854 entries for the 30 fps recording). Downsampled video variants (e.g. 6 fps) contain fewer frames but index into the same SRT file, producing a frame-to-telemetry mismatch that silently corrupts altitude and GPS readings. This was caught during development and documented as a mandatory constraint: only native-rate video may be used for quantitative analysis.

## M.10 Team Coordination Under Field Conditions

The field day was the first occasion on which the entire team worked simultaneously on the integrated system outside the lab. Several coordination patterns emerged that informed the workflow for subsequent sessions.

**Role distribution.** Three roles crystallised naturally within the first hour. The *software operator* (PuTTY Terminal 3) executed test scripts, interpreted output, and decided the test sequence. The *systems monitor* (Pi screen, Terminal 2) watched the diagnostics dashboard for anomalies—camera dropouts, Cube heartbeat loss, GPS satellite count fluctuations—and called “stop” if any indicator fell below threshold. The *hardware handler* repositioned the drone for camera calibration shots, swapped test targets, and managed power cycling. This division of labour allowed tests to proceed with minimal idle time: while the software operator analysed benchmark results, the hardware handler was already positioning the checkerboard for lens calibration.

**Communication protocol.** In the field, verbal communication replaced the asynchronous messaging used during remote development. The team adopted a simple protocol: the software operator announced each test before execution (“Running benchmark, 50 passes, do not touch the drone”), the systems monitor confirmed readiness (“Diagnostics green, go ahead”), and the hardware handler confirmed the physical setup (“Checkerboard at 0.5 metres, holding steady”). This call-and-response pattern, borrowed from flight crew resource management, prevented the camera-in-use conflict from recurring after the initial discovery and ensured that no test ran with the system in an unexpected state.

**Decision-making under uncertainty.** The weather cancellation decision was the most significant collective choice of the day. The flight-day checklist specified objective thresholds for wind and precipitation, which removed ambiguity from the decision. The pilot, who held formal authority for the go/no-go call, confirmed that conditions exceeded the checklist limits. The team agreed without debate and pivoted immediately to the bench-only programme. In retrospect, the speed of this pivot—from “we are not flying” to “here is the alternative test programme”—was enabled by having the calibration and benchmark scripts already written, tested in simulation, and loaded on the Pi. Had the team needed to improvise the bench programme on the spot, the afternoon would have been substantially less productive.

## M.11 Lessons for the Next Field Day

The 11 March experience produced a concrete list of procedural and technical improvements for subsequent sessions:

1. **Pre-stage the startup sequence.** The three-terminal workflow and field quick-reference sheet eliminated the ad-hoc debugging that consumed the first hour. Future sessions begin with a copy-paste startup, not improvisation.
2. **Calibrate before flying.** FOV and lens calibration should be the first activities at any new site, before propellers are fitted. The 22% focal-length error would have silently corrupted every autonomous GPS estimate had it not been caught.
3. **Carry colour and dimensional references.** A printed colour chart and a tape measure are minimal-weight additions to the field kit that enable in-situ validation of camera output and FOV.
4. **Prepare fallback test plans.** The pivot to bench-only testing was smooth because calibration and benchmark scripts were already written and tested. Future field days should always carry a “no-fly” test programme that maximises data collection from ground-based activities.
5. **Record everything for post-hoc analysis.** The DJI video, captured weeks earlier for a different purpose, became the most data-rich asset of the field day. Every future flight—even manual reconnaissance—should record video with telemetry for later replay through the detection pipeline.
6. **Respect the go/no-go checklist.** The decision to cancel was the right one. The checklist exists to remove ego from the decision; following it without negotiation is a team discipline, not a failure.

The field day demonstrated that the progressive testing philosophy is not merely a risk-management strategy but an information-maximisation strategy. By treating each tier as independently valuable, the team extracted six quantitative outputs from a session that produced zero seconds of flight time. Every one of those outputs improved the system’s readiness for the flight that weather denied.

## N Contingency Planning and Deliverable Tiers

This appendix defines the project’s risk management strategy for flight day: a three-tier deliverable structure (minimum viable, target, and stretch) with explicit fallback paths between them, paired with pre-rehearsed contingency plans for eight identified failure scenarios. Autonomous UAV missions introduce uncertainty at every layer—hardware, software, environmental, and regulatory—and the consequences of an unanticipated failure during a time-limited university flight slot are severe: lost data, potential hardware damage, and no opportunity to reschedule. The tiered approach ensures that the team always has a known-good configuration ready to demonstrate, even if the most ambitious capabilities cannot be proven in the field. Each contingency is mapped to a specific on-site test script that can verify the fallback path before committing to autonomous flight.

### N.1 Minimum Viable Deliverable

The minimum viable deliverable (MVD) satisfies the core brief requirements using the simplest operational mode: the human pilot retains manual control of the aircraft while the onboard system provides passive computer vision and telemetry logging. Table 69 summarises the MVD components and their requirement coverage.

The MVD does not satisfy R06 (PLB focus area redirect) or R07 (payload delivery near casualty), as both require autonomous decision-making beyond what Mission Planner AUTO mode provides. However, it demonstrates a working search-and-detect capability with full safety compliance, which is sufficient to evidence the core competencies assessed in the brief rubric: specialist skills (CV pipeline on embedded hardware), decision making (progressive test methodology), and communication (live telemetry and post-flight reporting).

### N.2 Target Deliverable

The target deliverable is the system described throughout this report: a fully autonomous mission orchestrated by the Python state machine on the Raspberry Pi companion computer. Table 71 maps each capability to its requirement.

The target deliverable achieves full compliance with R01–R05 and R08–R12. Partial compliance with R06 is achieved through the focus area support module (Section V.7), which re-generates the search pattern over a smaller polygon when PLB coordinates are received. R07 compliance depends on the Tarot payload release mechanism integration,

which is implemented in software (`SERVO_CHANNEL`, `SERVO_FULL_PWM` in `config.py`) but requires field calibration of the servo PWM values.

### N.3 Stretch Goals

If the target deliverable is achieved with margin, several enhancements would extend the system’s operational capability:

- **Multi-target detection and classification.** The detection queue already supports multiple concurrent targets with spatial clustering and inverse-variance weighting. The stretch goal adds automatic classification labels (dummy, item of interest, false positive) based on detection size, aspect ratio, and persistence across frames, reducing operator workload.
- **Live PLB focus area redirect (R06).** Mid-flight injection of a new focus area polygon via the ground station, triggering immediate pattern regeneration from the drone’s current position. The `focus_area.json` file is re-read on each PLB event, so the ground station need only write updated coordinates and send a trigger command.
- **Payload release mechanism (R07).** Full integration of the Tarot double-throw servo for first aid kit delivery, including a two-stage release sequence (partial open at 1300  $\mu$ s PWM, full release at 1100  $\mu$ s) with altitude and position guards to prevent premature deployment.
- **NCNN inference backend.** The NCNN export of the retrained YOLOv8n model is available in `cv_models/sar_v2.1088/n` and is expected to achieve  $\sim 15$  FPS on the Pi 5 CPU, tripling the current TFLite throughput of  $\sim 5$  FPS and enabling detection during faster transit legs.
- **Real-time GPS estimation on dashboard.** Streaming the Kalman-filtered target position estimate to the web ground station as a live scatter plot, giving the operator spatial context for the detection cluster before the verify prompt.

### N.4 How the Tier Structure Evolved

The three-tier deliverable structure was not defined at project inception; it emerged through iterative refinement as the team gained experience with the gap between simulation readiness and flight readiness. Understanding this evolution is important because it illustrates a general principle in safety-critical robotics: the deliverable scope must be calibrated to the verification evidence available, not to the ambition of the design.

**Initial ambition (weeks 1–4).** The project began with a single target: a fully autonomous SAR mission including search, detection, autonomous confirmation, landing, and payload delivery. At this stage, the team had not yet assembled the hardware or conducted any simulation beyond basic MAVLink tutorials. The assumption was that if the software worked in simulation, it would work on the aircraft with minor configuration changes.

**Realism check after simulation (weeks 5–8).** Once the SITL simulation was running end-to-end, two observations forced a reassessment. First, the number of parameters that required real-world calibration—FOV, GPS noise floor, detection altitude ceiling, inference latency, wind tolerance—was far larger than anticipated. Each unknown parameter represented a potential failure mode that could not be resolved without flight data. Second, the team’s access to flight time was limited to scheduled university slots, weather-permitting, with shared airspace and strict safety oversight. A single ambitious flight that failed at the first obstacle would waste the entire slot.

These observations motivated the introduction of the MVD tier: a configuration that requires *no custom flight code* in the autopilot loop, relying entirely on Mission Planner AUTO mode for waypoint execution and limiting the Pi’s role to passive observation. The MVD could be demonstrated even if every autonomous software feature failed, because it depended only on commercial autopilot firmware that the team did not write and could not break.

**Hardware integration reality (weeks 9–12).** Bench testing on the assembled Pi–Cube–camera stack revealed further integration challenges: Python 3.13 compatibility issues with `tflite-runtime` and `pyserial` (requiring the `ai-edge-litert` package and a mavproxy UDP bridge, respectively), a 22% focal length calibration error invisible to simulation, and the IMX296 camera’s BGR colour channel ordering mislabelled as RGB by the `picamera2` driver. Each of these issues was resolved within hours of discovery, but their existence confirmed that the gap between “works in simulation” and “works on hardware” contained numerous small failures that could only be found through physical testing.

This experience formalised the target deliverable as a distinct tier above the MVD: the full autonomous mission, but gated behind a progressive testing framework (Section E.2) that required each integration level to pass before the next was attempted. The stretch goals were then defined as capabilities that could be layered on top of the target if flight time permitted, without requiring changes to the core state machine.

**The resulting structure.** Table 73 summarises the relationship between the tiers, the verification evidence required for each, and the fallback path if that evidence is not available.

The key insight is that each tier is independently valuable. The MVD produces detection logs, GPS-tagged imagery, and a verification report that satisfy the brief’s core competency requirements. The target adds autonomous decision-making and payload delivery. The stretch adds operational efficiency and reduced operator workload. At no point does a failure at a higher tier invalidate the evidence collected at lower tiers—the progressive structure ensures that the team always has demonstrable results, regardless of how far up the ladder conditions allow them to climb.

## N.5 Contingency Plans

Table 75 defines the primary failure scenarios, their detection criteria, and the pre-planned response. Each contingency was rehearsed in simulation and, where applicable, during bench testing with the assembled hardware.

## N.6 Test Script Mapping

Each contingency scenario has a corresponding test script that can be executed on-site to verify the fallback path before committing to autonomous flight. Table 77 maps scenarios to scripts.

The progressive flight test methodology (Section E) ensures that each tier is validated before the next is attempted. If any test step fails, the system falls back to the last validated tier rather than proceeding with an unproven configuration. This approach—inspired by standard practice in aerospace verification and validation [55]—means that flight day always has a known-good fallback ready, even if the most ambitious capability is not yet proven in the field.

## O Test Script Reference

Systematic verification of a safety-critical autonomous system requires a structured test harness that covers every subsystem in isolation before exercising integrated behaviours. The 74 test scripts described in Section E.3 constitute a purpose-built verification harness designed around the progressive trust-building philosophy. This section provides a complete reference: what each script does, when it is used, and how it maps to the mission requirements and the five-tier progressive framework (Section E.2).

### O.1 Organisation by Category

Scripts are grouped into six directories, each targeting a distinct verification concern. The naming convention reflects execution order: lower numbers within the `flight/` directory carry lower risk and must pass before higher-numbered scripts are attempted.

**tests/hardware/** *Component-level checks* (8 scripts). Each script tests a single peripheral—camera, AI model (TFLite and NCNN backends), GPS receiver, or buzzer—in isolation. All issue zero flight commands and are safe to run on the bench at any time.

**tests/flight/** *Progressive flight tests* (14 scripts, numbered 0a–5). These form the critical path from bench to autonomous flight. Six bench scripts (0a–0f) validate the command pipeline, planning, geofence, and servo without propellers; five airborne scripts (1–5) build trust incrementally from manual RC with passive logging through to full autonomous search-and-land; plus geofence verification and map utilities.

**tests/diagnostics/** *Runtime monitoring* (5 scripts). Dashboards and camera streams used during preflight checks and live flight to monitor system health. The multi-panel diagnostics dashboard serves as the final go/no-go gate.

**tests/calibration/** *Parameter measurement* (7 scripts). FOV calibration (bench and altitude), lens distortion characterisation, GPS ground-truth comparison, and two focal-length calibration tools (terminal-based and OpenCV GUI). Results feed directly into `config.py` parameters that govern the coordinate-transform accuracy of the entire pipeline.

**tests/day\_1.experiments/** *Structured data collection* (6 scripts). All passive (zero commands), all output CSV. Designed to run during Tier 4 manual flights to collect quantitative performance data—detection rate versus altitude, GPS estimation error, model comparison—before autonomy is enabled.

**tests/laptop/** *Development-only tools* (21 scripts). Model inspection, TFLite and NCNN debugging, DJI video replay with detection overlay, A/B model comparison, path visualisation, NFZ manual testing, and approach/descent simulation. Not deployed to the Pi or used on flight day.

**tests/unit/** *Unit tests* (6 files, 127 test functions). Pytest-based tests covering configuration, states, GPS transforms, planning, vision, and GPS utilities.

**tests/automated/** *Automated integration tests* (7 scripts, 33 test functions). Extended coverage of planning, utilities, vision, geofence, and SITL smoke tests.

## O.2 Flight Test Progression

The flight scripts implement the incremental trust-building strategy. Each script is designed so that its failure mode is benign at its tier: bench scripts cannot cause vehicle motion; passive scripts issue zero commands; waypoint scripts operate without CV decision-making. Table 79 details the progression.

**Table 79:** Flight test progression. Each script must pass before the next is attempted. “Commands” indicates whether the script sends flight commands to the autopilot.

| Script              | Tier        | What It Proves                                                                                                                                                                                                  | Commands           | Risk     |
|---------------------|-------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|----------|
| 0a_cube_commands    | 3 (Bench)   | Individual MAVLink commands reach the Cube and return correct ACKs: mode changes, parameter reads, arming pre-checks.                                                                                           | Mode set           | Low      |
| 0b_bench_mission    | 3 (Bench)   | Full mission command sequence (arm, takeoff, waypoint, land) executes without error on the bench with propellers removed. Every <code>COMMAND_ACK</code> is verified.                                           | Arm, goto, land    | Low      |
| 0c_feedback_test    | 3 (Bench)   | Vision→GPS pipeline: operator carries drone past a dummy; camera captures frames, AI detects, pixel coordinates are transformed to GPS. Validates the entire sensing chain.                                     | Zero               | Low      |
| 1_passive_flight    | 4 (Passive) | Real outdoor CV performance: pilot flies manually on RC while Pi runs camera + AI + GPS estimation. Logs every detection with confidence, altitude, and estimated position. Buzzer confirms detections audibly. | Zero               | Medium   |
| 2_waypoint_test     | 5 (Auto)    | Autonomous GPS navigation: arms, takes off to 10 m, flies 3–4 GPS waypoints in GUIDED mode, lands. No camera, no detection—proves that MAVLink arm/takeoff/goto/land commands work on real hardware.            | Arm, goto, land    | High     |
| 3_auto_detect       | 5 (Auto)    | AUTO waypoint pattern + AI detection. If a target is detected, the drone transitions to GUIDED hover over the detection point. Does not descend or land autonomously.                                           | Mode switch, hover | High     |
| 4_detect_and_center | 5 (Auto)    | Full autonomous mission: search pattern, detect, centre on target using velocity commands, descend, verify, approach, land. This is the complete system under test.                                             | Velocity, land     | Critical |

Three laptop-only drawing utilities support the flight scripts without running on the Pi: `draw_search_area.py` produces a `search_area.json` polygon, `draw_waypoints.py` produces a `waypoints.json` file, and `live_map.py` provides a real-time position viewer during flight. These separate the interactive map-drawing step (laptop with display) from the headless execution step (Pi over SSH).

## O.3 Hardware Tests

**Table 80:** Hardware test scripts. All run on the Pi bench with zero flight commands.

| Script                                  | Purpose                                                                                                                                                                                                     | Output          |
|-----------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------|
| <code>benchmark.py</code>               | Runs 50 inference cycles on the TFLite model, reports mean/min/max latency, FPS, detection count, and confidence statistics. Quantifies whether the Pi meets the real-time requirement (<250 ms per frame). | Terminal report |
| <code>cv_benchmark.py</code>            | Live detection rate and speed with the camera active, including simulated motion blur at configurable levels. Tests the full capture→inference→decode pipeline.                                             | Terminal report |
| <code>gps_test.py</code>                | GPS diagnostics: satellite count, fix type (2D/3D/DGPS), HDOP, position, and time-to-fix. Confirms the GPS receiver is functional and tracks enough satellites.                                             | Terminal table  |
| <code>gps_health.py</code>              | Step-by-step GPS verification guide: walks the operator through satellite count, fix quality, and position stability checks with pass/fail criteria at each step.                                           | Interactive     |
| <code>buzzer_test.py</code>             | Sends MAVLink tune commands to the Cube’s buzzer. Validates the MAVLink command path and provides audible confirmation that the Cube is receiving commands.                                                 | Audible tones   |
| <code>detection_snapshot_test.py</code> | Captures a single camera frame, runs inference, and saves the annotated detection image. Quick smoke test for the camera→model pipeline.                                                                    | Saved image     |
| <code>benchmark_full.py</code>          | Comprehensive multi-model benchmark: system info, preprocessing timing, all available TFLite/NCNN models, threading overhead, and memory profiling.                                                         | Terminal report |
| <code>ncnn_benchmark.py</code>          | NCNN backend inference timing (50 runs). Benchmarks the faster alternative backend (~72 ms vs 206 ms TFLite).                                                                                               | Terminal report |

## O.4 Calibration Tests

Calibration scripts measure physical parameters that the coordinate-transform pipeline (Section 2.6) depends on. Errors in these parameters propagate directly into GPS estimation accuracy.

**Table 81:** Calibration test scripts.

| Script                           | Method                                                                                                                                                          | Result                                                                        |
|----------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| <code>fov_calibrate.py</code>    | Bench FOV measurement: place ruler at known distance, capture frame, measure visible width. Comprehensive multi-sample calibration.                             | <code>FOCAL_LENGTH_MM</code> for <code>config.py</code> (calibrated: 5.46 mm) |
| <code>fov_test_simple.py</code>  | Quick single-measurement FOV check. Lighter-weight version of the full calibration.                                                                             | Quick HFOV estimate                                                           |
| <code>alt_test.py</code>         | FOV measurement at real hover altitude. Run during a manual flight to validate bench calibration against in-flight conditions.                                  | Altitude-corrected FOV                                                        |
| <code>lens_calibrate.py</code>   | Checkerboard lens distortion calibration (13×8 inner corners, 28 mm squares). Produces undistortion remap matrices loaded by <code>vision.py</code> at startup. | <code>calibration_data.npz</code> (RMS = 0.399)                               |
| <code>gps_ground_truth.py</code> | Post-flight comparison: CV-estimated GPS position vs. surveyed ground-truth position of the dummy. Quantifies end-to-end localisation accuracy.                 | Position error in metres                                                      |

## O.5 Experiment Scripts

Six structured experiment scripts are designed to run passively during Tier 4 manual flights. Each connects to the MAVLink telemetry stream (read-only) and the camera, but issues zero flight commands. All output timestamped CSV files suitable for direct statistical analysis.

**Table 82:** Day 1 experiment scripts. All passive (zero commands), all output CSV.

| Script                         | What It Measures                                                                                                                                                   | Output                          |
|--------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------|
| <code>altitude_sweep.py</code> | Detection rate vs. altitude in 10–30 m buckets. Establishes the detection ceiling for the current model and camera.                                                | <code>altitude_sweep.csv</code> |
| <code>speed_sweep.py</code>    | Detection rate vs. ground speed, plus a Laplacian blur metric per frame. Determines the maximum search speed that maintains acceptable detection performance.      | <code>speed_sweep.csv</code>    |
| <code>gps_accuracy.py</code>   | GPS estimation error: compares the CV-derived target position against a known dummy location, logging per-frame error in metres.                                   | <code>gps_accuracy.csv</code>   |
| <code>gps_drift.py</code>      | GPS noise floor: records raw GPS position while the drone hovers stationary, computing CEP50 and CEP95 circular error statistics.                                  | <code>gps_drift.csv</code>      |
| <code>model_compare.py</code>  | Benchmarks all available <code>.tflite</code> models side-by-side: inference speed, detection rate, confidence distribution.                                       | Terminal comparison             |
| <code>detection_log.py</code>  | General-purpose flight logger: records every detection (timestamp, pixel coordinates, confidence, altitude, GPS estimate) as a catch-all for post-flight analysis. | <code>detection_log.csv</code>  |

## O.6 Diagnostics

Diagnostics scripts provide real-time system health monitoring. The multi-panel dashboard aggregates all communication links, camera status, GPS fix quality, and battery level into a single screen, serving as the final go/no-go gate before arming.

- **`diagnostics.py`** — Multi-view connectivity dashboard: camera preview, MAVLink heartbeat, GPS fix, battery voltage, signal strength. Run as the last preflight check.
- **`cube_monitor.py`** — Live Cube telemetry: attitude (roll/pitch/yaw), GPS position, altitude, battery, flight mode. Used for debugging during bench and flight tests.
- **`camera_stream.py`** — Basic MJPEG stream to a web browser at `http://PI_IP:8090`. Enables remote camera viewing from the ground station laptop.
- **`camera_stream_fast.py`** — Threaded MJPEG stream with reduced latency. Preferred over the basic version for flight-day streaming.
- **`camera_stream_h264.py`** — H.264/HLS stream via FFmpeg. Higher compression but increased latency; not used in practice.

## O.7 Script-to-Requirement Mapping

Table 83 maps each test category to the mission requirements it verifies and the testing tier at which it operates.

**Table 83:** Test script categories mapped to mission requirements and verification tiers.

| Category     | Requirement Verified                                              | Tier | Key Scripts                                     |
|--------------|-------------------------------------------------------------------|------|-------------------------------------------------|
| Hardware     | AI inference <250 ms; GPS fix; camera functional                  | 3    | benchmark, gps_health                           |
| Calibration  | FOV accuracy; lens correction; localisation error                 | 3–4  | fov_calibrate, lens_calibrate, gps_ground_truth |
| Flight 0a–0c | MAVLink command path; sensing pipeline; bench safety              | 3    | cube_commands, bench_mission, feed-back_test    |
| Flight 1     | CV performance under real conditions; detection altitude          | 4    | passive_flight                                  |
| Flight 2     | Autonomous GPS navigation; arm/takeoff/land                       | 5    | waypoint_test                                   |
| Flight 3–4   | Full mission: search, detect, centre, verify, land                | 5    | auto_detect, detect_and_center                  |
| Experiments  | Detection rate vs. altitude/speed; GPS accuracy; model selection  | 4    | altitude_sweep, speed_sweep, gps_accuracy       |
| Diagnostics  | System health; communication integrity; preflight gate            | 3–5  | diagnostics, cube_monitor                       |
| Laptop/Video | Model validation on real footage; FOV calibration; A/B comparison | 1    | video_test, video_test_compare                  |

## O.8 What Was Tested on Real Hardware

The following scripts were executed on the Raspberry Pi 5 with real peripherals during bench and field sessions:

- **Bench (Pi + Cube + camera, no props):** `benchmark.py` (206.5 ms mean inference, 4.8 FPS), `gps_health.py`, `0a_cube_commands.py`, `0b_bench_mission.py`, `fov_calibrate.py` (5.46 mm focal length), `lens_calibrate.py` (RMS = 0.399), `diagnostics.py`, `camera_stream.py`.
- **Field (Pi assembled on drone):** `passive_watch.py` (passive detection during manual carry, GPS estimation validated at ground level), `capture_training.py` (16 real training images collected).
- **Laptop simulation (SITL):** All flight scripts 0a–4, full state machine, geofence enforcement, 20-scenario stress test (Table 31), dry-run validation, DJI video analysis pipeline.

## O.9 Remaining Test Gaps

Table 84 documents the tests that have not yet been executed on real hardware, mapped to the flight test progression. These represent honest gaps between simulation validation and full operational verification.

**Table 84:** Tests not yet completed on real hardware. All have passed in SITL simulation.

| Flight Step             | What Remains                                                                                              | Blocking Factor                    |
|-------------------------|-----------------------------------------------------------------------------------------------------------|------------------------------------|
| Step 1: Passive flight  | Full airborne passive detection (flown, not hand-carried)                                                 | Weather cancelled first flight day |
| Step 2: Waypoint test   | Autonomous GPS waypoint navigation on real hardware                                                       | Requires Step 1 pass               |
| Step 3: Auto-detect     | AUTO pattern + AI-triggered GUIDED hover                                                                  | Requires Step 2 pass               |
| Step 4: Detect & centre | Full autonomous mission with velocity centering and landing                                               | Requires Step 3 pass               |
| Day 1 experiments       | Altitude sweep, speed sweep, GPS accuracy/drift CSV collection                                            | Requires airborne flight           |
| Calibration             | <code>alt_test.py</code> (FOV at real altitude), <code>gps_ground_truth.py</code> (post-landing accuracy) | Requires airborne flight           |

All remaining tests have passed in SITL simulation and the code is deployed to the Pi. The gap is not software readiness but flight opportunity: the first outdoor session was cancelled due to weather, and the progressive framework prohibits skipping tiers. The bench and simulation results provide confidence that the system will perform correctly when flight conditions permit, but real-world validation of detection altitude, GPS accuracy under wind, and motion blur effects remains outstanding.

## P Sensor Calibration

Accurate target localisation from an aerial platform depends on the fidelity of every link in the imaging chain: if the camera’s field of view, lens geometry, or colour representation deviates from the values assumed by the software, errors propagate multiplicatively through the GPS estimation pipeline. This section documents the four calibration campaigns conducted during development, each of which corrected a specific source of systematic error.

### P.1 Field-of-View Calibration

The ground footprint of a pinhole camera at altitude  $h$  is governed by

$$W_{\text{ground}} = \frac{w_{\text{sensor}} \cdot h}{f} \quad (38)$$

where  $w_{\text{sensor}}$  is the physical sensor width and  $f$  is the effective focal length. This relationship is used twice in the system: first by `planning.py` to compute the scan-strip spacing (with a 20% overlap margin), and second by the GPS estimation module to convert pixel offsets from the image centre into metric displacements on the ground. An error of  $\Delta f/f$  in the focal length therefore produces an identical fractional error in every GPS estimate and in the strip spacing, meaning the search pattern either leaves uncovered gaps or wastes flight time on excessive overlap.

The IMX296 global-shutter sensor has a nominal pixel pitch of 3.45  $\mu\text{m}$  and an active area width of 5.02 mm. The lens datasheet quoted a focal length of 7.0 mm, which was used as the initial default in `config.py`. To verify this value, a bench calibration was performed on the assembled camera module:

1. The camera was mounted vertically at a measured height of 1.000(5) m above a flat surface, using the `capture_training.py` MJPEG stream for live preview.
2. A tape measure was placed horizontally across the full width of the frame.
3. The visible extent was read directly from the tape as 0.92(1) m.

Substituting into Equation 38 and solving for  $f$ :

$$f = \frac{w_{\text{sensor}} \cdot h}{W_{\text{ground}}} = \frac{5.02 \times 1.00}{0.92} = 5.46 \text{ mm} \quad (39)$$

The calibrated value is 22% lower than the datasheet nominal. The corresponding horizontal field of view is

$$\theta_{\text{HFOV}} = 2 \arctan\left(\frac{w_{\text{sensor}}}{2f}\right) = 2 \arctan\left(\frac{5.02}{2 \times 5.46}\right) \approx 49.3^\circ \quad (40)$$

Had the uncalibrated value of 7.0 mm been retained, every ground-distance calculation would have been scaled by a factor of  $5.46/7.0 = 0.78$ —i.e. the system would have *underestimated* the camera footprint by 22%, producing GPS target estimates displaced from the true position by the same fraction of the pixel offset. At a search altitude of 35 m, the ground footprint is 32.2 m (calibrated) versus 25.1 m (uncalibrated); a detection at the frame edge would yield a GPS error of  $\approx 1.5$  m purely from the focal length mismatch. The strip spacing in the search pattern would also have been 22% too narrow, wasting approximately four additional scan lines over the survey area.

### P.2 Lens Distortion Calibration

Wide-angle lenses introduce radial and tangential distortion that displaces pixel coordinates from their ideal pinhole-model positions. For a target near the frame edge, uncorrected barrel distortion can shift the detected bounding-box centre by several pixels, degrading the pixel-to-GPS conversion. The standard Brown–Conrady model [56] parameterises radial distortion as

$$r' = r(1 + k_1 r^2 + k_2 r^4 + k_3 r^6) \quad (41)$$

where  $r$  is the normalised radial distance from the principal point and  $k_1, k_2, k_3$  are the radial distortion coefficients. Tangential terms ( $p_1, p_2$ ) account for misalignment between the lens and sensor planes.

Calibration was performed using a printed checkerboard target (calib.io pattern,  $14 \times 9$  squares at 28 mm pitch, yielding  $13 \times 8$  inner corners). Approximately 20 images were captured at varying orientations using `tests/calibration/lens_calibrate.py`, which calls `cv2.findChessboardCornersSB` (the sector-based detector, more robust to lighting variation) with a fallback

to the standard `cv2.findChessboardCorners` with adaptive thresholding and corner normalisation flags. Sub-pixel refinement was applied with `cv2.cornerSubPix` using a (5,5) search window and termination criteria of 30 iterations or  $\epsilon = 0.001$ .

The resulting reprojection RMS was **0.399 px**, indicating an excellent fit. The five distortion coefficients and  $3 \times 3$  camera matrix were saved to `calibration_data.npz`. At runtime, `vision.py` loads this file during initialisation and precomputes a pair of remap lookup tables via `cv2.initUndistortRectifyMap`:

```
new_mtx, _ = cv2.getOptimalNewCameraMatrix(
    mtx, dist, (cam_w, cam_h), 0, (cam_w, cam_h))
map1, map2 = cv2.initUndistortRectifyMap(
    mtx, dist, None, new_mtx, (cam_w, cam_h), cv2.CV_16SC2)
```

Every frame is then undistorted with a single call to `cv2.remap(frame, map1, map2, cv2.INTER_LINEAR)` before being passed to the detection model. The precomputed integer remap (CV\_16SC2) adds only 1.5 ms per frame—negligible compared to the 207 ms TFLite inference time—so the correction is effectively free. All downstream modules (passive observer, ground station, autonomous mission) benefit transparently because undistortion is applied inside `detect_in_image()`.

### P.3 Camera Colour-Space Discovery

The IMX296 sensor is configured via `picamera2` with the `RGB888` pixel format, which the library documents as returning 8-bit RGB data. However, during integration testing on the Raspberry Pi 5, colour reproduction appeared incorrect: the detection model (trained on BGR images, the OpenCV default) produced anomalously low confidence scores, and visual inspection of the stream showed a red/blue channel swap.

To diagnose the issue, all six permutations of the three colour channels were tested by displaying each alongside the raw frame:

```
for perm in [(0,1,2), (0,2,1), (1,0,2), (1,2,0), (2,0,1), (2,1,0)]:
    cv2.imshow(str(perm), frame[:, :, list(perm)])
```

The identity permutation (0,1,2)—i.e. treating the data as BGR without any conversion—produced correct colours. This confirms that the IMX296 driver delivers data in **BGR** byte order despite the `RGB888` format label. Applying the naïve `cv2.cvtColor(frame, cv2.COLOR_RGB2BGR)` conversion that the label would suggest actually *swaps* the channels into an incorrect order, resulting in the red and blue planes being exchanged.

The fix was to remove the incorrect `cvtColor` call entirely. The practical impact is twofold. First, the YOLO model—which expects BGR input—now receives correctly ordered data, restoring detection confidence to its trained level ( $> 0.95$  on the bench-test dummy). Second, the MJPEG stream served to the ground-station browser displays natural colours without post-processing, reducing per-frame latency. This discovery was documented as a permanent warning in the project documentation and `config.py` to prevent future regressions.

### P.4 DJI Video FOV Calibration

Separate from the onboard IMX296 camera, the team analysed recorded flight video from a DJI Mini 3 Pro to evaluate detection performance at real altitudes before the first autonomous flight. The DJI footage was centre-cropped from  $3840 \times 2160$  to  $1456 \times 1088$  to match the onboard camera’s aspect ratio, but the crop changes the effective field of view, which must be known to convert pixel detections into metric ground coordinates.

A calibration tool (`tools/fov_calibrate_video.py`) was developed. The operator pauses the video at frames where the mannequin is visible and clicks the top and bottom of the dummy; the tool computes the apparent height in pixels and, using the known real height of 1.8 m and the SRT-encoded altitude, solves for the focal length in pixels:

$$f_{\text{px}} = \frac{h_{\text{px}} \cdot h_{\text{alt}}}{h_{\text{real}}} \quad (42)$$

Nine samples were collected across altitudes from 15 m to 50 m, yielding a mean focal length of  $1416 \pm 41$  px. The corresponding horizontal FOV for the 1456-pixel crop width is

$$\theta_{\text{HFOV}} = 2 \arctan\left(\frac{1456}{2 \times 1416}\right) = 54.4^\circ \quad (43)$$

This DJI video HFOV ( $54.4^\circ$ ) is distinct from the onboard Pi IMX296 camera HFOV ( $49.3^\circ$ , Section F.2); the two values correspond to different cameras with different optics.

As a consistency check, the dummy height was back-calculated at each sample altitude using the calibrated focal length; all nine measurements fell within 1.7 m to 1.9 m, confirming the calibration against the known 1.8 m ground truth. This calibrated value replaced a hardcoded  $73^\circ$  assumption, which would have underestimated the focal length by 34% and corrupted every GPS estimate derived from the DJI analysis footage.

## P.5 Calibration Impact Summary

Table 85 summarises each calibration, its method, and the quantified impact on system accuracy.

The four calibrations are complementary: the FOV calibration corrects the dominant scale factor in the pinhole model, lens undistortion corrects the nonlinear residual, the colour-space fix ensures the detection model receives data in its trained domain, and the DJI FOV calibration enables accurate offline analysis of flight footage. Together, they reduce the systematic component of the target localisation error to the level set by GPS receiver noise ( $\sim 2$  m CEP) and atmospheric effects, rather than being dominated by avoidable modelling errors.

## P.6 Pre-Flight Calibration Checklist

Before every flight, the following calibration and readiness checks must be completed in order. Table 86 presents the checklist with acceptance criteria and the consequence of skipping each item.

## P.7 GPS Error Measurement Procedure

The GPS receiver introduces an irreducible random error that sets the floor of target localisation accuracy. Quantifying this error before flight allows the operator to judge whether conditions are adequate and provides a baseline against which CV-derived estimates are compared.

**Equipment.** Drone with GPS antenna, known reference position (e.g. the designated Take-Off Location at  $51.423406^\circ$  N,  $-2.671446^\circ$  W).

### Procedure.

1. Place the drone at the known reference GPS position with a clear sky view. Allow the GPS receiver to acquire a stable fix (fix type  $\geq 3$ ,  $\geq 10$  satellites).
2. Run the automated drift measurement script:

```
python tests/day_1_experiments/gps_drift.py
```

The script records GPS latitude, longitude, altitude, satellite count, and HDOP at 1 Hz for 60 s, writing results to a timestamped CSV file.

3. Compute the following metrics from the recorded data:
  - **CEP50:** the radius of a circle centred on the mean position that contains 50% of readings. This is the standard measure of GPS accuracy.
  - **CEP95:** the radius containing 95% of readings.
  - **Max deviation:** the largest single-sample displacement from the mean.
4. Record the satellite count and HDOP at the time of measurement for traceability.

**Acceptance criteria.**  $\text{CEP50} < 3$  m with  $\geq 10$  satellites and  $\text{HDOP} < 1.5$ . With a Here 3+ GNSS receiver in open sky, typical values are  $\text{CEP50} \approx 1.5$  m to 2.5 m.

**Impact.** GPS position error adds directly to target estimation error—it cannot be reduced by improving the vision pipeline. If  $\text{CEP50} = 2$  m, the target estimate can never be better than 2 m even with perfect FOV calibration and zero lens distortion. For a comprehensive decomposition of all error sources, see the error budget in Table 45 (Appendix K.6).

## P.8 Altitude Error Measurement Procedure

ArduCopter reports altitude relative to the arming point using a barometric pressure sensor. Because barometric pressure varies with temperature, humidity, and weather, the reported altitude drifts over time. Since altitude enters directly into the ground sample distance (GSD) calculation— $\text{GSD} = w_s \cdot h / (f \cdot W_{\text{px}})$ —any altitude error produces a proportional error in every pixel-to-metre conversion.

**Procedure.**

- 1. Place the drone on flat, level ground.
- 2. Arm the flight controller (propellers removed for bench testing) and note the initial altitude reading. It should read 0.0 m ( $\pm 0.1$  m).
- 3. Wait 5 min, monitoring the altitude readout via:

```
python tests/hardware/gps_test.py
```

- 4. Record the altitude at the end of the waiting period.
- 5. If the drift exceeds 0.5 m, recalibrate the barometer in Mission Planner (*Initial Setup*  $\rightarrow$  *Mandatory*  $\rightarrow$  *Accel Calibration*).

**Quantified impact.** Table 88 shows the position error introduced by altitude measurement error at the frame edge, where the pixel-to-GPS conversion is most sensitive.

**Acceptance criterion.** Altitude drift  $< 0.5$  m over 5 min on the ground. If conditions are marginal (hot day, rapidly changing pressure), re-arm immediately before takeoff to reset the altitude reference.

**P.9 Error Budget Cross-Reference**

The calibrations documented in this appendix address only the *systematic* (correctable) components of the target localisation error. The full error budget—including random GPS noise, GPS timing lag, attitude-induced projection shift, and motion blur—is presented in Table 45 (Appendix K.6). Table 89 maps each calibration in this appendix to the error source it mitigates, providing a unified view of which calibrations reduce which budget entries.

Calibration reduces the systematic error contribution from 4.8 m to 2.1 m RSS, bringing the total localisation uncertainty to a level dominated by the irreducible GPS receiver noise rather than avoidable modelling errors. The 2.1 m calibrated RSS is consistent with the CEP50  $\approx 2.3$  m measured from DJI flight video analysis (Section P.4).

**Q Simulation Validation and Sim-to-Real Gap**

A simulation framework is only useful if the behaviours it validates transfer reliably to hardware. This section examines what the multi-fidelity simulation (Section 2.8) proved, what it could not test, and the quantitative evidence underpinning confidence in the sim-to-real transfer.

**Q.1 Three Simulation Levels and What Each Validated**

Each simulation level targets a different subset of the system. Table 90 summarises the coverage.

**Table 90:** Validation coverage by simulation level.  $\checkmark$  = validated,  $\sim$  = partially validated,  $—$  = not testable.

| System Property                                        | Interactive  | SITL + Sim Cam | SITL + Webcam | Video Replay |
|--------------------------------------------------------|--------------|----------------|---------------|--------------|
| State machine logic (19 states, 32 transitions)        | $\checkmark$ | $\checkmark$   | $\checkmark$  | $—$          |
| MAVLink command interface (arm, takeoff, guided, land) | $—$          | $\checkmark$   | $\checkmark$  | $—$          |
| Geofence enforcement and SSSI avoidance                | $\sim$       | $\checkmark$   | $\checkmark$  | $—$          |
| Lawnmower pattern coverage                             | $\checkmark$ | $\checkmark$   | $\checkmark$  | $—$          |
| GPS estimation pipeline (pixel $\rightarrow$ WGS 84)   | $\checkmark$ | $\checkmark$   | $\checkmark$  | $\checkmark$ |
| Multi-target clustering and classification             | $\checkmark$ | $—$            | $—$           | $\sim$       |
| Visual servoing convergence                            | $\checkmark$ | $\checkmark$   | $\sim$        | $—$          |
| Operator verification workflow                         | $\checkmark$ | $\checkmark$   | $\checkmark$  | $—$          |
| Detection under real lighting/texture                  | $—$          | $—$            | $\sim$        | $\checkmark$ |
| Motion blur and vibration effects                      | $—$          | $—$            | $—$           | $\sim$       |
| Real GPS noise ( $\sim 2$ – $3$ m)                     | $—$          | $\sim$         | $\sim$        | $\checkmark$ |
| Wind disturbance on flight stability                   | $—$          | $—$            | $—$           | $—$          |

The interactive simulator (Level 1) was the primary tool for validating the GPS estimation mathematics and target approach algorithms. Over several hundred runs with configurable error injection, the three concurrent estimators—rolling average, cumulative average, and Kalman filter—were benchmarked against ground truth. The SITL level (Levels 2–3) validated the complete MAVLink command pipeline against ArduPilot’s C++ firmware, confirming that arming sequences, guided waypoint navigation, mode transitions, and landing commands executed correctly with realistic vehicle dynamics. Video replay (Level 4) bridged the gap to real conditions by exercising the detection model on genuine aerial imagery from a DJI flight.

## Q.2 Error Injection as a Validation Tool

The interactive simulator provides six configurable noise parameters that allow the estimation pipeline to be stress-tested against degraded sensor conditions:

1. **GPS drift** ( $\sigma = 0\text{--}5\text{ m}$ ): Ornstein–Uhlenbeck random walk matching real receiver autocorrelation. At  $\sigma = 3\text{ m}$  (typical open-sky conditions), the Kalman filter converges to sub-metre accuracy after 15–20 observations, while the rolling average achieves  $\sim 1.5\text{ m}$ .
2. **Camera shake** (0–15 px): altitude-dependent angular jitter modelling motor vibration. Detection confidence degrades below 0.5 at shake levels above 10 px (equivalent to  $\sim 40\text{ cm}$  ground displacement at 10 m altitude), indicating the boundary at which the detector’s performance is compromised.
3. **Inference throttle** (1–30 FPS): limiting CV throughput to the Pi’s measured 4 FPS reveals that at a search speed of 8 m/s and 35 m altitude, the camera footprint advances  $\sim 1.67\text{ m}$  between frames—well within the  $\sim 32\text{ m}$  ground footprint width—providing approximately 14 frames of target visibility per flyover pass.
4. **Altitude noise, yaw jitter, and FOV calibration error**: compound noise conditions that stress the coordinate transform under worst-case assumptions. A 10% FOV error produces a proportional 10% bias in all GPS estimates, confirming the importance of the bench calibration (Section F.2).

Real-time sliders allow these parameters to be swept mid-flight, producing sensitivity curves without restarting the simulator. This systematic approach identified the Kalman filter as the most robust estimator across the noise parameter space: it consistently converged within 1 m of ground truth under all tested noise combinations, whereas the rolling average was sensitive to bursts of high-noise observations at altitude.

## Q.3 DJI Video Replay as a Sim-to-Real Proxy

Recorded DJI flight footage provides the highest-fidelity offline validation available without flying the project drone. The 30 fps video ( $1456 \times 1088$  cropped) was paired frame-by-frame with SRT telemetry (GPS, altitude, heading) and processed through the identical detection and estimation pipeline used in all other simulation levels. This exercise validated three properties that no synthetic simulation can test:

- **Detection under real conditions.** The retrained YOLOv8n model (Section 2.3) detected the dummy in 87% of frames where the target occupied more than  $20 \times 20\text{ px}$ , with a mean confidence of 0.82. Missed detections correlated with sun glare, deep shadows, and partial occlusion by terrain features—failure modes absent from synthetic backgrounds.
- **GPS estimation accuracy.** Across 6854 frames from a single flight pass, the Kalman filter converged to a circular error probable (CEP50) of 2.3 m, with a maximum single-observation error of 16.5 m. The CEP50 is consistent with the compound uncertainty of the GPS receiver ( $\sim 2\text{ m}$ ), FOV calibration ( $\sim 5\%$ ), and the discovered 100–200 ms GPS timing lag that introduces a speed-dependent directional bias of  $\sim 1.6\text{ m}$  at 8 m/s.
- **GPS timing lag discovery.** The systematic along-track bias visible in the video replay scatter plots—a diagonal elongation of the observation cluster in the direction of travel—was traced to GPS receiver latency. This finding, which would have been nearly impossible to diagnose during a live autonomous flight, motivated the multi-pass averaging strategy and informed the error budget for the landing offset.

## Q.4 Validation Results Summary

Table 91 summarises the quantitative validation outcomes across all simulation levels.

**Table 91:** Quantitative simulation validation results.

| Metric                                              | Simulation Result                      | Real / Video Result        |
|-----------------------------------------------------|----------------------------------------|----------------------------|
| State machine end-to-end (19 states)                | All transitions correct                | Bench-verified (no flight) |
| Search pattern coverage                             | 100% of polygon area                   | Pattern executed in SITL   |
| GPS estimate convergence (Kalman, 20+ obs)          | $< 0.5\text{ m}$ from ground truth     | CEP50 = 2.3 m (DJI video)  |
| Detection rate (target $> 20 \times 20\text{ px}$ ) | 100% (synthetic background)            | 87% (real video)           |
| Mean detection confidence                           | 0.97 (synthetic)                       | 0.82 (real video)          |
| Inference latency (Pi 5, TFLite)                    | N/A (laptop)                           | 207 ms (4.8 FPS)           |
| Offset landing accuracy                             | $< 1\text{ m}$ (no wind, no GPS noise) | Not yet flight-tested      |

## Q.5 Progressive Build-Up: From Pure Software to Real Hardware

The simulation framework was not designed as a monolithic tool; it grew organically as each development phase revealed limitations in the previous level of testing. This subsection traces the chronological progression and explains why each simulation level was introduced.

**Phase 1: Software-only development (weeks 1–4).** All initial development occurred without any hardware. The state machine, lawnmower planner, and GPS estimation mathematics were implemented and tested entirely in the interactive simulator (`simple_simulator.py`), which models a drone as a point mass moving over a georeferenced satellite image. This environment was sufficient to validate the core algorithms—boustrophedon coverage, pixel-to-GPS coordinate transforms, Kalman filter convergence—but provided no confidence in the MAVLink command interface or the physical sensing pipeline. The key metric from this phase was that the Kalman filter converged to sub-metre accuracy from ground truth after 15–20 observations, establishing a baseline that later real-world testing could compare against.

**Phase 2: SITL integration (weeks 3–6).** ArduPilot’s Software-in-the-Loop simulator replaced the point-mass model with a full autopilot firmware running on the development laptop. This was the first time the code issued real MAVLink commands—`MAV_CMD_COMPONENT_ARM_DISARM`, `MAV_CMD_NAV_TAKEOFF`, `SET_POSITION_TARGET_GLOBAL_INT`—and received real telemetry responses. Several integration issues surfaced immediately: the `DISARM_DELAY` parameter (defaulting to 10 seconds) caused auto-disarm before the takeoff command was issued; `SET_POSITION_TARGET` messages sent during the climb phase cancelled the `NAV_TAKEOFF` command; and the state machine’s mode-change acknowledgement logic contained a race condition where a mode was requested and checked in the same heartbeat cycle. All three issues were resolved at this stage, at zero hardware risk. The SITL level also enabled the 22-test simulation checklist (Section E.6), which systematically exercised every state transition, CLI flag, and edge case.

**Phase 3: Docker Pi environment (week 7).** Before deploying code to the physical Raspberry Pi, a Docker container (`Dockerfile.pi-test`) was built to replicate the Pi’s ARM64 Linux environment with the production dependency set. This caught two critical compatibility issues. First, `tf-lite-runtime` could not be installed on Python 3.13; the replacement package `ai-edge-litert` was identified and verified in the container. Second, the `cv2.CAP_DSHOW` flag used for webcam access on Windows caused an import crash on Linux; a platform guard was added. The Docker test runs in under 30 seconds and was integrated into the pre-deployment checklist.

**Phase 4: Bench hardware integration (weeks 8–10).** With the Pi, Cube Orange, and IMX296 camera physically assembled, bench testing validated the complete sensing and command pipeline with propellers removed. The three bench scripts (`0a_cube_commands.py`, `0b_bench_mission.py`, `0c_feedback_test.py`) ran through individual commands, the full mission sequence, and the vision-to-GPS pipeline respectively. This phase discovered the FOV calibration error (22% deviation from the datasheet value, Section M.3), the IMX296 BGR colour channel issue, and the TFLite bounding box coordinate format inconsistency. Each was a defect invisible to all prior simulation levels.

**Phase 5: DJI video replay (weeks 10–12).** To bridge the gap between bench testing and actual flight before any flight time was available, recorded DJI footage from a reconnaissance overflight of the test site was replayed through the full detection pipeline. This was the first time the retrained model encountered genuine aerial imagery with real lighting, shadows, and terrain texture. The 87% detection rate and 2.3 m CEP50 GPS estimation accuracy established the most realistic performance benchmarks available without flying the project drone, and revealed the GPS timing lag phenomenon that synthetic imagery could never have exposed.

**What this progression demonstrates.** Each level caught defects that were invisible to all preceding levels. SITL caught MAVLink command sequencing issues. Docker caught Python 3.13 dependency failures. Bench testing caught sensor calibration errors and colour channel mismatches. Video replay caught real-world detection degradation and GPS timing bias. The cumulative effect is that by the time the system reaches outdoor flight, the majority of integration failures have already been found and resolved in controlled, zero-risk environments. This approach is consistent with the V-model verification philosophy (Section E.1) but adapted to the practical constraints of a university project with limited flight access.

## Q.6 Known Sim-to-Real Gaps

Despite the multi-fidelity approach, several properties remain untestable in simulation and represent the primary sources of residual risk:

- **Wind and turbulence.** No simulation level models wind effects on the airframe. Wind-induced position drift during visual servoing and GPS lock phases will degrade estimation accuracy and may trigger geofence warnings near boundary edges. SITL offers a wind parameter, but the interaction with camera stabilisation and detection reliability is unvalidated.
- **Motion blur at speed.** The interactive simulator and SITL produce sharp synthetic frames regardless of drone velocity. Real flight at 8 m/s with a 200 ms exposure could produce motion blur of  $\sim 1.6$  m ground equivalent,

potentially reducing detection confidence. The DJI video provides partial evidence (recorded at a different airspeed and camera system), but the IMX296 global shutter should mitigate this for the project hardware.

- **Real GPS accuracy.** SITL GPS noise ( $\sim 0.5$  m) is optimistic. Real open-sky receivers typically exhibit 2–3 m CEP, with occasional multipath excursions to 5–10 m near buildings or terrain. The DJI video analysis (CEP50 = 2.3 m) provides the best available estimate, but was recorded with a different GPS receiver.
- **RF interference and MAVLink dropouts.** The simulation assumes a perfect MAVLink link. In the field, WiFi, radio control, and telemetry transmitters share the 2.4 GHz band, potentially causing packet loss or latency spikes. The link-lost timeout (5 s  $\rightarrow$  RTL) is implemented but has not been tested under real interference conditions.
- **Sunlight and thermal effects.** Detection performance under direct sunlight, long shadows, and thermal shimmer is only partially captured by the DJI video replay, which was recorded under overcast conditions. The model has not been validated against harsh midday lighting or low sun angles.

## Q.7 Confidence Assessment

Based on the evidence gathered across all four simulation levels and the DJI video proxy, the system properties can be categorised into three confidence tiers:

**High confidence** (validated at multiple levels with consistent results): the finite state machine and all 32 transitions; the MAVLink command interface (arm, takeoff, guided navigation, land); the lawnmower search pattern geometry; the GPS estimation mathematics and Kalman filter convergence; the operator verification workflow; and geofence enforcement including SSSI avoidance.

**Moderate confidence** (validated in simulation but with known real-world degradation): detection reliability at altitude (87% on real video vs 100% synthetic); GPS estimation accuracy (CEP50 = 2.3 m on video, sub-metre in simulation); inference throughput on the Pi (4.8 FPS measured on bench, but not under thermal load during sustained flight).

**Low confidence** (untested or only indirectly tested): offset landing accuracy under wind; visual servoing stability in gusty conditions; detection performance under harsh lighting; system behaviour under RF interference or GPS multipath; and thermal throttling during extended missions. These properties require validation through the progressive flight testing framework (Section E.2), advancing from passive observation to full autonomy only after each preceding tier passes.

## Q.8 Lessons from the Sim-to-Real Transfer

The multi-fidelity simulation campaign produced several generalisable lessons about the value and limitations of simulation in safety-critical robotics development.

**Simulation catches logic errors; hardware catches calibration errors.** Of the 20 failure scenarios tested in SITL (Section E.6), nine required code fixes—all were logic or sequencing errors (missing timeouts, hardcoded parameters, race conditions). Of the six defects discovered during bench testing, all were calibration or compatibility issues (FOV error, colour channel mismatch, coordinate format inconsistency, Python 3.13 package incompatibility). The two categories of failure are almost perfectly disjoint: simulation cannot catch calibration errors because it uses ground-truth geometry, and bench testing cannot catch logic errors in the state machine because it lacks the dynamic vehicle context needed to exercise complex state transitions. This complementarity validates the multi-tier approach: both tiers are necessary, and neither is sufficient alone.

**Video replay is the highest-value offline test.** The DJI video analysis produced the most operationally relevant data of any non-flight activity. It was the only test that exposed the GPS timing lag phenomenon, the only test that measured detection performance under real lighting and texture, and the only test that produced realistic GPS estimation statistics. The cost was low: recording the reconnaissance video required approximately 15 minutes of manual flight time with a consumer drone, and the replay pipeline reused the identical detection and estimation code used in all other simulation levels. For projects with limited flight access, video replay from a survey drone represents an efficient way to bridge the sim-to-real gap before committing to autonomous operations.

**Configuration defaults from datasheets are unreliable.** The 22% focal length error (datasheet: 7.0 mm; measured: 5.46 mm) would have produced a systematic 6 m GPS estimation bias at search altitude. This error was undetectable in simulation and would have degraded every autonomous GPS estimate, potentially causing the drone to centre on a location 6 m from the actual target. The lesson is that any parameter derived from a datasheet or nominal specification should be treated as an initial guess and calibrated against physical measurement before use in a safety-critical pipeline. The project's calibration scripts (`tests/calibration/`) are designed to make this verification rapid and reproducible at any new site.

**The sim-to-real gap is not a single gap; it is many small gaps.** The phrase “sim-to-real gap” suggests a single chasm between simulation and reality. In practice, the gap is composed of dozens of small discrepancies—sensor quirks, driver bugs, environmental conditions, timing latencies—each of which is individually minor but collectively significant. The progressive testing framework addresses this by providing multiple opportunities to discover and resolve these discrepancies before they compound during autonomous flight. Each tier narrows the gap incrementally, rather than attempting to cross it in a single step.

## R Mission Flow

Understanding the system’s operational behaviour requires tracing a complete mission from power-on through search, detection, payload delivery, and return-to-launch. This section describes that end-to-end sequence, following the state machine (Figure 9) and the mission timeline (Figure 11), with explicit reference to the flight parameters summarised in Table 124 and the requirement identifiers from the project brief.

### R.1 Pre-Flight

Before every flight the crew executes a structured start-up sequence:

1. The pilot powers the Hexsoon EDU-450 and verifies that the Cube Orange obtains a 3D GPS fix with at least six satellites.
2. On the Raspberry Pi 5 companion computer, the operator starts `mavproxy` over the serial link to the Cube (921 600 baud), with two UDP outputs: port 14550 for the mission script and port 14551 for a diagnostics monitor. A TCP bridge on port 5762 allows Mission Planner on the ground-station laptop to connect simultaneously.
3. The operator runs `preflight.py`, which verifies camera frames, AI model loading, Cube heartbeat, and GPS lock in a single pass.
4. The operator opens the browser-based ground station at `http://PI_IP:8090`, which provides a live MJPEG video stream, a 2D GPS grid, detection cluster markers, and command buttons (Section 2.7).
5. The mission script (`main.py`) is launched. It loads the survey polygon, flight boundary, SSSI no-fly zone, and take-off location from the course KML file, generates the lawnmower search pattern, and enters the INIT state.

### R.2 Takeoff

The pilot arms the motors via the RC transmitter and switches to GUIDED mode. The state machine confirms motor arm status and issues `MAV_CMD_NAV_TAKEOFF` to the search altitude of 35 m (R04: maximum 50 m). The drone climbs vertically from the Take-Off Location (R03: within 5 m of the defined TOL). A 60-second timeout alerts the operator if the target altitude is not reached; on real hardware, an unexpected disarm during climb transitions the system to DONE rather than silently retrying, ensuring the pilot retains full awareness.

### R.3 Transit to Search Area

Once 90% of the target altitude is reached, the planner generates the lawnmower waypoints from the drone’s current position. If operator-defined transit waypoints exist (loaded from `waypoints.json`), the drone follows them first at  $15\text{ m s}^{-1}$  transit speed, providing a controlled ingress path that avoids the SSSI. The NFZ geofence is active throughout transit: any waypoint within 30 m of the SSSI boundary was already filtered at plan time, and the runtime speed ramp (Section R.11) provides a continuous safety margin.

### R.4 Search

The drone executes the boustrophedon search pattern at 35 m altitude. Speed follows an altitude-dependent schedule:  $6\text{ m s}^{-1}$  at 20 m, linearly interpolating to  $10\text{ m s}^{-1}$  at 50 m, yielding approximately  $8\text{ m s}^{-1}$  at the nominal search altitude. The computer vision pipeline runs continuously at 4.8 FPS on the Pi 5 (TFLite, XNNPACK backend). Each frame undergoes the following processing chain:

1. **Capture** — the IMX296 global-shutter camera acquires a  $1456 \times 1088$  frame.
2. **Undistort** — a precomputed lens-distortion remap corrects barrel distortion (1.5 ms cost).
3. **Resize and normalise** — the frame is resized to the model’s  $640 \times 640$  input tensor and normalised to  $[0, 1]$ .
4. **Inference** — the YOLOv8n TFLite model produces bounding-box predictions (206 ms average).
5. **Threshold** — detections with confidence below 0.2 are discarded. This deliberately low threshold maximises recall; the operator filters false positives via the dashboard.

The live MJPEG stream on the ground-station dashboard shows each frame with a detection overlay (bounding box, confidence score, crosshair-to-target line), the current state, altitude, speed, GPS coordinates, and waypoint progress. The operator monitors the stream passively during the search phase.

## R.5 Detection and Queuing

When the vision system detects a target, the pixel coordinates are converted to a GPS estimate using the drone's current position, altitude, and heading (Section 2.6). The estimate is then filtered through three validity checks:

- **NFZ check** — detections whose estimated GPS falls inside the SSSI polygon are discarded.
- **Search boundary check** — detections outside the survey polygon are discarded.
- **De-duplication** — detections within 5 m of a previously rejected target, a logged item of interest, or an already-queued detection are suppressed.

Valid detections are appended to a bounded queue (maximum 20 entries; oldest dropped on overflow). Critically, the drone does *not* stop or deviate on the first detection. It continues the search pattern, collecting multiple observations of the same target from different viewing angles, which improves the GPS estimate through spatial averaging. Only when a queued detection is ready for investigation does the state machine transition.

## R.6 PLB Focus Area Redirect

Between 5 and 15 minutes into the search, the operator receives simulated Personal Locator Beacon (PLB) coordinates (R06). The operator presses the B key on the dashboard or terminal, and the system loads a focus-area polygon from a JSON file. A new, tighter lawnmower pattern is generated over the focus polygon from the drone's current position, with reduced search speed ( $5 \text{ m s}^{-1}$ ) to maximise detection probability in the high-priority zone. The waypoint index and rescan counter reset; previously rejected false positives are preserved so they are not re-investigated.

## R.7 Investigation

When the operator observes a detection cluster on the dashboard and judges it worth investigating, they press the N key (investigate). The state machine pops the highest-priority queued detection, records the drone's current pattern position as a departure point, and flies to the cluster centroid at transit speed. This operator-in-the-loop decision (R08) prevents the drone from chasing every low-confidence detection autonomously, conserving flight time and battery.

## R.8 Verification

Upon arrival within 1 m of the estimated target position, the drone hovers and the state machine enters VERIFY. The operator sees a close-up live video feed on the dashboard and has 120 seconds to make a classification decision:

- **Y (confirm casualty)** — the target GPS is refined using a 10-second GPS position average. The operator is then prompted to select a landing side (N/E/S/W).
- **I (item of interest)** — the target is logged with its GPS coordinates and imagery (R05, R10) but no landing is attempted. The drone returns to its departure point and resumes the search pattern.
- **X / N (false positive / reject)** — the target is added to the rejected list (suppressing future detections within 5 m) and the drone returns to its departure point.

A countdown timer is displayed on both the HUD and the dashboard, with audible warnings at 60, 90, and 100 seconds. If no response is received within 120 seconds, the target is auto-rejected and the drone resumes searching.

## R.9 Approach and Landing

If the operator confirms the casualty, the drone calculates a landing point offset 7.5 m from the target in the operator-selected direction, satisfying the requirement to land within 10 m but not within 5 m of the casualty (R07). The choice of 7.5 m provides margin for GPS estimation error ( $\text{CEP}_{50} \approx 2.3 \text{ m}$  from video analysis) while remaining well within the 10 m outer bound.

The drone descends to 3 m above the offset landing point and enters the HOVER\_TARGET state for a 15-second payload deployment sequence:

1.  $t = 0\text{--}3 \text{ s}$ : stabilise hover at 3 m.
2.  $t = 3 \text{ s}$ : Tarot servo partial release (stage 1, PWM 1300) — parachute or strap slips.
3.  $t = 6 \text{ s}$ : Tarot servo full release (stage 2, PWM 1100) — first-aid kit separates.
4.  $t = 15 \text{ s}$ : servo retracts (PWM 1500), hover complete.

Direct offset landing was chosen over a visual-servoing descent for three reasons. First, the GPS estimate accuracy ( $\text{CEP}_{50} \approx 2.3 \text{ m}$ ) is sufficient for a 7.5 m offset, where the absolute position of the landing point matters less than its distance from the target. Second, visual servoing at low altitude introduces risk: the target may leave the camera's field of view during descent, and corrective manoeuvres close to the ground increase the probability of a crash. Third, the simplified approach reduces the number of states and transitions that must be validated, improving system reliability within the available development time.

## R.10 Return to Launch

After payload deployment, the drone climbs back to search altitude and retraces the transit waypoints in reverse order at  $15\text{ m s}^{-1}$ . Upon reaching the launch vicinity ( $< 2\text{ m}$  horizontal), it issues `MAV_CMD_NAV_LAND`. Touchdown is confirmed by three independent mechanisms: ArduPilot’s accelerometer-based auto-disarm, an altimeter threshold ( $0.5\text{ m}$  for five consecutive ticks), and an absolute 90-second landing timeout with forced disarm. The mission terminates in `DONE`, reporting the landing error relative to the home position (R03).

## R.11 Safety Layers During Flight

Four safety mechanisms operate concurrently throughout the mission, independent of the current state.

**NFZ Geofence (R01, R02).** The SSSI no-fly zone is enforced by a five-layer geofence (Section D.1):

1. **Plan-time filtering** — waypoints within  $30\text{ m}$  of the SSSI boundary are removed during pattern generation.
2. **Speed ramp** — within  $20\text{ m}$  of the boundary, a linear speed scalar reduces the drone’s velocity from  $3.0\text{ m s}^{-1}$  at the zone edge to  $0.3\text{ m s}^{-1}$  at the boundary, providing a smooth deceleration that prevents overshoot.
3. **Repulsive field** — within  $23\text{ m}$  of an inner offset polygon, a constant-magnitude ( $3.0\text{ m s}^{-1}$ ) velocity vector pushes the drone away from the boundary.
4. **Hard boundary** — if the drone enters the SSSI polygon, the system immediately transitions to `MANUAL` mode and halts all autonomous commands, returning control to the pilot.
5. **Firmware geofence** — an independent geofence on the Cube Orange autopilot triggers `LOITER` or `RTL` if all software layers fail.

**Manual Override.** At any point during the mission, the pilot can press the `M` key (via terminal, dashboard, or a mapped RC switch) to enter `MANUAL` mode. The state machine saves the current position and active state, then ceases all autonomous commands. The pilot flies the drone manually. Pressing `M` again transitions to `RETURN_FROM_MANUAL`, which navigates back to the saved departure point and resumes the interrupted state. This provides a non-destructive override that preserves mission progress (R09).

**RC Kill Switch (R09).** The RC transmitter’s mode switch can be set to `STABILIZE` at any time, bypassing the companion computer entirely. This hardware-level override is always available regardless of software state and constitutes the ultimate failsafe. The Safety Pilot retains authority to take this action at any point during the flight (R11).

**Link Loss.** The state machine monitors MAVLink heartbeats continuously. If no heartbeat is received for five seconds, the system triggers an automatic Return-to-Launch via ArduPilot’s built-in `FS_GCS_ENABLE` failsafe parameter, which is configured before flight. This ensures the drone returns home even if the companion computer loses communication with the autopilot.

## R.12 How the Mission Was Built Up Incrementally

The mission flow described above represents the final integrated system, but it was not designed or implemented as a single unit. Each subsection of the mission was developed, tested, and validated independently before being composed into the full sequence. This incremental approach was essential for managing complexity and maintaining confidence at each integration step.

**Stage 1: Search pattern in isolation.** The lawnmower planner (`planning.py`) was the first module developed, because it has no hardware dependencies. The module takes a polygon (the survey area), a strip width (derived from the camera FOV at the search altitude), and an edge margin, and generates a boustrophedon waypoint sequence. It was tested entirely offline using `python main.py --dry-run`, which visualises the pattern on a map without requiring a Cube connection, GPS fix, or even SITL. The output image (`dry_run_pattern.jpg`) provided immediate visual confirmation that the pattern covered the survey area, respected the SSSI exclusion zone, and connected waypoints in a flyable order. Over 50 dry-run iterations refined the strip width, turn radius, and edge margin before the planner was ever connected to an autopilot.

**Stage 2: Detection pipeline in isolation.** The vision module (`vision.py`) was developed as a fully independent component with a single public interface: `detect_in_image(frame)` returns `(found, x, y, conf)`. This interface decouples the detection algorithm from the mission logic entirely—`main.py` never imports TensorFlow, never handles model loading, and never processes raw bounding boxes. The vision module was tested first on static images (`tests/hardware/benchmark.py`), then on live camera feeds (`tests/laptop/test_cv.py --camera`), and finally on

recorded DJI flight video (`tests/laptop/video_test.py`). Each testing stage expanded the conditions under which the module was validated without requiring changes to its interface.

**Stage 3: MAVLink commands in SITL.** The arming, takeoff, waypoint navigation, and landing commands were tested against ArduPilot SITL on the development laptop. The state machine was connected to the simulated autopilot and the entire mission flow—from INIT through SEARCH to LANDING and DONE—was executed repeatedly. This stage caught the `DISARM_DELAY` issue, the `SET_POSITION_TARGET` conflict with `NAV_TAKEOFF`, and the mode-change acknowledgement race condition, all of which would have been dangerous if encountered during a real flight.

**Stage 4: Operator interface development.** The browser-based ground station (`pi_flight.py`) was developed and tested on the laptop in SIMULATION mode before being deployed to the Pi. The MJPEG video stream, GPS grid overlay, detection markers, and command buttons (N/Y/I/X/L) were all validated against the SITL simulator, allowing the full operator workflow—observe detection, press investigate, classify target, confirm landing—to be rehearsed dozens of times without hardware risk. The headless architecture (no `cv2.imshow` calls; all UI through the browser) was a deliberate design choice to ensure the same code would run on the Pi over SSH without a display server.

**Stage 5: Integration on hardware.** Only after each component had been validated independently was the full system assembled on the Pi with the Cube Orange. The bench test scripts (`0a_cube_commands.py`, `0b_bench_mission.py`, `0c_feedback_test.py`) verified the integrated pipeline—camera capture through AI inference through GPS estimation through MAVLink command—in a controlled, grounded environment. The final step before flight would be passive observation during manual RC flight (Tier 4), where the Pi runs the full pipeline but issues zero flight commands, collecting ground-truth detection data that informs the configuration parameters used when autonomy is finally enabled. This staged approach means that by the time the system reaches autonomous flight, every component has been validated at multiple integration levels, and the only remaining unknowns are the environmental factors—wind, lighting, GPS multipath, thermal effects—that no ground-based test can replicate. The progressive testing framework (Section E.2) formalises this philosophy: each tier must pass before the next is attempted, and failures at any tier are resolved at that level before proceeding.

## S Centering vs. Direct Offset Landing

A critical design decision in the landing pipeline is whether to visually servo the drone above the target before commanding a landing offset, or to proceed directly from a GPS-estimated position. The target localisation pipeline is structured as a two-step process:

1. **Step 1 — GPS estimation from search altitude.** During the lawnmower search pattern at 35 m, each detection is projected from pixel coordinates to a ground-level GPS position using the calibrated camera model (Section 2.6). Multiple detections are fused using inverse-variance weighting and Kalman filtering to produce a consolidated target estimate with  $CEP_{50} \approx 2.3$  m.
2. **Step 2 — Visual servo centering (optional).** The state machine includes a dedicated CENTERING state that flies the drone directly above the estimated target and drives the bounding-box centroid toward the image centre, eliminating off-axis projection error. A subsequent DESCENDING state would reduce altitude from 35 m to 15 m while maintaining visual lock, further reducing the ground sampling distance and improving localisation to  $CEP_{50} \approx 1.5$  m.

In the current implementation, Step 2 is bypassed: the drone centres above the target at search altitude (35 m), enters VERIFY for operator confirmation and GPS averaging, then proceeds to an offset landing using only the Step 1 estimate. This section analyses the centering-vs-direct-offset trade-off and justifies why Step 2 was deferred.

### S.1 The Centering Option

In the full centering pipeline (Step 2 above), the drone transitions from SEARCH to CENTERING, where it flies toward the GPS-estimated target position and attempts to place the detection bounding box within a pixel-tolerance of the image centre. The DESCENDING state then reduces altitude from 35 m to 15 m in increments while maintaining visual lock. In the baseline configuration, the descent phase is bypassed: the drone transitions directly from CENTERING to VERIFY at search altitude (35 m). Only after the operator confirms the target does the system command a landing offset.

This pipeline improves localisation accuracy because lower-altitude frames have a smaller ground sampling distance (GSD), and centering eliminates the pixel-projection error from off-axis detections. However, it introduces three additional risk factors:

1. **Low-altitude manoeuvring.** Descending from 35 m to 5 m while simultaneously correcting lateral position requires sustained GUIDED-mode velocity commands. Any interruption—communication dropout, momentary detection loss, or wind gust—can leave the drone at low altitude with stale position commands.
2. **Visual servo fragility.** At altitudes below 10 m, the target may exceed the camera field of view, motion blur increases with ground-relative speed, and the narrow depth of field of the IMX296 global shutter sensor reduces detection confidence. During simulation testing, three of twenty centering runs exhibited oscillatory lateral corrections at 8 m altitude.
3. **Increased flight time over the casualty.** The centering and descent phases add 30–60 s of hover time directly above the target, which in a real SAR scenario places the casualty at greater risk from rotor downwash and potential mechanical failure.

## S.2 Why Direct Offset Landing Is Sufficient

The offset landing approach bypasses centering: once the operator confirms the target, the drone flies to a GPS point displaced 7.5 m from the estimated target position (perpendicular to the wind vector when available, otherwise due south) and executes a standard `MAV_CMD_NAV_LAND`. The key insight is that the system requirement (R07) specifies landing *within* 5 m to 10 m of the casualty—not directly on top of it. A 7.5 m offset places the nominal landing point at the centre of this annular band, and the question reduces to whether GPS-based localisation is accurate enough to keep the actual touchdown within the band.

## S.3 Probability Analysis

The GPS-estimated target position was characterised in simulation with realistic noise parameters calibrated against DJI flight-video telemetry (Section 2.6). The circular error probable at the 50th percentile was measured as  $\text{CEP}_{50} = 2.3$  m. Assuming the radial error follows a Rayleigh distribution, the scale parameter is

$$\sigma = \frac{\text{CEP}_{50}}{\sqrt{2 \ln 2}} \approx \frac{2.3}{1.177} \approx 1.95 \text{ m} \quad (44)$$

The actual landing point is offset by  $d_0 = 7.5$  m from the estimated target. The distance from the *true* target to the landing point is therefore a random variable whose distribution depends on the vector sum of the deterministic offset and the stochastic GPS error. In the worst case (all error components aligned radially), the landing distance from the true target is  $r = d_0 + \epsilon$ , where  $\epsilon \sim \mathcal{N}(0, \sigma^2)$  projected onto the offset axis. The probability that the landing falls within the 5 m to 10 m band is

$$P(5 < r < 10) = \Phi\left(\frac{10 - 7.5}{1.95}\right) - \Phi\left(\frac{5 - 7.5}{1.95}\right) = \Phi(1.28) - \Phi(-1.28) \approx 0.900 - 0.100 = 0.800 \quad (45)$$

where  $\Phi$  is the standard normal CDF. This conservative uni-axial projection gives a lower bound of 80%. In practice the error is two-dimensional: the lateral component (perpendicular to the offset direction) shifts the landing point around the annular band rather than outside it, and the multi-detection fusion (inverse-variance weighting over  $\sim 10$  observations) reduces  $\sigma$  by  $\sqrt{N_{\text{eff}}}$ . With an effective sample size of  $N_{\text{eff}} = 6$ , the fused  $\sigma_{\text{fused}} \approx 1.95/\sqrt{6} \approx 0.80$  m, giving

$$P(5 < r < 10) = \Phi\left(\frac{2.5}{0.80}\right) - \Phi\left(\frac{-2.5}{0.80}\right) = \Phi(3.13) - \Phi(-3.13) \approx 0.999 - 0.001 > 0.99 \quad (46)$$

confirming that the fused GPS estimate with a 7.5 m offset meets the R07 band requirement with over 99% probability.

## S.4 Simulation Comparison

To validate the analysis, twenty simulated missions were executed for each approach using the SITL environment with GPS noise ( $\sigma_{\text{GPS}} = 1.5$  m), camera shake, and  $3 \text{ m s}^{-1}$  wind. Table 92 summarises the results.

The “mean error” column reports the absolute deviation of the actual landing distance from the intended 7.5 m offset—i.e. how far the touchdown was from the centre of the acceptable band. The direct offset approach achieved 95% compliance with R07, with a single outlier caused by a GPS multipath spike during the final descent. The centering approach improved mean accuracy to 1.8 m but produced three runs where the visual servo exhibited lateral oscillation below 10 m altitude, requiring the safety pilot to intervene (RC override to LOITER). In all three cases the drone recovered, but the behaviour would be unacceptable in an operational SAR deployment.

## S.5 Design Decision

Given that (a) the direct offset approach meets R07 with >95% probability under realistic noise, (b) centering introduces low-altitude instability risks that are difficult to mitigate without extensive tuning, and (c) the project

timeline prioritises a reliable first demonstration over optimal accuracy, the direct offset landing was selected as the baseline configuration. The centering states remain implemented and tested in simulation; enabling them requires only a single configuration flag (`ENABLE_CENTERING = True` in `config.py`), making this a low-risk upgrade path for future iterations.

## S.6 Safety Considerations of Hovering Above the Casualty

Visual centering requires the drone to hover directly above the detected casualty for 10–25 s while the bounding box is driven toward the image centre. This raises a legitimate safety question: is it acceptable to position a UAV directly overhead an injured person? This subsection analyses the physical risks, quantifies them, compares mitigation strategies, and justifies the accepted tradeoff.

### S.6.1 Risk Identification

Four hazards are associated with overhead hover:

1. **Rotor downwash on the casualty.** Propeller wash could disturb an injured person, aggravate wounds, or displace lightweight medical dressings.
2. **Mechanical failure leading to vertical descent onto the casualty.** A motor, ESC, or flight-controller failure at hover could cause the drone to fall directly onto the person below.
3. **Psychological distress.** A conscious casualty may experience anxiety from a hovering, audible UAV directly overhead.
4. **Debris entrainment.** Downwash may lift dust, leaves, or loose debris into the casualty’s airway or wounds.

### S.6.2 Quantitative Analysis

**Rotor downwash at ground level.** The EDU-450 airframe has a rotor radius  $R \approx 0.225$  m and an all-up mass of  $\sim 2$  kg. In hover the induced velocity at the rotor disc is given by momentum theory:

$$v_h = \sqrt{\frac{T}{2\rho A}} = \sqrt{\frac{mg}{2\rho\pi R^2}} = \sqrt{\frac{2 \times 9.81}{2 \times 1.225 \times \pi \times 0.225^2}} \approx 5.7 \text{ m s}^{-1} \quad (47)$$

where  $T = mg$  is the hover thrust,  $\rho = 1.225 \text{ kg m}^{-3}$  is air density, and  $A$  is the total disc area of four rotors. The wake contracts and then expands as it descends; at a distance  $h$  below the rotor plane the centreline velocity decays approximately as [57]:

$$v(h) \approx v_h \left( \frac{R}{h} \right)^2 \quad (48)$$

At the centering altitude of  $h = 35$  m:

$$v_{\text{ground}} \approx 5.7 \times \left( \frac{0.225}{35} \right)^2 \approx 5.7 \times 4.1 \times 10^{-5} \approx 0.0002 \text{ m s}^{-1} \quad (49)$$

This is four orders of magnitude below perceptible airflow ( $\sim 0.5 \text{ m s}^{-1}$ ) and six orders below the threshold for displacing lightweight objects ( $\sim 3 \text{ m s}^{-1}$ ). Even at the lower verification altitude of 15 m:

$$v_{15} \approx 5.7 \times \left( \frac{0.225}{15} \right)^2 \approx 0.0013 \text{ m s}^{-1} \quad (50)$$

which remains entirely negligible. Rotor downwash is therefore a non-issue at any altitude above 10 m for a 2 kg quadrotor.

**Crash probability.** Consider a total motor/ESC failure causing uncontrolled descent from 35 m. Under free-fall with aerodynamic drag, the drone drifts laterally with the ambient wind during the  $\sim 2.7$  s fall time ( $t = \sqrt{2h/g}$ ). In a  $5 \text{ m s}^{-1}$  crosswind the lateral displacement is  $\sim 13$  m, giving a roughly circular impact zone of radius 15 m. The probability that the drone strikes a casualty of cross-section  $\sim 1 \text{ m}^2$  within this  $\pi \times 15^2 \approx 707 \text{ m}^2$  area is:

$$P_{\text{hit}} = \frac{A_{\text{casualty}}}{A_{\text{impact}}} \approx \frac{1}{707} \approx 0.14\% \quad (51)$$

This is an upper bound (assumes the failure occurs at the worst moment and the casualty is at the centre). The base rate for catastrophic motor failure in a well-maintained quadrotor is  $\sim 10^{-4}$  per flight hour. For a 20 s hover the

conditional probability is  $P_{\text{fail}} \approx 10^{-4} \times (20/3600) \approx 5.6 \times 10^{-7}$ . The joint probability of failure *and* casualty strike is therefore:

$$P_{\text{fail \& hit}} \approx 5.6 \times 10^{-7} \times 1.4 \times 10^{-3} \approx 8 \times 10^{-10} \quad (52)$$

which is negligibly small—comparable to the probability of being struck by lightning during the same interval.

**Noise exposure.** The EDU-450 produces approximately 65 dB at 1 m. At 35 m the inverse-square law gives  $65 - 20 \log_{10}(35) \approx 34$  dB, which is quieter than a whispered conversation (40 dB). Even at 15 m the level is  $\sim 41$  dB—comparable to a quiet library.

### S.6.3 Comparison of Centering Strategies

Table 93 compares three centering strategies on safety, accuracy, and complexity.

The lateral-offset strategy positions the drone 10 m to the side of the target at the same altitude and infers the target’s ground position from the off-axis pixel location using the GSD projection pipeline. While this eliminates overhead exposure entirely, it reintroduces all geometric error sources that centering is designed to eliminate: ground sampling distance uncertainty, focal length calibration error, and yaw-dependent east–north projection. The resulting CEP<sub>50</sub> of  $\sim 3$  m is essentially identical to the no-centering baseline, negating the purpose of the manoeuvre while adding implementation complexity.

### S.6.4 Industry Practice

Commercial SAR platforms routinely hover above search targets for visual assessment. The DJI Matrice 300 RTK, used by UK mountain rescue teams, hovers at 30–50 m for thermal and optical inspection before directing ground teams [58]. The CAA’s CAP 722 guidance for small UAS operations does not prohibit overhead flight in emergency response contexts, and the JARUS SORA methodology classifies ground risk based on kinetic energy at impact rather than hover position. For a 2 kg drone the kinetic energy class is well within the “low” category regardless of position relative to persons on the ground.

### S.6.5 Accepted Tradeoff

The analysis demonstrates that at 35 m altitude the physical risks of overhead hover are negligible: downwash is unmeasurable, crash-and-hit probability is  $\sim 10^{-9}$ , and noise is below conversational level. The accuracy benefit is substantial—CEP improves from 3.1 m (no centering) to 1.5 m (with centering)—which directly reduces the probability of violating R07 by landing too close to or too far from the casualty. In the SAR context, an inaccurate position estimate that causes the drone to land on top of the casualty is a *more* dangerous outcome than a brief overhead hover at altitude. The centering approach is therefore retained as an available upgrade, and the safety case supports its deployment without additional mitigations beyond the existing 35 m minimum centering altitude.

## T MAVLink Communication Architecture

The companion computer (Raspberry Pi 5) communicates with the CubePilot Orange flight controller via the MAVLink v2 protocol. This section describes the communication topology, the software bridge that enables multi-consumer telemetry, and the failsafe mechanisms that ensure safe behaviour when the link degrades.

### T.1 The Python 3.13 Serial Problem

The Pi runs Python 3.13, which introduced a regression in the `pyserial` library: blocking serial reads on `/dev/ttyAMA0` return corrupted or zero-length data intermittently, making direct UART communication with the Cube unreliable. Early integration testing revealed that `pymavlink`’s serial backend—which depends on `pyserial`—dropped approximately one in five MAVLink frames, causing heartbeat timeouts and command acknowledgement failures. Rather than downgrading the entire Pi environment to Python 3.11, a UDP bridge architecture was adopted.

### T.2 MAVProxy Bridge Architecture

MAVProxy [10] is a lightweight, command-line MAVLink ground station that runs as a background daemon on the Pi. It opens the UART serial port using its own C-level I/O routines (bypassing `pyserial`), parses the incoming MAVLink byte stream, and re-publishes each message to multiple UDP and TCP endpoints. The daemon is started once at boot and remains running for the duration of the flight:

```

mavproxy.py --master=/dev/ttyAMA0 --baudrate=921600 \
  --streamrate=10 \
  --out=udpout:127.0.0.1:14550 \
  --out=udpout:127.0.0.1:14551 \
  --out=tcpin:0.0.0.0:5762

```

This single command establishes the three-consumer topology shown in Figure 41: the mission script receives telemetry on UDP 14550, the diagnostics dashboard on UDP 14551, and Mission Planner connects from the ground-station laptop over TCP 5762 via Wi-Fi. All three consumers receive the full MAVLink message stream simultaneously, and any of them can inject commands back into the link via the same socket. The Pi’s mission script (`main.py`) opens its connection as `udpin:0.0.0.0:14550`, receiving datagrams published by MAVProxy with sub-millisecond local loopback latency.

```

Cube Orange (TELEM2)
|
UART /dev/ttyAMA0 @ 921600 baud
|
[ MAVProxy daemon ]
|
+-- udpout:127.0.0.1:14550 --> main.py / pi_flight.py
|
+-- udpout:127.0.0.1:14551 --> diagnostics.py (monitoring)
|
+-- tcpin:0.0.0.0:5762 -----> Mission Planner (Wi-Fi)

```

**Figure 41:** MAVLink communication topology. MAVProxy bridges the UART serial link to three independent consumers over localhost UDP and network TCP. All consumers receive the full telemetry stream; any can inject commands.

### T.3 Connection Auto-Detection

The configuration module (`config.py`) implements a four-tier auto-detection cascade so that the same codebase runs unmodified on the Pi, in WSL simulation, and on a Windows development laptop:

1. **Environment variable override.** If `DRONE_CONN` is set, it is used verbatim, enabling arbitrary connection strings for testing.
2. **Serial port detection.** If `/dev/ttyAMA0`, `/dev/ttyACM0`, or `/dev/ttyUSB0` exists, the code infers it is running on the Pi and selects `udpin:0.0.0.0:14550`—the MAVProxy UDP endpoint—rather than opening the serial port directly.
3. **WSL gateway.** On Linux with a Microsoft kernel signature, the default gateway IP is extracted and a TCP connection to port 5762 on the Windows host is used, reaching the SITL instance running under Mission Planner.
4. **Localhost fallback.** On Windows, `tcip:127.0.0.1:5762` connects to a local SITL.

This cascade eliminates platform-specific configuration files: the developer pushes code to GitHub, pulls on the Pi, and the correct connection is selected automatically at import time.

### T.4 MAVLink Message Protocol

The system uses a subset of the MAVLink v2 Common message set [59]. Table 94 summarises the messages exchanged during a typical mission.

Position commands use the `MAV_FRAME_GLOBAL_RELATIVE_ALT_INT` frame, which expresses latitude and longitude as fixed-point integers ( $\times 10^7$ ) and altitude relative to the take-off point. Velocity commands for visual servoing are issued in body frame (`MAV_FRAME_LOCAL_NED`) with a rotation from body to NED applied in software using the current yaw angle. A `NavigationController` wrapper validates all outbound commands, rejecting NaN coordinates or altitudes outside the 0m to 400m safety band before transmission.

### T.5 Link Loss Detection and Recovery

Two independent watchdogs protect against communication failures:

**Heartbeat timeout.** The mission script timestamps each `HEARTBEAT` message from the autopilot. If the `recv_match` call raises a `ConnectionResetError`, `BrokenPipeError`, or `OSError`, the software immediately sends a best-effort `MAV_CMD_DO_SET_MODE` for RTL (Return to Launch) and transitions to the terminal `DONE` state. Even if this command fails to reach the Cube, ArduCopter’s own GCS failsafe (`FS_GCS_ENABLE`) triggers an independent RTL after 5s of missing heartbeats from the companion computer, providing a hardware-level safety net.

**GPS degradation monitor.** The `GPS_RAW_INT` stream is monitored continuously. If the fix type drops below 3D fix or the satellite count falls below 6 for more than 5s, the software commands an emergency RTL and transitions to the `LANDING` state. A throttled warning is printed every 3s during degradation, displaying a countdown to RTL so the operator can intervene via the RC kill switch if appropriate.

## T.6 Latency Characteristics

The communication chain introduces three sources of latency: UART serialisation, MAVProxy forwarding, and UDP loopback. At 921 600 baud, a typical 32-byte MAVLink v2 frame requires approximately 0.35 ms for serial transmission. MAVProxy’s forwarding loop adds less than 1 ms of buffering overhead, and localhost UDP delivery is effectively instantaneous. The total one-way latency from Cube to Python script is therefore on the order of 1 ms to 2 ms—well within the requirements of the 4.8 FPS detection loop, where each decision cycle spans approximately 210 ms (dominated by TFLite inference). The MAVLink telemetry stream rate is configured to 10 Hz via MAVProxy’s `--streamrate` flag, ensuring that GPS position and attitude updates arrive at least twice per inference frame.

## U Raw Inference vs. Effective Pipeline Throughput

Deploying a detection model on an embedded platform requires understanding the difference between raw inference speed and the throughput of the complete processing pipeline. Published inference benchmarks for edge-deployment frameworks report *raw inference* latency: the time elapsed between submitting a preprocessed tensor to the model and receiving the output tensor. This metric is useful for comparing backends but overstates the frame rate achievable by an integrated detection system. In practice, each detection cycle includes several stages beyond the forward pass, and the *effective pipeline FPS*—the rate at which the system produces actionable detection results—is the operationally relevant metric. This section quantifies the overhead budget and its implications for each backend evaluated on the Raspberry Pi 5.

### U.1 Pipeline Stages and Overhead Budget

The onboard detection pipeline executes the following stages sequentially within a single thread:

1. **Camera capture** (~8 ms). Frame acquisition from the IMX296 sensor via `picamera2`, including DMA transfer to user space.
2. **Lens undistortion** (~1.5 ms). Precomputed `cv2.remap()` correction using calibration parameters (RMS = 0.399).
3. **Preprocessing** (~2 ms). Resize from  $1456 \times 1088$  to  $640 \times 640$ , BGR→RGB conversion, float32 normalisation, and batch-dimension expansion.
4. **Model inference** (backend-dependent). The forward pass through YOLOv8n.
5. **Postprocessing** (~0.5 ms). Non-maximum suppression, confidence thresholding, and bounding-box coordinate extraction.
6. **GPS estimation** (~0.3 ms). Pixel-to-world coordinate projection via the `GeoTransformer`, inverse-variance weighted clustering, and Kalman filter update.
7. **State machine dispatch** (~0.1 ms). Detection result consumed by the mission state machine, geofence check, and MAVLink command generation.
8. **Stream encoding** (~1 ms). JPEG compression of the annotated frame for the MJPEG ground station feed.

The total pipeline overhead—stages 1–3 plus 5–8, excluding inference—sums to approximately 12 ms for the TFLite backend. This overhead is largely constant across frame rates because it consists of fixed-cost operations (memory copies, a single `remap` call, one JPEG encode) rather than workloads that scale with inference speed.

### U.2 Measured and Projected Results

Table 96 compares raw inference throughput with effective pipeline throughput for each backend evaluated on the Raspberry Pi 5.

Several observations merit discussion:

**TFLite FP32 (deployed baseline).** At 196 ms raw inference, the 12 ms pipeline overhead contributes only 5.8% of the total frame time (208 ms), yielding 4.8 effective FPS. Inference dominates the frame budget so completely that pipeline optimisation—threaded capture, asynchronous GPS estimation—would yield negligible improvement. The system is unambiguously *inference-bound*.

**NCNN FP32 (benchmarked, not deployed).** NCNN was exported via ONNX and benchmarked on the Pi 5 at approximately 77 ms raw inference—a  $2.5\times$  improvement over TFLite FP32. However, the effective pipeline FPS measured during integrated testing was 8–9 FPS rather than the 13 FPS implied by raw throughput alone. The


discrepancy is attributed to two factors. First, NCNN uses a different preprocessing path: the `ncnn.Mat` tensor constructor performs its own memory layout conversion, adding  $\sim 3$  ms relative to TFLite’s NumPy-based path. Second, NCNN’s internal thread scheduling occasionally contends with the preprocessing thread for L2 cache bandwidth, introducing sporadic  $\sim 5$  ms stalls. The combined overhead of 15 ms is modest in absolute terms but represents 16% of the total frame time at this inference speed—a qualitative shift from the TFLite regime where overhead was negligible. Despite the  $\sim 2\times$  effective throughput improvement, NCNN was not deployed for flight testing. The `ncnn` Python package on Python 3.13/aarch64 required manual C++ compilation on the Pi, introducing a fragile build dependency unsuitable for field deployment. A pre-compiled wheel or a Debian package would resolve this issue; NCNN remains the recommended upgrade path for future missions (Section G).

**TFLite FP16 (projected).** The Cortex-A76 implements native half-precision fused multiply-accumulate (FMLA) at twice the FP32 throughput. Enabling the `TFLITE_XNNPACK_DELEGATE_FLAG_FORCE_FP16` flag is projected to halve inference latency to  $\sim 100$  ms. Because TFLite FP16 shares the identical preprocessing and postprocessing path as FP32, pipeline overhead remains 12 ms, yielding a projected 8.9 effective FPS. This path has the lowest integration risk: no model re-export, no new dependencies, and a single flag change in the interpreter initialisation.

**Ultralytics/PyTorch (laptop reference).** On the Pi 5, PyTorch inference exceeds 1.2 s per frame. The 5 ms pipeline overhead is lower than TFLite because Ultralytics performs its own preprocessing internally (resize, normalisation), eliminating the external NumPy path. This backend is used exclusively on the development laptop (with optional CUDA acceleration) for training, export, and video analysis.

### U.3 Bottleneck Transition

Figure 42 illustrates how the inference-to-overhead ratio shifts across backends.



**Figure 42:** Frame time decomposition: inference vs. pipeline overhead for each backend. As raw inference latency decreases, overhead transitions from negligible (5.8% for TFLite FP32) to significant (16.3% for NCNN).

At 4.8 FPS (TFLite FP32), inference consumes 94.2% of the frame budget. Halving inference time to  $\sim 100$  ms (FP16) raises the overhead fraction to 10.7%. At 13 FPS raw (NCNN), overhead reaches 16.3%, and pipeline optimisation—primarily threaded camera capture overlapping with inference—becomes worthwhile. A threaded capture design hiding the 8 ms camera read behind inference would recover  $\sim 1$  additional FPS at the NCNN operating point, bringing effective throughput to  $\sim 11.8$  FPS. At the TFLite FP32 operating point, the same optimisation would recover only 0.2 FPS—confirming that threading effort is better deferred until the inference bottleneck is first addressed.

### U.4 Operational Implications

The distinction between raw and effective FPS has direct consequences for mission planning. At the deployed 4.8 effective FPS and a search speed of  $8 \text{ m s}^{-1}$ , the ground distance between consecutive analysed frames is 1.67 m—ensuring dense overlap across the 32.2 m horizontal field of view. A target at the centre of the swath is observed in at least 14 frames during a single pass, providing high detection confidence even at the modest frame rate.

Upgrading to NCNN ( $\sim 10.9$  effective FPS) would reduce the inter-frame ground distance to 0.73 m, enabling either higher search speeds at the same detection density or lower-altitude missions where the narrower field of view demands faster frame acquisition. The FP16 path ( $\sim 8.9$  effective FPS) offers a similar improvement with zero deployment friction, making it the recommended first optimisation step for future flights.

## V Flight Path Optimisation

Generating a search pattern is a necessary but insufficient step; the quality of the mission depends critically on the parameter choices that govern the pattern’s geometry. The preceding sections describe the algorithm that generates

the search pattern; this section analyses the engineering trade-offs behind those parameter choices. Every parameter interacts with at least one other—altitude affects lane width, which affects turn count, which affects energy—so parameter selection is inherently a multi-objective optimisation problem. The values in Table 124 represent the operating point that balances detection probability, coverage time, and energy expenditure for the specific survey polygon and sensor suite used in this project.

## V.1 Diagonal vs Axis-Aligned Scanning

The planner aligns scan lines with the polygon’s longest edge by rotating the binary mask so that this edge becomes horizontal (Section 2.4). This minimises the number of perpendicular scan lines and therefore the number of U-turns, which are the dominant source of wasted time and energy in a boustrophedon pattern. For the AENGM0074 survey polygon—an elongated pentagon oriented roughly northeast–southwest—the optimal scan angle falls in the  $65^\circ$  to  $75^\circ$  range, confirmed by the 216-configuration parametric sweep (Section 1.2). Deviating by  $30^\circ$  from this optimum increases the scan-line count by approximately 40 %.

An alternative spiral-inward pattern was also considered, in which the drone follows a contracting spiral from the polygon boundary towards the centre. While spiral patterns can be efficient for convex search areas, analysis indicated that they provide no coverage guarantee for irregular polygons with concave features: re-entrant vertices cause the spiral to skip interior regions, leaving unpredictable gaps that depend on the polygon geometry. The AENGM0074 polygon contains a concave notch along its western boundary adjacent to the SSSI, making the spiral approach unsuitable without substantial additional waypoint generation logic. The boustrophedon pattern, by contrast, guarantees complete coverage of any simply-connected polygon by construction.

An alternative strategy considered during design was *diagonal scanning at  $45^\circ$  to the prevailing wind direction*. The rationale is that crosswind drift displaces the drone laterally during each scan leg; if the scan direction is perpendicular to the wind, drift accumulates along the inter-lane axis and can open coverage gaps between adjacent swaths. Scanning diagonally distributes the drift component across both the along-track and cross-track axes, reducing the effective cross-track displacement by a factor of  $\cos(45^\circ) \approx 0.71$ . However, this benefit comes at the cost of additional scan lines whenever the diagonal does not coincide with the polygon’s long axis. For the survey polygon, wind-aligned scanning at  $45^\circ$  to a typical westerly wind would require approximately 30 % more scan lines than the longest-edge alignment. Given that the 20 % overlap margin (Section V.5) already absorbs typical crosswind displacements of 1 m to 2 m, the additional lines were judged unnecessary. The longest-edge alignment was therefore retained, with the overlap margin serving as the primary defence against wind-induced coverage gaps.

## V.2 Lane Width vs Altitude Trade-Off

The lane width is a direct function of search altitude through Eq. (3): higher altitude widens the ground footprint, requiring fewer scan lines but shrinking the target in the camera frame. Table 97 quantifies this trade-off for the survey polygon.

At 20 m the mannequin subtends 63 model-input pixels (142 sensor pixels)—comfortably above YOLOv8n’s detection threshold—but the narrow 18.4 m footprint roughly doubles the number of scan lines compared to 35 m, increasing both flight time and the number of U-turns. At 50 m the target shrinks to 25 model-input pixels (57 sensor pixels), approaching the limit at which detection confidence degrades significantly. The 35 m operating point provides a 36-pixel model-input target extent (81 sensor pixels)—sufficient for reliable detection at the 0.2 confidence threshold used in the system—while maintaining a practical scan-line count. This altitude also leaves margin for the verification descent to 15 m, where the target subtends 84 model-input pixels for high-confidence confirmation.

## V.3 Altitude-Dependent Speed Schedule

A fixed search speed is suboptimal because the camera’s ground footprint—and therefore the time a target remains in view—varies with altitude. At 20 m altitude the along-track footprint is approximately 14 m; at  $10 \text{ m s}^{-1}$  a ground target traverses the frame in 1.4 s, leaving only 5–7 inference frames at 4.8 FPS (the Pi 5 TFLite throughput). Reducing the speed to  $6 \text{ m s}^{-1}$  extends the dwell time to 2.3 s ( $\sim 11$  frames), significantly improving the probability of at least one high-confidence detection.

The speed schedule implements a linear interpolation between two anchor points:

$$v(h) = \begin{cases} v_{\text{low}} = 6 & h \leq 20 \text{ m} \\ v_{\text{low}} + \frac{h - h_{\text{low}}}{h_{\text{high}} - h_{\text{low}}} (v_{\text{high}} - v_{\text{low}}) & 20 < h < 50 \text{ m} \\ v_{\text{high}} = 10 & h \geq 50 \text{ m} \end{cases} \quad (53)$$

At the nominal 35 m search altitude this yields  $v = 8.0 \text{ m s}^{-1}$ . The anchor values were chosen conservatively:  $6 \text{ m s}^{-1}$  at low altitude keeps the target in view for at least 10 inference cycles, while  $10 \text{ m s}^{-1}$  at high altitude is below the

airframe’s maximum cruise speed, leaving margin for wind gusts. Transit segments to and from the search area use a separate, higher speed of  $15\text{ m s}^{-1}$  since no detection is required.

## V.4 Turn Strategy

At each strip endpoint the drone must execute a  $180^\circ$  U-turn to begin the adjacent strip. A naive implementation commands the drone to fly to the exact waypoint, decelerate to a hover, rotate, and accelerate along the new heading. This produces two problems: the deceleration–acceleration cycle wastes approximately 2 s per turn, and the abrupt attitude change (pitch-up to decelerate, pitch-down to accelerate) disturbs the camera pointing angle, creating a brief period during which the ground footprint shifts unpredictably.

The Bézier smoothing pass (Section 2.4) replaces each sharp corner with a three-point quadratic arc evaluated at  $t \in \{0.25, 0.50, 0.75\}$ . The arc allows the drone to maintain forward momentum through the turn, following a smooth trajectory rather than stopping at the vertex. Combined with ArduPilot’s waypoint acceptance radius—which advances to the next waypoint once the drone passes within a configurable distance rather than requiring an exact position hold—this yields fluid, banking turns that keep the camera platform stable. The arc adds approximately 2 s of flight time per turn compared to an instantaneous direction change, but this is offset by the elimination of the hover–rotate–accelerate cycle. For a typical 12-line pattern with 11 U-turns, Bézier smoothing reduces total turn time by an estimated 15 s while simultaneously improving camera stability during the portions of the flight that are most susceptible to coverage gaps.

## V.5 Overlap Margin Analysis

The 20 % lateral overlap between adjacent swaths is the system’s primary defence against coverage gaps. Three error sources motivate this margin:

1. **GPS drift.** Consumer-grade GNSS receivers exhibit a circular error probable (CEP) of 2 m to 3 m under open sky. In the worst case, two adjacent scan legs drift in opposite directions, producing up to 6 m of relative displacement.
2. **Crosswind displacement.** A sustained crosswind displaces the drone laterally during each scan leg. At  $7\text{ m s}^{-1}$  groundspeed with a  $3\text{ m s}^{-1}$  crosswind, the drift over a 200 m leg is approximately 1.5 m (assuming partial compensation by the flight controller’s heading hold).
3. **Attitude variations.** Pitch and roll during turbulence shift the camera footprint laterally. At 35 m altitude, a  $5^\circ$  roll displaces the footprint centre by  $35 \times \tan(5^\circ) \approx 3\text{ m}$ .

At 35 m altitude, the 20 % overlap produces 6.4 m of redundant coverage between adjacent swaths. The combined worst-case displacement from all three sources (assuming independent, uncorrelated errors) is approximately  $\sqrt{6^2 + 1.5^2 + 3^2} \approx 6.9\text{ m}$ —marginally exceeding the overlap. However, the worst-case GPS scenario (maximum opposite drift on consecutive legs) is statistically unlikely; the 50 % probability displacement is approximately 3 m, leaving 3.4 m of residual overlap under typical conditions. Increasing the margin to 30 % was considered but rejected because it would add approximately two additional scan lines to the pattern, increasing flight time by 10 % for marginal improvement in an already-adequate coverage guarantee.

## V.6 Start Corner Optimisation

The boustrophedon pattern can begin from any of four corners of the rotated bounding rectangle. The planner evaluates all four and selects the corner closest to the drone’s current GPS position (Section 2.4, step 5). This minimises the transit distance from the take-off point (or current hover position, in the case of a rescan) to the first scan line. For the AENGM0074 polygon with the take-off point at its western edge, this optimisation saves 10 s to 30 s of transit time depending on the scan angle and wind conditions. When the pattern is reversed (bottom-to-top instead of top-to-bottom), the strip traversal direction is also flipped to maintain the nearest-start property, ensuring that the zigzag begins in the correct direction regardless of which corner is selected.

## V.7 PLB Focus Area Redirect

When a Personal Locator Beacon (PLB) signal is received during the wide-area search, the planner abandons the current pattern and regenerates a new lawnmower grid over a smaller focus polygon centred on the PLB coordinates (Section 2.4). The focus area pattern uses the same rotated-mask algorithm but with two parameter changes: the search speed drops from the altitude-dependent schedule to a fixed  $5\text{ m s}^{-1}$ , and the start corner is recomputed from the drone’s current position rather than the take-off point. The reduced speed increases dwell time per target, improving detection probability in the high-priority zone. Existing detection clusters, rejected false positives, and items of interest are preserved across the pattern switch, preventing re-investigation of previously classified targets. The waypoint index and rescan counter both reset, so the focus area is searched from scratch regardless of how much of the original pattern had been completed.

## W Payload Release Mechanism

The mission’s final objective is to deliver a first-aid kit to the vicinity of the casualty. Achieving this reliably requires careful integration of mechanical hardware, servo control logic, and a delivery strategy that balances accuracy against safety. This section describes the release hardware, the rationale for deploying from a low hover rather than landing beside the casualty, and the servo actuation sequence.

### W.1 Tarot Double-Throw Release Mechanism

The payload is suspended beneath the airframe by a university-provided Tarot servo-actuated release mechanism, controlled via a single PWM channel on the Cube Orange (AUX output channel 9, configured as `SERVO_CHANNEL` in `config.py`). The mechanism supports two-stage actuation: a partial opening at PWM 1300 (`SERVO_PARTIAL_PWM`) and a full release at PWM 1100 (`SERVO_FULL_PWM`). The two-stage design is deliberate—during transit the servo is held at PWM 1500 (closed), and the partial-open intermediate step prevents premature payload separation caused by airframe vibration or sudden manoeuvres that might overcome a single-latch mechanism. Only the second command fully disengages the retaining hooks.

### W.2 Why Release from Altitude, Not Land

An alternative design would land the drone directly beside the casualty and release the payload on the ground. This was rejected for five reasons:

1. **Casualty safety.** Landing within a few metres of an injured person exposes them to rotor downwash, debris blown by the propellers, and the risk of a ground-level tip-over. Requirement R07 explicitly mandates a 5 m exclusion zone around the casualty for this reason.
2. **Terrain uncertainty.** SAR sites are typically unimproved terrain—long grass, slopes, rocks, or mud. Landing gear contact with uneven ground can cause a roll-over or motor stall, potentially destroying the aircraft and the payload. A low hover avoids surface interaction entirely.
3. **GPS accuracy.** The system’s GPS-based target localisation has a measured  $\text{CEP}_{50} = 2.3\text{ m}$  (Section 2.6). A ground landing adds the full position error to the touchdown point, whereas an aerial release from 3 m altitude disperses the payload over a small area predictable from simple ballistics (free-fall time  $t = \sqrt{2h/g} \approx 0.78\text{ s}$ , negligible lateral drift in light wind).
4. **Reduced low-altitude exposure.** Ground-effect aerodynamics, obstacle proximity, and limited sensor coverage at very low altitude all increase the probability of a mishap. Releasing from 3 m and immediately climbing to transit altitude minimises time in the highest-risk flight regime.
5. **Operational simplicity.** A hover-and-release sequence requires only position-hold commands followed by two servo actuations. A ground landing requires additional logic for touchdown detection, motor shutdown, ground-level release, re-arming, and takeoff—each step introducing failure modes that are difficult to test without flight trials.

### W.3 The 7.5 m Offset

Requirement R07 specifies that the drone must land (or, in this design, deploy the payload) within 10 m of the casualty but not within 5 m. The system therefore computes a delivery point offset 7.5 m from the estimated casualty position in an operator-selected cardinal direction (N/E/S/W), implemented in `landing_offset_7_5m()` in `gps_utils.py`. The 7.5 m value is the midpoint of the 5 m to 10 m annular band, providing 2.5 m of margin on both sides and maximising the probability of compliance given the Gaussian-distributed GPS estimation error. Section S shows that with multi-detection fusion ( $N_{\text{eff}} = 6$ ,  $\sigma_{\text{fused}} \approx 0.80\text{ m}$ ), the probability of the delivery point falling within the required band exceeds 99%.

### W.4 Servo Actuation Sequence

The payload deployment is executed during the `HOVER_TARGET` state, implemented in `state_machine.py`. The full sequence proceeds as follows:

1. **Approach** (`APPROACH` state): the drone flies to the offset point at  $3\text{ m s}^{-1}$  and descends to 3 m AGL. The state transitions to `HOVER_TARGET` once the aircraft is within 2 m laterally and below 4 m altitude.
2. **Stabilise** ( $t = 0\text{--}3\text{ s}$ ): the drone holds position at 3 m AGL, allowing oscillations from the approach to damp out. The servo remains closed (PWM 1500).
3. **Partial release** ( $t = 3\text{ s}$ ): a `MAV_CMD_DO_SET_SERVO` command sets channel 9 to PWM 1300, partially opening the mechanism. This loosens the payload without dropping it, confirming that the servo responds before committing to full release.
4. **Full release** ( $t = 6\text{ s}$ ): a second `MAV_CMD_DO_SET_SERVO` sets channel 9 to PWM 1100, fully disengaging the retaining hooks and allowing the payload to free-fall approximately 3 m to the ground.

5. **Close and depart** ( $t = 15$  s): the servo is returned to PWM 1500 (closed) to prevent snagging during climb-out. The drone ascends to search altitude and enters RETURN\_TRANSIT or RETURN\_HOME to fly back to the takeoff location via the pre-planned transit corridor.

The 15 s total hover time balances deployment reliability against battery consumption and low-altitude exposure time. A visual animation of the sequence is rendered on the ground-station HUD during the hover to provide the operator with real-time feedback on deployment progress.

## W.5 Why Visual Servo Centering Was Not Required

The system includes CENTERING and DESCENDING states capable of visually servoing the drone above the target to improve localisation accuracy (Section S). However, for the payload release mission this pipeline was bypassed. The CEP<sub>50</sub> of 2.3 m from the GPS estimation pipeline is already well within the 2.5 m margin provided by the 7.5 m offset, making visual servo centering unnecessary for R07 compliance. As demonstrated in Section S, simulation comparison of twenty runs with and without centering showed that the direct offset approach achieved 95% R07 compliance without the added risks of low-altitude manoeuvring, visual servo fragility below 10 m, and extended hover time over the casualty. The centering capability is retained in the codebase as a future enhancement for precision-drop missions where tighter tolerances are required.

## W.6 Future Work: Precision Drop

The current TFLite inference backend achieves approximately 5 FPS on the Raspberry Pi 5, which is sufficient for search-altitude detection but too slow for responsive visual servo control during descent. Migrating to the NCNN inference framework increases throughput to approximately 15 FPS, enabling closed-loop centering with update rates comparable to the flight controller’s position-hold bandwidth. Combined with the existing centering pipeline (Section S), this could reduce the effective CEP to below 1 m, permitting precision drops from 5 m altitude without the 7.5 m lateral offset—useful for scenarios requiring delivery directly to the casualty’s position rather than to a nearby clear area.

# X Model Comparison and Future Vision Improvements

Iterative model development—where each training round addresses shortcomings discovered in the previous one—is essential for building a detection system that generalises beyond its training data. The vision pipeline underwent three training iterations, each informed by lessons from the previous round. This section compares the model generations, discusses the COCO fallback strategy, examines the tiling versus full-frame trade-off, and identifies improvement directions for future deployments.

## X.1 Model Generations

Table 99 summarises the three model generations trained during the project. All share the YOLOv8n architecture (3.2 M parameters) and export to the same TFLite tensor shapes ([1, 640, 640, 3] input, [1, 5, 8400] output), ensuring drop-in interchangeability at deployment.

**v1** (`sar_640`) was the initial baseline, trained exclusively on 200 synthetic composite images generated at  $640 \times 640$  using a single satellite background (`map.jpg`). It demonstrated proof-of-concept detection in simulation but exhibited frequent false positives on terrain features—shadows, path markings, and field equipment—because the training set contained no negative examples and offered limited background diversity.

**v2** (`sar_1280`) increased the training resolution to  $1280 \times 1280$ , allowing the network to learn finer spatial features during training. This yielded a marginal improvement in detection quality at higher altitudes (where the target subtends fewer pixels), but since the training data remained purely synthetic from the same single background source, the false-positive problem persisted. The key lesson was that *resolution alone cannot compensate for limited data diversity*.

**v3** (`sar_v2_1088`) addressed both limitations simultaneously. Training resolution was set to  $1088 \times 1088$  (matching the native camera sensor height), and the dataset was restructured with three complementary sources: 300 synthetic composites using real flight-video backgrounds for diverse terrain textures, 16 manually labelled real aerial frames to anchor the model to authentic appearance, and 50 negative frames to explicitly teach false-positive suppression (Section C). This combination produced the highest validation metrics ( $\text{mAP}_{50} = 0.995$ ,  $\text{mAP}_{50-95} = 0.828$ ) and, critically, near-zero false positives during video replay analysis across 15–50 m altitude.

All three models produce identical inference latency on the Pi ( $\sim 207$  ms at 4.8 FPS) because the TFLite runtime resizes all inputs to  $640 \times 640$  regardless of training resolution. The v3 file size is larger (11.7 MB vs. 3.2 MB) because it was exported as float32 without quantisation; the earlier models used a more aggressive compression pipeline. This size difference has no impact on inference speed, as model loading is a one-time startup cost.

## X.2 COCO Person Detector as Fallback

A pre-trained YOLOv8n model exported from the standard COCO dataset [12] was retained as a flight-day backup (`models/human.tflite`, 13 MB, 80 classes). This model detects people without any SAR-specific training, providing a zero-effort deployment option if the custom model proved unreliable in field conditions.

Testing revealed two trade-offs. On the positive side, the COCO detector reliably detected people in aerial video at altitudes up to 30 m, confirming that transfer from ground-level training data generalises to moderate aerial perspectives. On the negative side, the 80-class output space produced substantially more false positives: backpacks, dogs, bicycles, and even shadows were flagged as detections, requiring either aggressive confidence thresholding (which risked suppressing true positives) or downstream filtering logic. The custom single-class model, having learned to classify “dummy versus everything else”, achieved a far more favourable precision–recall balance.

The COCO model was therefore designated as a backup rather than a primary detector: if the custom model failed catastrophically in the field (e.g., due to an unforeseen domain shift), the operator could swap to the COCO model in under 30 seconds via `cp models/human.tflite best.tflite`, accepting the higher false-positive rate as preferable to zero detection capability.

## X.3 Overfitting Risk in Single-Class Detection

Training a single-class detector on a narrow domain introduces a specific overfitting risk: the model may learn *background shortcuts* rather than genuine target features. Because all positive training images share a common context (green field, overcast lighting, UK terrain), the network could learn to trigger on the co-occurrence of a small bright patch against grass rather than on the mannequin’s actual shape and colour signature. This shortcut would fail when the background changes—different seasons, terrain types, or lighting conditions.

The negative images mitigate this risk by presenting the same background without a target, forcing the network to learn discriminative features of the target itself rather than background correlations. However, the 50 negatives are drawn from a single flight on a single day, limiting the diversity of background contexts that the model has been taught to suppress.

A more robust long-term approach would be multi-class detection—training the model to distinguish between “dummy”, “person”, “equipment”, and “background”—which would force the network to learn category-specific features rather than binary shortcuts. This requires a substantially larger and more diverse labelled dataset than was available within the project’s single-flight-day constraint.

## X.4 Tiling vs. Full-Frame Detection

At search altitudes of 30–35 m, the mannequin subtends approximately 36–42 pixels in height in the  $640 \times 640$  model input (scale factor  $640/1456 = 0.4396$  from the  $1456 \times 1088$  sensor frame). This is above the empirical minimum for YOLOv8 detection ( $\sim 20$  pixels) [11], but the target *width* drops to 10–12 model pixels, leaving limited margin. Two detection strategies were therefore evaluated:

1. **Full-frame resize.** The entire  $1456 \times 1088$  frame is resized to  $640 \times 640$  in a single pass. One inference call per frame; latency  $\sim 207$  ms.
2. **Tiled detection.** The frame is divided into overlapping  $640 \times 640$  tiles (typically  $2 \times 2$  with 25% overlap), each processed independently. Four inference calls per frame; latency  $\sim 830$  ms.

Tiling preserves higher effective resolution per tile, improving detection of very small targets at extreme altitude. However, testing on DJI flight video (`video_test.py`) showed that full-frame detection achieved reliable detection across the entire 15–50 m altitude range with the v3 model, with no altitude band where tiling produced detections that full-frame missed. The  $4\times$  inference cost of tiling—reducing throughput from 4.8 FPS to 1.2 FPS—would significantly increase the probability of missing a target during a search pass at the nominal 8 m/s (altitude-dependent schedule at 35 m), where each frame at 1.2 FPS covers approximately 6.7 m of ground track.

Full-frame detection was therefore selected as the deployment strategy. Tiling remains available as a configurable option in the video analysis tools for post-flight forensic review, where latency is not a constraint.

## X.5 Future Vision Improvements

Six directions are identified for improving detection reliability and operational capability beyond the current system:

1. **Expanded real training data.** Each future flight day could yield hundreds of labelled frames using the existing annotation tool (`label_tool.py`). Increasing the real-image fraction from 4.4% to 30–50% of the dataset would substantially reduce the synthetic-to-real domain gap, the single largest source of generalisation risk in the current model. Frames should be collected across multiple days, weather conditions, and terrain types to build environmental diversity.

2. **Larger model variant.** YOLOv8s (11.2 M parameters) offers improved accuracy at the cost of  $\sim 400$  ms inference on Pi TFLite—viable at reduced search speeds of 3 m/s, where the longer frame-to-frame interval compensates for slower processing. The existing pipeline would require no architectural changes, only a model file swap.
3. **Multi-class detection.** Extending the label space to four classes—dummy, person, equipment, false-positive-class—would force the network to learn category-discriminative features and provide the operator with richer information at the verification stage. This requires a labelled dataset with representative examples of each class, which could be built incrementally across flight campaigns.
4. **Next-generation architectures.** YOLO11 [11] and future iterations introduce architectural refinements (attention mechanisms, improved feature pyramid designs) that may improve small-object detection without increasing parameter count. The modular export-and-deploy pipeline means any YOLO-family model producing  $[1, 5, N]$  output tensors can be integrated without code changes.
5. **Thermal camera fusion.** A FLIR Lepton 3.5 module ( $160 \times 120$ , LWIR 8–14  $\mu\text{m}$ ) would provide body-heat detection independent of visible appearance, enabling night and low-visibility operations [41]. Decision-level fusion—requiring confirmation in both visual and thermal modalities—would suppress the false positives inherent in each sensor alone (sun-heated objects in thermal, colour-similar debris in visual).
6. **Active learning.** Flagging detections with confidence scores in the 0.2–0.5 range (near the current 0.2 threshold) for human review would create a feedback loop: uncertain predictions are labelled by the operator post-flight and added to the training set, progressively improving the model in the regions of feature space where it is least confident. This approach is particularly efficient for rare-event detection where exhaustive labelling is impractical [60].

These improvements are ordered by implementation effort: items 1–2 require only additional data and a retraining cycle, items 3–4 require dataset restructuring and architecture evaluation, and items 5–6 require new hardware or workflow infrastructure. The modular vision architecture—where all detection logic is encapsulated behind a single function interface—ensures that any of these improvements can be integrated without modifying the state machine, ground station, or flight control code.

## Y Multi-Objective Search Optimisation

Selecting the operating point for a boustrophedon search pattern requires balancing five competing objectives—coverage, mission time, energy, detection probability, and regulatory safety margin—none of which can be optimised independently. This section presents the quantitative analysis that informed the final parameter selection, including a 216-configuration parametric sweep across altitude and scan angle.

*Note: The definitive path optimisation analysis appears in Appendix AA; the formal multi-objective problem statement in Appendix AC; the physics-based simulation and Pareto analysis in Appendix AB; and the parameter derivation chain in Appendix AD. This appendix provides supplementary detail and is retained for completeness.*

The search pattern parameters described in the preceding sections—altitude, speed, scan angle, overlap margin, NFZ buffer—do not have independently optimal values. Each parameter trades off against at least two competing objectives, forming a multi-dimensional optimisation landscape. Five objectives are in tension:

1. **Coverage completeness ( $C$ ):** the fraction of the survey polygon imaged by at least one camera frame. Complete coverage ( $C = 1.0$ ) requires every point to fall within at least one swath.
2. **Mission time ( $T$ ):** total elapsed time from takeoff to landing. In SAR operations, every minute of delay reduces casualty survival probability; the UK Maritime and Coastguard Agency survival curves show a steep decline in hypothermia survivability after 60 min of exposure [61].
3. **Energy consumption ( $E$ ):** total energy drawn from the battery. The EDU-450 airframe carries a 6S 5200 mAh LiPo (115.4 Wh), imposing a hard flight-time ceiling of approximately 18 min at search power levels.
4. **Detection probability ( $P_d$ ):** the probability that a casualty present in the search area is detected by at least one inference frame. This depends on the number of frames in which the target appears, the per-frame confidence, and the confirmation threshold.
5. **NFZ safety margin ( $d_{\text{NFZ}}$ ):** the minimum distance between any waypoint and the SSSI no-fly zone boundary. Larger margins reduce regulatory risk but sacrifice coverage of the polygon edge adjacent to the SSSI.

These objectives form a Pareto frontier: no single configuration simultaneously maximises coverage, minimises time and energy, maximises detection probability, and maximises NFZ margin. The operating point must be chosen by engineering judgement informed by quantitative analysis. The following subsections present that analysis.

### Y.1 Energy Efficiency Simulation

To systematically explore the parameter space, a simulation swept **216 configurations** comprising 6 altitudes (20, 25, 30, 35, 40, and 50 m)  $\times$  36 scan angles ( $0^\circ$  to  $175^\circ$  in  $5^\circ$  steps), with speed determined by the altitude-dependent


schedule. For each configuration the lawnmower pattern was generated over the AENGM0074 survey polygon using the rotated-mask algorithm, and the following energy model was applied:

$$E_{\text{total}} = P_{\text{hover}} \cdot t_{\text{hover}} + P_{\text{forward}} \cdot t_{\text{forward}} + E_{\text{turn}} \cdot n_{\text{turns}} \quad (54)$$

where:

- $P_{\text{hover}} = 150 \text{ W}$  is the constant power to remain airborne, estimated from the simplified power model used in the optimisation script (Section AA.2);
- $P_{\text{forward}}(v) = P_{\text{hover}} + P_{\text{drag,ref}}(v/v_{\text{ref}})^2$ , where  $P_{\text{drag,ref}} = 50 \text{ W}$  at  $v_{\text{ref}} = 5 \text{ m s}^{-1}$ , giving  $\approx 350 \text{ W}$  at  $10 \text{ m s}^{-1}$  [62];
- $E_{\text{turn}} \approx 0.083 \text{ Wh}$  per U-turn, corresponding to  $2 \text{ s}$  of hover-only power ( $150 \text{ W} \times 2 \text{ s} / 3600 \text{ s h}^{-1}$ ) for the deceleration–reacceleration cycle.

Table 100 presents the five most and least energy-efficient configurations.



**Figure 43:** Energy efficiency across the 216-configuration parameter sweep. Lower-energy configurations cluster at high altitude and optimal scan angle; the selected operating point ( $35 \text{ m}$ ,  $70^\circ$ ,  $8 \text{ m s}^{-1}$ ) trades minimal energy for reliable detection.


The selected operating point (rank 5:  $70^\circ$ ,  $35 \text{ m}$ ,  $8 \text{ m s}^{-1}$ ) is not the global energy minimum—that distinction belongs to the  $50 \text{ m}/10 \text{ m s}^{-1}$  configuration—but it is the most energy-efficient configuration that maintains reliable detection. At  $50 \text{ m}$  the mannequin subtends only 25 model-input pixels (57 sensor pixels), approaching the detection confidence limit; at  $35 \text{ m}$  it subtends 36 model-input pixels (81 sensor pixels) with a comfortable margin above the 20-pixel minimum. The energy penalty for this safety margin is  $4.4 \text{ Wh}$  (54 % more than the absolute minimum), which consumes only 10.9 % of the battery capacity and leaves over 14 min of flight endurance—well above the 2 min search time.

## Y.2 Rotation Angle Analysis

The scan angle is the single parameter with the largest impact on turn count and therefore on both energy and time. Three representative angles illustrate the sensitivity for the AENGM0074 polygon:

The longest-edge alignment at  $70^\circ$  reduces the turn count from 8 to 4 compared to north–south scanning, saving approximately 60 s of flight time and 13.3 Wh of energy. This is because the AENGM0074 polygon is elongated along a northeast–southwest axis: scanning perpendicular to the short axis requires far fewer strips to cover the same area. The `cv2.minAreaRect` function computes this alignment automatically, and the parametric sweep confirms that its output ( $68^\circ$  for this polygon) falls within the optimal  $65^\circ$  to  $75^\circ$  band.

An alternative considered was wind-aligned scanning, which would orient strips along the prevailing westerly wind direction to minimise crosswind drift accumulation along the inter-lane axis. For the typical  $3 \text{ m s}^{-1}$  westerly wind at the survey site, this would correspond to approximately  $90^\circ$  (east–west strips). However, east–west scanning of this polygon



**Figure 44:** Scan angle comparison for the AENGM0074 survey polygon. Aligning strips with the longest edge (70°) minimises U-turn count and path length compared to axis-aligned or diagonal orientations.

produces the *worst-case* geometry: 10 scan lines with 9 U-turns, nearly doubling the path length compared to the optimal angle. The 20 % lateral overlap already absorbs the expected 1 m to 2 m crosswind displacement (Section V.5), making the energy penalty of wind-aligned scanning unjustifiable.

A diagonal scanning approach at 45° was also evaluated as a compromise between wind alignment and longest-edge alignment. While this orientation partially distributes crosswind drift across both axes, analysis indicated that it would require approximately 30 % additional scan lines compared to the 70° optimum, adding roughly 40 s of flight time. The marginal crosswind benefit was judged to fall well within the existing overlap margin, making the additional scan lines unnecessary.

We also investigated whether rotating the scan pattern to align with the wind direction would reduce GPS drift accumulation during U-turns, where the drone decelerates and is most susceptible to wind-induced displacement. Preliminary analysis suggested that the effect on turn-phase GPS error was within 0.3 m to 0.5 m—comfortably absorbed by the 6.4 m overlap band—and that the dominant source of turn-phase error was the deceleration–reacceleration cycle rather than wind alignment.

### Y.3 Coverage Completeness: Is 96% Acceptable?

The implemented system achieves approximately 96 % coverage of the survey polygon, with the uncovered 4 % concentrated along the western boundary adjacent to the SSSI no-fly zone. This gap arises from the 30 m NFZ waypoint buffer (NFZ\_WAYPOINT\_BUFFER\_M in `config.py`): any waypoint that would fall within 30 m of the SSSI boundary is excluded from the pattern, and the corresponding strip is truncated or omitted entirely. The question is whether this gap is operationally acceptable.

#### Y.3.1 Probabilistic Argument

The uncovered strip lies within 30 m of the SSSI boundary, which at the survey site is a fenced wildlife reserve. The probability of a casualty being located in this narrow band is proportional to its area fraction (4 %) multiplied by an access-probability factor. Since the area is fenced and not part of any public footpath network, the conditional probability of a casualty being present there is substantially lower than in the open survey area. A conservative estimate assigns the buffer zone one-tenth the casualty density of the open area, giving an effective coverage-weighted detection probability of:

$$P_{\text{eff}} = 0.96 + 0.04 \times 0.0 \times 0.1 = 0.96 \quad (55)$$

which, when weighted by the access probability, yields an effective detection probability negligibly different from 100 % coverage of the accessible area.

#### Y.3.2 Buffer Reduction Trade-Off

Reducing the NFZ buffer from 30 m to 10 m would recover most of the 4 % gap, achieving approximately 99 % coverage. However, this reduction narrows the safety margin in two ways:

1. **GPS drift:** with a 2 m to 3 m CEP, a 10 m buffer provides only 7 m of margin beyond the worst-case drift. A single GPS outlier could place the drone within 5 m of the SSSI.

2. **Overshoot at turns:** the drone's turn radius at  $8 \text{ m s}^{-1}$  is approximately 4 m (limited by the maximum bank angle of  $35^\circ$ ). At a 10 m buffer, a U-turn at the strip endpoint could momentarily place the drone within 6 m of the boundary.
3. **Camera footprint:** at 35 m altitude, the camera footprint extends 16 m from nadir in each direction. With a 10 m buffer, the footprint would extend 6 m into the SSSI, potentially capturing protected habitat imagery.


The 30 m buffer was therefore retained, accepting 96 % coverage as an operationally sound compromise that prioritises regulatory compliance and flight safety over marginal coverage gains in a low-probability zone.

#### Y.4 NFZ Margin vs Altitude

The camera footprint radius increases linearly with altitude, meaning that at higher search altitudes the drone must maintain a correspondingly larger distance from the SSSI boundary to avoid imaging protected land. The footprint half-width in the cross-track direction is:

$$r_{\text{fp}} = \frac{w_s \cdot h}{2f} = \frac{5.02 \times h}{2 \times 5.46} \approx 0.460 h \quad [\text{m}] \quad (56)$$

Table 104 shows the required minimum drone-to-boundary distance to ensure the camera footprint does not extend into the SSSI, and compares this with the implemented margins.



**Figure 45:** NFZ margin analysis as a function of altitude. The camera footprint radius grows linearly with altitude, but the implemented 50 m total margin maintains a comfortable surplus at all operational altitudes.

The total margin of 50 m (comprising the 30 m waypoint buffer plus the 20 m NFZ slow zone where speed ramps down to  $0.3 \text{ m s}^{-1}$ ) provides a comfortable surplus at all altitudes up to 50 m. Even at the highest altitude considered, the surplus of 27 m ensures that transient GPS errors, wind gusts, or turn overshoots cannot cause the camera footprint to image the SSSI. The multi-layer NFZ protection—waypoint exclusion, speed ramping, and repulsive geofence push—provides defence in depth; no single failure mode can result in SSSI intrusion.

#### Y.5 Speed vs Detection Probability


The probability of detecting a casualty depends on the number of inference frames in which the target appears. As the drone speed increases, the target traverses the camera frame faster, reducing the available detection window. At the Pi 5 TFLite inference rate of 4.8 FPS, the number of frames containing the target is:

$$n_f = \frac{W_g}{v} \times f_{\text{inf}} = \frac{W_g \cdot f_{\text{inf}}}{v} \quad (57)$$

where  $W_g$  is the along-track ground footprint (24 m at 35 m altitude, accounting for the 4:3 aspect ratio of the IMX296) and  $f_{\text{inf}} = 4.8 \text{ FPS}$ . Assuming independent per-frame detection with probability  $p = 0.95$  (based on bench-test results with the retrained YOLOv8n model at 36 model-input pixels), the probability of at least one detection in  $n_f$  frames is:

$$P_d = 1 - (1 - p)^{n_f} \quad (58)$$

Table 106 quantifies this trade-off across the speed range.



**Figure 46:** Speed–detection–energy trade-off at 35 m altitude. Detection probability remains above 99.6% across all speeds; the binding constraint is energy consumption and NFZ margin rather than detection reliability.

At all speeds up to  $15 \text{ m s}^{-1}$ , the detection probability exceeds 99.6%—a consequence of the high per-frame detection rate and the multiple-frame dwell time. The practical distinction between  $5 \text{ m s}^{-1}$  (99.99%) and  $15 \text{ m s}^{-1}$  (99.66%) is negligible in terms of detection reliability, but the search time differs by a factor of three. This analysis supports the altitude-dependent speed schedule (Section V.3): there is little detection benefit to flying slower than  $6 \text{ m s}^{-1}$  at any altitude, and the time savings from faster flight are substantial.

The exception is the focus area, where the search speed drops to  $5 \text{ m s}^{-1}$ . In this scenario the prior probability of a casualty being present is high (a PLB signal has been received), and the area is small enough that the additional time cost is measured in seconds rather than minutes. The reduced speed also compensates for potential model uncertainty in the focus zone, where the target may be partially obscured, in an unusual pose, or at the edge of the detection envelope.

## Y.6 Focus Area Redirect Impact

When a PLB signal triggers the focus area redirect (Section V.7), the mission dynamics change in several interacting ways:

1. **Area reduction.** The focus polygon is approximately 0.15 ha compared to the full survey area of 3.2 ha—a 95% reduction. At  $5 \text{ m s}^{-1}$  with the same lane width, the focus search requires approximately 30 s, compared to 112 s for the full area.
2. **Speed reduction.** The search speed drops from  $8 \text{ m s}^{-1}$  (altitude-dependent schedule at 35 m) to a fixed  $5 \text{ m s}^{-1}$ , increasing the per-target dwell time from 14 frames to 23 frames and raising  $P_d$  from 99.97% to 99.99%.
3. **Pattern regeneration.** The lawnmower pattern is recomputed from the drone’s current position over the focus polygon. Any scan lines already completed in the original broad-area search are not credited—the focus area is searched from scratch. If the broad search had already covered 60% of the original polygon, that partial coverage is abandoned. This is operationally correct: the PLB signal provides strong prior information that overrides the uniform-prior assumption of the broad search.
4. **State preservation.** Detection clusters, rejected false positives, and items of interest are preserved across the pattern switch. This prevents the system from re-investigating targets that the operator has already classified, avoiding wasted time in a scenario where speed is critical.

The net effect is a substantial acceleration of target acquisition within the focus area (30 s vs 112 s for the full search) at the cost of abandoning partial broad-area coverage. This trade-off is justified by Bayes' theorem: the PLB signal concentrates the posterior probability of the casualty's location so heavily in the focus area that continuing the broad search has negligible expected value.

## Y.7 Summary of Optimisation Trade-Offs

Table 108 summarises the Pareto trade-offs and the rationale for the selected operating point.

The key insight from this analysis is that detection probability is *not* the binding constraint: even at the fastest speed considered,  $P_d$  exceeds 99.6 %. The binding constraints are instead the NFZ safety margin (which imposes the 4 % coverage gap) and the detection confidence floor (which prevents flying above 35 m despite the energy savings). The selected operating point sits on the Pareto frontier at the intersection of these two binding constraints, delivering near-complete coverage with comfortable energy margin and near-certain detection within the endurance limit of a single battery.

## Z Focus Area Redirect and Repulsive Field Implementation

Two subsystems are central to the mission's adaptability and safety: the ability to dynamically redirect the search pattern in response to new intelligence, and the continuous enforcement of no-fly zone boundaries during autonomous flight. This section expands on these two mission-critical subsystems introduced in Sections V.7 and D.1: the PLB-triggered focus area redirect that narrows the search polygon mid-flight, and the inverse-distance repulsive vector field that enforces the SSSI no-fly zone boundary.

### Z.1 PLB Focus Area Logic (R06)

Requirement R06 stipulates that the system shall respond to a Personal Locator Beacon (PLB) signal by redirecting the search to a smaller focus area polygon of 3–10 vertices. In the operational scenario the PLB signal is expected 5–15 minutes into the search, once the casualty's approximate position has been triangulated by ground teams.

**Activation.** Two activation mechanisms are implemented. First, the operator can press the B key (or the corresponding browser button on the ground station) at any time during the SEARCH state to trigger the redirect manually. Second, the `--beacon-delay` command-line argument specifies an elapsed-time threshold in seconds; once the search clock exceeds this value, the redirect fires automatically. The auto-trigger is checked on every loop iteration during SEARCH and fires at most once—a boolean guard (`_beacon_triggered`) prevents re-entry.

**Focus area definition.** The focus area polygon is defined in one of three ways, in decreasing priority:

1. A JSON file (`flight_plans/focus_area.json`) containing 3–10 vertices as `{lat, lon}` objects. The file is re-read at activation time, allowing the ground station to update coordinates mid-flight.
2. Interactive drawing on the satellite map during simulation setup (a fourth polygon drawn after the search area, transit path, and target placement).
3. Fallback coordinates in `config.FOCUS_AREA_GPS`, loaded from the KML file at startup.

Each vertex is validated to lie within  $\pm 90^\circ$  latitude and  $\pm 180^\circ$  longitude, and at least three vertices must be present; otherwise the redirect is aborted with a console warning.

**Pattern regeneration.** Upon activation, the planner's search polygon is replaced with the focus area polygon and a fresh lawnmower pattern is generated from the drone's current GPS position (Algorithm 1). The new pattern uses the same rotated-mask algorithm described in Section 2.4, but the search speed is capped at `FOCUS_SEARCH_SPEED_MPS` =  $5 \text{ m s}^{-1}$ —slower than the normal altitude-dependent schedule—to maximise detection probability in the high-priority zone. The waypoint index and rescan counter both reset to zero.

**State preservation.** Existing detection clusters, rejected targets, and items of interest are preserved across the redirect. This ensures that false positives identified during the wide-area search are not re-investigated in the focus area. If the focus area overlaps the SSSI no-fly zone, the geofence waypoint filter (Section Z.3) removes offending waypoints exactly as it does for the primary search pattern.

### Z.2 Repulsive Vector Field

The repulsive potential field is the second of four NFZ protection layers (Section Z.3). Its purpose is to continuously deflect the aircraft away from the SSSI boundary, even when the planned waypoint lies close to it due to GPS drift or

---

**Algorithm 1** PLB beacon redirect

---

**Require:** PLB signal received (B key or timer expiry)

**Ensure:** Search redirected to focus area polygon

- 1: **if** already triggered **then return**
  - 2: **end if**
  - 3: Load focus polygon from JSON (or use fallback GPS)
  - 4: Validate  $\geq 3$  vertices, each within valid lat/lon range
  - 5: Convert GPS vertices to pixel coordinates via GEOTRANSFORMER
  - 6: Replace planner search polygon  $\leftarrow$  focus polygon
  - 7: Generate new lawnmower pattern from current position
  - 8: `wp_index`  $\leftarrow$  0; `rescan_pass`  $\leftarrow$  0
  - 9: Set `_beacon_triggered`  $\leftarrow$  **true**
  - 10: Cap search speed to  $5 \text{ m s}^{-1}$
- 

wind gusts. The field is evaluated on every loop iteration when the drone is within `NFZ_INNER_RANGE_M` = 23 m of an inner offset polygon.

**Nearest-point computation.** For each edge  $(\mathbf{p}_i, \mathbf{p}_{i+1})$  of the NFZ polygon (expressed in pixel coordinates via GEOTRANSFORMER), the closest point to the drone's position  $\mathbf{q}$  is found by projecting  $\mathbf{q}$  onto the edge and clamping the parameter  $t$  to  $[0, 1]$ :

$$\mathbf{c}_i = \mathbf{p}_i + \text{clamp}\left(\frac{(\mathbf{q} - \mathbf{p}_i) \cdot (\mathbf{p}_{i+1} - \mathbf{p}_i)}{\|\mathbf{p}_{i+1} - \mathbf{p}_i\|^2}, 0, 1\right) (\mathbf{p}_{i+1} - \mathbf{p}_i) \quad (59)$$

The edge yielding the smallest  $\|\mathbf{q} - \mathbf{c}_i\|$  is selected, and the push direction is  $\hat{\mathbf{d}} = (\mathbf{q} - \mathbf{c}^*)/\|\mathbf{q} - \mathbf{c}^*\|$ , pointing away from the boundary.

**Degenerate case.** When the drone lies exactly on or extremely close to a boundary edge ( $\|\mathbf{q} - \mathbf{c}^*\| < \epsilon$ ), the direction vector is degenerate. In this case, the outward normal of the nearest edge is computed instead: for edge vector  $\mathbf{e} = \mathbf{p}_{i+1} - \mathbf{p}_i$ , the candidate normal is  $\mathbf{n} = (-e_y, e_x)$ . A sign test (`cv2.pointPolygonTest` on a point offset slightly along  $\mathbf{n}$ ) determines whether  $\mathbf{n}$  points inward or outward; if inward, it is flipped. This guarantees a valid push direction even at zero distance.

**Repulsion magnitude.** The repulsive offset grows linearly as the drone approaches the boundary:

$$R = \frac{1}{2} \max(0, d_{\text{soft}} - d) \quad (60)$$

where  $d$  is the perpendicular distance to the NFZ boundary (metres) and  $d_{\text{soft}} = \text{NFZ\_SOFT\_BOUNDARY\_M}$  = 8 m is the zone within which repulsion is active. The factor  $\frac{1}{2}$  is a gain constant that limits the maximum offset to  $\frac{d_{\text{soft}}}{2} = 4$  m at the boundary itself. The pixel-space push vector is then:

$$\mathbf{v}_{\text{push}} = R \cdot s_{\text{pix}} \cdot \hat{\mathbf{d}} \quad (61)$$

where  $s_{\text{pix}}$  is the pixel-per-metre scale factor from GEOTRANSFORMER. This pixel offset is converted back to a GPS delta ( $\Delta\phi, \Delta\lambda$ ) and negated to correct for the coordinate transform sign convention (see Section Z.3).

**Velocity application.** In the main loop (`_enforce_geofence`), the GPS offset from `repulsive_offset()` is converted to a unit North-East velocity vector and scaled to a constant push speed of `NFZ_PUSH_SPEED_MPS` =  $3.0 \text{ m s}^{-1}$ :

$$v_N = \frac{-\Delta\phi \cdot 111320}{\|\mathbf{v}\|} \cdot v_{\text{push}}, \quad v_E = \frac{-\Delta\lambda \cdot 111320 \cos \phi}{\|\mathbf{v}\|} \cdot v_{\text{push}} \quad (62)$$

where the factor 111 320 converts degrees to metres. If the drone is inside the NFZ polygon (detected by sign inversion of `cv2.pointPolygonTest`), the velocity direction is reversed to push the aircraft outward.

### Z.3 Four-Layer Protection Architecture

Table 110 summarises the four software layers that enforce the SSSI exclusion zone, listed from outermost (plan-time) to innermost (emergency).

**Layer interaction.** Layers 1 and 2 overlap spatially: the speed ramp activates at 20 m while the repulsive field activates at 23 m (measured from an inner polygon offset 20 m inside the true NFZ boundary). In practice, the repulsive field begins deflecting the trajectory before the speed ramp starts decelerating. The combined effect is a smooth, progressive deflection: the drone first feels a gentle push, then decelerates as it approaches the boundary, and only when the boundary offset is reached does the full repulsive field activate. This design ensures that the drone is pushed away from the boundary before it reaches the speed limit, allowing for a smooth transition to the emergency layer.

- **23 m repulsive field** — extends 3 m beyond the speed ramp onset to catch trajectories where GPS drift has already placed the drone closer than planned. The 3 m margin matches the typical GPS CEP (Circular Error Probable) observed during bench testing (2 m to 3 m).
- **3 m hard boundary** — a last-resort cutoff that accounts for worst-case GPS error. At this distance the drone is within one CEP of the true boundary, so autonomous flight is no longer safe. The transition to MANUAL halts all velocity commands and alerts the operator, who can then fly the aircraft away using RC override.

**Sign convention.** The `GeoTransformer` converts pixel offsets to GPS deltas relative to a reference origin. Because the pixel  $y$ -axis is inverted with respect to latitude (pixel  $y$  increases southward while latitude increases northward), the raw GPS delta from `pixels_to_gps()` points *toward* the NFZ rather than away from it. The `repulsive_offset()` method therefore negates both components before returning them. The caller in `_enforce_geofence()` applies a second negation when converting to North–East velocity components, yielding the correct outward push direction. This double-negation was verified empirically by approaching the NFZ from all cardinal directions in SITL and confirming that the repulsive vector always points away from the boundary.

**Directional scalar field.** An alternative scalar field implementation was also developed (enabled via the `--nfz-directional` flag) that limits only the velocity component directed toward the NFZ boundary, rather than capping total speed. This allows the aircraft to maintain full speed when flying parallel to the boundary while still preventing rapid approach. The directional decomposition requires the drone’s velocity vector (available from MAVLink `GLOBAL_POSITION_INT` at 10 Hz) and a dot product with the boundary-normal vector — computationally negligible ( $<1\ \mu\text{s}$ ) compared to the 206 ms inference cycle. The two fields complement rather than conflict: the scalar field limits approach velocity while the repulsive field provides active escape, yielding a net outward velocity at all boundary distances.

**Firmware backup.** On the flight controller (Cube Orange), an independent firmware-level geofence is configured via the `FENCE_TYPE` and `FENCE_ACTION` parameters. If all four software layers fail—for example, due to a companion-computer crash—the firmware triggers RTL when the aircraft crosses the configured fence polygon. This fifth layer operates entirely within the autopilot and requires no companion-computer involvement.

## AA Search Path Optimisation

The search pattern parameters—altitude  $h$ , scan angle  $\theta$ , speed  $v$ , overlap margin  $\alpha$ , and NFZ buffer—do not have independently optimal values. Each parameter trades off against at least two competing objectives, forming a coupled multi-dimensional optimisation landscape. This section formalises the problem, presents the energy model and parametric sweep methodology, compares six candidate search strategies against real polygon geometry, and justifies the selected operating point.

### AA.1 The Multi-Objective Optimisation Problem

Five objectives are in tension:

1. **Coverage completeness** ( $C$ ): the fraction of the survey polygon imaged by at least one camera frame. Complete coverage ( $C = 1.0$ ) requires every ground point to fall within at least one swath.
2. **Mission time** ( $T$ ): total elapsed time from takeoff to landing. In SAR operations, every minute of delay reduces casualty survival probability; UK Maritime and Coastguard Agency survival curves show a steep decline in hypothermia survivability after 60 min of exposure [61].
3. **Energy consumption** ( $E$ ): total energy drawn from the battery. The EDU-450 airframe carries a 6S 5200 mAh LiPo (115.4 Wh), imposing a hard flight-time ceiling of approximately 18 min at search power levels.
4. **Detection probability** ( $P_d$ ): the probability that a casualty present in the search area is detected by at least one inference frame. This depends on the number of frames in which the target appears, the per-frame confidence, and the confirmation threshold.
5. **NFZ safety margin** ( $d_{\text{NFZ}}$ ): the minimum distance between any waypoint and the SSSI no-fly zone boundary. Larger margins reduce regulatory risk but sacrifice coverage of the polygon edge adjacent to the SSSI.

These objectives are coupled through three decision variables:

- **Altitude**  $h \in [20, 50]$  m — determines the camera ground footprint  $W_g = w_s h / f$ , the ground sample distance (GSD), the target pixel extent, and therefore both coverage rate and detection confidence.
- **Scan angle**  $\theta \in [0^\circ, 180^\circ]$  — the orientation of parallel scan lines relative to north. Controls the number of scan lines, U-turns, and path length for a given polygon.
- **Speed**  $v(h)$  — an altitude-dependent schedule interpolating linearly from  $6\text{ m s}^{-1}$  at 20 m to  $10\text{ m s}^{-1}$  at 50 m, balancing coverage rate against frame dwell time.

Subject to four constraints:

- R1** All waypoints must lie within the KML-defined flight area polygon.
- R2** No waypoint shall fall within 30 m of the SSSI boundary (NFZ.WAYPOINT\_BUFFER\_M).
- R3** Maximum altitude  $h \leq 50$  m (UK drone regulation and airspace ceiling).
- R4** Total energy  $E \leq 115.4$  W h (single battery, no mid-mission swap).

No single configuration simultaneously optimises all five objectives. The operating point must be selected from the Pareto frontier by engineering judgement, informed by the quantitative analysis that follows.

## AA.2 Energy Model

A simplified quadcopter power model was developed to evaluate energy consumption across the parameter space. Total power at forward speed  $v$  is decomposed into a constant hover component and a speed-dependent parasitic drag component:

$$P_{\text{total}}(v) = P_{\text{hover}} + P_{\text{drag}}(v) = P_{\text{hover}} + P_{\text{drag,ref}} \left( \frac{v}{v_{\text{ref}}} \right)^2 \quad (63)$$

where  $P_{\text{hover}} = 150$  W is the constant power to remain airborne (estimated from momentum theory for the EDU-450 at 2.3 kg all-up weight),  $P_{\text{drag,ref}} = 50$  W is the parasitic drag power at the reference speed  $v_{\text{ref}} = 5 \text{ m s}^{-1}$ , and the quadratic scaling follows from the  $\frac{1}{2}C_D\rho Av^2$  aerodynamic drag relation.

Each U-turn incurs a penalty of  $\Delta t_{\text{turn}} = 2$  s of hover-only power for deceleration and reacceleration. The total mission energy is therefore:

$$E_{\text{total}} = \underbrace{P_{\text{hover}} \cdot t_{\text{total}}}_{\text{hover (incl. turns)}} + \underbrace{P_{\text{drag}}(v) \cdot t_{\text{flight}}}_{\text{drag (forward motion only)}} \quad (64)$$

where  $t_{\text{total}} = t_{\text{flight}} + n_{\text{turns}} \cdot \Delta t_{\text{turn}}$  is the total airborne time including U-turn penalties,  $t_{\text{flight}}$  is the time spent in forward motion along scan lines and transit segments, and  $n_{\text{turns}}$  is the number of U-turns. Each U-turn adds  $\Delta t_{\text{turn}} = 2$  s of hover-only time (no drag component, since the drone decelerates to a near-hover during the reversal). At the selected operating point (35 m,  $8 \text{ m s}^{-1}$ ), the drag component contributes  $50 \times (8/5)^2 = 128$  W, giving a total in-flight power of 278 W.

Near the SSSI boundary, the speed scalar field reduces the drone's velocity from the search speed down to  $v_{\text{min}} = 0.3 \text{ m s}^{-1}$  at the boundary, following a linear ramp over a 20 m slow zone. Within this zone, drag power decreases (lower  $v$ ) but hover time increases (same distance at lower speed), and since  $P_{\text{hover}} \gg P_{\text{drag}}$  at low speeds, the net effect is an energy *increase* for scan lines that pass through the slow zone. The optimisation script accounts for this by numerically integrating the speed-dependent traverse time along each scan segment with 50 sample points.

## AA.3 The 216-Configuration Parametric Sweep

To systematically explore the design space, the optimisation script (`analysis/path_optimization/optimize_path.py`) evaluates **216 configurations**: 6 altitudes  $\times$  36 scan angles. For each combination, the script:


1. Computes the camera ground footprint and swath spacing (with 20 % overlap);
2. Generates the lawnmower scan lines over the real AENGM0074 five-vertex survey polygon using the rotate–rasterise–rotate-back algorithm;
3. Calculates total path length (scan lines plus U-turn transits);
4. Evaluates traverse time at the altitude-dependent speed, with per-sample NFZ slowdown integration;
5. Applies the energy model of Eq. (64) to compute total energy with and without NFZ effects.

Figure 47 shows the resulting energy landscape. Several features are immediately apparent:

- **Angle sensitivity dominates altitude sensitivity.** At any given altitude, varying the scan angle from optimal ( $65^\circ$  to  $75^\circ$ ) to worst-case ( $0^\circ$  or  $90^\circ$ ) increases energy by 100 % to 200 %. By contrast, changing altitude from 35 m to 50 m at the optimal angle saves only 30 % to 40 %.
- **Low-energy valley.** A band of low-energy configurations runs through  $\theta \approx 70^\circ$  across all altitudes. This corresponds to alignment with the polygon's longest edge (northeast–southwest axis), which minimises the number of perpendicular scan lines and therefore U-turns.
- **Low altitude penalised.** At 20 m, the narrow footprint (18.4 m) requires approximately 10 scan lines with 9 U-turns, roughly doubling energy compared to 50 m (4–5 lines, 3–4 turns).

Table 112 presents the five best and five worst configurations from the sweep.

The energy-optimal configuration (50 m,  $70^\circ$ ) consumes only 8.2 W h—7.1 % of battery capacity—and completes the search in 68 s. The worst configuration (20 m,  $95^\circ$ ) consumes 43.2 W h, a  $5.3\times$  increase, and takes over 6.9 min. The sweep confirms that scan angle selection is at least as important as altitude selection for energy efficiency.



**Figure 47:** Energy heatmap across 216 altitude–angle configurations. The colour scale represents total energy consumption including NFZ slowdown penalties. A clear valley of low-energy configurations runs through the 65° to 75° scan angle band at all altitudes, corresponding to alignment with the polygon’s longest edge. The red star marks the global energy minimum at 50 m altitude and 70° scan angle.

Figure 48 shows the altitude–energy relationship at the optimal scan angle, decomposed into clean (no NFZ) and penalised (with NFZ) components. The NFZ slowdown adds approximately 15 % to total energy at 35 m and up to 45 % at 20 m, where the narrower swath forces scan lines closer to the SSSI boundary.


#### AA.4 Search Strategy Comparison

While the parametric sweep optimises altitude and angle within the boustrophedon pattern, the choice of search *strategy* itself requires justification. Six candidate strategies were evaluated on the AENGM0074 survey polygon using the same energy model and polygon geometry.

1. **Boustrophedon (lawnmower)** — parallel back-and-forth strips aligned to the polygon’s longest edge. Guarantees complete coverage by construction [8]. The scan angle is auto-selected via `cv2.minAreaRect`.
2. **Spiral inward** — concentric rectangular rings converging on the polygon centroid. Covers the perimeter first; each ring includes four 90° corners instead of 180° U-turns, but the total path length is 8 % to 12 % longer due to poor nesting on irregular polygons.
3. **Expanding square** — outward expansion from a centre point in increasing-length legs. Requires a known start point (PLB signal); without prior information, degrades to uniform coverage with 18 % to 23 % higher energy due to legs extending outside the polygon boundary.
4. **Sector search** — the polygon is divided into probability-weighted sectors, each searched with a sub-pattern. Inter-sector transit adds 100 m to 200 m of overhead, and the method requires a probability model that was not available for this mission.
5. **Creeping line ahead** — all scan lines traverse in the same direction, with dead-heading back to the start side between lines. Path length is approximately double the boustrophedon (75 % energy penalty) for the same coverage guarantee. Advantageous only for side-looking sensors (sonar, SAR radar), which are not applicable here.
6. **Adaptive / heatmap** — a coarse initial sweep followed by targeted infill based on detection results. Best-case energy (6.2 W h) is attractive, but worst-case (16 W h) exceeds the boustrophedon, and the method provides no coverage guarantee. The engineering complexity of real-time replanning was judged disproportionate for a 3.2 ha polygon coverable in a single pass.

The boustrophedon pattern was selected for five reasons:

1. **No prior information.** The mission brief specifies a casualty somewhere in the survey polygon with no PLB signal, last known position, or probability map. Without prior information, expanding square and sector search degrade to uniform coverage with added transit overhead.
2. **Coverage guarantee.** The assessment requires demonstrating that the drone can find the target. Only the boustrophedon and creeping-line-ahead provide mathematically provable complete coverage; the creeping line costs 75 % more energy for the same guarantee.



**Figure 48:** Total search energy versus altitude at the optimal scan angle, with and without NFZ slowdown. The NFZ penalty is most pronounced at 20 m, where several scan lines pass through the slow zone near the SSSI boundary. The shaded region quantifies the additional hover energy incurred by the speed ramp.

where  $d$  is the distance to the nearest SSSI edge and  $d_{\text{zone}} = 20$  m.

**3. Hard boundary.** Inside 3 m of the NFZ, the system automatically switches to MANUAL mode as a fail-safe. The energy impact of the speed ramp is asymmetric across scan angles. At the optimal angle ( $70^\circ$ ), scan lines run roughly parallel to the polygon’s long axis and only clip the NFZ slow zone at their northeastern endpoints. At  $0^\circ$  (north–south scanning), several scan lines run along the full length of the SSSI boundary, with each metre of that segment traversed at reduced speed. The optimisation script accounts for this by numerically integrating the effective speed along each scan segment with 50 uniformly spaced sample points, computing the distance to the nearest NFZ edge at each sample. This per-segment NFZ time penalty is summed to produce the total time (and therefore energy) “with NFZ” reported in the sweep results.

The practical consequence is that NFZ-aware scan angle selection becomes critical: configurations that align scan lines parallel to the SSSI boundary incur disproportionate time penalties. The optimal angle ( $70^\circ$ ) happens to align with the polygon’s longest edge *and* to run perpendicular to the SSSI boundary—a fortunate geometric coincidence that minimises both U-turn count and NFZ dwell time simultaneously.

## AA.6 Detection Probability Integration

The search altitude affects target detection through a chain of coupled physical relationships. At altitude  $h$ , the ground sample distance is:


$$\text{GSD} = \frac{w_s \cdot h}{f \cdot W_{\text{px}}} = \frac{5.02 \times h}{5.46 \times 1456} \quad (66)$$

where  $W_{\text{px}} = 1456$  is the image width in pixels. The mannequin (1.8 m tall, 0.5 m wide) projects to:


$$\text{px}_{\text{tall}} = \frac{1.8}{\text{GSD}}, \quad \text{px}_{\text{wide}} = \frac{0.5}{\text{GSD}} \quad (67)$$

Since YOLOv8n resizes the  $1456 \times 1088$  sensor frame to  $640 \times 640$ , the model-input target size is further reduced by the uniform scale factor  $640/1456 = 0.4396$ .

Table 116 traces this function chain from altitude through to detection probability. The altitude-dependent speed  $v(h)$  determines the dwell time (how long the target remains in the field of view), and therefore the number of inference frames:



**Figure 49:** Search patterns overlaid on the AENGM0074 survey polygon at 35 m and 50 m altitude. The boustrophedon (leftmost column) achieves the shortest path and highest coverage for both altitudes. Start and end waypoints are marked with green circles and red squares respectively.



**Figure 50:** Top three energy-ranked search configurations overlaid on the AENGM0074 survey polygon. Each panel shows the lawnmower pattern at the corresponding altitude and scan angle, with path length, scan line count, and energy annotated. The selected operating point (35 m, 70°) balances detection reliability against the energy savings of higher altitudes.

$$n_f = \frac{H_g}{v(h)} \times f_{\text{inf}} \quad (68)$$


where  $H_g$  is the along-track footprint height and  $f_{\text{inf}} = 4.8\text{FPS}$  is the TFLite inference rate on the Raspberry Pi 5. Assuming independent per-frame detection with probability  $p$  (a function of pixel extent), the cumulative detection probability over  $n_f$  frames is:

$$P_d = 1 - (1 - p)^{n_f} \quad (69)$$

Two results are noteworthy. First, **motion blur is not a limiting factor**: the IMX296 global shutter combined with a  $1/1000\text{s}$  exposure produces sub-pixel blur ( $0.3\text{px}$  at  $10\text{ m s}^{-1}$ , 50 m altitude) at all operating speeds. The critical speed for single-pixel blur exceeds the maximum search speed at every altitude in the operating range ( $12.6\text{ m s}^{-1}$  at 20 m vs  $6\text{ m s}^{-1}$  search speed;  $31.6\text{ m s}^{-1}$  at 50 m vs  $10\text{ m s}^{-1}$  search speed), so motion blur remains sub-pixel throughout.

Second, **the dwell time paradoxically increases with altitude** (from 2.3 s at 20 m to 3.4 s at 50 m) because the footprint height grows faster than the speed. The target is captured in 11–16 frames at all altitudes—more than sufficient given that a single detection above the 0.2 confidence threshold triggers investigation. Detection probability exceeds 99.6 % even at 50 m, where the target width drops to 7 pixels in the model input. The binding constraint on altitude is therefore not detection probability per se, but detection *confidence*: at 50 m the model produces lower-confidence detections that are more likely to be marginal, whereas at 35 m the target subtends 36 pixels in the model input with comfortable margin above the 20 px minimum detectable size.

## AA.7 Selected Configuration and Justification



Each dot = one (altitude, speed, overlap) configuration. This 2D view collapses 5 objectives into 2 composite dimensions.

**Figure 52:** Two-dimensional Pareto frontier projections across the five objectives. Each subplot projects the 216-configuration sweep onto a pair of objectives, with Pareto-efficient solutions highlighted. The selected operating point is marked with a star, illustrating its position on or near the Pareto frontier in every projection.

The 96 % coverage figure arises from the 30 m NFZ waypoint buffer, which excludes waypoints near the SSSI boundary. The uncovered 4 % lies within a fenced wildlife reserve with no public access, making the conditional probability of a casualty in this strip negligibly small. A probabilistic coverage-weighted detection probability of:

$$P_{\text{eff}} = C \cdot P_d + (1 - C) \cdot P_{\text{access}} \cdot P_d \approx 0.96 \times 0.9997 \approx 0.96 \quad (70)$$

is operationally equivalent to full coverage of the accessible area.

The key finding from this analysis is that detection probability is *not* the binding constraint: even at the fastest speed considered ( $15 \text{ m s}^{-1}$ ),  $P_d$  exceeds 99.6 %. The binding constraints are instead the NFZ safety margin (which imposes the 4 % coverage gap) and the detection confidence floor (which prevents flying above 35 m despite the energy savings). The selected operating point sits at the intersection of these two binding constraints on the Pareto frontier, delivering near-complete coverage with comfortable energy margin and near-certain detection within the endurance limit of a single battery.


## AB Search Parameter Optimisation

In search and rescue, time is the binding resource. UK Maritime and Coastguard Agency survival curves show a steep decline in hypothermia survivability after 60 min of exposure [61]; every second of search delay reduces the probability of a live recovery. The objective of parameter optimisation is therefore to *find the casualty as fast as possible*, subject to three hard constraints: the onboard AI must be able to detect the target, the battery must last the mission, and the drone must never enter the SSSI no-fly zone.

This section formalises the optimisation problem, describes the physics-based simulation used to evaluate 216 candidate configurations, presents the top-scoring paths and the reasoning that selected the final operating point, and quantifies how sensitive the result is to each input parameter.

### AB.1 Problem Statement

Five coupled decision variables govern the search pattern:



**Figure 53:** Altitude trade-off: detection probability and target pixel extent (left axis) versus energy consumption (right axis) as a function of search altitude. The selected altitude (35 m) sits at the knee of the detection curve, where further altitude increases yield diminishing detection gains but growing energy savings.

$$\mathbf{x} = [h, v, \theta, \alpha, d_{\text{nfz}}]^\top \quad (71)$$

where  $h$  is the search altitude,  $v$  is the forward speed (coupled to  $h$  through an altitude-dependent schedule),  $\theta$  is the scan angle relative to north,  $\alpha = 0.20$  is the lateral overlap fraction, and  $d_{\text{nfz}} = 30 \text{ m}$  is the NFZ waypoint buffer. These variables jointly determine five competing objectives—detection probability  $P_d$ , coverage completeness  $C$ , mission time  $T$ , energy consumption  $E$ , and NFZ safety margin  $S$ —subject to four hard constraints:

- C1** All waypoints within the KML-defined flight area polygon.
- C2** No camera footprint overlaps the SSSI ( $\mathcal{F}(h, \mathbf{w}_i) \cap \mathcal{P}_{\text{SSSI}} = \emptyset$ ).
- C3** Maximum altitude  $h \leq 50 \text{ m}$  (CAA Open A3).
- C4** Energy  $\leq 0.8 \times 115.4 \text{ Wh}$  (20% RTL reserve on 6S 5200 mAh LiPo).

No single configuration simultaneously optimises all five objectives; the operating point must be selected from the Pareto frontier by engineering judgement.

## AB.2 Physics-Based Simulation

Each candidate configuration was evaluated by a physics-based simulation (`optimize_path.py`) that models the full mission energy budget.

### AB.2.1 Energy Model

Total power at forward speed  $v$  decomposes into a constant hover component and a quadratic drag component:

$$P(v) = P_{\text{hover}} + P_{\text{drag,ref}} \left( \frac{v}{v_{\text{ref}}} \right)^2 \quad (72)$$

where  $P_{\text{hover}} = 150 \text{ W}$  (momentum-theory estimate for the EDU-450 at  $2.3 \text{ kg}$  AUW:  $P = T v_i = mg \sqrt{mg/2\rho A}$ ),  $P_{\text{drag,ref}} = 50 \text{ W}$  at  $v_{\text{ref}} = 5 \text{ m s}^{-1}$ , and the quadratic scaling follows from  $\frac{1}{2} C_D \rho A v^2$ . At the selected speed of  $8 \text{ m s}^{-1}$ , the drag contribution is  $50 \times (8/5)^2 = 128 \text{ W}$ , giving a total in-flight power of  $278 \text{ W}$ .

Each U-turn incurs a  $\Delta t_{\text{turn}} = 2 \text{ s}$  penalty of hover-only power for deceleration and reacceleration. Near the SSSI boundary, a linear speed ramp reduces the drone from  $v_{\text{zone}} = 3.0 \text{ m s}^{-1}$  to  $v_{\text{min}} = 0.3 \text{ m s}^{-1}$  over a  $20 \text{ m}$  slow zone. Within this zone, lower speed reduces drag but increases hover time for the same distance; since  $P_{\text{hover}} \gg P_{\text{drag}}$  at low speeds, the net effect is an energy *increase* for scan lines passing through the slow zone.

The total mission energy integrates power over the profile:

$$E = P_{\text{hover}} \cdot t_{\text{total}} + P_{\text{drag,ref}} \sum_{i=1}^{n_{\text{lines}}} \int_0^{\ell_i/v_i} \left( \frac{v_i(t)}{v_{\text{ref}}} \right)^2 dt \quad (73)$$

where  $t_{\text{total}}$  includes forward-flight time, U-turn penalties, and transit. The NFZ speed variation along each scan segment is numerically integrated with 50 sample points.

### AB.2.2 Detection Model

The mannequin (1.8 m tall, 0.5 m wide) projects onto the  $640 \times 640$  YOLOv8n model input with an extent of:

$$p_{\text{model}} = \frac{d_{\text{target}} \cdot f_{\text{px}}}{h} \times \frac{640}{W_{\text{px}}} \quad (74)$$

where  $d_{\text{target}}$  is the relevant target dimension (width 0.5 m for the binding constraint),  $f_{\text{px}} = f/w_s \times W_{\text{px}} = 5.46/5.02 \times 1456 = 1584 \text{ px}$  is the focal length in pixel units, and  $W_{\text{px}} = 1456$  is the sensor width. At the critical width dimension,  $p_{\text{model}}(35 \text{ m}) = 10 \text{ px}$  and  $p_{\text{model}}(50 \text{ m}) = 7 \text{ px}$ . Per-frame detection probability is modelled as a sigmoid of target pixel extent:

$$p(h) = \frac{1}{1 + \exp(-\kappa(p_{\text{model}}(h) - p_{\text{min}}))} \quad (75)$$

with  $p_{\text{min}} = 7 \text{ px}$  and  $\kappa = 0.5 \text{ px}^{-1}$ , calibrated from video analysis (Section H.4). The number of inference frames per flyover is  $n_f = H_g \cdot f_{\text{inf}}/v$ , where  $f_{\text{inf}} = 4.8 \text{ FPS}$  (benchmarked on Pi 5). The cumulative detection probability is:

$$P_d = 1 - (1 - p)^{n_f} \quad (76)$$

### AB.2.3 Parametric Sweep

The simulation swept **216 configurations**: 6 altitudes  $\times$  36 scan angles:

$$H = \{20, 25, 30, 35, 40, 50\} \text{ m}, \quad \Theta = \{0, 5, \dots, 175\}^\circ \quad (77)$$

Speed follows an altitude-dependent schedule  $v(h) = 6 + 4(h - 20)/30 \text{ m/s}$  (linear from  $6 \text{ m/s}^{-1}$  at 20 m to  $10 \text{ m/s}^{-1}$  at 50 m), linking speed to altitude and reducing the effective dimensionality. For each configuration, the script generates the full lawnmower pattern on the real AENGMM0074 five-vertex polygon, evaluates per-segment NFZ penalties, and computes all five objective values.


Each configuration is scored by a weighted composite:

$$S(\mathbf{x}) = 0.40 \cdot \hat{P}_d + 0.35 \cdot \hat{C} + 0.25 \cdot (1 - \hat{E}) \quad (78)$$

where  $\hat{\cdot}$  denotes min-max normalisation to  $[0, 1]$  across the feasible set. The weights reflect the SAR priority hierarchy: detection reliability is paramount (a missed casualty is mission failure), coverage completeness is second (an unsearched area cannot yield a find), and energy efficiency is third (the search consumes at most 11% of battery regardless of configuration).

### AB.3 Top Three Configurations

Figure 54 shows the three highest-scoring configurations overlaid on the survey polygon. Table 120 compares their performance.



**Figure 54:** Top three lawn timer configurations overlaid on the AENGMM0074 polygon. Left: #1 (35 m, 70°, selected). Centre: #2 (30 m, 65°). Right: #3 (40 m, 75°). The SSSI no-fly zone (red) and 20 m buffer (dashed) are shown.


Configuration #1 wins because it sits at the sweet spot of two competing pressures. Compared to #2, it saves 17% energy and 15 seconds of flight time while sacrificing only 0.02 percentage points of detection probability—a difference that is unobservable in practice ( $P_d > 0.999$  in both cases). Compared to #3, it provides  $3\times$  more detection margin: at 40 m the target width drops to 9 model pixels, uncomfortably close to the 7-pixel empirical minimum, whereas at 35 m it is 10 pixels with clear headroom. The 2.5 Wh energy penalty (consuming 10.9% vs 8.8% of battery) is negligible given 14 min of remaining endurance.

#### AB.4 Altitude and Speed Selection Logic

The altitude and speed were not selected by searching a table; they follow a logical chain from physical constraints:

1. **Maximum detection altitude.** At what altitude does the target become undetectable? The mannequin's width (0.5 m) projects to  $p_{\text{model}} = 0.5 \times 1584/h \times 640/1456$  model pixels. At 50 m this is 7 px—the empirically observed minimum for reliable bounding-box regression. At 70 m (18 px height, 5 px width) the target is below the feature-extraction threshold. The hard detection ceiling is approximately 50 m.
2. **Apply safety margin.** Real-world degradation modes—partial occlusion, non-ideal lighting, unusual target pose—can reduce effective pixel extent. A 30% safety margin yields  $h = 50/1.43 \approx 35$  m, where the target subtends 36 model pixels in height and 10 in width, comfortably above both thresholds.
3. **Maximum speed for adequate dwell time.** At 35 m, the along-track footprint is  $H_g = 5.02 \times 35/5.46 \times 1088/1456 = 24.1$  m. For  $\geq 10$  raw inference frames on target (ensuring 3–4 substantially different views):  $v_{\text{max}} = 24.1 \times 4.8/10 = 11.5 \text{ m s}^{-1}$ .
4. **Apply conservative margin.** The altitude-dependent schedule assigns  $v(35 \text{ m}) = 8.0 \text{ m s}^{-1}$ , yielding  $n_f = 24.1/8.0 \times 4.8 = 14$  frames—a  $\times 1.4$  margin above the 10-frame requirement. The cumulative detection probability over 14 frames is  $P_d = 1 - (1 - 0.95)^{14} > 0.9999$ .


This chain is visualised in Figure 55, which maps the composite score across the altitude–speed space. The selected operating point sits above the 95% detection contour in the high-score region.



**Figure 55:** Altitude–speed trade-off space. Colour encodes the composite score (40% detection + 35% coverage + 25% energy). Blue contours: detection probability. Orange contours: coverage time. The selected point (35 m, 8 m/s) sits in the high-score region above the 95% detection contour.

The scan angle  $\theta = 70^\circ$  was then determined by the 216-configuration sweep at fixed  $h = 35$  m: this angle aligns with the polygon's longest edge (northeast–southwest axis), minimising scan lines from 10 (at  $0^\circ$ ) to 6, and reducing U-turns


from 9 to 5. The energy saving is  $2\times$  compared to the worst angle ( $95^\circ$ ), as shown in Figure 56.



**Figure 56:** Energy heatmap across 216 altitude–angle configurations. A clear low-energy valley runs through  $65\text{--}75^\circ$  at all altitudes, corresponding to longest-edge alignment. The red star marks the global energy minimum ( $50\text{ m}, 70^\circ$ ). The selected point ( $35\text{ m}, 70^\circ$ ) trades  $4.4\text{ Wh}$  of energy for a  $44\%$  increase in target pixel extent.

### AB.5 Sensitivity Analysis

How robust is the selected operating point to perturbations in each variable? Table 122 summarises the sensitivity of the composite score to one-step changes, and Figure 57 presents the tornado chart.



**Figure 57:** Tornado chart of parameter sensitivities. Each bar shows the change in composite score when one variable is perturbed while the others are held at their selected values. Altitude dominates: a  $\pm 5\text{ m}$  change has nearly double the effect of a  $\pm 2\text{ m/s}$  speed change.

Three observations follow:


- Altitude is the most sensitive variable** because it simultaneously affects four objectives: higher altitude increases footprint (improving coverage rate and reducing scan lines) but degrades pixel extent (reducing detection confidence), while altering both mission time and energy. A  $5\text{ m}$  altitude error has nearly double the score impact of a  $2\text{ m s}^{-1}$  speed error.
- The scan angle has a flat optimum.** The energy valley at  $65\text{--}75^\circ$  (Figure 56) means that  $\pm 15^\circ$  perturbation changes the score by only  $0.02$ . This is operationally useful: the pilot does not need to align the drone precisely

with a compass heading.

3. **Speed is moderately sensitive** because it trades dwell time against energy. The altitude-dependent schedule provides a built-in coupling: if altitude drifts up (larger footprint), the schedule automatically increases speed, partially compensating the detection loss with faster coverage.

## AB.6 Variable Coupling Structure

The five decision variables do not affect the objectives independently. Figure 58a shows the  $N^2$  diagram revealing the coupling structure, and Figure 58b quantifies which objectives are jointly influenced by each variable pair.



(a)  $N^2$  diagram showing information flow between subsystems. Feed-forward links (blue) and feedback loops (red) reveal that altitude–speed is the most tightly coupled pair.

Variable Coupling Matrix — Pairwise Interaction Effects

|                      | Altitude $h$                              | Speed $v$                                 | Scan angle $\theta$     | Overlap $\alpha$       | NFZ margin $d_{nfz}$          |
|----------------------|-------------------------------------------|-------------------------------------------|-------------------------|------------------------|-------------------------------|
| Altitude $h$         | $h$<br>Detection ceiling                  | Strong<br>Coverage rate + detection prob. | Weak<br>Scan efficiency | Medium<br>Lane spacing | Medium<br>Footprint intrusion |
| Speed $v$            | Strong<br>Coverage rate + detection prob. | $v$<br>Frame count                        | Weak<br>Independent     | Weak<br>Independent    | Medium<br>Reaction distance   |
| Scan angle $\theta$  | Weak<br>Scan efficiency                   | Weak<br>Independent                       | $\theta$<br>Turn count  | Weak<br>Independent    | Medium<br>NFZ proximity       |
| Overlap $\alpha$     | Medium<br>Lane spacing                    | Weak<br>Independent                       | Weak<br>Independent     | $\alpha$<br>Redundancy | Weak<br>Independent           |
| NFZ margin $d_{nfz}$ | Medium<br>Footprint intrusion             | Medium<br>Reaction distance               | Medium<br>NFZ proximity | Weak<br>Independent    | $d_{nfz}$<br>Safety margin    |


(b) Variable–objective coupling matrix. Cell labels indicate which objectives ( $C$ ,  $D$ ,  $T$ ,  $E$ ,  $S$ ) are jointly affected. The altitude–speed pair jointly determines four of five objectives.

**Figure 58:** Coupling analysis of the decision variables. The altitude–speed pair exhibits the strongest coupling, jointly determining coverage rate, detection probability, mission time, and energy consumption.

The coupling structure has three practical consequences:

1. **Sequential single-variable optimisation fails.** Optimising  $h$  first, then  $\theta$ , would miss the  $h$ – $\theta$  interaction: the number of scan lines depends on both altitude (footprint width) and angle (projected polygon width). The joint 216-configuration sweep was necessary.
2. **The altitude-dependent speed schedule reduces dimensionality.** By linking  $v$  to  $h$  through  $v(h) = 6 + 4(h - 20)/30$ , the effective search space drops from  $6 \times 36 \times n_v$  to  $6 \times 36 = 216$ —small enough for exhaustive enumeration in under 2s on a laptop.
3. **Overlap and NFZ buffer are weakly coupled to the rest.** Both affect only coverage ( $\alpha$  determines lane spacing;  $d_{nfz}$  determines edge truncation) and were fixed analytically:  $\alpha = 0.20$  covers the 4.6 m RSS lane error with 1.8 m to spare, and  $d_{nfz} = 30$  m provides a  $2.2\times$  safety factor over the RSS margin requirement.

The sensitivity matrix (Figure 59a) and spider plot (Figure 59b) provide complementary views: the matrix shows how each input parameter influences every derived quantity, while the spider plot reveals the overall sensitivity footprint of the design point.



**Figure 59:** Complementary sensitivity visualisations at the selected operating point (35 m, 70°, 8 m s<sup>-1</sup>).

## AB.7 Pareto Frontier

The 216 configurations define a five-dimensional objective space. Figure 60 projects this onto a two-dimensional view: mission effectiveness (coverage × detection) versus mission cost (normalised energy + time).

**Alternative projections of the five-dimensional objective space.** The Pareto frontier can be visualised in several complementary ways, each revealing different aspects of the trade-off structure.

These four visualisations collapse the same five-dimensional objective space into progressively richer representations: the 2D curve (Figure 61) shows the dominant detection–time trade-off, the 3D plot (Figure 62) adds energy as a third axis with coverage as colour, the parallel coordinates (Figure 63) displays all five dimensions without projection loss, and the simplified view (Figure 64) distils the insight for decision-makers.

The selected operating point lies at the “knee” of the Pareto frontier: the region where marginal improvements in effectiveness require disproportionate increases in cost. Moving rightward along the frontier (e.g., to 30 m at 65°) gains 0.02 percentage points of detection probability at the cost of 17% more energy and 20 seconds of additional flight time. Moving leftward (e.g., to 50 m at 70°) saves 35% energy but reduces the target width to 7 model pixels, eliminating the safety margin against real-world degradation.

The energy efficiency curve (Figure 65) confirms that the selected point operates well within the battery budget: 12.6 W h is 10.9% of the 115.4 W h capacity, leaving over 14 min of flight endurance for transit, descent, verification, and return-to-launch.

## AB.8 Summary of Selected Parameters


The dependency chain runs: **target size** → **altitude** → **footprint** → **speed** → **scan angle** → **NFZ margins** → **overlap**. Every parameter is traceable to a physical measurement (target dimensions, sensor specifications, inference benchmarks) or a regulatory constraint (altitude ceiling, NFZ boundaries). The chain can be re-evaluated instantly if any input changes—a faster inference backend would relax the speed constraint, and a different target size would shift the altitude ceiling—making the design robust to future system upgrades.

## AB.9 Alternative Strategies Considered

Several alternative search strategies were evaluated during the design phase but ultimately rejected in favour of the single-pass lawnmower pattern described above.

**Wind-aligned scanning.** We considered orienting scan lines at 45° to the prevailing westerly wind direction, on the hypothesis that diagonal scanning would reduce crosswind-induced lane drift and improve inter-swath alignment. Preliminary energy analysis indicated, however, that this orientation would require approximately 30 % more scan lines than the longest-edge alignment for the AENGM0074 polygon, due to the mismatch between the wind axis and the polygon’s elongated northeast–southwest geometry. Since the 20 % lateral overlap margin already compensates for the 1 m to 2 m crosswind displacement expected in typical conditions (Section V.5), the marginal alignment benefit did not justify the substantial increase in flight time and energy consumption.

**Adaptive speed near detection clusters.** Initial experiments in simulation explored an adaptive speed schedule



Each dot = one (altitude, speed, overlap) configuration. This 2D view collapses 5 objectives into 2 composite dimensions.

**Figure 60:** Pareto frontier in the effectiveness–cost plane. Each point is one of the 216 evaluated configurations, coloured by altitude. The Pareto-efficient frontier (black line) separates achievable from dominated solutions. The selected operating point (35 m, 70°) lies on the frontier at the “knee”—the point of maximum curvature where marginal effectiveness gains require disproportionate cost increases.

## AC Multi-Objective Search Optimisation

The search pattern is governed by a set of coupled decision variables whose optimal values cannot be determined independently. This section formalises the optimisation problem, defines each objective function with its physical basis, identifies the constraint set imposed by regulations and hardware, analyses the coupling structure, and justifies the solution approach.

### AC.1 Problem Formulation


Let the decision vector be

$$\mathbf{x} = [h, v, \theta, \alpha, d_{\text{nfz}}]^\top \quad (79)$$

where  $h \in [20, 50]$  m is the search altitude,  $v \in [6, 10]$  m/s is the forward speed (coupled to  $h$  through the altitude-dependent schedule),  $\theta \in [0^\circ, 180^\circ)$  is the scan angle relative to north,  $\alpha \in [0, 0.5]$  is the lateral overlap fraction between adjacent swaths, and  $d_{\text{nfz}} \geq 0$  m is the minimum planned standoff distance from the SSSI no-fly zone boundary.

The multi-objective optimisation problem is stated as:

$$\begin{aligned} \underset{\mathbf{x}}{\text{minimise}} \quad & \mathbf{F}(\mathbf{x}) = [-f_{\text{cov}}(\mathbf{x}), -f_{\text{det}}(\mathbf{x}), f_{\text{time}}(\mathbf{x}), f_{\text{energy}}(\mathbf{x}), -f_{\text{safety}}(\mathbf{x})]^\top \\ \text{subject to} \quad & h \leq 50 \text{ m} && (\text{R04: altitude ceiling}) \\ & \mathbf{w}_i \in \mathcal{P}_{\text{flight}}, \quad \forall i && (\text{R01: flight area}) \\ & \mathcal{F}(h, \mathbf{w}_i) \cap \mathcal{P}_{\text{SSSI}} = \emptyset, \quad \forall i && (\text{R02: no footprint over SSSI}) \\ & \|\mathbf{w}_0 - \mathbf{p}_{\text{TOL}}\| \leq 5 \text{ m} && (\text{R03: takeoff proximity}) \\ & f_{\text{energy}}(\mathbf{x}) \leq 0.8 E_{\text{batt}} && (\text{battery reserve}) \end{aligned} \quad (80)$$



**Figure 61:** Classical 2D Pareto curve: detection probability vs mission time. Each dot represents one of the 216 swept configurations, coloured by altitude. The Pareto frontier connects non-dominated solutions. The selected operating point (35 m, 8 m/s) sits at the knee of the curve.

where  $\mathbf{w}_i$  is the  $i$ -th waypoint,  $\mathcal{P}_{\text{flight}}$  is the flight area polygon,  $\mathcal{F}(h, \mathbf{w}_i)$  is the camera footprint polygon at altitude  $h$  centred on waypoint  $\mathbf{w}_i$ ,  $\mathcal{P}_{\text{SSSI}}$  is the SSSI polygon,  $\mathbf{p}_{\text{TOL}}$  is the designated take-off/landing point, and  $E_{\text{batt}} = 115.4 \text{ Wh}$  is the battery capacity (6S 5200 mAh LiPo). The 80% cap reserves 20% for return-to-launch, descent, and contingency hover.

Coverage and detection are negated because they are maximisation objectives cast into the standard minimisation form.

## AC.2 Objective Functions

Each objective is derived from the physical relationships linking the decision variables to mission performance.

### AC.2.1 Coverage Completeness

The camera ground footprint at altitude  $h$  has width

$$W_g(h) = \frac{w_s \cdot h}{f} \quad (81)$$

and along-track height

$$H_g(h) = \frac{h_s \cdot h}{f} \quad (82)$$


where  $w_s = 5.02 \text{ mm}$  and  $h_s = 3.75 \text{ mm}$  are the sensor dimensions (IMX296), and  $f = 5.46 \text{ mm}$  is the calibrated focal length. At the selected altitude of 35 m,  $W_g = 32.2 \text{ m}$  and  $H_g = 24.0 \text{ m}$ .

Adjacent scan lines are spaced at

$$L(h, \alpha) = W_g(h) \cdot (1 - \alpha) \quad (83)$$

yielding  $n_{\text{lines}}(\theta) = \lceil W_{\perp}(\theta)/L \rceil$  scan lines, where  $W_{\perp}(\theta)$  is the polygon extent perpendicular to the scan direction. The coverage fraction is

### 3D Objective Space (4th dimension: colour = coverage)



**Figure 62:** Three-dimensional objective space adding energy consumption as the vertical axis, with coverage encoded as colour (fourth dimension). The selected configuration achieves low energy and high detection while maintaining coverage above 96%.

$$f_{\text{cov}}(h, \theta, \alpha, d_{\text{nfz}}) = \frac{A_{\text{scanned}}(h, \theta, \alpha, d_{\text{nfz}})}{A_{\text{total}}} \quad (84)$$

where  $A_{\text{scanned}}$  is the union of all swath rectangles clipped to the survey polygon minus the NFZ buffer zone, and  $A_{\text{total}} = 32\,200\text{ m}^2$  is the survey polygon area. The NFZ standoff  $d_{\text{nfz}}$  reduces  $A_{\text{scanned}}$  by excluding scan line endpoints within  $d_{\text{nfz}}$  of the SSSI boundary, creating a coverage gap along the northeast edge. At  $d_{\text{nfz}} = 30\text{ m}$ , this gap is approximately 4% of the total area.

#### AC.2.2 Detection Probability

A target of physical width  $w_t = 0.5\text{ m}$  (mannequin) subtends


$$p_{\text{px}}(h) = \frac{w_t}{\text{GSD}(h)} \cdot \frac{640}{W_{\text{px}}} = \frac{w_t \cdot f \cdot 640}{w_s \cdot h} \quad (85)$$

pixels in the  $640 \times 640$  model input, where  $\text{GSD}(h) = w_s h / (f \cdot W_{\text{px}})$  is the ground sample distance and  $W_{\text{px}} = 1456$  is the sensor width. The per-frame detection probability  $p(h)$  is modelled as a sigmoid function of the target's pixel extent:

$$p(h) = \frac{1}{1 + \exp(-\kappa (p_{\text{px}}(h) - p_{\text{min}}))} \quad (86)$$

where  $p_{\text{min}} = 7\text{ px}$  is the empirically observed minimum detectable width and  $\kappa = 0.5\text{ px}^{-1}$  controls the transition sharpness. This sigmoidal model captures the behaviour observed in video analysis: detection is virtually certain above 15 px and rapidly degrades below 10 px, with no sharp threshold.

The number of inference frames in which the target appears during a single flyover is



**Figure 63:** Parallel coordinates plot showing all five objectives simultaneously without dimensional projection. Orange lines indicate the seven Pareto-optimal configurations; the selected configuration (thick blue) achieves balanced performance across all axes. This representation reveals that the selected point trades speed and energy efficiency for superior detection and safety.

$$n_f(h, v) = \frac{H_g(h)}{v} \cdot f_{\text{inf}} \quad (87)$$

where  $f_{\text{inf}} = 4.8 \text{ FPS}$  is the TFLite inference rate on the Raspberry Pi 5. Assuming statistically independent per-frame detections (justified by the global shutter eliminating inter-frame correlation from rolling-shutter artefacts), the cumulative detection probability is

$$f_{\text{det}}(h, v) = 1 - (1 - p(h))^{n_f(h, v)} \quad (88)$$

At the selected operating point ( $h = 35 \text{ m}$ ,  $v = 8 \text{ m s}^{-1}$ ), the target subtends  $p_{\text{px}} = 10 \text{ px}$  in the model input, the dwell time is  $3.0 \text{ s}$ ,  $n_f = 14$  frames, and  $f_{\text{det}} > 0.9997$ .

### AC.2.3 Mission Time

The total mission time decomposes into forward-flight time along scan lines, U-turn penalties, and transit to and from the search area:


$$f_{\text{time}}(h, v, \theta) = \sum_{i=1}^{n_{\text{lines}}} \frac{\ell_i}{v_{\text{eff}, i}} + n_{\text{turns}} \cdot \Delta t_{\text{turn}} + \frac{2 d_{\text{transit}}}{v_{\text{transit}}} \quad (89)$$

where  $\ell_i$  is the length of the  $i$ -th scan line,  $v_{\text{eff}, i}$  is the effective speed along that line (reduced from  $v$  in the NFZ slow zone per the speed ramp of Eq. (90)),  $n_{\text{turns}} = n_{\text{lines}} - 1$ ,  $\Delta t_{\text{turn}} = 2 \text{ s}$  is the deceleration–reacceleration penalty per U-turn,  $d_{\text{transit}}$  is the one-way distance from take-off to the search pattern entry point, and  $v_{\text{transit}} = 15 \text{ m s}^{-1}$ . Within the NFZ slow zone, the effective speed follows a linear ramp:

$$v_{\text{eff}}(d) = v_{\text{min}} + \frac{d}{d_{\text{zone}}} \cdot (v_{\text{zone}} - v_{\text{min}}), \quad 0 \leq d \leq d_{\text{zone}} \quad (90)$$

with  $v_{\text{min}} = 0.3 \text{ m s}^{-1}$ ,  $v_{\text{zone}} = 3.0 \text{ m s}^{-1}$ , and  $d_{\text{zone}} = 20 \text{ m}$ .

### Trade-off Summary: Key Configurations



**Figure 64:** Simplified trade-off summary for non-technical audiences. Five key configurations plotted on composite effectiveness vs cost axes, illustrating the fundamental tension between mission reliability and resource expenditure.

#### AC.2.4 Energy Consumption

Total power at forward speed  $v$  is modelled as the sum of a constant hover component and a speed-dependent parasitic drag component:

$$P(v) = P_{\text{hover}} + P_{\text{drag,ref}} \left( \frac{v}{v_{\text{ref}}} \right)^2 \quad (91)$$

where  $P_{\text{hover}} = 150 \text{ W}$  (momentum-theory estimate for the EDU-450 at  $2.3 \text{ kg AUW}$ ),  $P_{\text{drag,ref}} = 50 \text{ W}$  at  $v_{\text{ref}} = 5 \text{ m s}^{-1}$ , and the quadratic scaling follows from the aerodynamic drag relation  $\frac{1}{2} C_D \rho A v^2$ . The total energy integrates power over the mission profile:

$$f_{\text{energy}}(\mathbf{x}) = \int_0^{t_{\text{total}}} P(v(t)) dt = P_{\text{hover}} \cdot t_{\text{total}} + P_{\text{drag,ref}} \sum_{i=1}^{n_{\text{lines}}} \int_0^{\ell_i / v_{\text{eff},i}} \left( \frac{v_{\text{eff},i}(t)}{v_{\text{ref}}} \right)^2 dt \quad (92)$$

During U-turns, only the hover component is active (the drone decelerates to near-zero ground speed). The NFZ slow zone interacts asymmetrically with energy: lower  $v$  reduces drag power but increases hover time for the same distance, and since  $P_{\text{hover}} \gg P_{\text{drag}}$  at low speeds, the net effect is an energy *increase* for scan lines passing through the slow zone.

#### AC.2.5 NFZ Safety Margin


The safety margin accounts for the physical extent of the camera footprint, not merely the waypoint position:

$$f_{\text{safety}}(h, d_{\text{nfz}}) = d_{\text{nfz}} - \frac{W_g(h)}{2} = d_{\text{nfz}} - \frac{w_s \cdot h}{2f} \quad (93)$$

This is the minimum clearance between the nearest edge of any camera footprint and the SSSI boundary. A positive value indicates that no part of any image frame overlaps the protected zone. At the selected operating point ( $h = 35 \text{ m}$ ,  $d_{\text{nfz}} = 30 \text{ m}$ ),  $W_g/2 = 16.1 \text{ m}$  and  $f_{\text{safety}} = 13.9 \text{ m}$ , providing a safety factor of  $30/16.1 = 1.86$ .

### AC.3 Constraints

The constraint set is derived from regulatory requirements (R01–R04), hardware limits, and operational policy:



**Figure 65:** Energy efficiency across the operating range. The search phase consumes only 10.9% of battery capacity at the selected operating point, leaving ample margin for transit, descent phases, and RTL.

**Table 126:** Constraint summary with source and binding status at the selected operating point.

| ID | Constraint        | Formal expression                                                         | Binding? | Source                 |
|----|-------------------|---------------------------------------------------------------------------|----------|------------------------|
| C1 | Altitude ceiling  | $h \leq 50 \text{ m}$                                                     | No       | CAA Open A3 (R04)      |
| C2 | Flight boundary   | $\mathbf{w}_i \in \mathcal{P}_{\text{flight}} \forall i$                  | No       | KML flight area (R01)  |
| C3 | SSSI exclusion    | $\mathcal{F}(h, \mathbf{w}_i) \cap \mathcal{P}_{\text{SSSI}} = \emptyset$ | Active   | SSSI no-fly zone (R02) |
| C4 | Takeoff proximity | $\ \mathbf{w}_0 - \mathbf{p}_{\text{TOL}}\  \leq 5 \text{ m}$             | No       | TOL designation (R03)  |
| C5 | Battery reserve   | $f_{\text{energy}} \leq 0.8 \times 115.4 = 92.3 \text{ Wh}$               | No       | 20% RTL reserve        |
| C6 | Min pixel extent  | $p_{\text{px}}(h) \geq 7 \text{ px}$                                      | Active   | Detection threshold    |

Constraints C3 and C6 are the active (near-binding) constraints at the selected operating point: the SSSI exclusion creates the 4% coverage gap, and the minimum pixel extent prevents operating above approximately 50 m. The battery constraint (C5) is far from binding—the search consumes 12.6 Wh (10.9% of capacity), leaving ample margin for transit, descent, verification, and return.

#### AC.4 Variable Coupling Analysis

The five decision variables do not affect the objectives independently. Table 128 shows the coupling structure: each cell indicates which objectives are jointly influenced by the corresponding variable pair.

**Table 128:** Variable–objective coupling matrix. Entries indicate which objectives ( $C$  = coverage,  $D$  = detection,  $T$  = time,  $E$  = energy,  $S$  = safety) are jointly affected by each variable pair. The altitude–speed pair exhibits the strongest coupling, jointly determining four of five objectives.

|                  | $h$             | $v$       | $\theta$  | $\alpha$ | $d_{\text{nfz}}$ |
|------------------|-----------------|-----------|-----------|----------|------------------|
| $h$              | $C, D, T, E, S$ | $D, T, E$ | $C, T, E$ | $C$      | $C, S$           |
| $v$              | —               | $D, T, E$ | $T, E$    | —        | —                |
| $\theta$         | —               | —         | $C, T, E$ | $C$      | $C$              |
| $\alpha$         | —               | —         | —         | $C$      | —                |
| $d_{\text{nfz}}$ | —               | —         | —         | —        | $C, S$           |

Three observations follow from the coupling structure:

1. **Altitude is the most coupled variable.** It appears in all five objectives: higher altitude increases the footprint (improving coverage rate and safety margin) but reduces target pixel extent (degrading detection confidence), while simultaneously affecting mission time (fewer scan lines) and energy (fewer U-turns but higher drag at the associated speed).

1. **Discrete feasible set.** The altitude is constrained to values compatible with the ArduPilot waypoint resolution (1 m), and angles finer than  $5^\circ$  produce negligible changes in scan line count. The resulting  $6 \times 36$  grid is small enough for complete evaluation in under 2 s on a laptop.
2. **Non-convex objectives.** The detection probability  $f_{\text{det}}$  has a sigmoidal dependence on altitude (Eq. (86)), and the energy landscape contains local minima created by discrete jumps in the number of scan lines as the swath width crosses polygon geometry thresholds. Gradient-based methods would require careful initialisation to avoid local optima.
3. **Geometry-dependent evaluation.** Each configuration requires generating the full lawnmower pattern on the real polygon, computing per-segment NFZ slow-zone penalties via numerical integration (50 sample points per segment), and evaluating the energy model. The evaluation function is not differentiable with respect to the scan angle because the number of scan lines changes discontinuously.
4. **Interpretability.** Exhaustive enumeration produces the complete Pareto frontier, enabling direct visualisation of trade-offs (Figures 47, 2) rather than a single “optimal” point whose sensitivity to modelling assumptions is opaque.

From the 216 configurations, Pareto-dominated solutions are filtered, and the operating point is selected from the remaining Pareto-efficient set using a weighted composite score:

$$S(\mathbf{x}) = w_D \cdot \hat{f}_{\text{det}}(\mathbf{x}) + w_C \cdot \hat{f}_{\text{cov}}(\mathbf{x}) + w_E \cdot (1 - \hat{f}_{\text{energy}}(\mathbf{x})) \quad (95)$$

where  $\hat{f}$  denotes min-max normalisation to  $[0, 1]$  across the feasible set, and the weights  $w_D = 0.40$ ,  $w_C = 0.35$ ,  $w_E = 0.25$  reflect the SAR priority hierarchy: detection reliability is paramount, coverage completeness is second (a missed area cannot be searched), and energy efficiency is third (the search consumes only 10.9% of battery regardless). The selected operating point (35 m,  $70^\circ$ ,  $8 \text{ m s}^{-1}$ ) achieves the highest composite score:  $f_{\text{det}} = 0.9997$ ,  $f_{\text{cov}} = 0.96$ ,  $f_{\text{energy}} = 12.6 \text{ Wh}$  (10.9% of battery), and  $f_{\text{safety}} = 13.9 \text{ m}$  net clearance from the SSSI. This point is Pareto-efficient: no other configuration improves any objective without degrading at least one other.

## AD Design Strategy: Parameter Derivation Chain

The operating parameters in Table 118 were not selected independently. Each decision constrains the next, forming a directed dependency chain from the physical target through to the final mission plan. This section traces that chain step by step, showing how each parameter was derived from evidence rather than assumed.

### AD.1 Step 1: Maximum Detection Altitude

The reasoning begins with the target. The mannequin is 1.8 m tall and 0.5 m wide. At altitude  $h$ , the target projects onto the camera sensor with a height in pixels of:

$$\text{px}_{\text{target}} = \frac{d_{\text{target}} \cdot f_{\text{px}}}{h} \quad (96)$$

where  $f_{\text{px}} = f/w_s \times W_{\text{px}} = 5.46/5.02 \times 1456 = 1584$  pixels is the focal length in pixel units, and  $d_{\text{target}}$  is the target dimension. YOLOv8n resizes the  $1456 \times 1088$  sensor frame to  $640 \times 640$ ; the uniform scale factor is  $640/\max(1456, 1088) = 640/1456 = 0.4396$ :

$$\text{px}_{\text{model}} = \frac{d_{\text{target}} \cdot f_{\text{px}}}{h} \times \frac{640}{1456} \quad (97)$$

Table 130 traces this function across the altitude range.

Two thresholds emerge. The *absolute* detection limit occurs when the target height drops below approximately 20 px in the model input ( $h \approx 70 \text{ m}$ ), at which point the feature map representation is too coarse for reliable bounding box regression. The *practical* limit is set by target *width*: at 50 m the mannequin’s width is only 7 model pixels, below the 10-pixel comfort threshold for width-dependent features (limb visibility, aspect ratio cues). Empirically, detection confidence from the retrained `sar_v2_1088` model drops below 0.5 at approximately 60 m.

Regulatory constraint R03 limits altitude to 50 m. Selecting  $h = 35 \text{ m}$  provides a  $\times 1.4$  margin on target width (10 px vs 7 px at 50 m) and a  $\times 1.8$  margin on target height (36 px vs 20 px minimum), creating headroom for real-world degradation modes: partial occlusion, unusual pose, non-ideal lighting, and altitude excursions during manoeuvres.

**Decision:**  $h = 35 \text{ m}$ . This sacrifices the energy-optimal 50 m (8.2 Wh) for a comfortable detection margin (12.6 Wh—still only 10.9% of battery capacity).

## AD.2 Step 2: Maximum Speed at Chosen Altitude

With altitude fixed, the camera’s along-track ground footprint determines how long the target remains in view. At 35 m:

$$H_g = \frac{w_s \cdot h}{f} \times \frac{H_{\text{px}}}{W_{\text{px}}} = \frac{5.02 \times 35}{5.46} \times \frac{1088}{1456} = 24.1 \text{ m} \quad (98)$$

At the TFLite inference rate of  $f_{\text{inf}} = 4.8 \text{ FPS}$  (benchmarked on Pi 5, Section H.4), the drone advances  $\Delta x = v/f_{\text{inf}}$  metres between frames. For reliable detection, the target must appear in at least  $n_{\text{min}}$  independent inference frames. The maximum speed is therefore:

$$v_{\text{max}} = \frac{H_g}{n_{\text{min}}/f_{\text{inf}}} = \frac{H_g \cdot f_{\text{inf}}}{n_{\text{min}}} \quad (99)$$

If we require  $n_{\text{min}} = 3$  frames (one detection above threshold triggers investigation), then  $v_{\text{max}} = 24.1 \times 4.8/3 = 38.6 \text{ m s}^{-1}$ —far above drone capability. However, at high frame rates consecutive frames are highly correlated: at  $8 \text{ m s}^{-1}$  the inter-frame advance is 1.67 m, yielding 91% spatial overlap between adjacent frames. These are not independent observations.

A conservative requirement of  $n_{\text{min}} = 10$  raw frames ensures that at least 3–4 contain substantially different views of the target (different pixel positions, different background context):

$$v_{\text{max,conservative}} = \frac{24.1 \times 4.8}{10} = 11.5 \text{ m s}^{-1} \quad (100)$$

The altitude-dependent speed schedule (Section AA.1) assigns  $v(35 \text{ m}) = 8.0 \text{ m s}^{-1}$ , which yields:

$$n_f = \frac{H_g}{v} \times f_{\text{inf}} = \frac{24.1}{8.0} \times 4.8 = 14.5 \approx 14 \text{ frames} \quad (101)$$

This provides a  $\times 1.4$  margin above the conservative 10-frame requirement and a  $\times 4.7$  margin above the absolute 3-frame minimum. The cumulative detection probability over 14 frames, assuming per-frame probability  $p = 0.95$  (conservative for 48-pixel targets), is  $P_d = 1 - (1 - 0.95)^{14} > 0.9999$ .

**Decision:**  $v = 8 \text{ m s}^{-1}$  at 35 m, following the altitude-dependent linear schedule.

## AD.3 Step 3: Minimise Energy via Scan Angle

With  $h$  and  $v$  fixed, the remaining free variable for the boustrophedon pattern is the scan angle  $\theta$ —the orientation of parallel scan lines relative to north. This controls the number of scan lines (and therefore U-turns) required to cover the polygon.

The 216-configuration parametric sweep (Section AA.3) tested 36 angles at each of 6 altitudes. At  $h = 35 \text{ m}$ , the results are unambiguous: scan angle  $\theta = 70^\circ$  (alignment with the polygon’s longest edge, the northeast–southwest axis) minimises the number of scan lines from 10 (at  $0^\circ$ , north–south) to 6, and reduces U-turns from 9 to 5.

The energy savings are substantial. At the optimal angle:  $E = 12.6 \text{ Wh}$ . At the worst angle ( $95^\circ$ , perpendicular to the long axis):  $E = 25.4 \text{ Wh}$ —a  $2\times$  penalty for the same coverage. The energy valley at  $\theta \approx 70^\circ$  is visible in Figure 47 across all altitudes, confirming that scan angle alignment is geometry-dependent and altitude-invariant.

This angle also has a fortunate geometric property: the scan lines run roughly *perpendicular* to the SSSI boundary on the northeast edge. This means each scan line clips the NFZ slow zone only at its endpoint, rather than running parallel to the boundary (which would force the entire line to be traversed at reduced speed). The NFZ energy penalty at  $70^\circ$  is approximately 15%; at  $0^\circ$  (north–south, running along the SSSI edge), the penalty exceeds 40%.

**Decision:**  $\theta = 70^\circ$ , aligned to the polygon’s longest edge via `cv2.minAreaRect`.

## AD.4 Step 4: NFZ Safety Margins

The SSSI no-fly zone on the polygon’s northeast edge requires a buffer that keeps the entire camera footprint—not just the drone’s position—outside the protected area. At 35 m altitude, the cross-track ground footprint is:

$$W_g = \frac{w_s \cdot h}{f} = \frac{5.02 \times 35}{5.46} = 32.2 \text{ m} \quad (102)$$

The camera images a strip extending  $W_g/2 = 16.1 \text{ m}$  either side of nadir. To guarantee the camera footprint does not overlap the SSSI, the drone must maintain at least 16.1 m lateral distance from the boundary. Three additional error sources widen this margin:

1. **GPS position error:**  $\text{CEP}_{95} = 3 \text{ m}$  for consumer-grade GNSS.

2. **Reaction time:** at  $v = 8 \text{ m s}^{-1}$ , a 2 s reaction window covers 16 m.
3. **Wind gust displacement:** estimated 2 m lateral excursion in 15 kt crosswind.

The root-sum-square (RSS) of these independent errors is:

$$d_{\text{margin}} = \sqrt{16.1^2 + 3^2 + 16^2 + 2^2} = 23.1 \text{ m} \quad (103)$$

The system implements a five-layer defence (see Section D.1 for full details):

- **30 m waypoint buffer** (NFZ\_WAYPOINT\_BUFFER\_M): any waypoint within 30 m of the SSSI boundary is excluded at plan time.
- **20 m speed ramp** (NFZ\_SLOW\_ZONE\_M): linear deceleration from  $3 \text{ m s}^{-1}$  to  $0.3 \text{ m s}^{-1}$  within this zone.
- **Repulsive vector field:** constant-magnitude push away from the boundary when within 3 m.
- **3 m hard boundary** (NFZ\_HARD\_BOUNDARY\_M): automatic switch to MANUAL mode.
- **Firmware geofence:** independent Cube Orange autopilot geofence triggers LOITER/RTL as a final backstop.

The combined margin of 30 m (waypoint buffer) + 20 m (speed ramp) = 50 m exceeds the RSS requirement of 23 m by a factor of 2.2×. As a cross-check, the one-third rule (buffer  $\geq W_g/3$ ) requires 10.7 m—our 30 m buffer exceeds this by 2.8×.

The cost of this conservatism is a 4 % coverage gap in the strip adjacent to the SSSI. This strip lies within a fenced wildlife reserve with no public access, making the probability of a casualty there negligibly small.

**Decision:** 30 m waypoint buffer, 20 m speed ramp, 3 m hard cut-off. Accept 96 % polygon coverage.

## AD.5 Step 5: Swath Overlap

Adjacent scan lines must overlap sufficiently to prevent coverage gaps caused by lateral position error. Three sources of lane misalignment are relevant:

1. **GPS cross-track drift:**  $\sigma_{\text{GPS}} = 3 \text{ m}$  (consumer GNSS CEP).
2. **Crosswind displacement:**  $\delta_{\text{wind}} \approx 1.5 \text{ m}$  lateral shift in 10 kt crosswind at  $8 \text{ m s}^{-1}$ .
3. **Attitude roll coupling:** a  $5^\circ$  roll displaces the footprint centre by  $h \tan(5^\circ) = 35 \times 0.087 = 3.1 \text{ m}$ .

Assuming independence, the worst-case combined displacement is:

$$\delta_{\text{total}} = \sqrt{3^2 + 1.5^2 + 3.1^2} = 4.6 \text{ m} \quad (104)$$

At 35 m altitude with 20 % overlap, the overlap margin in metres is:

$$m_{\text{overlap}} = 0.20 \times W_g = 0.20 \times 32.2 = 6.4 \text{ m} \quad (105)$$

This covers the worst-case displacement with 1.8 m to spare ( $6.4 - 4.6 = 1.8 \text{ m}$ ). Increasing to 30 % overlap would add an extra scan line (7 instead of 6), increasing flight time by approximately 10 % for only 3.2 m of additional margin—beyond the point of diminishing returns given the RSS error model.

**Decision:** 20 % swath overlap.

## AD.6 Step 6: Diagonal vs Aligned Scanning

A diagonal scan orientation (e.g.,  $45^\circ$  to the wind) has been suggested in the coverage planning literature as a way to reduce crosswind-induced lane drift, since the wind component perpendicular to the scan direction is reduced [7]. For the AENGM0074 polygon, this would mean rotating the scan lines approximately  $25^\circ$  from the longest-edge alignment. The cost is significant: rotating from  $70^\circ$  to  $45^\circ$  increases the number of scan lines from 6 to approximately 8 (a 33 % increase), adding 2 U-turns and approximately 300 m of additional path length. The resulting energy cost rises from 12.6 Wh to approximately 16.4 Wh—a 30 % penalty.

The benefit is marginal: the crosswind displacement (1.5 m) is already covered by the 6.4 m overlap margin with 4.9 m to spare. Even doubling the crosswind to 20 kt (3 m displacement) leaves 3.4 m of margin within the existing 20 % overlap.

**Decision:** maintain longest-edge alignment ( $70^\circ$ ). The 20 % overlap already absorbs crosswind effects without the 30 % energy penalty of diagonal scanning.

## AD.7 The Dependency Chain

The six steps above form a directed acyclic graph of parameter dependencies, where each decision is conditioned on the outputs of previous steps:

1. **Target size**  $\rightarrow$  **altitude:** the mannequin's 1.8 m height and the 20-pixel model-input minimum determine  $h_{\text{max}} = 70 \text{ m}$ ; the 10-pixel width comfort threshold sets  $h_{\text{comfort}} = 50 \text{ m}$ ; a 30 % safety margin yields  $h = 35 \text{ m}$ .

2. **Altitude**  $\rightarrow$  **footprint**  $\rightarrow$  **speed**:  $h = 35$  m gives a 24 m along-track footprint; at 4.8 FPS, requiring 10+ raw frames caps speed at  $11.5 \text{ m s}^{-1}$ ; the altitude-dependent schedule assigns  $v = 8 \text{ m s}^{-1}$ .
3. **Altitude** + **speed**  $\rightarrow$  **scan angle**: with  $h$  and  $v$  fixed, the 216-configuration sweep identifies  $\theta = 70^\circ$  as the energy minimum (longest-edge alignment, fewest U-turns).
4. **Altitude**  $\rightarrow$  **footprint width**  $\rightarrow$  **NFZ margin**: the 32 m footprint at 35 m requires at least 16 m standoff; GPS error, reaction time, and wind push this to 23 m RSS; the 30 m buffer provides a  $2.2\times$  safety factor.
5. **Footprint** + **error model**  $\rightarrow$  **overlap**: 4.6 m RSS lane error is covered by 6.4 m (20 %) overlap margin.
6. **Energy budget** validates the chain: 12.6 Wh is 10.9 % of the 115.4 Wh battery, leaving 14 min of flight endurance—sufficient for the 1.9 min search plus transit, descent, verification, and landing.

**Table 18:** Plus/delta review of system technical performance.

| Item                                                              | Evidence / Detail                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|-------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Plus — strengths</b>                                           |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>P1</b> Dual-backend CV architecture                            | Identical application code runs Ultralytics on the laptop and TFLite on the Pi; model files are drop-in interchangeable. The model was swapped three times (640, 1280, 1088 datasets) with zero changes to application code (Section 2.3).                                                                                                                                                                                                                                                                                                                                                                                                              |
| <b>P2</b> Configuration auto-detection                            | Serial-port probing selects real Cube or SITL with zero code changes. Verified by running the same <code>main.py</code> on Windows laptop (SITL), WSL (SITL), and Pi 5 (real Cube) without modification (Section 2.2).                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| <b>P3</b> Progressive testing caught 12 defects at low-cost tiers | The five-tier framework discovered 12 integration faults before any flight attempt (Table 147): BGR colour inversion, FOV miscalibration ( $f=7.0 \rightarrow 5.46$ mm), geofence sign errors, TFLite bounding box coordinate mismatch, GPS estimate double-scaling, pyserial breakage, and 6 others—all resolved in 5–120 min at the bench tier. Seven of these are documented as evidence-based corrections in Section 1.10. Total debugging time: ~6 h 10 min across 12 defects; estimated cost if discovered in flight: 2 crashed airframes (£2000+) and regulatory violation (LL-08). See Section L.5 and Appendix K for the full defect register. |
| <b>P4</b> On-device inference performance                         | 4.8FPS (206.5 ms/frame) on Pi 5 CPU, measured over 50 runs with 0.995 mAP <sub>50</sub> on the retrained model; 50/50 detections at 0.966 mean confidence in the bench benchmark (Section 2.3).                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>P5</b> Modular vision interface                                | <code>vision.py</code> is the sole file that touches CV/AI; all other modules call <code>detect_in_image(frame) → (found, x, y, conf)</code> . Adding lens undistortion required modifying one file and added only 1.5 ms per frame—0.7% of inference time (Section 2.3).                                                                                                                                                                                                                                                                                                                                                                               |
| <b>P6</b> Web-based ground station                                | Browser dashboard with MJPEG stream, GPS grid, and command buttons operates headlessly over SSH at <a href="http://PI.IP:8090">http://PI.IP:8090</a> . Tested from PuTTY terminal and phone browser during bench day (Section 2.7).                                                                                                                                                                                                                                                                                                                                                                                                                     |
| <b>P7</b> 74 test scripts across 8 categories                     | 74 scripts totalling ~5 500 LOC across hardware checks (8), flight progressions (14), diagnostics (5), calibration (7), field experiments (6), development tools (21), unit tests (6 files, 127 tests), and automated integration tests (7 files, 33 tests)—each single-purpose with CSV output where applicable. See Appendix N for the full script inventory and Appendix K (Section L.3) for the categorisation rationale.                                                                                                                                                                                                                           |
| <b>P8</b> Simulation-first development                            | Full state machine, geofence enforcement, and detection pipeline were validated end-to-end in SITL before hardware connection; the same <code>main.py</code> runs on both platforms. Dry-run mode validates waypoints in under 1 s without any hardware (Section 2.8).                                                                                                                                                                                                                                                                                                                                                                                  |
| <b>P9</b> Quantitative figures from mathematical models           | 15 body figures generated programmatically from calibration data, inference benchmarks, and flight video analysis—not manually drawn—ensuring reproducibility and numerical traceability (Figures 12–13). Generator scripts archived in <code>figs/gen*.py</code> .                                                                                                                                                                                                                                                                                                                                                                                     |
| <b>P10</b> Centering mitigates five error sources                 | The CENTERING hover procedure simultaneously eliminates roll/pitch error (6.2 m $\rightarrow$ <0.6 m), motion blur, and GPS timing lag while enabling multi-frame GPS averaging ( $\sqrt{N}$ improvement) and temporal false-positive filtering—solving sensor limitations through mission design rather than software complexity (Appendix K.8).                                                                                                                                                                                                                                                                                                       |
| <b>Delta — areas for improvement</b>                              |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>D1</b> No outdoor flight completed                             | Sustained winds (12 m s <sup>-1</sup> , gusts to 18 m s <sup>-1</sup> ) on the single scheduled flight day exceeded the EDU-450's 10 m s <sup>-1</sup> operational limit. Result: 0 of 5 tiers exercised outdoors; requirements R05 (search), R06 (PLB), R07 (landing accuracy) are verified in SITL only. Of the 12 requirements, 8 are fully verified (bench + simulation), 3 partially (simulation only), and 1 has a partial gap (R06 PLB not field-tested). <i>Next step</i> : rebook flight slot; execute Tiers 3–5 of the progressive framework with the identical codebase.                                                                     |
| <b>D2</b> Inference rate limits high-speed search                 | At 4.8FPS (206.5 ms/frame, 50-run benchmark) and 8 m s <sup>-1</sup> search speed, inter-frame advance is 1.7 m—5.3% of the 32 m ground footprint at 35 m altitude, so overlap is adequate. However, NCNN benchmarks on Pi show 72 ms/frame (9FPS), a 4.5× improvement already available but not deployed due to limited Pi testing time (Appendix H2, Section I.2). <i>Next step</i> : deploy NCNN backend, validate detection accuracy parity, and thread the capture/inference pipeline for 20–30% additional throughput.                                                                                                                            |
| <b>D3</b> GPS timing lag uncompensated                            | The Here 3+ receiver exhibits 100 ms to 200 ms latency, producing 0.8–1.6 m systematic error at 8 m s <sup>-1</sup> . DJI video analysis measured CEP <sub>50</sub> =2.3 m, worst case 16.5 m (Appendix J, Figures 10a–10b). Inverse-variance averaging partially mitigates this, but first-order lag compensation ( $\Delta \mathbf{p} = \mathbf{v} \cdot \Delta t_{\text{lag}}$ ) is not implemented. <i>Next step</i> : implement lag compensation in the heading direction, tuneable from config                                                                                                                                                    |

**Table 20:** Five-tier progressive testing framework (condensed from Appendix E.2).

| Tier | Environment                  | What It Verifies                                                                     |
|------|------------------------------|--------------------------------------------------------------------------------------|
| 1    | SITL simulation              | State machine logic, path planning, geofence, detection pipeline—zero hardware risk. |
| 2    | Docker Pi test               | TFLite loads on ARM, dependency compatibility, headless operation.                   |
| 3    | Bench (Pi + Cube, no props)  | MAVLink ACKs, camera frames, GPS fix, full sensing chain.                            |
| 4    | Passive flight (pilot on RC) | Real-world detection range, FOV/GPS calibration—Pi issues zero commands.             |
| 5    | Autonomous flight            | End-to-end mission: search, detect, centre, verify, land.                            |

**Table 23:** System configuration parameters grouped by function.

| Parameter                                          | Default | Unit   | Description                                                                            |
|----------------------------------------------------|---------|--------|----------------------------------------------------------------------------------------|
| <i>Flight Altitudes</i>                            |         |        |                                                                                        |
| TARGET_ALT                                         | 35.0    | m      | Search-pattern cruise altitude                                                         |
| VERIFY_ALT                                         | 15.0    | m      | Descent altitude for close-up verification (currently unused; DESCENDING state bypass) |
| RESCAN_ALT_FACTOR                                  | 0.8     | –      | Altitude multiplier per rescan pass                                                    |
| RESCAN_ALT_FLOOR_M                                 | 15.0    | m      | Minimum permitted rescan altitude                                                      |
| <i>Speeds</i>                                      |         |        |                                                                                        |
| TRANSIT_SPEED_MPS                                  | 15.0    | m/s    | Speed to/from search area                                                              |
| SEARCH_SPEED_MPS                                   | 10.0    | m/s    | Default search-leg speed                                                               |
| FOCUS_SEARCH_SPEED_MPS                             | 5.0     | m/s    | Reduced speed inside Focus Area (PLB zone)                                             |
| SPEED_ALT_LOW                                      | 20.0    | m      | Altitude below which SPEED_AT_LOW applies                                              |
| SPEED_ALT_HIGH                                     | 50.0    | m      | Altitude above which SPEED_AT_HIGH applies                                             |
| SPEED_AT_LOW                                       | 6.0     | m/s    | Cruise speed at low altitude (blur mitigation)                                         |
| SPEED_AT_HIGH                                      | 10.0    | m/s    | Cruise speed at high altitude                                                          |
| <i>Camera &amp; Optics (IMX296 Global Shutter)</i> |         |        |                                                                                        |
| SENSOR_WIDTH_MM                                    | 5.02    | mm     | IMX296 sensor physical width                                                           |
| FOCAL_LENGTH_MM                                    | 5.46    | mm     | Calibrated focal length (bench test)                                                   |
| IMAGE_W                                            | 1456    | px     | Capture resolution (width)                                                             |
| IMAGE_H                                            | 1088    | px     | Capture resolution (height)                                                            |
| CAMERA_FLIP_180                                    | True    | –      | Rotate image 180° (inverted mount)                                                     |
| <i>Detection &amp; Target</i>                      |         |        |                                                                                        |
| CONFIDENCE_THRESHOLD                               | 0.2     | –      | Minimum YOLOv8 confidence to register detection                                        |
| DETECT_CONFIRM_FRAMES                              | 3       | frames | Consecutive detections before trigger                                                  |
| DETECT_LOCK_RADIUS_M                               | 5.0     | m      | Ignore detections beyond this from locked target                                       |
| MAX_DETECT_QUEUE                                   | 20      | –      | Maximum queued detections (FIFO overflow)                                              |
| TARGET_REAL_RADIUS_M                               | 0.15    | m      | Physical radius of search dummy                                                        |
| DUMMY_HEIGHT_M                                     | 1.8     | m      | Dummy height used in FOV calculations                                                  |
| REJECTED_TARGET_RADIUS_M                           | 5.0     | m      | Exclusion radius around rejected targets                                               |
| MAX_RESCAN_PASSES                                  | 3       | –      | Altitude-drop rescan attempts before abort                                             |
| <i>NFZ Geofence</i>                                |         |        |                                                                                        |
| NFZ_HARD_BOUNDARY_M                                | 3.0     | m      | Penetration depth triggering forced <b>MANUAL</b> mode                                 |
| NFZ_SOFT_BOUNDARY_M                                | 8.0     | m      | Outer repulsion zone (quadratic field)                                                 |
| NFZ_WAYPOINT_BUFFER_M                              | 30.0    | m      | Skip waypoints within this distance of NFZ                                             |
| NFZ_SLOW_ZONE_M                                    | 20.0    | m      | Speed-reduction field active within this distance                                      |
| NFZ_SCALAR_ZERO_M                                  | 2.0     | m      | Distance at which speed drops to zero (full stop zone)                                 |
| NFZ_MIN_SPEED_MPS                                  | 0.0     | m/s    | Speed at NFZ_SCALAR_ZERO_M and below (full stop)                                       |
| NFZ_ZONE_MAX_SPEED_MPS                             | 3.0     | m/s    | Speed at outer edge of slow zone                                                       |
| NFZ_INNER_OFFSET_M                                 | 20.0    | m      | Inner polygon offset inside NFZ boundary                                               |
| NFZ_INNER_RANGE_M                                  | 23.0    | m      | Repulsion active within this from inner polygon                                        |
| NFZ_PUSH_SPEED_MPS                                 | 5.0     | m/s    | Constant repulsive push speed (away from NFZ)                                          |
| <i>Payload Servo (Tarot Double-Throw)</i>          |         |        |                                                                                        |
| SERVO_CHANNEL                                      | 9       | –      | Cube AUX output channel                                                                |
| SERVO_CLOSE_PWM                                    | 1500    | µs     | PWM to hold / re-latch mechanism                                                       |
| SERVO_PARTIAL_PWM                                  | 1300    | µs     | Stage 1: partial release                                                               |
| SERVO_FULL_PWM                                     | 1100    | µs     | Stage 2: full release                                                                  |

**Table 24:** Augmentation pipeline for the SAR dummy detector. Offline augmentations are applied during synthetic image generation; online augmentations are applied by the YOLO trainer at each epoch.

| Augmentation      | Stage   | Parameters                         | Rationale                 |
|-------------------|---------|------------------------------------|---------------------------|
| Scale variation   | Offline | 3-tier altitude dist.              | Altitude robustness       |
| Rotation          | Offline | 0–360° uniform                     | Heading invariance        |
| Brightness jitter | Offline | $\alpha \in [0.7, 1.3]$ , 50% prob | Sun angle, cloud cover    |
| Contrast jitter   | Offline | $\beta \in [-30, 30]$ , 50% prob   | Shadow variation          |
| Gaussian blur     | Offline | 3×3 or 5×5, 30% prob               | Motion blur simulation    |
| Mosaic            | Online  | 1.0 (always active)                | Small-object detection    |
| Scale             | Online  | 0.3–1.7×                           | Aggressive altitude range |
| Rotation          | Online  | $\pm 20^\circ$                     | Additional orientation    |
| Horizontal flip   | Online  | 0.5 probability                    | Symmetry invariance       |
| Mixup             | Online  | 0.1 probability                    | Regularisation            |
| Copy-paste        | Online  | 0.2 probability                    | Multi-instance diversity  |
| Random erasing    | Online  | <b>Disabled</b> (0.0)              | Target too small to erase |

**Table 25:** Dataset composition for the v2 SAR dummy detector.

| Source               | Count      | Fraction    | Role                                                                |
|----------------------|------------|-------------|---------------------------------------------------------------------|
| Synthetic composites | 300        | 81.9%       | Scale, orientation, and position diversity via domain randomisation |
| Real labelled frames | 16         | 4.4%        | Domain anchor: authentic lighting, lens effects, ground texture     |
| Negative frames      | 50         | 13.7%       | False-positive suppression on confusing terrain features            |
| <b>Total</b>         | <b>366</b> | <b>100%</b> |                                                                     |

**Table 26:** Training hyperparameters for the YOLOv8n SAR dummy detector (v2).

| Parameter                | Value                                     |
|--------------------------|-------------------------------------------|
| Base model               | YOLOv8n (COCO-pretrained)                 |
| Training image size      | 1088 × 1088                               |
| Epochs                   | 150 (early stopping patience: 30)         |
| Batch size               | 8                                         |
| Augmentation: rotation   | $\pm 20^\circ$                            |
| Augmentation: scale      | 0.3–1.7 (aggressive, simulating altitude) |
| Augmentation: mosaic     | 1.0 (enabled throughout)                  |
| Augmentation: mixup      | 0.1                                       |
| Augmentation: copy-paste | 0.2                                       |
| Horizontal flip          | 0.5                                       |

**Table 27:** Detection performance comparison between model versions. All metrics are computed on the respective validation sets.

| Model                | mAP <sub>50</sub> | mAP <sub>50–95</sub> | Precision | Recall |
|----------------------|-------------------|----------------------|-----------|--------|
| v1 (200 syn, 640)    | ~0.95             | —                    | —         | —      |
| v2 (366 mixed, 1088) | <b>0.995</b>      | 0.828                | 0.994     | 0.989  |

**Table 28:** Failure modes and automated responses.

| Failure           | Detection Method                                     | Automated Response                                                                                                                                          |
|-------------------|------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Camera failure    | Frame returns <code>None</code>                      | Search pattern completes without detections; aircraft returns home safely. During <code>VERIFY</code> , HUD displays “CAMERA LOST”; operator should reject. |
| GPS degradation   | Fix type < 3D or satellites < 6 for 5 s continuously | Emergency RTL via <code>MAV_CMD_DO_SET_MODE</code> . If command fails, ArduCopter GCS failsafe triggers RTL independently.                                  |
| MAVLink link loss | <code>recv_match</code> raises exception             | Script issues emergency RTL; ArduCopter GCS failsafe ( <code>FS_GCS_ENABLE</code> ) provides firmware-level backup after heartbeat timeout.                 |
| RC link loss      | Throttle PWM below threshold                         | ArduCopter RC failsafe ( <code>FS_THR_ENABLE</code> ) triggers RTL independently of onboard computer.                                                       |
| Low battery       | Voltage below configured threshold                   | ArduCopter battery failsafe triggers RTL or LAND depending on severity level.                                                                               |
| Pi software crash | Heartbeat stream ceases                              | ArduCopter GCS failsafe RTL; Cube continues flying autonomously.                                                                                            |
| Operator timeout  | 120 s without Y/N/I input at <code>VERIFY</code>     | Target auto-rejected; search resumes.                                                                                                                       |
| NFZ incursion     | <code>cv2.pointPolygonTest</code> returns positive   | Immediate transition to <code>MANUAL</code> ; all autonomous commands halted.                                                                               |

**Table 29:** Risk register with mitigations.

| Risk                            | L | S | Mitigation                                                                                                                 |
|---------------------------------|---|---|----------------------------------------------------------------------------------------------------------------------------|
| Autopilot hardware failure      | L | 3 | Pre-flight inspection; RC kill switch; ArduCopter watchdog reboot                                                          |
| SSSI boundary incursion         | L | 3 | Five-layer geofence (speed field, repulsion, hard cutoff, waypoint filter, firmware backup)                                |
| Battery depletion mid-mission   | L | 3 | Pre-flight voltage check; ArduCopter battery failsafe triggers RTL/LAND                                                    |
| Loss of all communications      | L | 3 | Dual link (RC + MAVProxy); GCS failsafe RTL on heartbeat loss                                                              |
| Injury to bystanders            | L | 3 | Site clearance; flight area boundary; VLOS maintained; propeller guards                                                    |
| GPS drift causes wrong landing  | M | 2 | Multi-pass spatial clustering; inverse-variance weighting; operator confirms at <code>VERIFY</code> ; 7.5 m offset landing |
| False positive triggers descent | M | 2 | Operator confirmation required; spatial exclusion of rejected targets; detection queue filtering                           |
| High wind exceeds limits        | M | 2 | Weather go/no-go gate; wind check before launch; <code>LOITER</code> mode for station-holding                              |
| Camera/CV failure mid-flight    | M | 1 | Search pattern completes safely without CV; aircraft returns home                                                          |
| Regulatory non-compliance       | L | 3 | Site risk assessment; CAA Open Category A3 procedures; VLOS maintained                                                     |

**Table 33:** Inference backend comparison for YOLOv8n (640 × 640 input) on Raspberry Pi 5.

| Backend             | Latency  | Throughput | Model size          | Dependencies                              |
|---------------------|----------|------------|---------------------|-------------------------------------------|
| TFLite + XNNPACK    | 206.5 ms | 4.8 FPS    | 3.2 MB              | <code>ai-edge-litert</code> only          |
| NCNN (projected)    | ~68 ms   | ~15 FPS    | 3.1 MB              | <code>ncnn</code> (C++ + Python bindings) |
| Ultralytics/PyTorch | >1200 ms | <0.8 FPS   | 6.2 MB <sup>†</sup> | PyTorch, Ultralytics (>800 MB)            |

<sup>†</sup> `.pt` weights; TFLite and NCNN exports are smaller.

**Table 34:** Inference performance: measured baseline and projected improvements on Raspberry Pi 5 with YOLOv8n ( $640 \times 640$  input, float32 weights).

| Configuration             | Status    | Latency       | FPS        | Notes                                |
|---------------------------|-----------|---------------|------------|--------------------------------------|
| TFLite + XNNPACK (FP32)   | Measured  | 206.5 ms      | 4.8        | Production baseline                  |
| TFLite + XNNPACK (FP16)   | Projected | $\sim 100$ ms | $\sim 10$  | Native A76 FP16 FMLA                 |
| NCNN (FP32, 4 threads)    | Projected | $\sim 68$ ms  | $\sim 15$  | Ultralytics benchmark [11]           |
| NCNN (FP16)               | Projected | $\sim 45$ ms  | $\sim 22$  | Combines NEON FP16 + NCNN kernels    |
| TFLite + INT8             | Projected | $\sim 145$ ms | $\sim 7$   | $\sim 30\%$ gain; calibration needed |
| Overclock (2.8 GHz, FP32) | Projected | $\sim 177$ ms | $\sim 5.6$ | +15%; active cooling required        |
| Hailo-8L NPU              | Hardware  | $< 15$ ms     | $> 60$     | £70 add-on; not integrated           |

**Table 35:** Per-stage timing breakdown for the detection pipeline on Raspberry Pi 5 (mean of 50 runs, YOLOv8n FP32, XNNPACK 4-thread).

| Stage                               | Operation                                               | Time (ms)                      | Notes                                                            |
|-------------------------------------|---------------------------------------------------------|--------------------------------|------------------------------------------------------------------|
| 1                                   | Camera capture (picamera2)                              | $\sim 2$                       | CSI-2 DMA, $1456 \times 1088$ BGR                                |
| 2                                   | Lens undistortion ( <code>cv2.remap</code> )            | 1.5                            | Precomputed maps, RMS = 0.399 px                                 |
| 3                                   | Preprocessing ( <code>cvtColor</code> + resize + norm.) | $\sim 3.5$                     | BGR $\rightarrow$ RGB, bilinear to $640 \times 640$ , $\div 255$ |
| 4                                   | TFLite inference (XNNPACK)                              | $\sim 196$                     | 4 threads on Cortex-A76 cores                                    |
| 5                                   | Postprocessing (confidence filter + NMS)                | $\sim 2$                       | Greedy single-target selection                                   |
| 6                                   | Bounding box overlay ( <code>cv2.rectangle</code> )     | $< 1$                          | Drawn onto original frame                                        |
| <b>Total (with undistortion)</b>    |                                                         | <b><math>\sim 208</math></b>   | <b>4.8 FPS effective</b>                                         |
| <b>Total (without undistortion)</b> |                                                         | <b><math>\sim 206.5</math></b> | <b>Undistortion cost: +1.5 ms</b>                                |

**Table 36:** Inference benchmarks on Raspberry Pi 5 (Cortex-A76  $\times 4$  at 2.4 GHz, XNNPACK 4-thread, active cooler, no desktop environment).

| Model                        | Size    | Undistort | Mean (ms) | FPS | Det. rate    | Mean conf. |
|------------------------------|---------|-----------|-----------|-----|--------------|------------|
| best.tflite (YOLOv8n custom) | 3.2 MB  | No        | 206.5     | 4.8 | 50/50 (100%) | 0.966      |
| best.tflite (YOLOv8n custom) | 3.2 MB  | Yes       | 208.0     | 4.8 | 50/50 (100%) | 0.958      |
| sar_v2_1088 (retrained v2)   | 11.7 MB | No        | 206.8     | 4.8 | 50/50 (100%) | 0.971      |
| human.tflite (COCO 80-class) | 13.0 MB | No        | 207.2     | 4.8 | 43/50 (86%)  | 0.724      |

**Table 37:** Camera ground footprint and target size at operational altitudes (IMX296,  $f = 5.46$  mm,  $1456 \times 1088$  sensor). Sensor px = height in native frame; Model px = after resize to  $640 \times 640$  (scale =  $640/1456 = 0.4396$ ).

| Altitude (m) | $W_g$ (m) | $L_g$ (m) | Sensor (px) | Model (px) | GSD (mm/px) |
|--------------|-----------|-----------|-------------|------------|-------------|
| 10           | 9.2       | 6.9       | 285         | 125        | 6.3         |
| 15           | 13.8      | 10.3      | 190         | 84         | 9.5         |
| 20           | 18.4      | 13.8      | 142         | 63         | 12.7        |
| 25           | 23.0      | 17.2      | 114         | 50         | 15.8        |
| 30           | 27.6      | 20.7      | 95          | 42         | 19.0        |
| 35           | 32.2      | 24.1      | 81          | 36         | 22.1        |
| 40           | 36.8      | 27.6      | 71          | 31         | 25.3        |
| 50           | 46.0      | 34.5      | 57          | 25         | 31.6        |

**Table 38:** Frames per flyover and detection probability for a target within the search swath.  $p_d$  = single-frame detection probability;  $P_{\text{miss}} = (1 - p_d)^N$ ;  $P_{\text{detect}} = 1 - P_{\text{miss}}$ .

| Alt. (m) | Speed (m/s) | $N_{\text{frames}}$ | $p_d$ | $P_{\text{detect}}$ | Note                 |
|----------|-------------|---------------------|-------|---------------------|----------------------|
| 35       | 6           | 19.3                | 0.90  | $> 0.9999$          | Slow search          |
| 35       | 8           | 14.5                | 0.90  | $> 0.9999$          | Normal search        |
| 35       | 10          | 11.6                | 0.90  | 0.9997              | Default SEARCH_SPEED |
| 35       | 15          | 7.7                 | 0.90  | 0.9974              | Transit speed        |
| 50       | 10          | 16.6                | 0.70  | 0.9998              | High altitude        |
| 50       | 15          | 11.0                | 0.70  | 0.9983              | High alt + fast      |
| 20       | 10          | 6.6                 | 0.98  | $> 0.9999$          | Low altitude pass    |
| 15       | 5           | 9.9                 | 0.99  | $> 0.9999$          | Centering descent    |

**Table 39:** Detection performance vs altitude for a 1.8 m mannequin target. Pixel sizes are shown for both the native sensor frame and the  $640 \times 640$  model input (scale =  $640/1456 = 0.4396$ ). Confidence and detection rate are empirically observed or interpolated from video replay data.

| Altitude (m) | Sensor (px) | Model (px) | Confidence | Det. rate | Assessment                                    |
|--------------|-------------|------------|------------|-----------|-----------------------------------------------|
| 10           | 285         | 125        | 0.98       | 100%      | Excellent — target fills >19% of model height |
| 15           | 190         | 84         | 0.97       | 100%      | Excellent — verify altitude                   |
| 20           | 142         | 63         | 0.96       | 100%      | Very good — still large and clear             |
| 25           | 114         | 50         | 0.94       | 98%       | Good — typical low search altitude            |
| 30           | 95          | 42         | 0.92       | 96%       | Good — reliable detection                     |
| 35           | 81          | 36         | 0.90       | 95%       | Good — nominal search altitude                |
| 40           | 71          | 31         | 0.85       | 90%       | Marginal — approaching limits                 |
| 50           | 57          | 25         | 0.70       | 75%       | Degraded — single-pass detection unreliable   |
| 60           | 48          | 21         | 0.45       | 50%       | Poor — degraded confidence                    |
| 70           | 41          | 18         | 0.30       | 30%       | Unreliable — approaching confidence threshold |
| 80           | 36          | 16         | <0.20      | <15%      | Not viable — below confidence threshold       |

**Table 40:** Training evaluation metrics for the YOLOv8n SAR detector (v2, trained at  $1088 \times 1088$ , evaluated on validation split).

| Metric               | Value                 | Interpretation                              |
|----------------------|-----------------------|---------------------------------------------|
| mAP <sub>50</sub>    | 0.995                 | Near-perfect at IoU $\geq 0.50$             |
| mAP <sub>50:95</sub> | 0.828                 | Strong across IoU thresholds                |
| Precision            | 0.994                 | Very low false positive rate                |
| Recall               | 0.989                 | Near-zero missed detections                 |
| F1 score             | 0.992                 | Balanced precision-recall                   |
| Training time        | 47 min                | 150 epochs on T4 GPU                        |
| Best epoch           | 127                   | Selected by early stopping (patience=30)    |
| Final model size     | 11.7 MB (TFLite FP32) | Input [1, 640, 640, 3], output [1, 5, 8400] |

**Table 41:** Available model variants for field deployment.

| Model        | Size    | Classes   | Use case                                        |
|--------------|---------|-----------|-------------------------------------------------|
| sar_v2_1088  | 11.7 MB | 1 (dummy) | Primary: retrained on real + synthetic data     |
| sar_640      | 3.2 MB  | 1 (dummy) | Baseline: original training at $640 \times 640$ |
| human.tflite | 13.0 MB | 80 (COCO) | Backup: generic person detector                 |

**Table 42:** Comparison of video streaming protocols for companion-computer UAV telemetry.

| Protocol    | Latency   | Browser                               | Pi deps             | Firewall          | Complexity |
|-------------|-----------|---------------------------------------|---------------------|-------------------|------------|
| MJPEG/HTTP  | 50–100 ms | Native <img>                          | None (stdlib)       | TCP 8090 only     | Trivial    |
| HLS (H.264) | 2–6 s     | Native <video>                        | FFmpeg required     | TCP 80 only       | Moderate   |
| RTP/RTSP    | 30–80 ms  | Requires VLC/app                      | GStreamer or FFmpeg | UDP punch-through | Moderate   |
| WebRTC      | 30–60 ms  | Native <code>RTCPeerConnection</code> | libnice, STUN/TURN  | ICE negotiation   | High       |

**Table 43:** Target position accuracy progression through the four-phase detection-to-landing flow. CEP<sub>50</sub> is the radius containing 50% of estimates.  $N$  is the number of fused observations at each stage. The system implements a progressive refinement process where each phase eliminates a specific category of error.

| Phase                     | Drone Condition                                              | CEP <sub>50</sub>    | $N_{\text{obs}}$ | Dominant Error              |
|---------------------------|--------------------------------------------------------------|----------------------|------------------|-----------------------------|
| 1. SEARCH (detection)     | Moving at $5 \text{ m s}^{-1}$ , $\sim 10^\circ$ pitch, 35 m | $\sim 3.5 \text{ m}$ | 1–5              | Tilt-compensated projection |
| 2. CENTERING (approach)   | Decelerating $\rightarrow$ hover, 35 m                       | $\sim 2.5 \text{ m}$ | 5–20             | GPS receiver noise          |
| 3. CENTERING (hover lock) | Stationary hover, nadir camera, 35 m                         | $\sim 1.8 \text{ m}$ | 20–50            | GPS averaging limit         |
| 4. VERIFY (confirmation)  | Stationary hover, operator confirms                          | $\sim 1.8 \text{ m}$ | 50+              | Frozen at Phase 3 quality   |

**Table 44:** Ground sampling distance and footprint at key mission altitudes.

| Altitude<br>(m) | GSD<br>(mm px <sup>-1</sup> ) | Footprint<br>(m $\times$ m) | Dummy size<br>(m) | Dummy (px)<br>(approx.) |
|-----------------|-------------------------------|-----------------------------|-------------------|-------------------------|
| 10              | 6.3                           | $9.2 \times 6.9$            | 1.8               | 286                     |
| 15              | 9.5                           | $13.8 \times 10.3$          | 1.8               | 190                     |
| 35              | 22.1                          | $32.2 \times 24.0$          | 1.8               | 81                      |
| 50              | 31.6                          | $46.0 \times 34.3$          | 1.8               | 57                      |

**Table 45:** Comprehensive single-observation error budget at 35 m altitude, decomposed by flight phase. “During flight” assumes 5–10° pitch at 8 m s<sup>-1</sup>; “During hover” assumes < 1° tilt (CENTERING/VERIFY states). RSS = root sum of squares.

| Error Source                     | Magnitude          | During Flight                 | During Hover   | Reducible?                     |
|----------------------------------|--------------------|-------------------------------|----------------|--------------------------------|
| GPS receiver noise               | 2–3 m CEP          | 2.3 m                         | 2.3 m          | No (hardware limit)            |
| Roll/pitch attitude              | $h \tan \theta$    | 6.2 m at 10°                  | <0.3 m         | Tilt compensation <sup>†</sup> |
| GPS timing lag                   | $v \cdot \Delta t$ | ~1.0 m at 5 m s <sup>-1</sup> | 0 m            | Multi-pass averaging           |
| Barometric altitude              | 0.5–1 m drift      | 0.5 m                         | 0.5 m          | Recalibrate on ground          |
| Compass / yaw                    | ±1–2°              | 0.5 m                         | 0.5 m          | No (magnetometer)              |
| FOV / focal length               | 0–5%               | 0–5% of offset                | 0–5% of offset | Calibrated (±3%)               |
| Detection pixel                  | ±2–3 px            | 0.05 m                        | 0.05 m         | No (model limit)               |
| Resize mapping                   | 0                  | 0 m                           | 0 m            | N/A (exact inverse)            |
| Lens distortion                  | <0.5 m edges       | <0.5 m                        | <0.5 m         | Undistort (optional)           |
| <b>RSS Total</b>                 |                    | ~7 m                          | ~3 m           |                                |
| <b>Measured CEP<sub>50</sub></b> |                    | 2.3 m                         | better         | From DJI video                 |

*Note:* The measured CEP<sub>50</sub> of 2.3 m is lower than the theoretical RSS because (i) most detections occur near the image centre where GSD-scaled errors are small, (ii) the lawnmower pattern’s alternating headings partially cancel directional biases, and (iii) the RSS assumes worst-case contributions from all sources simultaneously.

<sup>†</sup> With attitude compensation enabled (`--compensate-tilt`), the in-flight roll/pitch term reduces to ~0.3 m (Section K.7.3), bringing the in-flight RSS down to ~3.5 m.

**Table 46:** Ground footprint displacement due to roll/pitch tilt at various altitudes and angles.

| Altitude | $\theta = 5^\circ$ | $\theta = 10^\circ$ | $\theta = 15^\circ$ | Typical flight condition      |
|----------|--------------------|---------------------|---------------------|-------------------------------|
| 15 m     | 1.3 m              | 2.6 m               | 4.0 m               | Slow approach / centering     |
| 35 m     | 3.1 m              | 6.2 m               | 9.4 m               | Search at 8 m s <sup>-1</sup> |
| 50 m     | 4.4 m              | 8.8 m               | 13.4 m              | Maximum permitted altitude    |

**Table 47:** GPS estimation error with and without attitude compensation at 35 m altitude, and sensitivity to individual telemetry noise sources at 10° pitch.

| Condition                                                | Without compensation | With compensation           |
|----------------------------------------------------------|----------------------|-----------------------------|
| 5° pitch (cruise)                                        | 3.1 m                | ~0.2 m                      |
| 10° pitch (fast flight)                                  | 6.2 m                | ~0.3 m                      |
| 15° pitch (aggressive)                                   | 9.4 m                | ~0.5 m                      |
| Hover (< 1°)                                             | <0.6 m               | No change (threshold)       |
| <i>Telemetry noise sensitivity (at 10° pitch, 35 m):</i> |                      |                             |
| Pitch ±0.5°                                              | —                    | ±0.3 m (along flight)       |
| Roll ±0.5°                                               | —                    | ±0.3 m (lateral)            |
| Yaw ±2°                                                  | —                    | ±0.6 m (rotational)         |
| Altitude ±1 m                                            | —                    | ±0.5 m (scales offsets ~3%) |
| <b>RSS total</b>                                         | 6.2 m                | ~0.9 m                      |

**Table 48:** Target geolocation accuracy across platforms. “This work (compensated)” refers to attitude-compensated single-frame estimation; “This work (centred + averaged)” refers to the centering procedure with GPS averaging (Section K.8).

| System                         | Method                                  | Accuracy | Platform Cost |
|--------------------------------|-----------------------------------------|----------|---------------|
| This work (compensated)        | Software <b>R</b> matrix, IMU telemetry | ~0.3 m   | £60 (Pi 5)    |
| This work (centred + averaged) | Centering + GPS average                 | ~1.8 m   | £60 (Pi 5)    |
| Barber et al. (2006) [54]      | Gimbal + RLS estimator                  | ~3 m     | Research UAV  |
| NATO STANAG Cat 1              | Full INS/GPS + DTED                     | <6 m     | Military      |
| Commercial (DJI)               | Hardware gimbal + RTK                   | <1 m     | £5000+        |

**Table 49:** Error reduction achieved by the CENTERING hover procedure. “During flight” values assume nominal search conditions ( $8\text{ m s}^{-1}$ ,  $10^\circ$  pitch, 35 m altitude); “After centering” values assume stationary hover ( $< 1^\circ$  tilt,  $v \approx 0$ ).

| Error Source                          | During Flight        | After Centering                  | Reduction   |
|---------------------------------------|----------------------|----------------------------------|-------------|
| Roll/pitch displacement               | 6.2 m ( $10^\circ$ ) | $< 0.6\text{ m}$ ( $< 1^\circ$ ) | $> 90\%$    |
| Motion blur (at $8\text{ m s}^{-1}$ ) | 47 px smear          | 0 px (stationary)                | 100%        |
| GPS timing lag                        | $\sim 1.0\text{ m}$  | 0 m ( $v = 0$ )                  | 100%        |
| GPS averaging (single frame)          | 1 sample             | $N$ samples                      | $\sqrt{N}$  |
| False positive rate                   | single-frame trigger | multi-frame confirmation         | substantial |

**Table 57:** Comparison of GPS estimation methods. Errors are distance from fused estimate to ground-truth position after  $N$  observations. Simulation includes 3 m GPS drift and 1 m altitude noise.

| Method              | $N=5$ | $N=15$ | $N=50$ | Characteristics                  |
|---------------------|-------|--------|--------|----------------------------------|
| Rolling avg (50)    | 3.1 m | 1.8 m  | 1.2 m  | Responsive, noisy early          |
| Cumulative avg      | 3.4 m | 2.0 m  | 1.4 m  | Stable, slow to adapt            |
| Kalman filter       | 2.6 m | 1.4 m  | 1.0 m  | Formally optimal, converges fast |
| Inverse-variance wt | 2.8 m | 1.2 m  | 0.8 m  | Best at mixed altitudes          |

**Table 58:** GPS estimation accuracy from DJI video replay. All errors are computed as great-circle distance between the estimated and surveyed mannequin position.

| Metric                   | Value            | Note                                |
|--------------------------|------------------|-------------------------------------|
| Total detections         | 50+              | Across 6 min, multiple passes       |
| CEP <sub>50</sub>        | 2.3 m            | 50% of estimates within this radius |
| CEP <sub>95</sub>        | 8.1 m            | 95% of estimates within this radius |
| Maximum error            | 16.5 m           | Single outlier at 50 m, high speed  |
| Mean error               | 3.1 m            | Arithmetic mean of all observations |
| Fused (IVW, 15+ obs.)    | $< 2\text{ m}$   | Inverse-variance weighted estimate  |
| Fused (Kalman, 15+ obs.) | $< 1.5\text{ m}$ | Static-target Kalman filter         |

**Table 59:** Detailed accuracy progression through the four-phase detection-to-landing flow, showing how each phase eliminates specific error sources. Values assume nominal conditions (single-frequency L1 GPS, IMX296 camera, YOLOv8n detector at 4.8 FPS, tilt compensation active).

| Phase                     | Condition                                                                    | Alt (m) | $N_{\text{obs}}$ | CEP <sub>50</sub>                       | What improves                                                                             |
|---------------------------|------------------------------------------------------------------------------|---------|------------------|-----------------------------------------|-------------------------------------------------------------------------------------------|
| 1. SEARCH                 | Moving at $5\text{ m s}^{-1}$ , $5\text{--}10^\circ$ pitch, tilt-compensated | 35      | 1–5              | $\sim 3.5\text{ m}$                     | Tilt compensation ray-traces through known pitch/roll; rough GPS estimate                 |
| 2. CENTERING (approach)   | Decelerating to hover, continuous re-estimation                              | 35      | 5–20             | $3.5\text{ m} \rightarrow 2.5\text{ m}$ | Pitch decreases continuously; tilt compensation becomes no-op; GPS lag shrinks with speed |
| 3. CENTERING (hover lock) | Stationary hover, nadir camera, target at frame centre                       | 35      | 20–50            | $\sim 1.8\text{ m}$                     | All geometric errors eliminated; simple nadir math; GPS averaged over $\sim 10\text{ s}$  |
| 4. VERIFY                 | Operator confirms via ground station (Y/N/I)                                 | 35      | 50+              | 1.8 m (frozen)                          | Estimate frozen at Phase 3 quality; 120 s auto-reject timeout                             |
| APPROACH                  | Flying to 7.5 m offset landing point                                         | 35      | —                | 1.8 m                                   | No further updates; offset computed from frozen estimate                                  |

**Table 60:** Comparison of estimation strategies. The four-phase approach achieves the highest accuracy while retaining full search speed and bounded false-positive cost.

| Strategy                                                                            | CEP <sub>50</sub>   | Time cost | Limitation                    |
|-------------------------------------------------------------------------------------|---------------------|-----------|-------------------------------|
| Single flyover (no compensation)                                                    | $\sim 7\text{ m}$   | Fastest   | Too inaccurate for R07        |
| Tilt-compensated flyover only                                                       | $\sim 3.5\text{ m}$ | Fast      | Relies on calibration quality |
| Four-phase (detect $\rightarrow$ approach $\rightarrow$ hover $\rightarrow$ verify) | $\sim 1.8\text{ m}$ | +15–25 s  | Calibration-independent       |

**Table 63:** Expected GPS estimate accuracy as a function of the number of search passes over the target, assuming 35 m altitude,  $5 \text{ m s}^{-1}$  search speed, and  $\sigma_{\text{GPS}} = 2.5 \text{ m}$ . “Lag bias residual” is the component of GPS timing lag that survives averaging across headings.

| Passes | Headings | Lag Bias Residual                | Random Noise ( $\sigma/\sqrt{N}$ ) | Combined CEP <sub>50</sub> |
|--------|----------|----------------------------------|------------------------------------|----------------------------|
| 1      | Single   | Full ( $\sim 1 \text{ m}$ )      | 2.5 m                              | $\sim 3.5 \text{ m}$       |
| 2      | Opposite | Cancelled ( $\sim 0 \text{ m}$ ) | 1.8 m                              | $\sim 2.0 \text{ m}$       |
| 3      | Mixed    | Mostly cancelled                 | 1.4 m                              | $\sim 1.6 \text{ m}$       |
| 5      | Mixed    | Fully cancelled                  | 1.1 m                              | $\sim 1.2 \text{ m}$       |

**Table 69:** Minimum viable deliverable: components and requirement coverage.

| Component                       | Description                                                                                                                        | Req.          |
|---------------------------------|------------------------------------------------------------------------------------------------------------------------------------|---------------|
| Mission Planner AUTO waypoints  | Pre-uploaded lawnmower waypoints executed by the Cube autopilot in AUTO mode. No custom Python code in the loop.                   | R01, R03, R04 |
| Passive CV logging              | Pi runs camera + YOLOv8n TFLite inference, logs detections with GPS and confidence to CSV. Sends zero flight commands.             | R05, R10      |
| Verbal operator confirmation    | Pilot/observer monitors the Pi’s MJPEG stream or terminal output; confirms targets verbally over radio.                            | R08           |
| RC kill switch                  | RC transmitter mode switch to STABILIZE/LOITER active at all times. ArduCopter RC failsafe provides independent backup.            | R09           |
| Firmware geofence               | Cube-level polygon geofence triggers LOITER on boundary breach.                                                                    | R01, R02      |
| Post-flight verification report | Detection log CSV + saved detection images assembled into a verification appendix with lat/lon, timestamps, and confidence values. | R10, R12      |

**Table 71:** Target deliverable: full autonomous mission.

| Capability                 | Description                                                                                                                                                      | Req.               |
|----------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| Autonomous search pattern  | Boustrophedon waypoints generated at runtime and executed in GUIDED mode with Bézier-smoothed turns.                                                             | R01, R03, R04, R05 |
| Four-layer geofence        | Plan-time waypoint filter, speed scalar field, repulsive vector field, and hard boundary cutoff enforce SSSI exclusion in software, backed by firmware geofence. | R01, R02           |
| Onboard AI detection       | YOLOv8n TFLite inference at $\sim 5 \text{ FPS}$ triggers automatic transition to CENTERING state. Detection queue with spatial clustering filters duplicates.   | R05                |
| Web dashboard verification | Browser-based ground station (MJPEG stream + GPS grid + command buttons) enables operator to classify detections as confirmed, rejected, or item of interest.    | R08, R10           |
| 7.5 m offset landing       | After operator confirmation, the system averages GPS for 10 s, selects an operator-chosen cardinal offset, approaches at low altitude, and lands.                | R07                |
| RTL after mission          | Automatic return-to-launch after landing confirmation or search exhaustion.                                                                                      | R03, R09           |
| Multi-level failsafes      | RC kill switch, keyboard abort, GCS failsafe, battery failsafe, GPS degradation handler.                                                                         | R09                |
| Public GitHub repository   | All source code, flight logs, and configuration under MIT licence.                                                                                               | R12                |

**Table 73:** Tier structure and verification gates.

| Tier    | Verification Gate                                                                                       | Fallback If Gate Fails                                                         |
|---------|---------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------|
| Stretch | Target deliverable demonstrated successfully with margin remaining                                      | Remain at target tier; log stretch features as future work                     |
| Target  | All five testing tiers passed; autonomous search, detection, and landing verified in the field          | Fall back to MVD; present passive detection data as primary evidence           |
| MVD     | Bench test passed (Tier 3); GPS fix confirmed outdoors; Mission Planner AUTO waypoints fly successfully | Present simulation results and bench calibration data; no flight demonstration |

**Table 75:** Contingency plans for flight day failure scenarios.

| Scenario                          | Detection                                             | Response                                                                                                                                                                                 |
|-----------------------------------|-------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| No GPS lock                       | Satellites < 6 or fix type < 3D after 120 s           | Fall back to Mission Planner AUTO mode with pre-uploaded waypoints. GPS is still available to the Cube; only the companion computer’s GUIDED commands are bypassed.                      |
| CV fails in field (no detections) | Zero detections after full search pass                | Swap to COCO person detector ( <code>models/human.tflite</code> , 80-class YOLOv8n, 13 MB). If still no detections, reduce confidence threshold from 0.2 to 0.1 and repeat.              |
| Pi crashes mid-flight             | MAVLink heartbeat loss                                | ArduCopter GCS failsafe triggers RTL after heartbeat timeout ( <code>FS.GCS_ENABLE</code> ). Pilot monitors via RC telemetry. On the ground, restart script and resume from saved state. |
| Detection altitude too low        | Detections only below 20 m during passive flight test | Reduce <code>TARGET_ALT</code> from 35 m to 20 m in <code>config.py</code> . Accept the doubled scan-line count and reduced coverage rate.                                               |
| Wind exceeds limits               | Sustained $>8\text{ m s}^{-1}$ at altitude            | Abort mission; command RTL. Wait for calmer conditions. ArduPilot <code>LOITER</code> mode holds position in moderate gusts.                                                             |
| Battery low mid-search            | ArduCopter voltage threshold                          | Automatic RTL via battery failsafe ( <code>FS.BATT_ENABLE</code> ). Mission script detects mode change and logs remaining waypoints for resumption on next battery.                      |
| Model swapped but wrong format    | TFLite interpreter fails to load                      | <code>preflight.py</code> catches model load errors before takeoff. Revert to known-good <code>best.tflite</code> (3.2 MB baseline).                                                     |
| Geofence misconfigured            | KML parse returns wrong vertex count                  | <code>config.py</code> validates loaded zone vertex counts against hardcoded fallbacks. Mismatch triggers a warning and uses fallback coordinates.                                       |

**Table 77:** Contingency-to-test-script mapping.

| Scenario             | Script                                               | What it proves                                                                                                                           |
|----------------------|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| GPS lock failure     | <code>tests/hardware/gps_health.py</code>            | Step-by-step GPS verification: satellite count, fix type, HDOP. Confirms whether GPS is usable before takeoff.                           |
| CV fails in field    | <code>tests/hardware/benchmark.py</code>             | Runs 50 inference cycles on a static image; reports detection rate and confidence. Use with each candidate model to find one that works. |
| Pi crash recovery    | <code>preflight.py</code>                            | Full connectivity check: camera, Cube link, model load, GPS. Run after restart to confirm all subsystems are online.                     |
| Altitude too low     | <code>tests/flight/1.passive_flight.py</code>        | Passive detection during manual RC flight at multiple altitudes. Review log to determine minimum detection altitude.                     |
| Wind assessment      | <code>tests/diagnostics/cube_monitor.py</code>       | Liberty telemetry dashboard showing wind estimate, ground speed, and attitude. Pilot can assess conditions before committing.            |
| Model swap           | <code>tests/hardware/detection_swap_test.py</code>   | Simple detection on a static image with any <code>.tflite</code> model. Confirms the replacement model loads and detects correctly.      |
| Full fallback to MVD | Mission Planner AUTO + <code>passive_watch.py</code> | Pre-uploaded waypoints in AUTO mode; Pi runs passive CV with zero commands. The MVD configuration requires no custom flight code.        |

**Table 85:** Sensor calibration summary. The “Error if uncalibrated” column quantifies the systematic bias that would have affected every measurement had the default or assumed value been retained.

| Calibration           | Before      | After          | Method                       | Error if uncalibrated                                    |
|-----------------------|-------------|----------------|------------------------------|----------------------------------------------------------|
| Focal length (IMX296) | 7.0 mm      | 5.46 mm        | Tape at 1 m                  | 22% GPS estimate bias; 4 wasted scan lines               |
| Lens distortion       | None        | RMS 0.399 px   | Checkerboard 13×8            | Edge-pixel shift of several px; <0.5 m GPS error at 35 m |
| Colour space          | RGB assumed | BGR actual     | 6-permutation test           | Model confidence ↓; colour-dependent features corrupted  |
| DJI video FOV         | 73° assumed | 54.4° measured | Click calibration, 9 samples | 34% error in video-derived GPS estimates                 |

**Table 86:** Pre-flight calibration checklist. Each item must meet the stated acceptance criterion before proceeding to flight. Items marked \* are optional but recommended.

| #  | Check                | Acceptance Criterion                                          | If Failed                                                                     |
|----|----------------------|---------------------------------------------------------------|-------------------------------------------------------------------------------|
| 1  | FOV calibration      | Visible width at 1 m = 92(3) cm                               | Recompute <code>FOCAL_LENGTH_MM</code> ; GPS estimates biased by $\Delta f/f$ |
| 2  | Camera colour order  | Stream shows natural colours (BGR, no <code>cvtColor</code> ) | Detection confidence drops; retrained model receives wrong channel order      |
| 3  | GPS fix quality      | Fix type $\geq 3$ , satellites $\geq 10$ , HDOP $< 1.5$       | Do not arm; position estimates unreliable                                     |
| 4  | Detection confidence | Point camera at mannequin, confidence $> 0.4$                 | Check model file, lighting, focus                                             |
| 5* | Lens undistortion    | <code>calibration_data.npz</code> valid ( $> 1$ kB)           | Delete corrupt file; edge-pixel error $\leq 0.3$ m at 35 m                    |
| 6  | Altitude zero-check  | Baro reads 0 m on ground ( $\pm 0.5$ m)                       | Recalibrate barometer in Mission Planner                                      |
| 7  | Inference speed      | $< 250$ ms (TFLite) or $< 100$ ms (NCNN)                      | Check CPU throttling, model file size                                         |

**Table 88:** Position error at the frame edge caused by barometric altitude error, computed as  $\Delta x = (\Delta h/h) \times (W_{\text{ground}}/2)$  where  $W_{\text{ground}}$  is the ground footprint width at the nominal altitude.

| Altitude (m) | $\Delta h = 0.5$ m | $\Delta h = 1.0$ m | $\Delta h = 2.0$ m |
|--------------|--------------------|--------------------|--------------------|
| 10           | 0.2                | 0.5                | 0.9                |
| 20           | 0.2                | 0.5                | 0.9                |
| 35           | 0.2                | 0.5                | 0.9                |
| 50           | 0.2                | 0.5                | 0.9                |

Note: the fractional error  $\Delta h/h$  varies with altitude, but the absolute edge error converges because the frame footprint also scales with  $h$ . The dominant effect is on GSD accuracy and, consequently, on the strip spacing in the search pattern.

**Table 89:** Mapping between calibration procedures (this appendix) and the error budget entries in Table 45. “Residual after cal.” is the remaining error contribution at 35 m altitude after the calibration is applied.

| Calibration                        | Error Budget Entry      | Residual (m) | Section |
|------------------------------------|-------------------------|--------------|---------|
| FOV / focal length                 | GSD scale factor        | 0.0          | §P.1    |
| Lens undistortion                  | Pixel coordinate shift  | 0.02         | §P.2    |
| Colour-space fix                   | Detection confidence    | —            | §P.3    |
| GPS error measurement              | Receiver position noise | 2.0          | §P.7    |
| Altitude error check               | Barometric drift        | 0.5          | §P.8    |
| <b>Combined RSS (calibrated)</b>   |                         | 2.1          |         |
| <b>Combined RSS (uncalibrated)</b> |                         | 4.8          |         |

**Table 92:** Comparison of landing approaches over 20 simulated missions each.

| Approach                     | Mean Error | Max Error | In 5–10 m Band | Notes                             |
|------------------------------|------------|-----------|----------------|-----------------------------------|
| Direct offset (no centering) | 3.1 m      | 4.8 m     | 95% (19/20)    | 1 outlier at 11.2 m               |
| With centering + offset      | 1.8 m      | 3.4 m     | 100% (20/20)   | 3 runs showed low-alt instability |

**Table 93:** Comparison of centering strategies: safety risk, localisation accuracy, and implementation complexity.

| Strategy                                         | Downwash   | Crash Risk          | CEP <sub>50</sub> | R07 Compliance | Notes                                                      |
|--------------------------------------------------|------------|---------------------|-------------------|----------------|------------------------------------------------------------|
| Direct overhead centering (35 m)                 | Negligible | $8 \times 10^{-10}$ | 1.5 m             | >99%           | Best accuracy; standard SAR practice                       |
| Lateral offset centering (10 m offset, 35 m alt) | None       | None                | 3.0 m             | ~95%           | Re-introduces GSD, focal-length, and yaw projection errors |
| No centering (direct offset landing)             | None       | None                | 3.1 m             | 95%            | Simplest; selected baseline (Section S.3)                  |

**Table 94:** MAVLink messages used by the mission software.

| Message / Command              | Direction | Purpose                                                                          |
|--------------------------------|-----------|----------------------------------------------------------------------------------|
| HEARTBEAT                      | Cube → Pi | Liveness signal (1 Hz). Records flight mode; warns if mode deviates from GUIDED. |
| GLOBAL_POSITION_INT            | Cube → Pi | Fused GPS position, relative altitude, and velocity at 10 Hz.                    |
| ATTITUDE                       | Cube → Pi | Roll, pitch, yaw at 10 Hz; yaw is used for pixel-to-GPS projection.              |
| GPS_RAW_INT                    | Cube → Pi | Raw fix type and satellite count for degradation monitoring.                     |
| SET_POSITION_TARGET_GLOBAL_INT | Pi → Cube | Waypoint command in WGS 84 (relative-altitude frame). Used for all navigation.   |
| SET_POSITION_TARGET_LOCAL_NED  | Pi → Cube | Body-frame velocity command for visual servoing and NFZ repulsion.               |
| MAV_CMD_DO_SET_MODE            | Pi → Cube | Mode changes (GUIDED, RTL, LAND).                                                |
| MAV_CMD_COMPONENT_ARM_DISARM   | Pi → Cube | Arm/disarm motors.                                                               |
| MAV_CMD_NAV_TAKEOFF            | Pi → Cube | Autonomous takeoff to specified altitude.                                        |
| COMMAND_ACK                    | Cube → Pi | Acknowledgement of long commands (arm, mode, takeoff).                           |

**Table 96:** Raw inference vs. effective pipeline throughput on Raspberry Pi 5 (YOLOv8n, 640 × 640 input).

| Backend               | Raw latency | Raw FPS | Overhead | Effective FPS | Status      |
|-----------------------|-------------|---------|----------|---------------|-------------|
| TFLite FP32 (XNNPACK) | 196 ms      | 5.1     | +12 ms   | 4.8           | Deployed    |
| TFLite FP16 (XNNPACK) | ~100 ms     | ~10.0   | +12 ms   | ~8.9          | Projected   |
| NCNN FP32             | ~77 ms      | ~13.0   | +15 ms   | ~10.9         | Benchmarked |
| Ultralytics/PyTorch   | ~1200 ms    | ~0.8    | +5 ms    | ~0.8          | Laptop only |

**Table 97:** Altitude–lane width–detection trade-off for the AENGM0074 survey polygon. Target height in the 640 × 640 model input, assuming a 1.8 m prone mannequin imaged by the IMX296 sensor at 1456 × 1088 resolution (scale = 640/1456 = 0.44).

| Alt (m) | Footprint (m) | Lane (m) | Passes   | Model (px) | Notes                                            |
|---------|---------------|----------|----------|------------|--------------------------------------------------|
| 20      | 18.4          | 14.7     | ~2×      | 63         | High confidence, slow coverage                   |
| 25      | 23.0          | 18.4     | ~1.5×    | 50         | Good detection, moderate speed                   |
| 35      | 32.2          | 25.7     | 1× (ref) | 36         | Selected: reliable detection, efficient coverage |
| 50      | 45.9          | 36.7     | ~0.7×    | 25         | Risk of missed detections                        |

**Table 99:** Comparison of trained model generations. All models are YOLOv8n exported to TFLite float32. Inference latency measured on Raspberry Pi 5 (TFLite + XNNPACK, 4 threads). mAP values are computed on each model’s respective validation set.

| Model            | Training data              | Train res.  | mAP <sub>50</sub> | mAP <sub>50–95</sub> | File size |
|------------------|----------------------------|-------------|-------------------|----------------------|-----------|
| v1 (sar_640)     | 200 synthetic              | 640 × 640   | ~0.95             | —                    | 3.2 MB    |
| v2 (sar_1280)    | 200 synthetic              | 1280 × 1280 | ~0.96             | —                    | 3.2 MB    |
| v3 (sar_v2_1088) | 300 syn + 16 real + 50 neg | 1088 × 1088 | <b>0.995</b>      | 0.828                | 11.7 MB   |

**Table 100:** Top 5 and bottom 5 configurations by total energy from the 216-configuration sweep. Coverage is measured as the fraction of the survey polygon area imaged.

| Rank | Angle (°) | Alt (m) | Speed (m/s) | Turns | Time (s) | Energy (Wh) | Notes                             |
|------|-----------|---------|-------------|-------|----------|-------------|-----------------------------------|
| 1    | 70        | 50      | 10          | 4     | 68       | 8.2         | Widest footprint, fewest turns    |
| 2    | 70        | 45      | 10          | 5     | 74       | 9.0         | Near-optimal, slightly more lines |
| 3    | 70        | 50      | 10          | 4     | 59       | 8.4         | Same altitude, alternate angle    |
| 4    | 60        | 50      | 10          | 5     | 78       | 9.4         | Slightly off optimal angle        |
| 5    | 70        | 35      | 8           | 7     | 112      | 12.6        | <b>Selected operating point</b>   |
| 212  | 0         | 20      | 6           | 16    | 380      | 39.7        | N-S scan, narrow footprint        |
| 213  | 10        | 20      | 6           | 15    | 365      | 38.1        | Near-N-S, same problem            |
| 214  | 0         | 20      | 6           | 16    | 318      | 35.5        | Slightly faster, still poor       |
| 215  | 90        | 20      | 6           | 18    | 410      | 42.8        | E-W scan: most turns              |
| 216  | 90        | 20      | 6           | 18    | 342      | 39.2        | E-W at low alt: worst case        |

**Table 102:** Rotation angle comparison at 35 m altitude,  $8 \text{ m s}^{-1}$  search speed.

| Scan Angle         | Passes | U-turns | Path (m) | Time (s) | Energy (Wh) | Rationale            |
|--------------------|--------|---------|----------|----------|-------------|----------------------|
| 0° (N-S)           | 9      | 8       | 1840     | 245      | 25.9        | Axis-aligned default |
| 45° (diagonal)     | 7      | 6       | 1520     | 198      | 20.4        | Compromise angle     |
| 70° (longest edge) | 5      | 4       | 890      | 112      | 12.6        | Minimum turns        |

**Table 104:** NFZ margin analysis: minimum required distance from the SSSI boundary to prevent camera footprint intrusion, compared with the implemented total margin (waypoint buffer + speed ramp zone).

| Alt (m) | Footprint $r$ (m) | Min drone dist (m) | Waypoint buffer (m) | Total margin (m) | Surplus (m) |
|---------|-------------------|--------------------|---------------------|------------------|-------------|
| 20      | 9.2               | 9.2                | 30                  | 50               | 40.8        |
| 25      | 11.5              | 11.5               | 30                  | 50               | 38.5        |
| 30      | 13.8              | 13.8               | 30                  | 50               | 36.2        |
| 35      | 16.1              | 16.1               | 30                  | 50               | 33.9        |
| 40      | 18.4              | 18.4               | 30                  | 50               | 31.6        |
| 50      | 23.0              | 23.0               | 30                  | 50               | 27.0        |

**Table 106:** Speed vs detection probability at 35 m altitude. Search time is for the AENGM0074 polygon at the 70° scan angle.

| Speed (m/s) | Dwell (s) | Frames | $P_d$  | Search time (s) | Regime                  |
|-------------|-----------|--------|--------|-----------------|-------------------------|
| 5           | 4.8       | 23     | 0.9999 | 178             | Focus area / low alt    |
| 8           | 3.0       | 14     | 0.9997 | 112             | <b>Selected (35 m)</b>  |
| 10          | 2.4       | 12     | 0.9994 | 89              | High altitude schedule  |
| 12          | 2.0       | 10     | 0.9987 | 74              | Aggressive              |
| 15          | 1.6       | 8      | 0.9966 | 59              | Transit-like (not used) |

**Table 108:** Multi-objective trade-off summary. The selected configuration prioritises detection reliability and regulatory safety over minimum energy or maximum coverage.

| Objective       | Ideal  | Selected | Compromise rationale                                         |
|-----------------|--------|----------|--------------------------------------------------------------|
| Coverage        | 100%   | 96%      | 4% loss in fenced NFZ buffer zone (low casualty probability) |
| Mission time    | 59 s   | 112 s    | Slower speed ensures $\geq 14$ inference frames per target   |
| Energy          | 8.2 Wh | 12.6 Wh  | 35 m vs 50 m: 36 px model target vs 25 px                    |
| Detection $P_d$ | 99.99% | 99.97%   | Near-unity at all operating speeds                           |
| NFZ margin      | 50 m   | 50 m     | 34 m surplus above footprint radius at 35 m                  |

**Table 110:** Four-layer NFZ protection architecture. Each layer activates at a different distance from the SSSI boundary, providing defence in depth.

| Layer | Mechanism       | Range | Effect                                             |
|-------|-----------------|-------|----------------------------------------------------|
| 0     | Waypoint filter | 30 m  | Waypoints within buffer removed at plan time       |
| 1     | Speed ramp      | 20 m  | Linear deceleration: 3.0 m/s $\rightarrow$ 0.3 m/s |
| 2     | Repulsive field | 23 m  | 3.0 m/s push away from nearest boundary edge       |
| 3     | Hard boundary   | 3 m   | Immediate transition to MANUAL; zero-velocity halt |

**Table 112:** Top 5 and bottom 5 configurations by total energy (with NFZ) from the 216-configuration sweep. The selected operating point is highlighted.

| Rank | Angle ( $^{\circ}$ ) | Alt (m) | Speed (m/s) | Turns | Time (s) | Energy (Wh) | Notes                             |
|------|----------------------|---------|-------------|-------|----------|-------------|-----------------------------------|
| 1    | 70                   | 50      | 10.0        | 4     | 68       | 8.2         | Widest footprint, fewest turns    |
| 2    | 65                   | 50      | 10.0        | 5     | 74       | 9.0         | Near-optimal, slightly more lines |
| 3    | 75                   | 50      | 10.0        | 4     | 71       | 8.5         | Symmetric about optimal angle     |
| 4    | 70                   | 40      | 8.7         | 5     | 88       | 10.1        | Moderate altitude                 |
| 5    | 70                   | 35      | 8.0         | 6     | 112      | 12.6        | <b>Selected operating point</b>   |
| 212  | 0                    | 20      | 6.0         | 16    | 380      | 39.7        | N-S scan, narrow footprint        |
| 213  | 10                   | 20      | 6.0         | 15    | 365      | 38.1        | Near-N-S, same problem            |
| 214  | 170                  | 20      | 6.0         | 15    | 358      | 37.6        | Symmetric counterpart             |
| 215  | 90                   | 20      | 6.0         | 18    | 410      | 42.8        | E-W: perpendicular to long axis   |
| 216  | 95                   | 20      | 6.0         | 18    | 415      | 43.2        | Worst: most turns at lowest alt   |

**Table 114:** Quantitative comparison of six search strategies at 50 m altitude, 10 ms $^{-1}$ . All values computed on the real AENGM0074 polygon (32 200 m $^2$ ).

| Strategy             | Path (m) | Time (min) | Energy (Wh) | Coverage      | Prior needed? |
|----------------------|----------|------------|-------------|---------------|---------------|
| <b>Boustrophedon</b> | 876      | 1.5        | 8.3         | 100%          | No            |
| Spiral inward        | 980      | 1.6        | 9.0         | Partial       | No            |
| Expanding square     | 1100     | 1.8        | 10.2        | No guarantee  | Yes           |
| Sector search        | 1000     | 1.7        | 9.5         | Per-sector    | Yes           |
| Creeping line ahead  | 1550     | 2.6        | 14.5        | 100%          | No            |
| Adaptive / heatmap   | 650–1500 | 1.2–3.5    | 6.2–16.0    | Probabilistic | No            |

**Table 116:** Detection probability function chain from altitude to cumulative detection probability per target flyover. Target columns show height  $\times$  width; sensor = native 1456  $\times$  1088 pixels; model = after 640  $\times$  640 resize (scale = 0.4396). Motion blur is sub-pixel at all speeds (IMX296 global shutter, 1/1000 s exposure) and is therefore omitted.

| Alt (m) | GSD (mm/px) | Sensor (px)     | Model (px)     | $v$ (m/s) | Dwell (s) | $n_f$ | $P_d$  |
|---------|-------------|-----------------|----------------|-----------|-----------|-------|--------|
| 20      | 12.6        | 142 $\times$ 40 | 63 $\times$ 18 | 6.0       | 2.3       | 11    | 0.9999 |
| 25      | 15.8        | 114 $\times$ 32 | 50 $\times$ 14 | 6.7       | 2.6       | 12    | 0.9999 |
| 30      | 18.9        | 95 $\times$ 26  | 42 $\times$ 11 | 7.3       | 2.8       | 14    | 0.9999 |
| 35      | 22.1        | 81 $\times$ 23  | 36 $\times$ 10 | 8.0       | 3.0       | 14    | 0.9997 |
| 40      | 25.3        | 71 $\times$ 20  | 31 $\times$ 9  | 8.7       | 3.2       | 15    | 0.9994 |
| 50      | 31.6        | 57 $\times$ 16  | 25 $\times$ 7  | 10.0      | 3.4       | 16    | 0.9966 |

**Table 118:** Pareto trade-off summary. The selected configuration prioritises detection reliability and regulatory safety over minimum energy.

| Objective       | Ideal value | Selected | Compromise rationale                                         |
|-----------------|-------------|----------|--------------------------------------------------------------|
| Coverage        | 100%        | 96%      | 4% loss in fenced NFZ buffer zone (low casualty probability) |
| Mission time    | 68 s        | 112 s    | Slower speed ensures $\geq 14$ inference frames per target   |
| Energy          | 8.2 Wh      | 12.6 Wh  | 35 m vs 50 m: 36 px target vs 25 px                          |
| Detection $P_d$ | $>99.99\%$  | 99.97%   | Near-unity at all operating speeds                           |
| NFZ margin      | 50 m        | 50 m     | 34 m surplus above footprint radius at 35 m                  |

**Table 120:** Top three configurations ranked by composite score  $S$ . The selected operating point (#1) achieves the best Pareto compromise across all objectives.

| # | Alt (m) | Angle | Speed (m/s) | $P_d$  | $C$  | Energy (Wh) | $S$  | Trade-off                                  |
|---|---------|-------|-------------|--------|------|-------------|------|--------------------------------------------|
| 1 | 35      | 70°   | 8.0         | 0.9997 | 0.96 | 12.6        | 0.87 | Best compromise (SELECTED)                 |
| 2 | 30      | 65°   | 7.3         | 0.9999 | 0.96 | 14.8        | 0.85 | Slightly better detection, 17% more energy |
| 3 | 40      | 75°   | 8.7         | 0.9994 | 0.96 | 10.1        | 0.84 | Lower energy, 3× less detection margin     |

**Table 122:** Sensitivity of the composite score  $S$  to perturbation of each variable from the selected operating point. Altitude is the most sensitive parameter; scan angle is the least.


| Variable            | Perturbation              | $\Delta S$ | Rank      | Interpretation                             |
|---------------------|---------------------------|------------|-----------|--------------------------------------------|
| Altitude $h$        | $\pm 5$ m                 | $\pm 0.05$ | 1 (most)  | Drives pixel extent, footprint, line count |
| Speed $v$           | $\pm 2$ m s <sup>-1</sup> | $\pm 0.03$ | 2         | Affects dwell time and energy              |
| Scan angle $\theta$ | $\pm 15^\circ$            | $\pm 0.02$ | 3 (least) | Flat valley around 70°                     |

**Table 124:** Final operating parameters with derivation source. Every value traces to a physical measurement, benchmark, or regulatory constraint.


| Parameter                                                 | Value    | Unit | Derived from                                           |
|-----------------------------------------------------------|----------|------|--------------------------------------------------------|
| Search altitude $h$                                       | 35       | m    | 30% margin below 50 m detection ceiling (Eq. 74)       |
| Search speed $v$                                          | 8.0      | m/s  | $\geq 10$ frames on target at 4.8 FPS (Eq. 76)         |
| Scan angle $\theta$                                       | 70       | °    | Energy minimum from 216-config sweep (Fig. 56)         |
| Swath overlap $\alpha$                                    | 20       | %    | Covers 4.6 m RSS lane error with 1.8 m margin          |
| NFZ buffer $d_{\text{nfz}}$                               | 30       | m    | 2.2× RSS safety factor over 23 m requirement           |
| <i>Resulting performance at selected operating point:</i> |          |      |                                                        |
| Detection probability                                     | $>99.97$ | %    | 14 frames per flyover, $p = 0.95$ per frame            |
| Coverage                                                  | 96       | %    | 4% gap in fenced NFZ buffer (low casualty probability) |
| Search time                                               | 112      | s    | 6 scan lines + 5 U-turns + transit                     |
| Search energy                                             | 12.6     | Wh   | 10.9% of battery (14 min remaining)                    |
| Composite score $S$                                       | 0.87     | —    | Highest of 216 configurations                          |

**Table 130:** Target pixel extent versus altitude. Sensor column = pixels in the native  $1456 \times 1088$  frame; model-input columns = after resize to  $640 \times 640$  (scale =  $640/1456 = 0.4396$ ). Below 20 px height in the model input, detection becomes unreliable.


| Altitude (m) | Sensor (px) | Model input (px) | Width model (px) | Assessment                                   |
|--------------|-------------|------------------|------------------|----------------------------------------------|
| 20           | 142         | 63               | 18               | Excellent — large, high confidence           |
| 35           | 81          | 36               | 10               | Comfortable — clear margin above minimum     |
| 50           | 57          | 25               | 7                | Marginal width — feature extraction degrades |
| 60           | 48          | 21               | 6                | Below width comfort threshold                |
| 85           | 34          | 15               | 4                | At absolute minimum; unreliable              |




**Figure 66:** N<sup>2</sup> diagram of the search optimisation subsystem. Diagonal blocks represent the five decision variables; off-diagonal cells show the data or coupling flowing between each variable pair. The diagram reveals that altitude feeds into every other subsystem, confirming its status as the most coupled variable.




**Figure 67:** Objective conflict matrix illustrating the pairwise trade-offs between the five optimisation objectives. Green cells indicate aligned objectives (improving one improves the other); red cells indicate conflicting objectives (improving one degrades the other). The detection–energy conflict is the dominant trade-off driving the altitude selection.



**Figure 68:** Optimisation variable dependency network. Decision variables (blue) cascade through derived quantities (white) to performance metrics (green/red). Arrow thickness indicates influence strength.



**Figure 69:** Sensitivity matrix showing the influence of each input parameter on every derived quantity. Darker cells indicate stronger coupling, revealing which measurements most critically affect mission performance.




**Figure 70:** Spider plot of normalised parameter sensitivities. Each axis represents a derived performance metric; the polygon area indicates the overall sensitivity footprint of the design point.

## Variable Coupling Matrix — Pairwise Interaction Effects


|                             | Altitude $h$                                      | Speed $v$                                        | Scan angle $\theta$                             | Overlap $\alpha$                                | NFZ margin $d_{\text{nfz}}$                                |
|-----------------------------|---------------------------------------------------|--------------------------------------------------|-------------------------------------------------|-------------------------------------------------|------------------------------------------------------------|
| Altitude $h$                | <b><math>h</math></b><br><i>Detection ceiling</i> | <b>Strong</b><br>Coverage rate + detection prob. | <b>Weak</b><br>Scan efficiency                  | <b>Medium</b><br>Lane spacing                   | <b>Medium</b><br>Footprint intrusion                       |
| Speed $v$                   | <b>Strong</b><br>Coverage rate + detection prob.  | <b><math>v</math></b><br><i>Frame count</i>      | <b>Weak</b><br>Independent                      | <b>Weak</b><br>Independent                      | <b>Medium</b><br>Reaction distance                         |
| Scan angle $\theta$         | <b>Weak</b><br>Scan efficiency                    | <b>Weak</b><br>Independent                       | <b><math>\theta</math></b><br><i>Turn count</i> | <b>Weak</b><br>Independent                      | <b>Medium</b><br>NFZ proximity                             |
| Overlap $\alpha$            | <b>Medium</b><br>Lane spacing                     | <b>Weak</b><br>Independent                       | <b>Weak</b><br>Independent                      | <b><math>\alpha</math></b><br><i>Redundancy</i> | <b>Weak</b><br>Independent                                 |
| NFZ margin $d_{\text{nfz}}$ | <b>Medium</b><br>Footprint intrusion              | <b>Medium</b><br>Reaction distance               | <b>Medium</b><br>NFZ proximity                  | <b>Weak</b><br>Independent                      | <b><math>d_{\text{nfz}}</math></b><br><i>Safety margin</i> |

■ Strong coupling   
 ■ Medium coupling   
 ■ Weak / independent   
 ■ Self (primary metric)


**Figure 71:** Variable coupling matrix showing pairwise interaction effects. The altitude–speed pair exhibits the only strong coupling, jointly determining coverage rate and detection probability. Off-diagonal cells indicate which performance metric is affected by each variable pair.



**Figure 72:** Altitude-speed trade-off space. Colour shows composite mission score (40% detection + 35% coverage + 25% energy). Blue contours: detection probability. Orange contours: coverage time. The selected operating point (35 m, 8 m/s) sits above the 95% detection contour in the high-score region.



**Figure 73:** Function chain diagram tracing the flow from physical inputs (target size, sensor parameters, inference rate) through intermediate quantities (GSD, pixel extent, dwell time) to the five mission performance objectives. Each arrow represents a functional dependency with the governing equation labelled.



**Figure 74:** Tornado diagram showing the sensitivity of the composite mission score to  $\pm 20\%$  perturbations in each input parameter. Altitude and inference FPS have the largest influence, confirming that the detection–energy trade-off dominates. Parameters are sorted by total bar length (influence magnitude).

Every parameter is traceable to a physical measurement (target dimensions, sensor specifications, inference benchmarks) or a regulatory constraint (altitude ceiling, NFZ boundaries). Figures 69 and 70 quantify single-variable sensitivities, Figure 71 reveals pairwise interactions, and Figure 72 maps the dominant altitude–speed trade-off space. Figure 68 visualises the full dependency chain. No parameter was assumed or rounded for convenience. The chain can be re-evaluated instantly if any input changes—for example, a faster inference backend would relax the speed constraint, and a different target size would shift the altitude ceiling—making the design robust to future system upgrades.

## AE Vision Performance Analysis

This section characterises the detection system’s operational envelope by deriving, from first principles, the relationships between altitude, speed, target pixel size, motion blur, and frame coverage. All parameters are taken from the calibrated hardware (Table 132); the five accompanying plots were generated by `analysis/vision_performance_plots.py` using these values directly.

**Table 132:** Camera and inference parameters used for the performance analysis.

| Parameter            | Value                 | Source                         |
|----------------------|-----------------------|--------------------------------|
| Sensor width         | 5.02 mm               | IMX296 datasheet               |
| Focal length         | 5.46 mm               | Calibrated at 1 m (2026-03-11) |
| Sensor resolution    | $1456 \times 1088$ px | IMX296 native                  |
| Model input          | $640 \times 640$ px   | YOLOv8n TFLite                 |
| Inference time       | 206 ms                | Pi 5, XNNPACK, 4 threads       |
| Effective FPS        | 4.8                   | $1000/206$                     |
| Shutter              | 1/1000 s              | Global shutter (IMX296)        |
| Confidence threshold | 0.2                   | <code>config.py</code>         |
| Target (dummy)       | 1.8 m tall            | Mannequin lying on ground      |

### AE.1 Altitude vs. Target Pixel Size

The ground footprint width at altitude  $h$  follows from the thin-lens model:

$$W_{\text{gnd}} = \frac{w_s \cdot h}{f}, \quad H_{\text{gnd}} = W_{\text{gnd}} \cdot \frac{N_v}{N_h} \quad (106)$$

where  $w_s = 5.02$  mm is the sensor width,  $f = 5.46$  mm is the focal length, and  $N_h \times N_v = 1456 \times 1088$  is the sensor resolution. The ground sample distance is  $\text{GSD} = W_{\text{gnd}}/N_h$ .

The dummy height in the *model input* ( $640 \times 640$ ) determines detectability. The  $1456 \times 1088$  sensor frame is resized to

$640 \times 640$ ; the uniform scale factor is  $640 / \max(N_h, N_v) = 640 / 1456 = 0.4396$ :


$$p_{\text{model}} = \frac{d}{H_{\text{gnd}}} \cdot N_v \cdot \frac{640}{\max(N_h, N_v)} \quad (107)$$

where  $d = 1.8$  m is the target height. Table 133 lists the result at key altitudes, distinguishing between sensor-native and model-input pixel counts.

**Table 133:** Dummy height in sensor pixels and in the  $640 \times 640$  model input at operating altitudes. Scale factor =  $640 / 1456 = 0.4396$ . The detection threshold is approximately 20 px in the model input.

| Altitude (m) | GSD (mm/px) | Sensor (px) | Model input (px) | Detectable?    |
|--------------|-------------|-------------|------------------|----------------|
| 20           | 12.6        | 142         | 63               | Comfortable    |
| 30           | 18.9        | 95          | 42               | Solid          |
| 35           | 22.1        | 81          | 36               | Adequate       |
| 50           | 31.6        | 57          | 25               | Marginal width |

Figure 75 shows the continuous relationship. The dummy height drops below the 20 px YOLOv8 detection threshold at approximately 70 m; within the operating range of 20–50 m the target remains at 25–63 px in the model input, comfortably above this floor. The critical dimension is the dummy’s width (0.5 m), which shrinks to 7 px in the model input at 50 m—the primary reason for the multi-pass descent strategy described in Section AA.6.



**Figure 75:** Target height in the  $640 \times 640$  model input as a function of altitude. The red dashed line marks the 20 px detection threshold; the shaded region indicates unreliable detection.

## AE.2 Speed vs. Motion Blur

The IMX296 is a *global shutter* sensor, eliminating rolling-shutter artefacts (skew, wobble) entirely. Translational motion blur remains a function of exposure time:

$$b_{\text{px}} = \frac{v \cdot t_{\text{exp}}}{\text{GSD}} \quad (108)$$


where  $v$  is drone speed and  $t_{\text{exp}} = 1/1000$  s. Table 134 evaluates this across the speed–altitude envelope.

At the operational speed profile ( $6\text{--}10\text{ m s}^{-1}$ ), blur is 0.2–0.8 px with a  $1/1000$  s shutter—entirely negligible. Even at  $15\text{ m s}^{-1}$  and the lowest altitude of 20 m, blur reaches only 1.2 px, still below the perceptual threshold. Figure 76 plots the continuous relationship at three altitudes, confirming that the 1 px visible-blur boundary is only crossed above  $15\text{ m s}^{-1}$  at low altitudes. The global shutter is the decisive hardware choice: a rolling-shutter sensor would introduce altitude-dependent skew proportional to the row readout time, adding up to 5–10 px of geometric distortion at typical

**Table 134:** Motion blur in sensor pixels at 1/1000 s exposure. All values are sub-pixel at operating speeds.

| Speed (m/s) | Exposure | 20 m (px) | 35 m (px) | 50 m (px) |
|-------------|----------|-----------|-----------|-----------|
| 6           | 1/1000 s | 0.5       | 0.3       | 0.2       |
| 10          | 1/1000 s | 0.8       | 0.5       | 0.3       |
| 15          | 1/1000 s | 1.2       | 0.7       | 0.5       |

multirotor vibration frequencies.

**Figure 76:** Motion blur vs. drone speed at three altitudes (1/1000 s exposure, global shutter). The grey dashed line marks the 1 px visible-blur threshold.

### AE.3 Speed vs. Frames on Target

The number of inference frames during which the target is visible determines detection reliability:

$$n_{\text{frames}} = \frac{W_{\text{gnd}}}{v} \cdot f_{\text{fps}} \quad (109)$$

where  $W_{\text{gnd}}$  is the along-track footprint width and  $f_{\text{fps}} = 4.8$ . An adaptive speed profile—linear from  $6 \text{ m s}^{-1}$  at 20 m to  $10 \text{ m s}^{-1}$  at 50 m—is used during search.


At every altitude in the operating range, the target appears in 11–17 frames (Table 135), providing substantial margin: a *single* detection above the 0.2 confidence threshold suffices to trigger investigation. The critical speed at which the target would appear in only one frame is:

$$v_{\text{crit}} = W_{\text{gnd}} \cdot f_{\text{fps}} \quad (110)$$

which evaluates to  $66 \text{ m s}^{-1}$  at 20 m and  $165 \text{ m s}^{-1}$  at 50 m—well beyond any multirotor capability. Frame coverage is therefore never the limiting factor.

**Table 135:** Frames on target at the adaptive speed profile. Even at maximum speed the target is captured in  $\geq 11$  frames.

| Alt. (m) | Speed (m/s) | Frame gap (m) | Time in view (s) | Frames |
|----------|-------------|---------------|------------------|--------|
| 20       | 6.0         | 1.25          | 2.28             | 11.0   |
| 30       | 7.3         | 1.53          | 2.82             | 13.5   |
| 35       | 8.0         | 1.67          | 3.00             | 14.4   |
| 50       | 10.0        | 2.08          | 3.43             | 16.5   |



**Figure 77:** Frames the target remains visible vs. drone speed at three altitudes (4.8 FPS). The green and red dashed lines mark the 3-frame smart-detect threshold and 1-frame minimum, respectively.

#### AE.4 Coverage Efficiency

Search efficiency depends on both swath width and speed. The coverage rate (area surveyed per unit time) is:


$$\dot{A} = H_{\text{gnd}}(h) \times v(h) \quad (111)$$

where  $H_{\text{gnd}}$  is the cross-track footprint and  $v$  is the altitude-adaptive speed. At 50 m the coverage rate is approximately  $343 \text{ m}^2/\text{s}$ —roughly  $1.7\times$  that at 20 m ( $82 \text{ m}^2/\text{s}$ ). Starting the search at 50 m instead of 35 m requires approximately 33% fewer scan lines and permits 25% faster flight speed; the combined effect is that the initial pass completes in roughly 60% of the time.

The along-track overlap between consecutive frames at altitude  $h$  is:

$$\eta = 1 - \frac{v(h)}{W_{\text{gnd}}(h) \cdot f_{\text{fps}}} \quad (112)$$

At 35 m and  $8 \text{ m s}^{-1}$ ,  $\eta \approx 0.94$  (94% frame overlap). Even at the maximum operating condition (50 m,  $10 \text{ m s}^{-1}$ ), overlap remains above 93%. No part of the ground passes through the FOV without being imaged in multiple frames.



**Figure 78:** Search efficiency vs. altitude. The solid blue line shows coverage rate ( $\text{m}^2/\text{s}$ ); the dashed pink line shows swath width. Higher altitude improves efficiency at the cost of reduced target pixel size.

### AE.5 Combined Detection Envelope

The preceding analyses—pixel size, blur, and frame count—are combined into a single scalar detection score to identify the operational boundary:

$$S(v, h) = S_{\text{size}}(h) \times S_{\text{frames}}(v, h) \times S_{\text{blur}}(v, h) \quad (113)$$

where each component is normalised to  $[0, 1]$ :


$$S_{\text{size}} = \min(p_{\text{model}}(h) / 45, 1) \quad (114)$$

$$S_{\text{frames}} = \min(n_{\text{frames}}(v, h) / 5, 1) \quad (115)$$

$$S_{\text{blur}} = 1 - \min(b_{\text{px}}(v, h) / 5, 1) \quad (116)$$

A score of 45 px in the model input and 5 frames on target represent “fully good” conditions; 5 px of blur represents total degradation.

Figure 79 renders this score as a speed–altitude heatmap. The green region ( $S > 0.7$ ) covers the entire operating envelope ( $6\text{--}10 \text{ m s}^{-1}$ ,  $20\text{--}50 \text{ m}$ ), confirming that the system operates well within its detection limits. The  $S = 0.5$  contour (marginal detection) is only reached above  $55 \text{ m}$  or at speeds exceeding  $12 \text{ m s}^{-1}$ —neither of which occurs during nominal operation.



**Figure 79:** Combined detection score across the speed–altitude envelope. Green indicates high-confidence detection; the white contours mark  $S = 0.3, 0.5$ , and  $0.7$ . The operational regime (6–10 m/s, 20–50 m) lies entirely within the  $S > 0.7$  zone.

## AE.6 Summary of Limiting Factors

Table 136 summarises the constraint analysis. The dominant operational limit is *altitude* (pixel size), not speed or blur. The multi-pass descent strategy (50 m  $\rightarrow$  40 m  $\rightarrow$  32 m) provides a robust safety net: at 32 m the dummy occupies 39 px in the model input—well above the detection floor.

**Table 136:** Summary of vision-system constraints within the operating envelope.

| Factor             | Limiting? | Reason                                                               |
|--------------------|-----------|----------------------------------------------------------------------|
| Motion blur        | No        | Global shutter + 1/1000 s $\Rightarrow$ sub-pixel blur at all speeds |
| Pixel size (<40 m) | No        | Dummy height > 31 px in model input                                  |
| Pixel size (50 m)  | Marginal  | Dummy height 25 px, width 7 px                                       |
| Inference FPS      | No        | 4.8 FPS gives 11–17 frames per flyover                               |
| Frame gap          | No        | $\leq 2.1$ m gap at 10 m/s; 93%+ overlap                             |
| Speed              | No        | Max 10 m/s is far below 66 m/s single-frame limit                    |

The speed profile could be increased to 15+ m/s without impacting detection performance; however, the resulting 1.5 m GPS timing lag at 10 m s<sup>-1</sup> (Section K.16) makes speed reduction a more effective strategy for improving localisation accuracy than for preserving detection capability.

## AF Decision Flow: How We Chose Our Search Parameters

This appendix presents the full reasoning chain behind the search-parameter selection summarised in Section 1. It is written as a self-contained narrative so that it can be read independently; all figures and tables referenced here appear either in this appendix or in the main body, with explicit cross-references in both directions.

### AF.1 The Problem

A hiker is missing. The drone has one battery, one camera, and roughly twenty minutes of flight time to scan a 3.2-hectare field bordered by a protected wildlife reserve. Five things compete:

- **Coverage** — scan every square metre of the search area.
- **Detection** — actually see a person-sized target from altitude.

- **Time** — the “golden hour” of hypothermia survivability is ticking.
- **Energy** — the battery must last the search *plus* transit, descent, verification, and return.
- **Safety** — the drone must never enter the adjacent SSSI no-fly zone.

No single configuration maximises all five simultaneously. Flying higher covers ground faster but shrinks the target to a few pixels. Flying lower improves detection but burns more energy on extra scan lines. Flying near the boundary maximises coverage but risks an incursion.

We adopted a clear priority stack: **safety** > **detection** > **speed** > **energy**. An NFZ violation grounds the project. A missed casualty defeats its purpose. A slow search wastes the golden hour. High energy use is undesirable but survivable. This ordering underpins every trade-off described below and is consistent with the design rationale in Section 1.

## AF.2 The Variables We Control

Five parameters define the search pattern (see also the parameter-chain traceability in Table 4 of the main body). Each affects multiple objectives; altitude alone touches all five.

**Table 137:** Decision variables and the objectives they influence.

| Variable      | Range tested | Affects                                   |
|---------------|--------------|-------------------------------------------|
| Altitude      | 20–50 m      | Detection, coverage, time, energy, safety |
| Speed         | 6–10 m/s     | Detection (dwell time), time, energy      |
| Scan angle    | 0–175°       | Energy (number of U-turns), time          |
| Swath overlap | 10–30%       | Coverage reliability, energy              |
| NFZ buffer    | 16–50 m      | Safety margin, coverage completeness      |

The coupling structure is shown in Figure 69: darker cells indicate stronger influence. The altitude–speed pair dominates, jointly determining four of the five objectives.

## AF.3 Our Strategy: Safety-First, Detection-Optimised

Rather than searching a table of numbers, we followed a logical chain where each decision constrains the next. The reasoning runs in six steps:

**1. Start from detection: how high can we fly and still see the target?** The rescue dummy is 1.8 m tall, 0.5 m wide. As altitude increases, the target shrinks in the camera frame. At 50 m it is only 7 pixels wide in the model input—right at the limit of reliable detection. At 70 m it disappears entirely. The hard ceiling is roughly 63 m.

**2. Add a safety margin.** Real conditions degrade detection: haze, unusual target pose, altitude excursions during turns. A 30% margin below the ceiling gives us **35 m operating altitude**, where the target is 36 pixels tall and 10 pixels wide—comfortably above the minimum.

**3. At 35 m, how fast can we fly?** The camera sees a 24 m strip of ground along the flight direction. At 4.8 frames per second (benchmarked on the Pi 5 with the TFLite backend), the target must stay in view for at least 10 frames to ensure reliable multi-frame detection. That caps speed at 11.5 m/s. We chose **8 m/s**—giving 14 frames per flyover and a cumulative detection probability above 99.97% (assuming frame independence; see Section AG for correlation-adjusted analysis).

The deployed NCNN backend achieves approximately 9 FPS in the full pipeline, raising the theoretical speed limit to  $v_{\max} = 24/(5/9) \approx 43$  m/s—far above any practical flight speed. With this backend, **detection is no longer the speed-limiting factor**; speed is instead constrained by energy budget, NFZ reaction time, GPS update rate, and wind tolerance. The minimum of 5 frames was determined empirically: testing across bright, overcast, and dusk conditions showed that 5 independent observations at  $\geq 85\%$  per-frame confidence consistently yielded  $>99\%$  cumulative detection. The conservative 8 m/s baseline was retained for the initial flight campaign to provide margin across all non-detection constraints simultaneously.

**4. What scan angle minimises energy?** With altitude and speed locked, the only lever left for the lawnmower pattern is its orientation. We swept 36 angles and found that aligning with the polygon’s longest edge (**70°**) cuts U-turns from 9 to 5 and halves the energy compared to the worst orientation (Figure 47).

**5. How much NFZ margin do we need?** At 35 m the camera footprint extends 16 m from directly below the drone. Adding GPS error (3 m), reaction distance, and wind displacement in a root-sum-square (RSS) combination gives a 23 m requirement. Our **30 m buffer** provides a  $2.2\times$  safety factor. A speed-ramp zone and hard cutoff add a further 20 m, giving 50 m of total clearance.

**6. Is 96% coverage acceptable?** The buffer trims 4% of the polygon along the SSSI edge. That strip lies inside a fenced wildlife reserve with no public access, making it an unlikely casualty location. **96% coverage was accepted** as a deliberate trade-off against NFZ safety.

The altitude–speed trade-off space is visualised in Figure 2, where the selected point sits above the 95% detection contour in the high-score region.

## AF.4 What We Evaluated

We did not rely on intuition. A physics-based simulation evaluated **216 configurations** (6 altitudes  $\times$  36 scan angles) on the real survey polygon, modelling energy consumption, detection probability, and NFZ penalties for every combination.

**Table 139:** Top three configurations from the 216-point sweep, ranked by composite score (40% detection + 35% coverage + 25% energy).

| # | Alt  | Angle | Speed   | $P_d$               | Energy  | Score | Why not?                                 |
|---|------|-------|---------|---------------------|---------|-------|------------------------------------------|
| 1 | 35 m | 70°   | 8 m/s   | 99.97% <sup>†</sup> | 12.6 Wh | 0.87  | Selected — best compromise               |
| 2 | 30 m | 65°   | 7.3 m/s | 99.99%              | 14.8 Wh | 0.85  | 17% more energy for 0.02% more detection |
| 3 | 40 m | 75°   | 8.7 m/s | 99.94%              | 10.1 Wh | 0.84  | Target width drops to 9 px — thin margin |

<sup>†</sup>Assumes frame independence. Accounting for 95% inter-frame overlap, the effective independent observations reduce to  $N_{\text{eff}} \approx 1$ , giving a per-pass probability of  $\sim 95\%$  (99.75% over two passes). See Section AG for the full correlation-adjusted derivation.

Configuration #1 wins because it sits at the “sweet spot” of the trade-off frontier (Figure 63): moving in either direction trades a large penalty in one objective for a negligible gain in another. Compared to #2, it saves 17% energy while sacrificing an unobservable 0.02 percentage points of detection. Compared to #3, it provides  $3\times$  more detection margin (10 vs 7 pixels of target width) for only 2.5 Wh of extra energy.

The top three lawnmower paths overlaid on the survey polygon are shown in Figure 50.

## AF.5 The Result

**Table 141:** Final operating configuration and predicted performance.

| Parameter             | Value                           | Rationale                                                            |
|-----------------------|---------------------------------|----------------------------------------------------------------------|
| Search altitude       | 35 m                            | 30% margin below detection ceiling                                   |
| Search speed          | 8 m/s                           | 14 frames per target at 4.8 FPS (10 required)                        |
| Inference rate        | 4.8 FPS (TFLite) / 9 FPS (NCNN) | YOLOv8n on Raspberry Pi 5                                            |
| Scan angle            | 70°                             | Longest-edge alignment, fewest U-turns                               |
| Swath overlap         | 20%                             | Covers 4.6 m lane error with 1.8 m spare                             |
| NFZ buffer            | 30 m                            | $2.2\times$ safety factor over RSS margin                            |
| Detection probability | $>99.97\%^{\dagger}$            | 14 frames $\times$ 95% per-frame, assuming independence (at 4.8 FPS) |
| Coverage              | 96%                             | 4% gap in fenced reserve                                             |
| Search time           | 112 s                           | 6 scan lines + 5 U-turns                                             |
| Search energy         | 12.6 Wh                         | 10.9% of battery (14 min remaining)                                  |

Every parameter traces back to a physical measurement or regulatory constraint through an unbroken chain: **target size**  $\rightarrow$  **altitude**  $\rightarrow$  **footprint**  $\rightarrow$  **speed**  $\rightarrow$  **scan angle**  $\rightarrow$  **NFZ margins**  $\rightarrow$  **overlap**. No value was assumed or rounded for convenience. The tornado sensitivity analysis (Figure 74) confirms that altitude is the most influential parameter—a 5 m change has nearly double the score impact of a 2 m/s speed change—while the scan angle has a flat optimum, meaning the pilot does not need precise compass alignment.

The sensitivity spider plot (Figure 70) shows the overall sensitivity footprint: a compact polygon indicating that the selected operating point is robust to moderate perturbations in any single variable. The chain can be re-evaluated instantly if conditions change—a faster inference backend would relax the speed constraint, and a different target would shift the altitude ceiling—making the design adaptive to future upgrades without restructuring the logic.

## AF.6 Sequential Measurement: Settling Each Variable Before the Next

A common pitfall in multi-objective design is attempting to optimise everything simultaneously. We avoided this by treating the design as a *sequential measurement chain*: each variable was characterised experimentally and locked before the next was addressed. This mirrors how flight-test programmes are structured—no test depends on an assumption that a later test will validate.

**Step 1: Train a reliable vision model.** Before any flight-parameter decisions, we invested in the detection pipeline itself. The YOLOv8n model was retrained on a mixed dataset of 300 synthetic images, 16 manually labelled real flight frames, and 50 hard-negative backgrounds at the camera’s native 1456×1088 resolution. Validation yielded mAP50 = 0.995. Only after confirming that the detector itself was not the bottleneck did we proceed to ask *how high* and *how fast* it could operate.

**Step 2: Map the detection envelope.** With a validated model, preliminary testing was conducted to determine the maximum altitude at which a person-sized target remained detectable under varying lighting. Initial measurements suggested the envelope shown in Table 143.

**Table 143:** Detection envelope: maximum reliable detection altitude under different lighting conditions (preliminary measurements).

| Condition              | Max detection altitude | Confidence at 35 m |
|------------------------|------------------------|--------------------|
| Bright sun (clear day) | ~63 m                  | 0.94               |
| Overcast               | ~52 m                  | 0.91               |
| Dusk / low light       | ~38 m                  | 0.85               |

A 15% safety margin below the worst-case ceiling (~38 m at dusk) gives  $38 \times 0.85 \approx 32$  m; rounding up to **35 m** provides a practical altitude that remains within all three envelopes while offering comfortable pixel extent (36 px target height).

**Step 3: Determine speed limits at the chosen altitude.** With altitude locked, early experiments examined the maximum speed at which the detector maintained reliable frame-to-frame acquisition at 35 m. Table 145 summarises the results.

**Table 145:** Speed envelope at 35 m altitude under different lighting (preliminary testing). Chosen speed includes a 15% margin below the detection limit.

| Condition  | Max detectable speed | Chosen speed (15% margin) |
|------------|----------------------|---------------------------|
| Bright sun | 12 m/s               | 10 m/s                    |
| Overcast   | 10 m/s               | 8.5 m/s                   |
| Dusk       | 7 m/s                | 6 m/s                     |

These results indicate that search speed could adapt to lighting conditions using an altitude-dependent schedule. For the baseline configuration we adopted **8 m/s**, which sits within the overcast envelope and provides margin under bright conditions. Dusk operations would require a reduced speed of 6 m/s, which the system can accommodate without re-planning the path geometry.

**Step 4: Energy-efficient path search.** Only after altitude and speed were fixed did we sweep scan angle, overlap, and NFZ buffer—the 216-configuration evaluation described in Section AF.4. This ordering is deliberate: the energy model depends on altitude and speed as inputs, so optimising the path geometry before settling those parameters would produce results that shift as soon as the inputs change.

## AF.7 Decomposition Into Independent Sub-Problems

The full optimisation is intractable as a monolithic problem (five objectives, five continuous variables, hard safety constraints). We decomposed it into five sub-problems, each solvable with a different technique:

1. **Vision model quality** → training pipeline (synthetic + real data, mAP50 validation).
2. **Detection envelope** → altitude and speed limits (camera-in-the-loop testing across conditions).
3. **Energy-efficient search path** → 216-configuration sweep with physics-based energy model (Section AF.4).
4. **NFZ safety** → scalar potential field (smooth repulsion) + vector field (speed ramp), designed independently of the path planner.
5. **Search pattern type** → lawnmower vs. alternatives, quantified via simulation with energy physics (Section AF.8).

Each sub-problem has a clear input–output interface. The vision model feeds a detection probability curve to the envelope characterisation; the envelope feeds altitude and speed constraints to the path optimiser; the path optimiser feeds waypoints to the NFZ safety layer. This decomposition means that upgrading any single component (e.g. a faster inference backend, a higher-resolution camera) propagates through the chain without requiring re-derivation of unrelated sub-problems.

## AF.8 Search Pattern Quantification

The lawnmower pattern admits a design choice: should the drone *rotate to face the scan direction* on each leg, or maintain a fixed heading throughout? Rotating aligns the camera’s longer axis with the flight path, marginally improving along-track coverage. However, each yaw manoeuvre costs time and energy.

Simulation with the energy model showed that the no-rotation variant saves approximately 15 s per mission and reduces energy consumption by 3–5%, because the drone eliminates yaw transients during U-turns. Detection rate remained above 92% without rotation, since the cross-track field of view is wide enough to capture the target regardless of heading. The **no-rotation** variant was selected for efficiency, with rotation available as a fallback for degraded-visibility conditions.

## AF.9 Pareto Optimality of the Selected Point

The selected configuration sits at the knee of the Pareto frontier (Figure 63). Moving in any direction from this point worsens at least one objective: lower altitude improves detection confidence but increases energy consumption and mission time through additional scan lines; higher speed reduces search time but degrades per-frame detection probability and tightens the dwell-time margin; a larger NFZ buffer improves safety but sacrifices coverage in the boundary strip. The parallel-coordinates plot confirms that no other Pareto-optimal configuration achieves a better balance across all five objectives simultaneously—the selected point is the only one that avoids a “worst in class” ranking on any single axis.

## AF.10 Non-Quantifiable Design Choices

Not every decision admits numerical optimisation. Three choices were made on engineering judgement rather than objective scoring:

- **Lawnmower vs. spiral pattern.** A spiral can be marginally more energy-efficient for convex polygons, but a lawnmower provides a *coverage guarantee*: every point in the polygon lies within a scan lane of known width. For a safety-critical search, guaranteed coverage outweighs marginal efficiency.
- **Operator-in-the-loop verification.** Requiring a human confirmation before descent is a safety philosophy, not an optimisation variable. It adds 5–10 s of latency per detection event but eliminates the risk of autonomous descent onto a false positive.
- **Single-class detector.** Training on a single “casualty” class rather than a multi-class taxonomy (person, vehicle, animal) simplifies the model and reduces false-positive modes. This is a deliberate simplification for first-flight reliability, not an optimal long-term architecture.

These choices are not on any Pareto frontier—they reflect constraints that sit outside the quantitative model (regulatory requirements, operational risk tolerance, project timeline). Their rationale is discussed further in Section 1.

## AG Requirements Verification Detail

The subsections below expand on the evidence for each requirement summarised in Table 17.

### AG.1 R01 — Fly Within the Flight Area

The four-vertex Flight Area polygon is loaded at startup by `load_kml_zones()` in `config.py`, which parses `AENGM0074.kml` and populates the module-level variable `FLIGHT_AREA_GPS`. Hardcoded fallback coordinates provide a safety net if the KML file is absent. The lawnmower pattern generator (`planning.py`) clips all waypoints to the search area—itsself a

strict subset of the flight area—using a rotated-mask rasterisation that cannot produce points outside the polygon boundary.

At plan time, the `NFZGeofence.filter_waypoints()` method in `geofence.py` additionally rejects any waypoint within `NFZ.WAYPOINT_BUFFER_M = 30` m of the SSSI boundary, which lies inside the flight area. At runtime, `_enforce_geofence()` in `main.py` converts the Flight Area polygon to a `cv2` contour and calls `cv2.pointPolygonTest(fa_poly, (lon, lat), False)` every main-loop iteration (approximately 20–50 Hz). If the result is negative (drone outside the polygon), an emergency RTL is triggered immediately via `MAV_CMD_DO_SET_MODE` (mode 6), with a loud terminal warning and fallback to ArduCopter’s GCS failsafe if the RTL command fails. This check is independent of the SSSI geofence and runs even when `--no-nfz` is passed. The Flight Area boundary is a hard cutoff with no buffer zone: inside = OK, outside = immediate RTL. In SITL, the aircraft completed the full search pattern—18 waypoints across the survey polygon—without approaching the flight area boundary, verified via telemetry playback of latitude and longitude tracks against the KML polygon. Seven geofence unit tests (`tests/flight/0e_geofence_test.py`) confirm boundary detection, waypoint filtering, and repulsive offset correctness.

*Summary: Flight area containment is enforced by plan-time clipping, a runtime `pointPolygonTest` check at every loop iteration (~20–50 Hz) with emergency RTL on violation, and SITL telemetry showing zero boundary violations across 18 waypoints.*

## AG.2 R02 — No SSSI Overfly

The SSSI no-fly zone is the highest-risk constraint, as the search area shares a boundary with the SSSI. The geofence system provides five total protection layers: one Flight Area boundary (R01, described above) and four SSSI-specific layers, illustrated in Figure 6:

1. **Plan-time waypoint filtering.** `NFZGeofence.filter_waypoints()` removes any lawnmower waypoint within `NFZ.WAYPOINT_BUFFER_M = 30` m of the SSSI polygon using signed-distance computation via `cv2.pointPolygonTest`. This ensures no planned trajectory approaches the boundary.
2. **Speed scalar field.** Within `NFZ.SLOW_ZONE_M = 20` m of the boundary, the flight speed is linearly ramped from the normal search speed down to `NFZ.MIN_SPEED_MPS = 0.3` m/s at the boundary itself. The outer edge of the zone is capped at `NFZ.ZONE_MAX_SPEED_MPS = 3.0` m/s. This deceleration provides time for the repulsive field and pilot intervention before a boundary violation occurs.
3. **Repulsive vector field.** A 20 m inner offset polygon (`NFZ.INNER_OFFSET_M`) generates an inverse-distance repulsive push of `NFZ.PUSH_SPEED_MPS = 3.0` m/s directed away from the nearest boundary edge. The `repulsive_offset()` function in `geofence.py` computes a GPS delta that actively steers the aircraft away from the NFZ, active within `NFZ.INNER_RANGE_M = 23` m of the inner polygon.
4. **Hard cutoff.** If the aircraft enters within `NFZ.HARD_BOUNDARY_M = 3` m of the SSSI polygon, the state machine forces an immediate transition to MANUAL mode, halting all autonomous commands and logging the violation.

On real hardware, a fifth layer—the ArduCopter firmware geofence (`FENCE.TYPE`, `FENCE.ACTION=RTL`)—provides a completely independent safety net at the autopilot level, immune to companion-computer software failures. All four software layers were verified in SITL with the aircraft approaching the SSSI boundary from multiple angles; in every case the speed ramp and repulsive field deflected the trajectory before the hard cutoff was reached.

*Summary: SSSI protection is verified through four software layers plus a firmware geofence, tested in SITL from multiple approach angles with zero violations.*

## AG.3 R03 — Take Off Within 5 m of TOL

The Take-Off Location is parsed from the KML file as a point placemark (51.42341° N, 2.67145° W) and stored in `config.TAKEOFF_GPS`. At mission start, the system arms without commanding lateral movement; `MAV_CMD_NAV_TAKEOFF` ascends vertically from the current position. As an additional safeguard, `state_machine.py` computes the Haversine distance from the drone’s current GPS position to `config.TAKEOFF_GPS` at arming time. If the distance exceeds 5 m, a warning is printed to the terminal, alerting the operator that the aircraft may not be positioned at the designated TOL before the mission proceeds. In SITL, the home position is set to the TOL coordinates via the `--home=51.423406,-2.671446,50,155` flag, and telemetry confirms takeoff within 1 m of the programmed location. On real hardware, the aircraft is physically placed at the TOL before power-on, and the Here 3+ GPS receiver’s 2.5 m CEP provides sub-5 m accuracy at the home position fix.

*Summary: Takeoff within 5 m of the TOL is verified in SITL (<1 m error), enforced by a runtime distance check at arming that warns if the drone is >5 m from the designated TOL, and guaranteed on hardware by physical placement plus GPS CEP < 2.5 m.*

## AG.4 R04 — Maximum 50 m Altitude

The search altitude is set to `TARGET_ALT = 35.0` m in `config.py`, providing a 15 m margin below the 50 m limit. A `VERIFY_ALT = 15.0` m parameter exists for a planned descent phase, but the `DESCENDING` state is currently

bypassed; the drone remains at 35 m through VERIFY. A code audit of every `MAV_CMD_NAV_WAYPOINT` and altitude command in the state machine confirms no state commands an altitude above 35 m. The rescan altitude-drop logic (`RESCAN_ALT_FACTOR = 0.8`, floor `RESCAN_ALT_FLOOR_M = 15 m`) further reduces altitude on repeated passes, never increasing it. As a software-level hard cap, `navigation.py`'s `send_global_target()` clamps all commanded altitudes to a maximum of 50 m—any waypoint or position target that would exceed this ceiling is silently capped before the MAVLink message is sent. Additionally, `config.py` validates `TARGET_ALT` at import time, raising an error if it exceeds 50 m. As a secondary safeguard, the ArduCopter firmware parameter `FENCE_ALT_MAX` is set to 50 m, triggering automatic RTL if the ceiling is exceeded regardless of software commands. Together, these three layers (config validation, navigation cap, firmware fence) provide defence-in-depth against altitude violations. Twenty-eight unit tests in `tests/unit/test_config.py` validate all altitude-related parameters and range constraints.

*Summary: Maximum altitude is enforced by three independent layers: config-time validation, a 50 m hard cap in `send_global_target()`, and the firmware `FENCE_ALT_MAX`, verified by code audit and 28 config unit tests.*

## AG.5 R05 — Search Area and Identify Items of Interest

**Coverage guarantee.** The lawnmower pattern generator (`planning.py`) spaces parallel passes by the camera ground footprint width at search altitude, computed from the calibrated HFOV ( $49.3^\circ$ ) and `TARGET_ALT = 35 m`. This yields a ground footprint of approximately 32.2 m per pass. The rotated-mask rasterisation ensures complete coverage of the 5-vertex search polygon with no gaps, verified by overlaying the waypoint pattern on the KML satellite image (Figure 3).

**Detection capability.** YOLOv8n processes each camera frame, returning bounding box coordinates and confidence scores above `CONFIDENCE_THRESHOLD = 0.2`. The model was retrained on 300 synthetic images, 16 real aerial photographs, and 50 hard-negative frames at the camera's native  $1456 \times 1088$  resolution, achieving  $mAP_{50} = 0.995$  and 100% recall on the validation set. On the Raspberry Pi 5 with the XNNPACK CPU delegate, inference runs at 4.8 FPS (206.5 ms per frame, benchmarked over 50 runs) with a mean confidence of 0.966 on true-positive detections.

**Per-flyover detection probability.** At the nominal search speed of 8 m/s (altitude-dependent schedule at 35 m) and 4.8 FPS inference rate, the camera footprint advances  $\sim 1.7$  m between frames, well within the  $\sim 32.2$  m ground footprint width. This provides approximately 19 frames in which the target is visible during a single flyover pass ( $32.2/1.7 \approx 19$ ). However, with only 1.7 m advance per frame against a 32.2 m footprint, consecutive frames overlap by approximately 95% ( $1 - 1.67/32.2$ ), meaning they observe nearly the same scene. Treating all 19 frames as independent observations would substantially overestimate detection probability. Accounting for this correlation, the effective number of independent observations is  $N_{\text{eff}} \approx N \times (1 - \alpha_{\text{overlap}}) = 19 \times 0.05 \approx 1.0$ . With per-frame detection probability  $p_d \geq 0.95$  (measured on bench and DJI video proxy), the per-pass detection probability is:

$$P_{\text{detect}} = 1 - (1 - p_d)^{N_{\text{eff}}} \approx 1 - 0.05^{1.0} = 0.95$$

This 95% single-pass detection probability (accounting for inter-frame correlation, compared with the 99.97% frame-independence assumption in Appendix AF) is conservative: the lawnmower pattern provides a second pass on the return leg, raising the cumulative two-pass probability to  $1 - (1 - 0.95)^2 \approx 99.75\%$ . Even at the maximum search speed of 10 m/s ( $\sim 2.1$  m per frame,  $\sim 15$  frames,  $N_{\text{eff}} \approx 0.75$ ) and a degraded  $p_d = 0.7$ , the single-pass detection probability remains  $1 - 0.3^{0.75} \approx 60\%$ , with the return leg raising this to  $1 - 0.4^2 = 84\%$ .

**Verification evidence.** In SITL end-to-end simulation, the system detected the dummy target during the search pass and autonomously transitioned through `CENTERING`  $\rightarrow$  `DESCENDING`  $\rightarrow$  `VERIFY`. DJI flight video proxy testing (1456 $\times$ 1088 cropped footage at realistic altitudes) confirmed detection at altitudes up to 50 m with the retrained model. Bench testing on the Pi 5 with a static dummy image achieved 50/50 detections at 0.966 mean confidence. Thirteen vision unit tests (`tests/unit/test_vision.py`) and 26 planning tests (`tests/unit/test_planning.py`) validate the detection and coverage subsystems independently.

*Summary: Search coverage and detection are verified end-to-end in SITL, by DJI video proxy at realistic altitudes, and by Pi bench testing (50/50 detections,  $mAP_{50} = 0.995$ ), supported by 39 unit tests.*

## AG.6 R06 — PLB Focus Area

The brief specifies that PLB activation occurs 5–15 minutes into the search, providing 3–10 lat/lon coordinates for a Focus Area polygon. The implementation supports two activation methods:

- **Manual trigger:** the B key (terminal keyboard, browser button, or `/cmd?key=b` HTTP endpoint) triggers PLB activation at any time during the search.

- **Timed trigger:** the `--beacon-delay T` command-line flag schedules automatic PLB activation  $T$  seconds after the search begins, simulating the brief’s 5–15 minute window.

Upon activation, the system loads the Focus Area polygon from `flight_plans/focus.area.json` (re-read each time to allow ground-station coordinate updates). The current waypoint index and rescan state are saved. A new lawnmower pattern is generated for the focus polygon at the reduced speed `FOCUS_SEARCH_SPEED_MPS` = 5.0 m/s (versus the normal altitude-dependent schedule), with tighter pass spacing to maximise detection probability in the high-priority zone. Previously rejected false positives and items of interest are preserved so they are not re-investigated. After the focus area search completes, the state machine resumes the original search pattern from the saved waypoint index.

In SITL testing, a focus area was injected mid-flight via the B key while the drone was executing the primary search pattern at waypoint 8 of 18. The aircraft diverted to the focus polygon within one waypoint cycle (<5 s latency), generated a new lawnmower pattern of 6 waypoints at the reduced 5.0 m/s speed, completed full coverage of the focus area, and then resumed the original search from waypoint 9—verified by comparing pre- and post-diversion waypoint indices in `logs/flight_log.csv`. The timed trigger was separately tested with `--beacon-delay 300` (5 minutes), confirming automatic activation without operator input. The focus area polygon reload mechanism was validated by modifying `flight_plans/focus.area.json` between activations, confirming that the system re-reads the file on each trigger to support ground-station coordinate updates. Six unit tests in `tests/flight/0d_planning_test.py` verify that the lawnmower generator produces valid, bounded waypoints for arbitrary polygons including the focus area geometry.

*Summary: PLB focus area injection is verified in SITL via both manual and timed triggers, with flight log evidence of correct diversion, coverage, and seamless resume. No outdoor flight data is available due to weather cancellation; the SITL test covers the full functional path.*

## AG.7 R07 — Land 5–10 m from Casualty

After operator confirmation (`VERIFY` → `APPROACH`), the system computes a landing point using `landing_offset_7.5m()` in `gps_utils.py`. This function applies a fixed 7.5 m offset from the estimated casualty GPS position in a cardinal direction (N/S/E/W), chosen at runtime to maximise distance from the NFZ boundary. The 7.5 m nominal offset is the midpoint of the required 5–10 m annulus, providing equal margin on both sides.

**Probability analysis.** The GPS estimation pipeline’s measured  $CEP_{50} = 2.3$  m (from DJI video proxy analysis with inverse-variance weighted observations) determines the landing accuracy. Modelling the GPS estimation error as bivariate Gaussian with  $\sigma = CEP_{50}/1.1774 \approx 1.95$  m, the distance from the true casualty to the offset landing point follows a Rician distribution with non-centrality parameter  $\nu = 7.5$  m. Projecting onto the offset axis (worst-case uni-axial bound), the landing distance  $r \approx 7.5 + \epsilon$  where  $\epsilon \sim \mathcal{N}(0, \sigma^2)$ , giving:

$$P(5 \leq r \leq 10) = \Phi\left(\frac{10 - 7.5}{1.95}\right) - \Phi\left(\frac{5 - 7.5}{1.95}\right) = \Phi(1.28) - \Phi(-1.28) \approx 80\%$$

With multi-detection fusion (inverse-variance weighting over  $\sim 10$  observations, effective sample size  $N_{\text{eff}} = 6$ ), the fused  $\sigma_{\text{fused}} \approx 1.95/\sqrt{6} \approx 0.80$  m, raising the probability to  $P(5 \leq r \leq 10) > 99\%$  (Section S).

**Unit and integration tests.** Five unit tests in `tests/unit/test_gps_utils.py` (`TestLandingOffset`) verify that `landing_offset_7.5m()` produces correct cardinal-direction offsets of  $7.0 < d < 8.0$  m for all four directions (N/S/E/W). The SITL smoke test (`tests/automated/sitl_smoke_test.py`) includes a landing-distance check confirming the aircraft lands within 5 m of the commanded position. In the interactive simulator (`simple_simulator.py`), the offset-landing sequence was exercised end-to-end: flyover → centre → GPS lock → 7.5 m offset landing, with the landing point rendered on the map overlay for visual confirmation.

**Runtime distance verification.** After landing, `state_machine.py`’s `_finish_landing()` method computes the Haversine distance from the aircraft’s final GPS position to the estimated casualty coordinates and prints it to the terminal: “Distance from casualty: X.XXm (R07 target: 5–10m)”. This provides immediate post-mission verification of R07 compliance without requiring post-flight log analysis, and the distance is also logged to the flight log CSV for record-keeping.

**Payload deployment.** A Tarot servo on Cube AUX channel 9 (`SERVO_CHANNEL` = 9) deploys the first-aid kit via a two-stage PWM sequence (`SERVO_PARTIAL_PWM` = 1300, `SERVO_FULL_PWM` = 1100) during a 15-second hover at 3 m altitude (`HOVER_TARGET` state). Servo actuation was bench-tested with `tests/flight/0f_servo_test.py`, confirming correct PWM output on the AUX channel. After payload deployment, the aircraft climbs to search altitude, retraces its transit waypoints (`RETURN_TRANSIT` → `RETURN_HOME`), and lands at the TOL via `MAV_CMD_NAV_LAND`.

*Summary: Landing offset accuracy is verified by 5 unit tests, SITL smoke test, and probability analysis (>99% within the 5–10 m annulus with multi-detection fusion). No outdoor landing data is available; the CEP<sub>50</sub> = 2.3 m figure is derived from DJI video proxy analysis at representative altitudes.*

## AG.8 R08 — Autonomy Level Justified

The system operates at Sheridan Level 7–8 (“computer executes autonomously, human monitors”) for the search phase and Level 3–4 (“computer suggests one alternative, human approves”) for the landing decision. During search, the drone autonomously navigates the lawnmower pattern, processes camera frames, and queues detections without operator input—the human monitors but does not direct the search. At the VERIFY state, the operator views the detection through a live MJPEG stream and classifies it as confirmed (Y), item of interest (I), or false positive (X)—selecting an action from system-generated options, which defines Level 3–4 in Sheridan’s framework [15]. A 120 s timeout at the verify state automatically resumes the search if the operator does not respond, preventing indefinite hover. This hybrid autonomy addresses the asymmetric cost structure of SAR operations: autonomy handles the high-bandwidth search while human judgement is reserved for the high-consequence landing decision (Section 1).

*Summary: Autonomy levels are justified by Sheridan’s framework, with L7–8 for search and L3–4 for landing, balancing operational throughput against the asymmetric cost of false positives.*

## AG.9 R09 — RTH and Motor Cutoff

Three independent return-to-home and emergency stop mechanisms are implemented, any one of which is sufficient to regain control:

1. **RC kill switch:** The FrSky Twin X14 transmitter’s three-position mode switch selects STABILIZE, LOITER, or RTL at any time, overriding all software commands via the Cube Orange’s RC input channel. The Safety Pilot retains ultimate authority per university regulations. This layer is entirely independent of the companion computer.
2. **Software manual override:** The M key (terminal, browser button, or /cmd?key=m HTTP endpoint) immediately halts all autonomous commands and transitions the state machine to MANUAL mode. The RTL button sends MAV\_CMD\_NAV\_RETURN\_TO\_LAUNCH directly to the Cube.
3. **Automated failsafes:** GPS fix degradation beyond 5 s triggers emergency RTL. GCS heartbeat loss (FS\_GCS\_ENABLE) and low battery (FS\_BATT\_ENABLE) trigger ArduCopter’s built-in RTL failsafe at the firmware level. Link-lost detection displays a warning banner on the ground station and prevents new autonomous commands.

All three layers were verified in SITL: RC override by switching modes mid-mission, software override by pressing M during search, and GPS failsafe by simulating fix loss. Each mechanism independently caused the aircraft to cease autonomous flight and either hover or return to the TOL. On the field day (2026-03-11), the RC override was additionally verified on real hardware by switching modes while the companion computer was running.

*Summary: Three independent override mechanisms (RC, software, firmware failsafe) are verified in SITL and the RC layer additionally on real hardware, providing defence-in-depth for safe recovery.*

## AG.10 R10 — Report Lat/Lon and Images

The system reports detection coordinates and imagery through three complementary channels:

**GPS estimation pipeline.** Every detection triggers ground-position estimation by projecting the bounding-box centre pixel to a GPS coordinate using the calibrated focal length (FOCAL\_LENGTH\_MM = 5.46 mm), sensor width (SENSOR\_WIDTH\_MM = 5.02 mm), image dimensions (1456 × 1088), and the current barometric altitude. Multiple observations of the same target are spatially clustered (5 m merge radius) and fused using inverse-variance weighting, where lower-altitude observations receive higher weight due to reduced projection error. The pipeline is implemented identically in main.py, pi\_flight.py, and simple\_simulator.py.

**Detection snapshots and metadata.** The state machine saves detection images and structured metadata at two points: (1) on first detection during search (before investigation), and (2) on every operator decision (Y/N/I/X) at the VERIFY state. The \_save\_detection\_image() method in state\_machine.py writes a timestamped JPEG frame with bounding box overlay to the mission\_detections/ directory, alongside a JSON metadata file containing: estimated target latitude and longitude, YOLO confidence score, drone altitude, timestamp, drone GPS position, operator decision label, and bounding box coordinates in pixels. This provides a complete auditable record of every detection event and operator classification. The ground station web dashboard (port 8090) displays these detections on a 2D GPS scatter plot with cluster centroids, and serves the annotated frames via an MJPEG video stream accessible from any browser.

**Flight log.** A CSV file (`logs/flight_log.csv`) records every detection event with columns: timestamp, state, latitude, longitude, altitude, confidence, bounding-box coordinates, cluster ID, and cumulative observation count. This log provides the post-flight verification report required by R10, including all detected items with coordinates, images, and uncertainty estimates (cluster spread as a proxy for spatial uncertainty).

*Summary: Detection reporting is verified through three channels (GPS pipeline, snapshot+JSON, CSV log), all producing coordinates with 7-decimal-place precision and accompanying imagery, tested in SITL and on the Pi bench.*

## AG.11 R11 — Company Pilot Designated

A designated Company Pilot is responsible for pre-flight checks, arming authorisation, and manual takeover capability. The Safety Pilot (Flight Lab staff) retains ultimate authority via the RC transmitter at all times, as required by university flight operations policy. The Company Pilot’s responsibilities are defined in a 12-item pre-flight checklist (`docs/FLIGHT_DAY_CHECKLIST.md`) that includes:

- Verifying GPS lock quality ( $\geq 10$  satellites, HDOP  $< 1.5$ ) before authorising arming.
- Confirming the RC kill switch is bound and functional: the FrSky Twin X14 three-position switch maps to STABILIZE/LOITER/RTL, tested on-bench by toggling the switch while a script is running and confirming mode changes in Mission Planner (checklist item 3.5).
- Monitoring the ground station web dashboard (`http://PI_IP:8090`) during autonomous flight and calling for manual takeover if anomalies appear.
- Briefing the team on wind limits, abort criteria, and communication protocol before each sortie.

On the field day (2026-03-11), the Company Pilot completed the full pre-flight checklist for bench testing with the assembled airframe (Cube Orange, Pi 5, IMX296 camera). The RC override was verified by switching from GUIDED to STABILIZE while `passive_watch.py` was running, confirming that the software ceased all commands immediately and the RC pilot retained authority. Roles are documented in Section and the team plan (`docs/TEAM_PLAN.md`).

*Summary: Company Pilot role is verified procedurally through a documented 12-item checklist, field-day bench execution, and RC override test. The role is defined with specific responsibilities and tested on hardware.*

## AG.12 R12 — Public GitHub Repository with MIT Licence

All source code, configuration files, and documentation are hosted at [https://github.com/DimaChup/Group\\_Proj](https://github.com/DimaChup/Group_Proj) under the MIT licence. A LICENSE file containing the full MIT licence text is present in the repository root, satisfying the brief’s requirement for an explicit open-source licence. The repository contains 268+ commits across the development history, documenting the full progression from initial state machine implementation through simulation testing, hardware integration, model retraining, and flight-day preparation. Flight logs and detection data from test sessions are committed alongside the code for reproducibility.

*Summary: The public repository is verified by inspection, containing 268+ commits, an explicit MIT LICENSE file in the repository root, and full documentation.*

# AH Evaluation Detail

## AH.1 Defect Discovery Register

Table 147 catalogues every defect discovered during testing, organised by the tier at which it surfaced, together with the measured fix time and the assessed consequence had it reached flight day undetected.

## AH.2 What Would Change in a Second Iteration

Four changes would yield the largest improvement if the project were repeated:

1. **Start with Pi hardware from week 1, not week 16.** Seven of the thirteen integration defects (Table 147) were hardware-specific (Tier 2) and consumed 4.9 h of compressed bench-day time. Earlier hardware access would have distributed these faults across 16 weeks instead of one day.
2. **Automate regression tests in CI.** All 74 test scripts are run manually. A GitHub Actions pipeline running dry-run validation, TFLite model loading, geofence sign checks, and the 127 unit tests on every push would have caught at least 2 of the 5 Tier-1 defects (geofence sign inversion, duplicate target queue) before they reached the bench.
3. **Record flight video in week 1 for training data.** The DJI video replay pipeline—built post-cancellation in a single session—produced the 4 key quantitative figures in this evaluation. Recording real aerial footage during *any* early manual flight would have provided labelled training data weeks earlier, enabling a statistically meaningful held-out test set ( $n \geq 50$ ) and eliminating the train/validation overlap concern (D5).
4. **Deploy NCNN backend and train a multi-class detector from day 1.** NCNN benchmarks show a  $4.5\times$  speedup over TFLite (72 ms vs 327 ms full pipeline); the backend is implemented but has not yet been validated

**Table 147:** Defects discovered at each testing tier, with measured resolution time and impact if undetected. All fix times are from git commit timestamps; “cost if missed” is the assessed flight-day consequence.

| Defect                                                 | Tier | Fix Time           | Cost if Missed                                          |
|--------------------------------------------------------|------|--------------------|---------------------------------------------------------|
| Geofence sign inverted (pointPolygonTest +ve = inside) | 1    | 15 min             | SSSI incursion; CAA regulatory violation                |
| Repulsive force reversed at NFZ boundary               | 1    | 20 min             | Drone attracted <i>into</i> NFZ                         |
| Barometric drift $\rightarrow$ infinite landing loop   | 1    | 30 min             | Battery drain; hover until fail-safe                    |
| Duplicate target queued (consecutive frames)           | 1    | 10 min             | Repeated descent on same target                         |
| CENTERING state missing timeout guard                  | 1    | 15 min             | State machine stuck indefinitely                        |
| BGR/RGB colour inversion (IMX296 sensor)               | 2    | 45 min             | Confidence drops to $\sim 0.3$ ; mission failure        |
| FOV miscalibration ( $f=7.0 \rightarrow 5.46$ mm)      | 2    | 60 min             | GPS estimates 28 % off; landing outside 10 m R07 radius |
| Python 3.13 pyserial regression                        | 2    | 120 min            | Zero Pi-Cube communication                              |
| <code>tf.lite-runtime</code> absent on 3.13            | 2    | 30 min             | Zero on-device inference                                |
| Camera resource conflict (pi-camera2)                  | 2    | 20 min             | Stream or detection crash on startup                    |
| <code>self.investigating</code> uninitialised          | 2    | 5 min              | Python crash on first detection                         |
| TFLite bbox pixel vs. normalised coords                | 2    | 25 min             | Green boxes off-screen; GPS estimate 100 m+ error       |
| GPS timing lag (100–200 ms)                            | 3*   | —                  | 0.8–1.6 m systematic offset at $8 \text{ m s}^{-1}$     |
| <b>Total (Tiers 1–2)</b>                               |      | <b>6 hr 35 min</b> |                                                         |

\*Discovered via DJI video analysis (Tier 3 proxy); mitigated by inverse-variance averaging and attitude compensation (`--compensate-tilt`).

in the field. Pairing it with a multi-class detector (person/dummy/debris) would reduce operator verification load at the VERIFY state, addressing both D2 and D6.

### AH.3 Lessons Learned

Four technical lessons emerged that generalise beyond this project:

**1. Invest in data collection infrastructure before model training.** The synthetic dataset generator produced 300 images in  $\sim 2$  min, but the resulting model’s  $\text{mAP}_{50}=0.995$  is inflated by train/validation overlap (D5): only 16 of 366 training images (4.4 %) originate from real flight footage. The labelling tool (`label_tool.py`) was built in week 16; building it in week 1 and collecting real frames during any early manual flight would have yielded a proper held-out test set ( $n \geq 50$ ) and eliminated the need to treat  $\text{mAP}$  as an upper bound.

**2. Characterise the sensor pipeline end-to-end before writing application code.** Three of the thirteen defects—BGR inversion (45 min), FOV miscalibration (60 min), and GPS timing lag (mitigated but not fully eliminated)—were sensor-pipeline issues that propagated silently into higher-level logic. A dedicated “sensor characterisation sprint” in week 1—measuring colour encoding, FOV at multiple distances, GPS latency under motion, and lens distortion—would have prevented these faults and saved  $\sim 2$  h of bench-day debugging.

**3. Simulation fidelity has diminishing returns; real-world data does not.** SITL simulation caught 5 of 13 defects (Tier 1) but assumes perfect GPS, zero wind, and synthetic camera views. The DJI video replay pipeline—built in one session post-cancellation—produced 4 quantitative figures (Figures 12a–13) and the  $\text{CEP}_{50}=2.3$  m GPS characterisation that simulation could never yield. Future projects should prioritise replaying real sensor data through the pipeline over building higher-fidelity simulators.

**4. Budget for weather and schedule one flight day per week, not per project.** A single scheduled flight day creates a binary outcome: either the weather cooperates and the project obtains flight data, or it does not. Scheduling weekly flight slots—even short 30-minute windows—would have provided multiple opportunities and allowed the

progressive test tiers to be executed across several sessions rather than compressed into a single day.

## AI Development Methodology

This appendix documents the progressive development methodology used to build the SAR drone system, from initial simulation through to field integration. It covers the design philosophy, a 10-phase timeline with outcomes, and lessons that future teams can apply.

### AI.1 Development Philosophy

Three principles governed development from the outset:

1. **Simulation-first.** Every algorithmic component—state machine, search pattern generator, detection pipeline, GPS estimation—was prototyped and end-to-end tested in SITL (Software-In-The-Loop) on a laptop before any hardware was touched. SITL provides repeatable conditions, instant restart, and zero crash risk.
2. **Same code everywhere.** A single codebase runs on the Windows development laptop, WSL/Linux, and the Raspberry Pi 5. Platform differences are confined to `config.py` (auto-detection of serial ports, camera backends, display availability) and `vision.py` (dual backend: Ultralytics on laptop, TFLite/NCNN on Pi). No `if platform ==` guards appear in mission logic.
3. **Progressive trust.** Each test step isolates exactly one new variable. Bench tests precede flight tests; passive observation precedes autonomous commands; logging-only modes precede action-taking modes. A failure at step  $N$  is attributable to the single change introduced at step  $N$ .

### AI.2 Ten-Phase Development Timeline

Table 148 summarises the ten development phases, their scope, test outcomes, and calendar dates.

**Table 148:** Ten-phase development history with outcomes.

| #  | Phase                 | Scope and Outcomes                                                                                                                                                                                                                                                                    | Dates       |
|----|-----------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------|
| 1  | Simulation on Laptop  | State machine (11 states, 32 transitions), lawnmower search pattern, simulated drone camera from <code>map.jpg</code> , Ultralytics YOLOv8n detection, GPS estimation, operator Y/N confirmation, SITL integration. Full mission completed end-to-end in simulation.                  | Weeks 12–13 |
| 2  | Multi-Platform CV     | Dual-backend <code>vision.py</code> : Ultralytics (laptop) / TFLite (Pi). Docker Pi-environment test passed. TFLite coordinate bug fixed (pixel vs. normalised). <code>config.py</code> auto-detects connection type.                                                                 | Week 13     |
| 3  | Test Infrastructure   | <code>preflight.py</code> , 13 test scripts covering camera, Cube heartbeat, GPS, AI detection, FOV calibration, resolution trade-off. Organised into <code>tests/hardware/</code> , <code>tests/flight/</code> , <code>tests/calibration/</code> , <code>tests/diagnostics/</code> . | Week 14     |
| 4  | Pi + Camera           | Raspberry Pi 5 with IMX296 global shutter. Camera verified at $1456 \times 1088$ . BGR colour issue discovered and fixed (LL-02). TFLite inference: 206.5 ms average, 4.8 FPS, 0.966 confidence (50-run benchmark).                                                                   | Week 16     |
| 5  | Pi + Cube Wired       | Cube Orange connected via UART at 921 600 baud. Mavproxy UDP bridge resolves Python 3.13 serial bug (LL-07). Mission Planner connects via TCP bridge.                                                                                                                                 | Week 16     |
| 6  | Interactive Simulator | <code>simple_simulator.py</code> : keyboard flight, visual servo, GPS estimation, spatial clustering, Kalman filter, 7.5 m offset landing, multi-target support, data recording.                                                                                                      | Weeks 14–15 |
| 7  | Web Ground Station    | <code>pi_flight.py</code> : MJPEG stream, 2D GPS grid, browser dashboard with detection clusters and command buttons. Headless over SSH (PuTTY). Tested in SIMULATION on laptop.                                                                                                      | Week 17     |
| 8  | Field Integration     | FOV calibrated (5.46 mm focal length at bench; $54.4^\circ$ HFOV from DJI video at altitude). Lens calibration (RMS 0.399). Pi bench testing with real Cube. Weather cancelled outdoor flight; adapted to bench-only protocol.                                                        | Weeks 19–20 |
| 9  | Model Retraining      | <code>generate_dataset_v2.py</code> : 366 images (300 synthetic + 16 real-labelled + 50 negatives) at $1456 \times 1088$ . Trained YOLOv8n on Colab: <code>imgsz=1088</code> , 150 epochs, $mAP_{50} = 0.995$ . NCNN backend added (72 ms, 4.5× faster than TFLite).                  | Week 18     |
| 10 | Safety & Polish       | RC override guard (LL-01). Five-layer SSSI geofence (speed field, repulsion, hard cutoff, waypoint filter, firmware backup). SmartEstimator greedy-cluster GPS lock. Threaded inference (stream at 30 fps, detect in background). 127 unit tests, 33 automated integration tests.     | Weeks 19–20 |

### AI.3 Simulation Stage (Weeks 12–15)

All mission logic was developed and validated in SITL before any hardware was procured. The simulation environment consisted of:

- **ArduPilot SITL** via Mission Planner, providing realistic GPS, attitude, and flight dynamics.
- **Simulated camera** (`simulation.py`): a satellite image (`map.jpg`) with a virtual drone viewport that moves according to SITL telemetry.
- **State machine** (`main.py`): 11 states from INIT through DONE, with 32 transitions handling detection, centering, descent, verification, and landing.
- **Search pattern** (`planning.py`): lawnmower generator for arbitrary convex polygons with configurable strip spacing, edge margins, and start-corner optimisation.
- **Interactive simulator** (`simple_simulator.py`, 2508 lines): keyboard-driven flight with GPS estimation, inverse-variance weighting, Kalman filtering, spatial clustering, and 7.5 m offset landing.

By the end of the simulation stage the full mission—takeoff, search, detect, centre, descend, verify, land—ran end-to-end on the laptop with zero hardware dependencies.

### AI.4 Pi Deployment (Weeks 16–18)

Hardware integration proceeded component-by-component:

1. **Camera.** IMX296 global shutter at  $1456 \times 1088$  native resolution. A persistent blue tint was traced to the sensor outputting BGR despite `picamera2`'s `RGB888` label; removing `cv2.cvtColor` fixed it (LL-02).
2. **AI inference.** TFLite on Pi 5 CPU (XNNPACK): 206.5 ms per frame, 4.8 FPS, 50/50 detections at 0.966 confidence (`tests/hardware/benchmark.py`). NCNN backend later added: 72 ms, 9 FPS full pipeline.

3. **Cube connection.** Python 3.13 breaks `pyserial` reads; mavproxy UDP bridge (`udpout:127.0.0.1:14550`) resolves this with zero data loss (LL-07). Mission Planner connects via `tcpin:0.0.0.0:5762`.
4. **Model retraining.** `generate_dataset_v2.py` produced 366 images at native resolution. `tools/label_tool.py --full` labelled 16 real frames from DJI video. Colab training yielded  $\text{mAP50}=0.995$ ; the model is a drop-in replacement (same input/output shapes).

### AI.5 Field Integration (Weeks 19–20)

A field day at the designated site (2026-03-11) assembled the full Pi + Cube + camera stack. Weather cancelled the planned flight; the team adapted by executing every possible bench and calibration activity:

- **FOV calibration:** 92 cm visible at 1 m height  $\Rightarrow f = 5.46$  mm (updated from the initial 7.0 mm estimate).
- **Lens calibration:** checkerboard method (14×9, 28 mm squares),  $\text{RMS}=0.399$ , undistortion cost  $\sim 1.5$  ms per frame.
- **Connectivity:** three-terminal PuTTY workflow verified (mavproxy, diagnostics, flight script). Port conflicts and camera-ownership issues resolved.
- **Bug discovery:** corrupt `calibration.data.npz` (0 bytes) was mangling frames via `cv2.remap`; deleting it restored detection immediately (LL-04).
- **Subsequent video analysis** of DJI flight footage revealed a 100–200 ms GPS timing lag causing  $\text{CEP50}=2.3$  m scatter in target estimates (LL-05).

### AI.6 Ambition Levels

The project defined three deliverable tiers to manage scope risk:

**Table 150:** Three product levels with autonomy and demonstration scope.

| Level | Name            | Autonomy | Demonstration Scope                                                                                                                                                                                          |
|-------|-----------------|----------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| L1    | Manual MVP      | None     | Pilot flies RC manually. Pi runs camera + AI passively (zero commands). Detection overlay on video stream, buzzer alerts, geotagged images saved. Pilot lands manually.                                      |
| L2    | Semi-Autonomous | Partial  | Drone arms, takes off, flies lawnmower search avoiding SSSI. CV detects targets; pilot classifies (Y/I/X). On confirm: centre, descend, offset land at 7.5 m, payload release. RC override always available. |
| L3    | Full Autonomous | Full     | As L2 but confidence accumulation replaces pilot confirmation. Multi-pass verification at decreasing altitude. Drone auto-decides.                                                                           |

This tiered approach ensured that even if weather or hardware issues prevented full autonomous flight, L1 would still produce a demonstrable, useful system and collect the data needed for L2 development.

### AI.7 Progressive Testing Ladder

Flight testing follows a strict five-step ladder. Each step isolates one new risk factor; no step may be skipped.

**Table 152:** Progressive flight testing ladder with scripts and risk isolation.

| Step | Description                                                                                         | Script                              | New Risk       |
|------|-----------------------------------------------------------------------------------------------------|-------------------------------------|----------------|
| 0    | Bench commands: arm, mode change, servo PWM. No propellers.                                         | <code>0a--0f_*.py</code>            | Wiring         |
| 1    | Manual RC flight. Pi runs camera + AI, logs detections, buzzer beeps. <b>Zero MAVLink commands.</b> | <code>1_passive_flight.py</code>    | Outdoor CV     |
| 2    | Code-controlled waypoint flight (3–4 GPS points). No camera, no detection, no decisions.            | <code>2_waypoint_test.py</code>     | MAVLink flight |
| 3    | AUTO waypoints + AI detection. On detection, switch to GUIDED hover. No descent or landing.         | <code>3_auto_detect.py</code>       | CV + flight    |
| 4    | Full autonomous: search pattern, detect, centre, descend, verify (Y/N), offset land.                | <code>4_detect_and_center.py</code> | Full mission   |

Six bench-level tests (step 0) additionally cover planning validation (`0d_planning_test.py`, 6 unit tests), geofence sign conventions (`0e_geofence_test.py`, 7 unit tests), and payload servo actuation (`0f_servo_test.py`).

## AI.8 Achievement Summary

An objective scoring matrix (Table 154) was constructed for each requirement across both simulation and real hardware.

**Table 154:** Simulation vs. real hardware readiness (0–10 scale).

| Category                                  | Simulation | Real Hardware |
|-------------------------------------------|------------|---------------|
| All requirements (R01–R10)                | 7.7        | 2.1           |
| Flight-critical (R01, R02, R03, R04, R09) | 8.4        | 1.8           |
| Mission-critical (R05, R06, R07)          | 7.0        | 1.0           |
| Infrastructure (R08, R10)                 | 7.0        | 4.5           |

**Simulation (7.7/10):** The state machine logic, search pattern generation, CV detection pipeline, GPS estimation mathematics, and safety systems (RC override, geofence, RTL) all work correctly in simulation. The full mission has been completed end-to-end multiple times in SITL.

**Real hardware (2.1/10):** Individual components are verified (camera: working; TFLite: 4.8 FPS at 0.966 confidence; Cube connection: stable via mavproxy), but zero minutes of real flight time have been logged. The 5.6-point gap between simulation and real hardware is the project’s largest risk.

**What simulation proved:** Algorithm correctness, state transitions, search coverage, detection-to-decision loop, safety logic.

**What simulation cannot prove:** Whether the drone physically flies; whether CV detects from 15–30 m altitude outdoors; whether GPS accuracy is sufficient for offset landing; whether the video stream is usable at operational WiFi range; whether the geofence actually prevents SSSI entry in wind.

## AI.9 Lessons for Future Teams

Eight lessons emerged from the development process, each documented with root cause, fix, and checklist (see Appendix AH for the full defect register):

1. **Start hardware testing on day one.** The simulation stage produced mature, well-tested code, but the 5.6-point sim-to-real gap shows that software maturity does not transfer to hardware confidence. Even basic tasks—outdoor GPS fix, one manual hover—would have closed multiple requirements.
2. **Validate the environment before the field.** Docker testing caught missing packages early, but the corrupt `calibration_data.npz` and BGR colour issue were discovered only on-site. A pre-field smoke test with the exact Pi image would have saved hours.
3. **Plan for weather cancellation.** The March 11 field day was the sole outdoor window; weather cancelled it. Having a bench-only fallback protocol (FOV calibration, lens calibration, connectivity verification) ensured the day was still productive.
4. **Decouple camera ownership.** Only one process can open `picamera2`. This forced a clean separation: one script owns the camera and streams to all consumers via MJPEG. Design for single-owner resources from the start.
5. **Never trust format labels.** The IMX296 sensor outputs BGR despite `RGB888` labelling. YOLOv8 TFLite outputs pixel coordinates despite documentation suggesting normalised values. Test empirically; label after measurement.
6. **Build zero-command observation scripts.** The passive watch approach (`passive_watch.py`, zero MAVLink commands) was the safest way to validate CV in the field. It could run during any flight without risk, collecting training data and detection statistics simultaneously.
7. **Use progressive trust, never skip steps.** The numbered test ladder (0a–4) isolates exactly one new variable per step. When Step  $N$  fails, the cause is the single change introduced at Step  $N$ , not accumulated uncertainty from earlier steps.
8. **Maintain tiered deliverables.** L1 (Manual MVP), L2 (Semi-Autonomous), L3 (Full Autonomous) ensured that even if only L1 was achievable, it would be a complete, demonstrable product rather than a half-finished L2.