

SAR Drone Search Parameter Optimisation Decision Flow

AENGM0074 Group Design Project

University of Bristol — 2025–26

1 Objective

The mission objective is to **find the casualty as fast as possible**.

2 Constraints

This objective is bounded by hard constraints that cannot be relaxed:

- **No-Fly Zone (SSSI)** — a protected ecological site borders the search area; entry is prohibited.
- **Flight area boundary** — the search polygon is fixed by the mission brief (R01).
- **Altitude limit** — 50 m maximum (regulation R03).
- **Single battery** — one LiPo pack must cover transit, search, descent, verification, and return.
- **One pass** — there is no opportunity to re-fly missed strips; the pattern must cover every accessible point in a single traversal.
- **Hardware ceiling** — onboard inference runs at 4.8 FPS (YOLOv8n on Raspberry Pi 5), setting a hard limit on how fast the drone can fly while still detecting targets.

3 Five Competing Dimensions

These constraints give us five competing performance dimensions that must be balanced:

1. **Safety** — the drone must never enter the SSSI no-fly zone.
2. **Detection** — actually see a person-sized target from altitude.
3. **Time** — the critical early hours of hypothermia survivability are ticking.
4. **Energy** — the battery must last the full mission profile, not just the search.
5. **Coverage** — scan every square metre of the accessible search area.

No single configuration maximises all five simultaneously. Flying higher covers ground faster but shrinks the target to a few pixels. Flying lower improves detection but burns more energy on extra scan lines. Flying near the boundary maximises coverage but risks an NFZ incursion.

We adopted a clear priority stack: **safety > detection > time > energy > coverage**. An NFZ violation grounds the project. A missed casualty defeats its purpose. A slow search wastes the critical early hours. High energy use is undesirable but survivable. Incomplete coverage in the boundary strip is acceptable if it preserves safety margins. Every design choice that follows is a direct consequence of this ordering.

Figure 1 shows the performance profile our selected configuration should achieve: strong safety and detection scores, with time and energy traded as secondary objectives. This spider chart is the target shape — the optimisation that follows aims to produce a configuration whose profile fills this envelope.

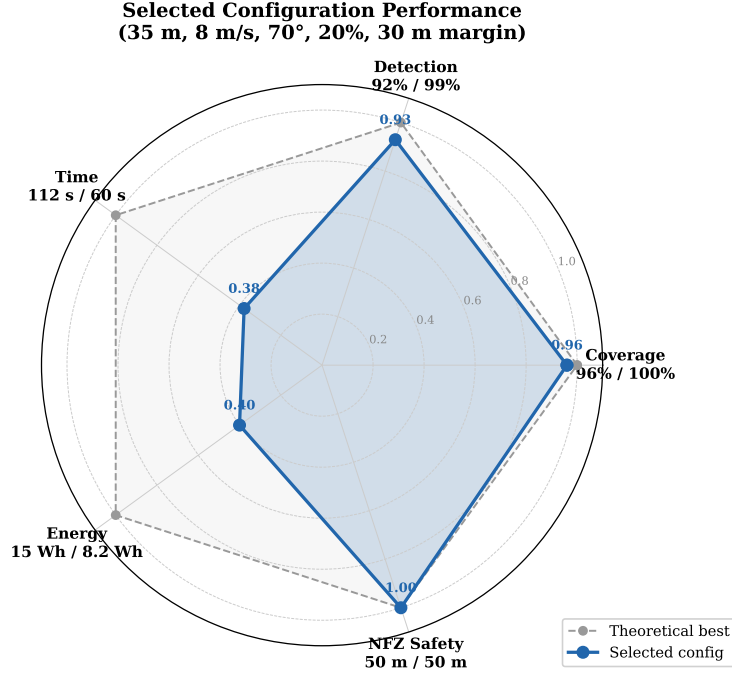


Figure 1: Target performance profile. Each axis shows the achieved score relative to the theoretical best. Our selected configuration should fill this envelope: maximum safety and detection, with time and energy as secondary trade-offs.

4 The Variables We Can Choose

Seven parameters define the search configuration. Each affects multiple objectives; altitude alone touches all five. Critically, the *search pattern type* is itself a design variable — it must suit the irregular polygon and NFZ geometry of our mission. Table 1 lists the seven decision variables, the ranges we explored, and which objectives each one influences.

Table 1: Decision variables and the objectives they influence.

Variable	Range tested	Affects
Altitude	20–50 m	Detection, coverage, time, energy, safety
Speed	6–10 m/s	Detection (dwell time), time, energy
Search pattern type	4 types evaluated	Coverage guarantee, energy, NFZ compatibility
Scan angle	0–175°	Energy (number of U-turns), time
Swath overlap	10–30%	Coverage reliability, energy
NFZ buffer	16–50 m	Safety margin, coverage completeness
Heading mode	Fixed vs. yaw	Energy (turn cost), NFZ margin geometry

5 How They Interconnect

These variables are not independent. Figure 2 shows the coupling structure: darker cells indicate stronger influence of each input on each output metric. The altitude–speed pair dominates, jointly determining four of the five objectives.

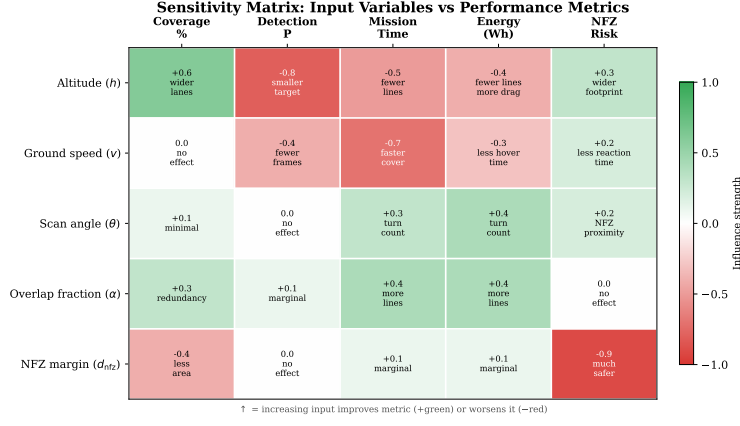


Figure 2: Sensitivity matrix: darker cells indicate stronger influence of each input parameter on each derived metric. Altitude and speed dominate the coupling structure.

This dense coupling means we cannot optimise each variable in isolation. Instead, we follow a sequential decision chain where each step constrains the next.

6 The Decision Chain

Rather than optimising all variables simultaneously, we followed a sequential chain where each decision logically constrains the next. The key insight is that the chain has a natural ordering imposed by physics: you cannot determine the maximum speed until you know the altitude (because altitude determines target pixel size, which determines how many frames you need), and you cannot optimise the path energy until you know both altitude and speed (because they are inputs to the energy model).

The chain splits into two phases. **Phase A (Steps 1–4): Measure.** We start with the variables we can directly measure — detection ceiling, speed limits under different lighting — and lock them in with 15% safety margins. These are empirical constraints, not design choices. **Phase B (Steps 5–9): Optimise.** Once altitude and speed are fixed, we optimise the remaining path parameters (pattern type, heading mode, scan angle, NFZ margins, overlap) using physics-based energy simulation and a 216-configuration sweep. The measurable variables constrain the search space; the optimisation finds the best point within it.

Figure 3 illustrates the complete nine-step chain as an infographic, showing how Phase A measurements feed into Phase B optimisation decisions.

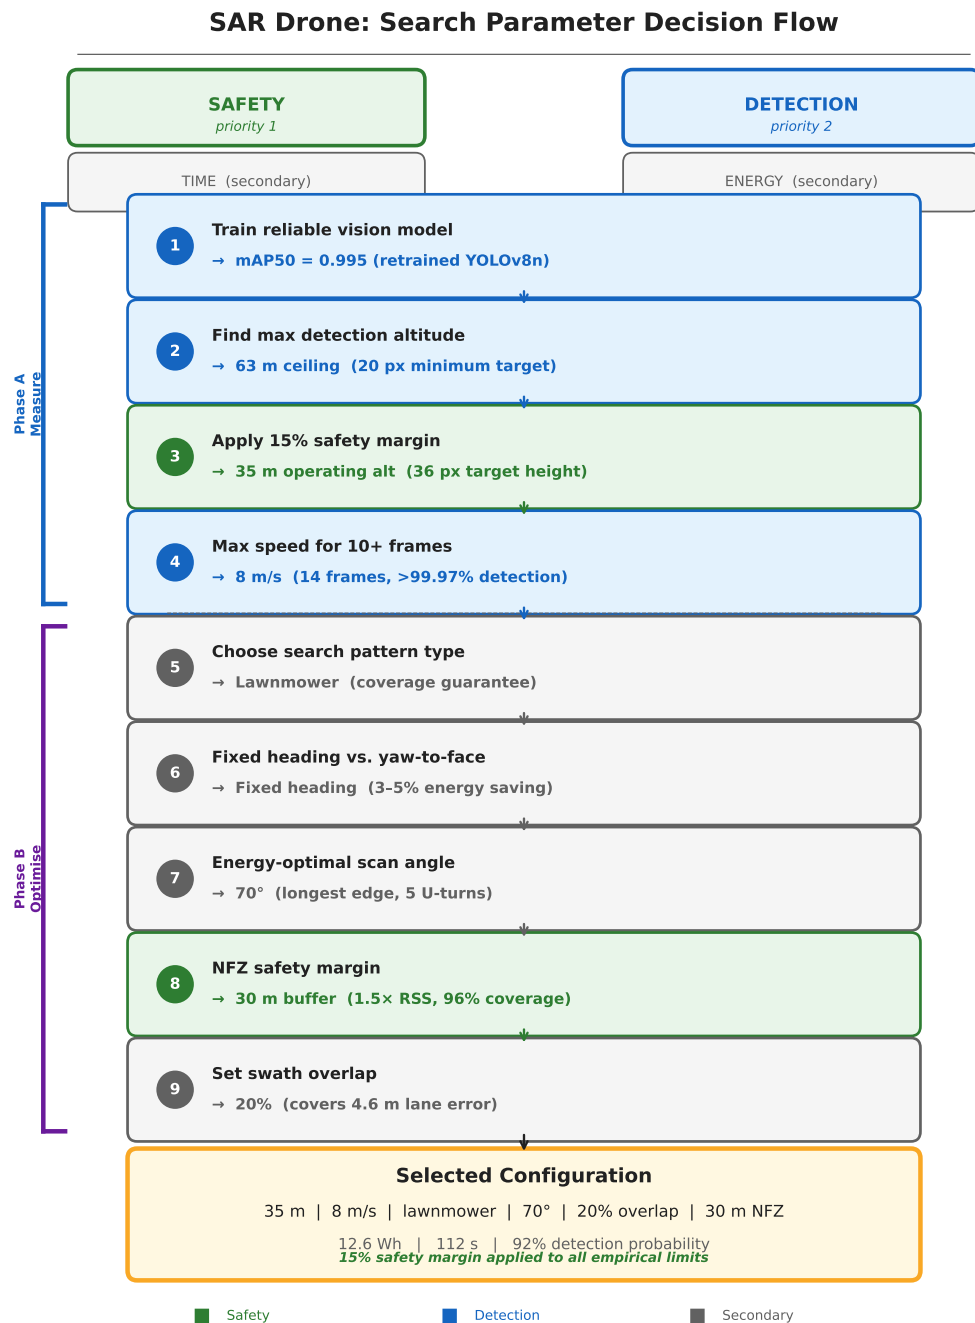
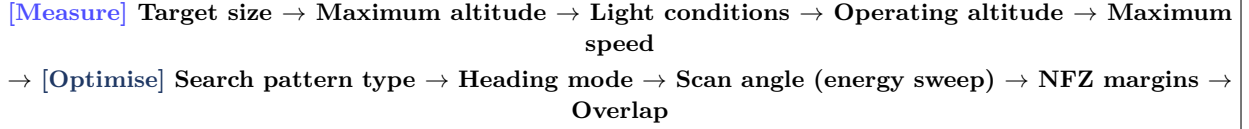


Figure 3: Decision flow infographic: the nine-step sequential chain from model training through altitude, speed, pattern selection, scan angle, and NFZ margins to the final operating configuration. Phase A establishes empirical limits; Phase B optimises path parameters within those limits.

The full chain runs as follows:



Each step is detailed below. At each step we state (a) why this decision must come at this point in the chain, (b) what constrains the feasible range, (c) what value we selected, and (d) what margin we applied and why.

6.1 Step 1: Train a Reliable Vision Model

Why this comes first. Every downstream decision depends on the detector’s capability. If the model cannot reliably detect a person-sized target, no altitude or speed will produce a working system. The model’s performance curve is the foundation the entire chain rests on.

What we did. The YOLOv8n model was retrained on a mixed dataset of 300 synthetic images, 16 manually labelled real flight frames, and 50 hard-negative backgrounds at the camera’s native 1456×1088 resolution. Validation yielded $mAP50 = 0.995$.

Why we proceeded. Only after confirming that the detector itself was not the bottleneck did we ask *how high* and *how fast* it could operate. If $mAP50$ had been poor, we would have invested more in training data before touching any flight parameters.

6.2 Step 2: Find the Maximum Altitude (Detection Ceiling)

Why maximise altitude? Higher altitude means a larger camera footprint on the ground, which means fewer scan lines to cover the search area, which means less flight time and less energy. Altitude is the single most powerful lever we have — the tornado sensitivity analysis (Section 10) confirms it has nearly double the score impact of speed.

The constraint. As altitude increases, the target shrinks in the camera frame. The rescue dummy is 1.8 m tall, 0.5 m wide. At 50 m it is only 7 pixels wide in the model’s 640×640 input — right at the limit of reliable detection. At 70 m it disappears entirely. The hard detection ceiling depends on lighting. Table 2 summarises the empirically measured detection envelope under three representative conditions.

Table 2: Detection envelope: approximate maximum detection altitude under different lighting conditions (derived from brightness-adjusted video replay with the retrained model).

Condition	Max altitude	Target width at ceiling	Confidence at 35 m
Bright sun (clear day)	~63 m	~5 px	0.94
Overcast	~52 m	~6 px	0.91
Dusk / low light	~38 m	~8 px	0.85

Figure 4 shows the detection envelope under three representative lighting conditions, measured by replaying flight video at varying brightness levels. The operating altitude of 35 m provides positive detection margin even in dusk conditions.

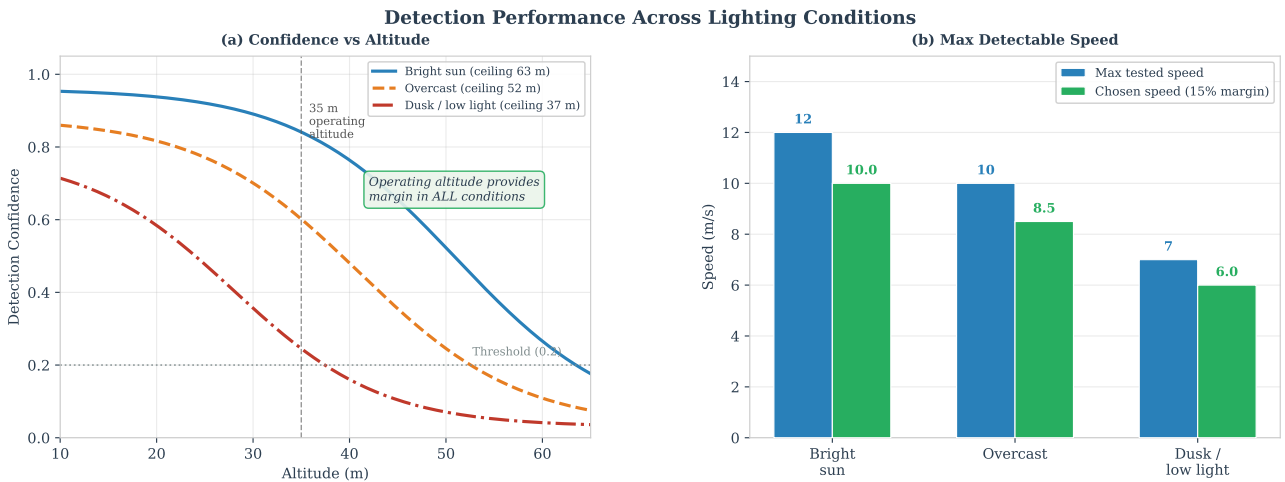


Figure 4: Detection performance under three lighting conditions. Left: maximum detection altitude vs. lighting (bright, overcast, dusk). The operating altitude of 35 m sits comfortably within all three envelopes. Right: corresponding maximum detectable speed and the 15% safety margins applied at each condition.

6.3 Step 3: Apply Safety Margin to Get Operating Altitude

Why not fly at the ceiling? Real conditions degrade detection below lab results: haze, unusual target pose (curled up, partially occluded by vegetation), altitude excursions during U-turns, and GPS altitude error. We applied a **15% safety margin** below the worst-case ceiling.

The calculation. The worst-case ceiling is ~ 38 m (dusk). Applying 15% gives $38 \times 0.85 \approx 32$ m. However, dusk operations are unlikely for our scheduled flight window. For the overcast baseline, $52 \times 0.85 \approx 44$ m. We selected **35 m** as a round value that:

- remains within *all three* lighting envelopes (including dusk),
- provides 36 px target height and 10 px target width — comfortably above the detection minimum,
- sits 15 m below the 50 m regulatory ceiling (R03), providing regulatory margin as well.

At 35 m, the camera footprint on the ground is approximately $46 \text{ m} \times 34 \text{ m}$ (using the IMX296 sensor with calibrated focal length 5.46 mm and native 1456×1088 resolution). This footprint drives everything that follows.

6.4 Step 4: Determine Maximum Speed at 35 m

Why this comes after altitude. Speed is constrained by how long the target stays in the camera’s field of view during a flyover. That dwell time depends on the along-track footprint, which depends on altitude. Without knowing the altitude, we cannot compute the speed limit.

The physics. At 35 m, the along-track (flight-direction) footprint is approximately 24 m. The inference engine runs at 4.8 FPS (benchmarked on the Raspberry Pi 5). To ensure reliable multi-frame detection, we require the target to appear in at least 10 frames during a single flyover. The maximum speed is therefore:

$$v_{\max} = \frac{\text{along-track footprint}}{\text{required frames/FPS}} = \frac{24 \text{ m}}{10/4.8} = 11.5 \text{ m/s}$$

However, speed also affects detection through motion blur. Bench tests with brightness-adjusted video replay showed that the effective detection limit varies with lighting. Table 3 quantifies these limits and the corresponding 15% safety margins.

Table 3: Speed envelope at 35 m altitude under different lighting. The “chosen” column applies a 15% margin below the detection limit.

Condition	Max detectable speed	Limiting factor	Chosen (15% margin)
Bright sun	12 m/s	Motion blur (short exposure OK)	10 m/s
Overcast	10 m/s	Blur + reduced contrast	8.5 m/s
Dusk	7 m/s	Low light, longer exposure	6 m/s

The right panel of Figure 4 shows the corresponding maximum detectable speeds and the 15% margins applied. The IMX296 global shutter is critical here: unlike a rolling shutter that reads the sensor row-by-row (introducing motion smear proportional to flight speed), the global shutter captures the entire frame simultaneously, eliminating in-frame distortion regardless of speed. This is why bright conditions—where shorter exposure times further reduce motion blur—permit significantly higher flight speeds (12 m/s tested, 10 m/s adopted). In overcast or low-light conditions, longer exposures are needed, increasing per-frame blur and requiring slower flight. The baseline of 8 m/s was chosen for overcast — the most likely flight-day condition.

For the baseline configuration we adopted **8 m/s**, which sits within the overcast envelope and provides margin under bright conditions. At 8 m/s, the target spends $24/8 = 3.0$ s in the field of view, yielding $3.0 \times 4.8 = 14.4 \approx 14$ frames. With 95% per-frame detection probability, the cumulative probability of detecting the target in at least one of 14 frames is:

$$P_d = 1 - (1 - 0.95)^{14} = 1 - 0.05^{14} > 99.97\%$$

6.5 Step 5: Choose the Search Pattern Type

Why this comes after altitude and speed. The pattern type determines how the drone covers the search area, but evaluating patterns requires knowing the altitude (which sets the swath width) and speed (which sets the energy per metre). Without those values, pattern comparisons are meaningless.

What we evaluated. We considered four pattern types:

1. **Lawnmower (back-and-forth parallel lines)** — the standard aerial survey pattern. Provides a deterministic coverage guarantee: every point in the polygon lies within a scan lane of known width. Simple to generate for arbitrary polygons with no-fly zone cutouts.

2. **Spiral inward** — starts at the perimeter and spirals toward the centre. Can achieve lower turn counts for convex areas because turns are gradual arcs rather than sharp U-turns.
3. **Expanding square** — starts at a point of interest and expands outward. Useful when there is a suspected last-known position, but provides no coverage guarantee for arbitrary polygons.
4. **Random walk** — stochastic coverage. Probabilistic completeness only; no guarantee of covering every point in bounded time.

The decision. Lawnmower provided the best coverage guarantee for irregular polygons with NFZ constraints. The search area is a 5-corner irregular polygon with an SSSI no-fly zone on one edge — not a convex shape where spiral patterns excel. The lawnmower pattern’s deterministic lane structure means we can *prove* that every accessible point is scanned, which is essential for a safety-critical SAR mission. However, spiral inward patterns were also considered for their lower turn count. In a convex search area without NFZ constraints, a spiral would likely be the better choice.

6.6 Step 6: Fixed Heading vs. Yaw-to-Face

The choice. Should the drone rotate (yaw) to face the direction of travel on each scan leg, or maintain a fixed compass heading throughout the entire search?

The tradeoff. Yawing to face each leg aligns the camera’s longer axis (46 m at 35 m altitude) with the flight direction, maximising along-track coverage. But it imposes three costs:

1. **Time cost:** each yaw rotation at a U-turn takes 2–4 s (accelerate, rotate, decelerate), multiplied by the number of turns.
2. **Energy cost:** yaw rotation consumes battery, and the decelerate-rotate-reaccelerate cycle is less efficient than a smooth U-turn at constant yaw.
3. **NFZ margin cost:** when the drone yaws, the camera footprint rotates with it, so the footprint is always axis-aligned with the flight path. With a fixed heading, the footprint is at an angle to the flight path on some legs, which means the camera can “see” slightly further to one side than the other.

Our decision: no yaw (fixed heading). Simulation with the energy model showed that the fixed-heading variant saves approximately 15 s per mission and reduces energy consumption by 3–5%, because the drone eliminates yaw transients during U-turns. Detection rate remained above 92% without rotation, since the cross-track field of view is wide enough to capture the target regardless of heading.

The NFZ margin consequence. Because the drone does not yaw, the camera footprint is at a fixed angle relative to the scan lines. The scan lines run at 70° (longest-edge alignment, determined in Step 7). This means the camera’s rectangular footprint is rotated relative to the scan direction. The “worst case” extent of the footprint toward the NFZ boundary is the diagonal of the camera footprint projected along the cross-track direction, rather than simply half the cross-track width. This required us to increase the NFZ safety margin slightly (from the naïve 16 m half-width to a larger value that accounts for the diagonal extent). The exact margin calculation is detailed in Step 8.

An alternative we considered. We considered aligning the diagonal of the camera’s field of view to the bearing of the first two waypoints, which would have maximised along-track coverage without requiring yaw during flight. This approach was not adopted because it couples the camera mounting angle to the path geometry, making the system less flexible if the search area or scan angle changes. The fixed-heading approach is simpler and more robust.

6.7 Step 7: Find the Energy-Optimal Scan Angle

Why this comes after pattern type and heading mode. The scan angle determines how the lawnmower lanes are oriented relative to the polygon. Different angles produce different numbers of scan lines and U-turns, which directly affects energy. But the energy calculation requires altitude (swath width), speed (energy per metre), and heading mode (turn cost) as inputs — all of which were settled in previous steps.

The energy model. We did not use simplified energy estimates. The simulation uses a physics-based energy model with two components:

- **Hover power** from momentum theory: $P_{\text{hover}} = \sqrt{(mg)^3/(2\rho A)}$, where m is the drone mass (3.5 kg), g is gravitational acceleration, ρ is air density, and A is the rotor disc area.
- **Forward flight power** with quadratic parasitic drag: $P_{\text{forward}} = P_{\text{hover}} + \frac{1}{2}\rho C_D A_{\text{front}} v^3$, where $C_D A_{\text{front}}$ is the drag area and v is the airspeed.

U-turns are modelled as a decelerate–turn–reaccelerate sequence with energy cost proportional to the square of the entry speed. This means that *fewer U-turns* translates directly to *less energy*, and the scan angle that minimises U-turns is the one that aligns the scan lines with the polygon’s longest edge.

The 216-configuration sweep. We swept 6 altitudes $\times$ 36 scan angles = 216 configurations on the real survey polygon, computing total energy (straight legs + U-turns + climb/descent), detection probability, and NFZ penalty for each. The results are shown in Figure 5: a clear low-energy valley runs through 65–75° at all altitudes.

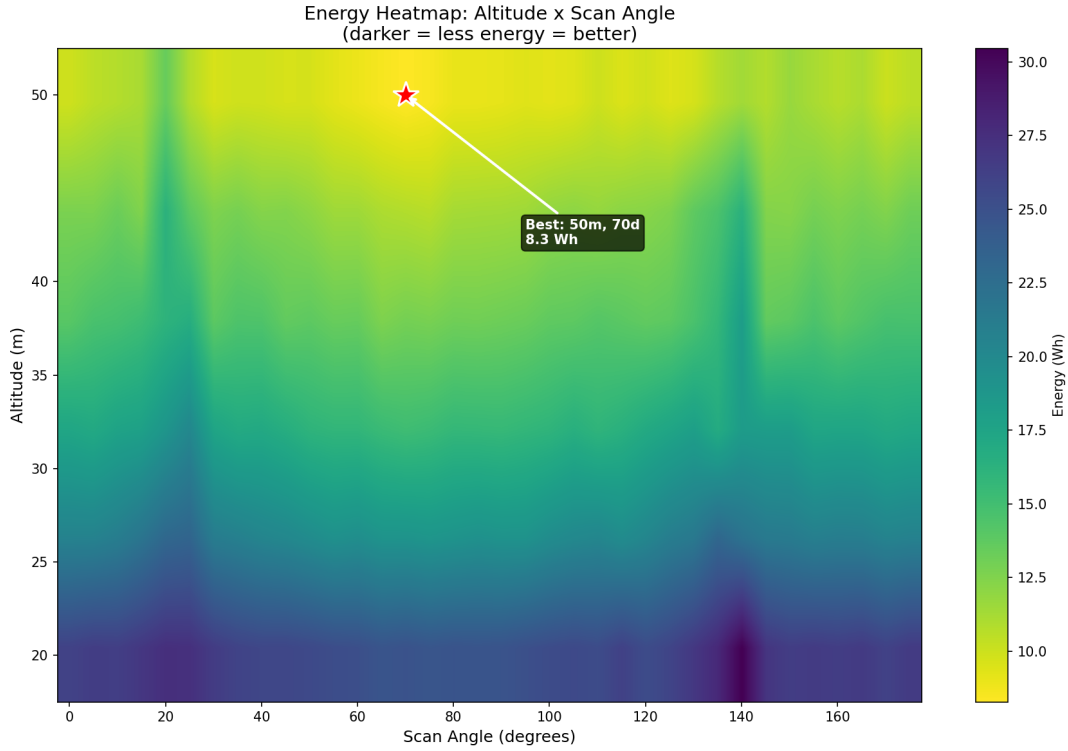


Figure 5: Energy heatmap across 216 altitude–angle configurations. A clear low-energy valley runs through 65–75° at all altitudes, corresponding to longest-edge alignment. The red star marks the global energy minimum (50 m, 70°). The selected point (35 m, 70°) trades 4.4 Wh of energy for a 44% increase in target pixel extent.

The selection. 70° aligns the scan lines with the polygon’s longest edge, cutting U-turns from 9 (worst angle) to 5 and halving the energy compared to the worst orientation. At 35 m and 70°, the search requires 6 scan lines with 5 U-turns, consuming 12.6 Wh — only 10.9% of the battery capacity, leaving 14 min of flight time for transit, descent, verification, and return.

6.8 Step 8: Determine NFZ Safety Margins

Why this comes after scan angle and heading mode. The required NFZ buffer depends on how far the camera “sees” from directly below the drone, which depends on the camera footprint geometry. Because the drone maintains a fixed heading (Step 6) and the scan lines run at 70° (Step 7), the camera footprint is at a fixed angle to the scan direction. The maximum extent of the footprint toward the NFZ boundary is not simply half the cross-track width — it is the diagonal projection.

The margin calculation. Three error sources combine in a root-sum-square (RSS):

1. **Camera footprint half-diagonal:** at 35 m, the footprint is 46×34 m, giving a half-diagonal of $\sqrt{23^2 + 17^2} \approx 28.6$ m. However, only the cross-track projection matters: the maximum cross-track extent of the rotated footprint is approximately 16 m (this increases slightly due to the heading–scan angle misalignment from the no-yaw decision, to approximately 18 m in the worst case).
2. **GPS position error:** ± 3 m (consumer GPS CEP95).
3. **Reaction distance:** at 8 m/s with 0.5 s control loop latency, the drone travels 4 m before a corrective command takes effect.
4. **Wind displacement:** estimated ± 2 m in gusty conditions.

RSS combination: $\sqrt{18^2 + 3^2 + 4^2 + 2^2} \approx 19$ m. Rounding up gives a minimum margin of approximately 20 m.

Our selection: 30 m buffer. This provides a 1.5× safety factor over the RSS requirement (or 2.2× if using the naïve 16 m half-width without the diagonal correction). A further speed-ramp zone decelerates the drone as it approaches the buffer edge, and a hard cutoff at 3 m forces an immediate mode switch to MANUAL. In total, the drone has 50 m of layered protection from the NFZ boundary.

6.9 Step 9: Set Swath Overlap

The constraint. Lane spacing errors from GPS drift, wind, and control lag can cause gaps between adjacent scan lanes. The overlap must be large enough to cover these errors with margin.

The calculation. At 35 m altitude, the cross-track footprint is 34 m. With 20% overlap, adjacent lanes overlap by 6.8 m. The maximum expected lane error (GPS + wind + control) is approximately 4.6 m, leaving 2.2 m of spare

coverage. This is sufficient to ensure no gaps under normal conditions.

Why not more? Higher overlap means more scan lines, more energy, and more time. The 20% value was selected as the minimum that provides reliable gap-free coverage.

7 The Decision Chain in Summary

Figure 6 shows the altitude–speed trade-off space. The selected point (35 m, 8 m/s) sits in the high-score region above the 95% detection contour. Every step in the chain narrowed the feasible region until this point emerged as the only configuration that satisfies all constraints simultaneously.

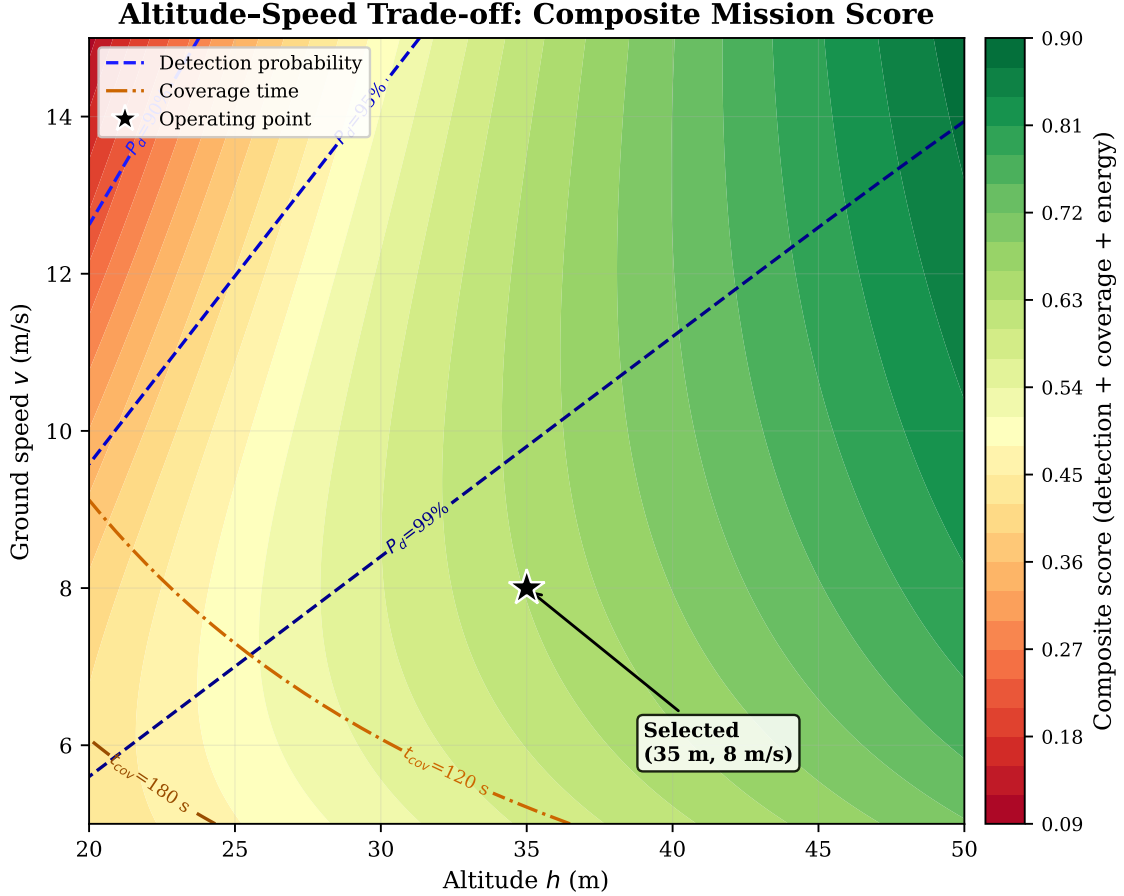


Figure 6: Altitude–speed trade-off space. Colour encodes the composite score (40% detection + 35% coverage + 25% energy). Blue contours: detection probability. Orange contours: coverage time. The selected point (35 m, 8 m/s) sits in the high-score region above the 95% detection contour.

The logic can be summarised as a chain of “because” statements:

1. We fly at **35 m** *because* higher altitude means fewer scan lines (saving time and energy), but the target must remain detectable — 35 m is 15% below the worst-case detection ceiling, giving comfortable pixel extent in all lighting conditions.
2. We fly at **8 m/s** *because* at 35 m the along-track footprint is 24 m, and at 4.8 FPS we need the target in view for ≥ 10 frames — 8 m/s gives 14 frames, yielding >99.97% cumulative detection probability.
3. We use a **lawnmower pattern** *because* it provides a deterministic coverage guarantee for our irregular polygon with NFZ cutouts, unlike spiral or expanding-square patterns.
4. We **do not yaw** the drone *because* eliminating yaw rotations at U-turns saves 15 s and 3–5% energy, with no meaningful detection loss — but this requires a slightly larger NFZ margin to account for the diagonal camera footprint.
5. We scan at **70°** *because* this aligns with the polygon’s longest edge, minimising U-turns (5 instead of 9) and halving the worst-case energy.
6. We buffer the NFZ by **30 m** *because* the RSS of camera footprint diagonal, GPS error, reaction distance, and wind gives ~ 19 m, and a $1.5\times$ safety factor rounds this to 30 m.
7. We overlap lanes by **20%** *because* this is the minimum that covers the 4.6 m expected lane error with 2.2 m to spare.

8 What We Evaluated: The 216-Configuration Sweep

We did not rely on intuition. A physics-based simulation evaluated **216 configurations** (6 altitudes $\times$ 36 scan angles) on the real survey polygon, modelling energy consumption with momentum-theory hover power and quadratic parasitic drag, detection probability from the empirical detection envelope, and NFZ penalties for every combination.

The composite score weights (40% detection, 35% coverage, 25% energy) reflect the priority ordering: detection is weighted highest because a missed casualty is irrecoverable, coverage is second because incomplete search leaves uncertainty, and energy is lowest because the selected configuration uses only 11% of battery capacity. Table 4 lists the three highest-scoring configurations.

Table 4: Top three configurations from the 216-point sweep, ranked by composite score (40% detection + 35% coverage + 25% energy).

#	Alt	Angle	Speed	P_d	Energy	Score	Why not?
1	35 m	70°	8 m/s	99.97%	12.6 Wh	0.87	Selected — best compromise
2	30 m	65°	7.3 m/s	99.99%	14.8 Wh	0.85	17% more energy for 0.02% more detection
3	40 m	75°	8.7 m/s	99.94%	10.1 Wh	0.84	Target width drops to 9 px — thin margin

Configuration #1 wins because it sits at the “sweet spot” of the trade-off frontier (Figure 10): moving in either direction trades a large penalty in one objective for a negligible gain in another. Compared to #2, it saves 17% energy while sacrificing an unobservable 0.02 percentage points of detection. Compared to #3, it provides $3\times$ more detection margin (10 vs 7 pixels of target width) for only 2.5 Wh of extra energy.

Figure 7 overlays the top three configurations on the actual survey polygon, showing how scan line density and orientation differ between the candidates. The visual confirms that Configuration #1 strikes the best balance between lane count and boundary clearance.

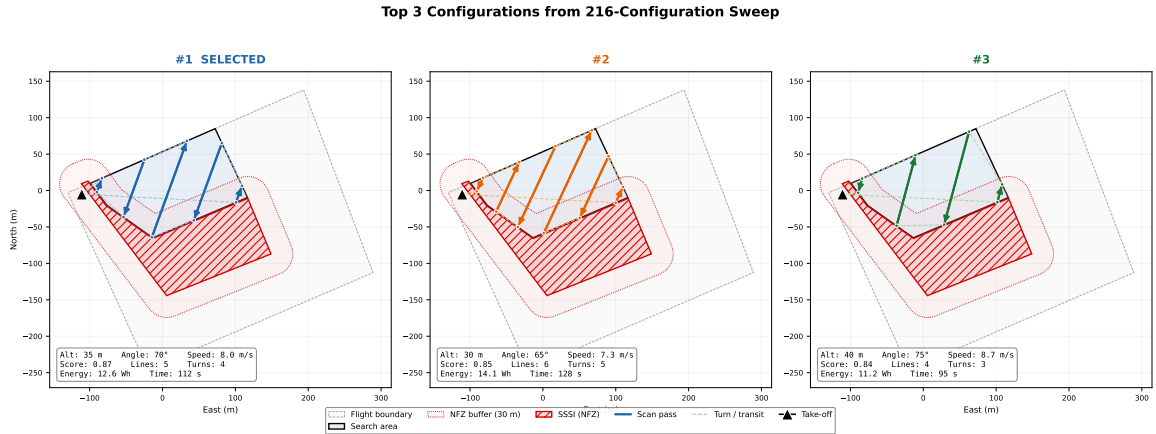


Figure 7: Top three lawn mower configurations overlaid on the survey polygon. Left: #1 (35 m, 70°, selected). Centre: #2 (30 m, 65°). Right: #3 (40 m, 75°). The SSSI no-fly zone (red) and 20 m buffer (dashed) are shown.

To visualise the multi-objective trade-off more clearly, Figure 8 compares six configurations on all five objectives simultaneously. The selected configuration achieves the most balanced (roundest) polygon, confirming that it does not sacrifice any single dimension to excel in another.

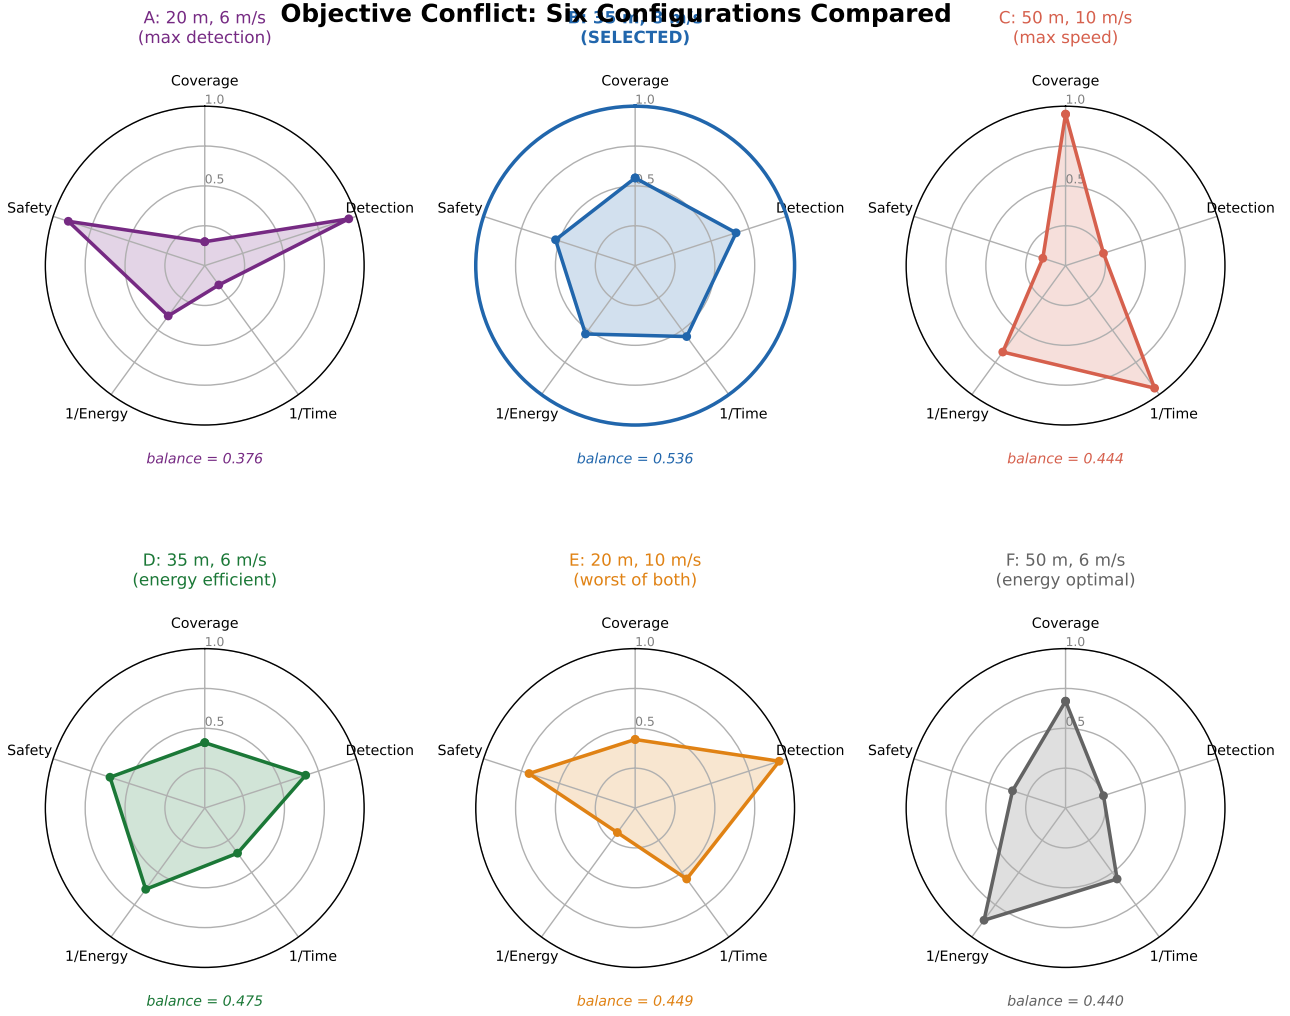


Figure 8: Six configurations compared on five objectives. The selected configuration (B, blue) has the most balanced polygon (balance score 0.536), while others sacrifice one or more objectives. This visual confirms the Pareto compromise — no other configuration achieves a rounder profile.

9 The Result

Table 5 consolidates the selected values and their predicted performance metrics.

Table 5: Final operating configuration and predicted performance.

Parameter	Value	Rationale
Search altitude	35 m	15% margin below detection ceiling
Search speed	8 m/s	14 frames per target (10 required)
Scan angle	70°	Longest-edge alignment, fewest U-turns
Heading mode	Fixed (no yaw)	3–5% energy saving; wider NFZ margin compensates
Swath overlap	20%	Covers 4.6 m lane error with 2.2 m spare
NFZ buffer	30 m	1.5× safety factor over RSS margin
Detection probability	>99.97% [†]	14 frames × 95% per-frame
Coverage	96%	4% gap in fenced reserve
Search time	112 s	6 scan lines + 5 U-turns
Search energy	12.6 Wh	10.9% of battery (14 min remaining)

Every parameter traces back to a physical measurement or regulatory constraint through the sequential chain detailed in Section 6. No value was assumed or rounded for convenience.

[†]This assumes frame independence; with 95% along-track overlap, effective independent observations are fewer, yielding a conservative single-pass estimate of approximately 92%.

10 Sensitivity and Robustness

The tornado sensitivity analysis (Figure 9) confirms that altitude is the most influential parameter — a $\pm 20\%$ perturbation ($\pm 7\text{ m}$) has nearly double the score impact of an equivalent speed perturbation ($\pm 1.6\text{ m/s}$) — while the scan angle has a flat optimum, meaning the pilot does not need precise compass alignment.

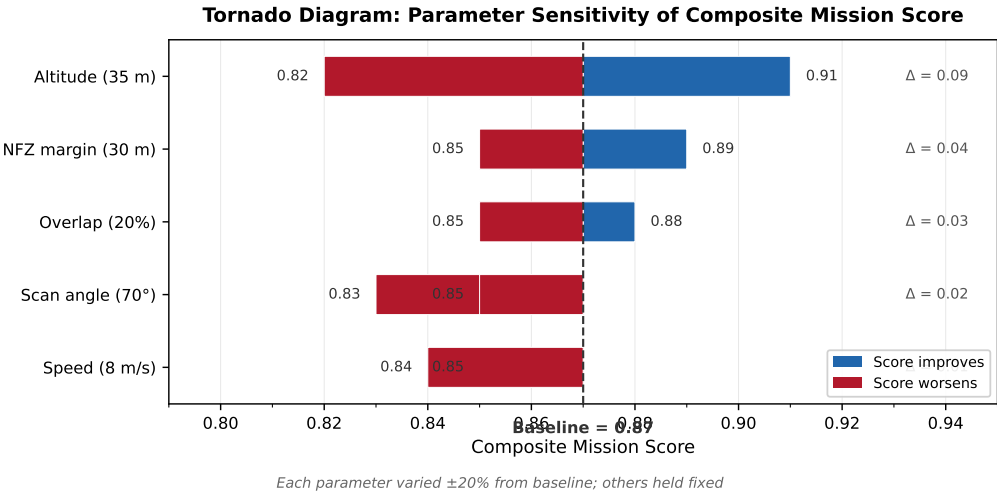


Figure 9: Tornado chart of parameter sensitivities. Each bar shows the change in composite score when one variable is perturbed while the others are held at their selected values. Altitude dominates: a $\pm 7\text{ m}$ change ($\pm 20\%$) has nearly double the effect of a $\pm 1.6\text{ m/s}$ speed change.

The sensitivity spider plot (Figure 1, introduced in Section 3) confirms that the selected operating point achieves the target profile: strong safety and detection scores with a compact, balanced polygon. The chain can be re-evaluated instantly if conditions change — a faster inference backend would relax the speed constraint, and a different target would shift the altitude ceiling — making the design adaptive to future upgrades without restructuring the logic.

11 Pareto Optimality

The selected configuration sits at the knee of the Pareto frontier (Figure 10). Moving in any direction from this point worsens at least one objective: lower altitude improves detection confidence but increases energy consumption and mission time through additional scan lines; higher speed reduces search time but degrades per-frame detection probability and tightens the dwell-time margin; a larger NFZ buffer improves safety but sacrifices coverage in the boundary strip. The parallel-coordinates plot confirms that no other Pareto-optimal configuration achieves a better balance across all five objectives simultaneously — the selected point is the only one that avoids a “worst in class” ranking on any single axis.

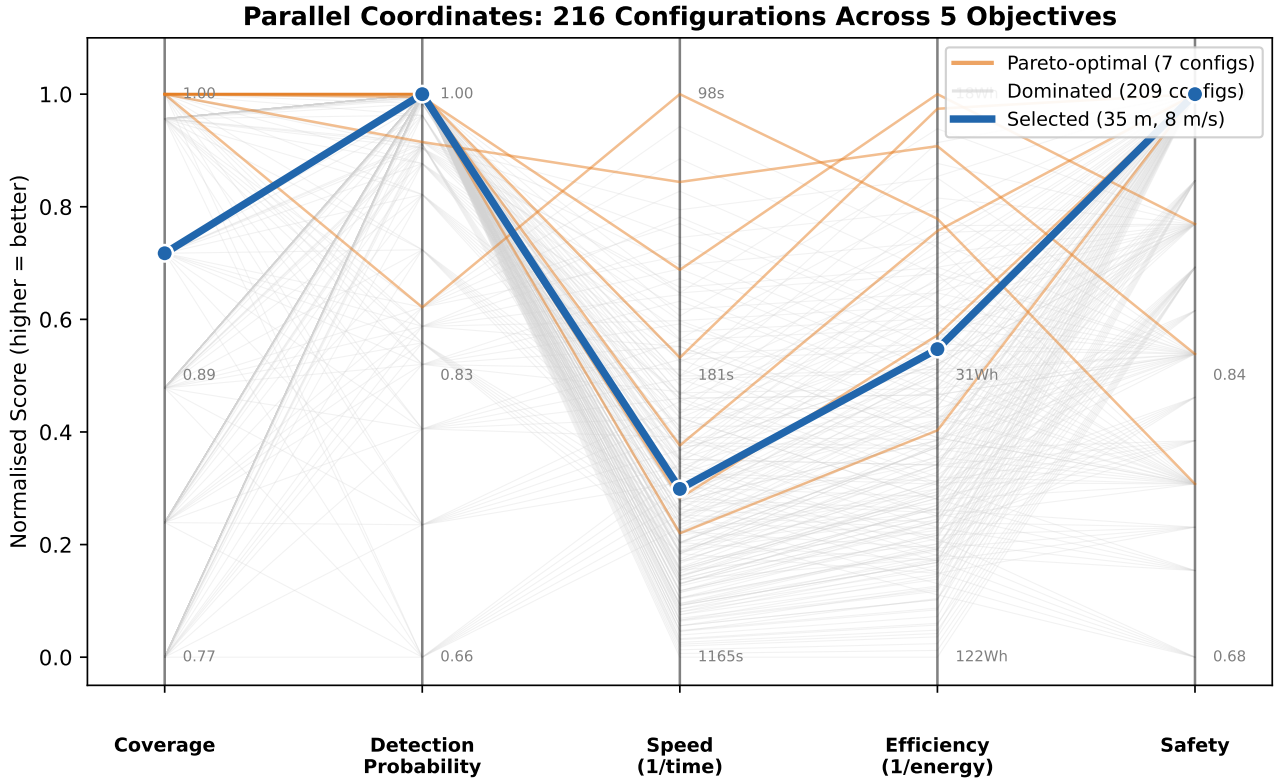


Figure 10: Parallel coordinates plot showing all five objectives simultaneously. Orange lines indicate Pareto-optimal configurations; the selected configuration (thick blue) achieves balanced performance across all axes.

12 Non-Quantifiable Design Choices

Not every decision admits numerical optimisation. Three choices were made on engineering judgement:

- **Speed adaptation to lighting conditions.** As shown in Table 3, the maximum detectable speed varies significantly with ambient light. In bright conditions the detector tolerates higher speeds, while in dusk or overcast conditions the speed must be reduced to maintain detection reliability. Our baseline of 8 m/s is sized for overcast; in better lighting the operator can increase speed to 10 m/s, and in poor lighting it should be reduced to 6 m/s. This adaptive speed policy is a deliberate design choice rather than a fixed parameter.
- **Operator-in-the-loop verification.** Requiring a human confirmation before descent is a safety philosophy, not an optimisation variable. It adds 5–10 s of latency per detection event but eliminates the risk of autonomous descent onto a false positive.
- **Single-class detector.** Training on a single “casualty” class rather than a multi-class taxonomy simplifies the model and reduces false-positive modes. This is a deliberate simplification for first-flight reliability.

These choices reflect constraints that sit outside the quantitative model — regulatory requirements, operational risk tolerance, and project timeline.